

# POWER Vector Library Manual

1.0.0

Generated by Doxygen 1.8.11



# Contents

<b>1</b>	<b>POWER Vector Library (pveclib)</b>	<b>1</b>
1.1	Notices	1
1.1.1	Reference Documentation	2
1.2	Rationale	2
1.2.1	POWER Vector Library Goals	3
1.2.1.1	POWER Vector Library Intrinsic headers	3
1.2.2	How pveclib is different from compiler vector built-ins	4
1.2.2.1	What can we do about this?	5
1.2.2.2	General Endian Issues	7
1.2.2.3	So what can the Power Vector Library project do?	8
1.2.2.4	Returning extended quadword results.	11
1.2.3	Background on the evolution of <altivec.h>	15
1.2.4	pveclib is not a vector math library	16
1.3	Performance data.	16
1.3.1	Additional analysis and tools.	20
<b>2</b>	<b>Todo List</b>	<b>23</b>
<b>3</b>	<b>Deprecated List</b>	<b>25</b>
<b>4</b>	<b>Class Index</b>	<b>27</b>
4.1	Class List	27
<b>5</b>	<b>File Index</b>	<b>29</b>
5.1	File List	29

<b>6</b>	<b>Class Documentation</b>	<b>31</b>
6.1	<code>__VEC_U_128</code> Union Reference	31
6.1.1	Detailed Description	32
6.2	<code>__VF_128</code> Union Reference	32
6.2.1	Detailed Description	32
<b>7</b>	<b>File Documentation</b>	<b>33</b>
7.1	<code>src/vec_bcd_ppc.h</code> File Reference	33
7.1.1	Detailed Description	33
7.1.2	Function Documentation	34
7.1.2.1	<code>bcd_qauntize0(_Decimal128 val)</code>	34
7.1.2.2	<code>vec_BCD2DFP(vui32_t val)</code>	34
7.1.2.3	<code>vec_bcdadd(vui32_t a, vui32_t b)</code>	34
7.1.2.4	<code>vec_bcdctb100s(vui8_t vra)</code>	35
7.1.2.5	<code>vec_bcddiv(vui32_t a, vui32_t b)</code>	35
7.1.2.6	<code>vec_bcdmul(vui32_t a, vui32_t b)</code>	35
7.1.2.7	<code>vec_bcdsub(vui32_t a, vui32_t b)</code>	36
7.1.2.8	<code>vec_DFP2BCD(_Decimal128 val)</code>	36
7.2	<code>src/vec_char_ppc.h</code> File Reference	36
7.2.1	Detailed Description	37
7.2.2	Endian problems with byte operations	38
7.2.3	Performance data.	39
7.2.3.1	More information.	39
7.2.4	Function Documentation	39
7.2.4.1	<code>vec_absdub(vui8_t vra, vui8_t vrb)</code>	39
7.2.4.2	<code>vec_clzb(vui8_t vra)</code>	39
7.2.4.3	<code>vec_isalnum(vui8_t vec_str)</code>	40
7.2.4.4	<code>vec_isalpha(vui8_t vec_str)</code>	40
7.2.4.5	<code>vec_isdigit(vui8_t vec_str)</code>	41
7.2.4.6	<code>vec_mrgahb(vui16_t vra, vui16_t vrb)</code>	41
7.2.4.7	<code>vec_mrgalb(vui16_t vra, vui16_t vrb)</code>	42

7.2.4.8	<a href="#">vec_mrgeb(vui8_t vra, vui8_t vrb)</a>	42
7.2.4.9	<a href="#">vec_mrgob(vui8_t vra, vui8_t vrb)</a>	43
7.2.4.10	<a href="#">vec_mulhsb(vi8_t vra, vi8_t vrb)</a>	43
7.2.4.11	<a href="#">vec_mulhub(vui8_t vra, vui8_t vrb)</a>	43
7.2.4.12	<a href="#">vec_mulubm(vui8_t vra, vui8_t vrb)</a>	44
7.2.4.13	<a href="#">vec_popcntb(vui8_t vra)</a>	44
7.2.4.14	<a href="#">vec_shift_leftdo(vui8_t vrw, vui8_t vrx, vui8_t vrb)</a>	46
7.2.4.15	<a href="#">vec_slbi(vui8_t vra, const unsigned int shb)</a>	46
7.2.4.16	<a href="#">vec_srabi(vi8_t vra, const unsigned int shb)</a>	47
7.2.4.17	<a href="#">vec_srbi(vui8_t vra, const unsigned int shb)</a>	47
7.2.4.18	<a href="#">vec_tolower(vui8_t vec_str)</a>	48
7.2.4.19	<a href="#">vec_toupper(vui8_t vec_str)</a>	48
7.2.4.20	<a href="#">vec_vmrgeb(vui8_t vra, vui8_t vrb)</a>	48
7.2.4.21	<a href="#">vec_vmrgob(vui8_t vra, vui8_t vrb)</a>	49
7.3	<a href="#">src/vec_common_ppc.h File Reference</a>	50
7.3.1	<a href="#">Detailed Description</a>	52
7.3.2	<a href="#">Consistent vector type naming</a>	53
7.3.3	<a href="#">Transferring 128-bit types</a>	53
7.3.4	<a href="#">Endian and vector constants</a>	54
7.4	<a href="#">src/vec_f128_ppc.h File Reference</a>	55
7.4.1	<a href="#">Detailed Description</a>	57
7.4.2	<a href="#">Examples</a>	58
7.4.3	<a href="#">Performance data</a>	59
7.4.4	<a href="#">Function Documentation</a>	59
7.4.4.1	<a href="#">vec_absf128(__binary128 f128)</a>	59
7.4.4.2	<a href="#">vec_all_isinftef128(__binary128 f128)</a>	59
7.4.4.3	<a href="#">vec_all_isinff128(__binary128 f128)</a>	60
7.4.4.4	<a href="#">vec_all_isnanf128(__binary128 f128)</a>	60
7.4.4.5	<a href="#">vec_all_isnormalf128(__binary128 f128)</a>	61
7.4.4.6	<a href="#">vec_all_issubnormalf128(__binary128 f128)</a>	61

7.4.4.7	<a href="#">vec_all_iszerof128(__binary128 f128)</a>	62
7.4.4.8	<a href="#">vec_const_huge_valf128()</a>	62
7.4.4.9	<a href="#">vec_const_inff128()</a>	62
7.4.4.10	<a href="#">vec_const_nanf128()</a>	63
7.4.4.11	<a href="#">vec_const_nansf128()</a>	63
7.4.4.12	<a href="#">vec_copysignf128(__binary128 f128x, __binary128 f128y)</a>	63
7.4.4.13	<a href="#">vec_isfinitef128(__binary128 f128)</a>	63
7.4.4.14	<a href="#">vec_isinf_signf128(__binary128 f128)</a>	64
7.4.4.15	<a href="#">vec_isinff128(__binary128 f128)</a>	64
7.4.4.16	<a href="#">vec_isnanf128(__binary128 f128)</a>	65
7.4.4.17	<a href="#">vec_isnormalf128(__binary128 f128)</a>	65
7.4.4.18	<a href="#">vec_issubnormalf128(__binary128 f128)</a>	66
7.4.4.19	<a href="#">vec_iszerof128(__binary128 f128)</a>	66
7.4.4.20	<a href="#">vec_setb_qp(__binary128 f128)</a>	67
7.4.4.21	<a href="#">vec_signbitf128(__binary128 f128)</a>	67
7.4.4.22	<a href="#">vec_xfer_bin128_2_vui128t(__binary128 f128)</a>	68
7.4.4.23	<a href="#">vec_xfer_bin128_2_vui16t(__binary128 f128)</a>	68
7.4.4.24	<a href="#">vec_xfer_bin128_2_vui32t(__binary128 f128)</a>	68
7.4.4.25	<a href="#">vec_xfer_bin128_2_vui64t(__binary128 f128)</a>	69
7.4.4.26	<a href="#">vec_xfer_bin128_2_vui8t(__binary128 f128)</a>	69
7.4.4.27	<a href="#">vec_xfer_vui128t_2_bin128(vui128_t f128)</a>	70
7.4.4.28	<a href="#">vec_xfer_vui16t_2_bin128(vui16_t f128)</a>	70
7.4.4.29	<a href="#">vec_xfer_vui32t_2_bin128(vui32_t f128)</a>	70
7.4.4.30	<a href="#">vec_xfer_vui64t_2_bin128(vui64_t f128)</a>	71
7.4.4.31	<a href="#">vec_xfer_vui8t_2_bin128(vui8_t f128)</a>	71
7.5	<a href="#">src/vec_f32_ppc.h File Reference</a>	72
7.5.1	<a href="#">Detailed Description</a>	73
7.5.2	<a href="#">Examples</a>	74
7.5.3	<a href="#">Performance data.</a>	75
7.5.4	<a href="#">Function Documentation</a>	75

7.5.4.1	<a href="#">vec_absf32(vf32_t vf32x)</a>	75
7.5.4.2	<a href="#">vec_all_isfinitef32(vf32_t vf32)</a>	75
7.5.4.3	<a href="#">vec_all_isinff32(vf32_t vf32)</a>	76
7.5.4.4	<a href="#">vec_all_isnanf32(vf32_t vf32)</a>	76
7.5.4.5	<a href="#">vec_all_isnormalf32(vf32_t vf32)</a>	77
7.5.4.6	<a href="#">vec_all_issubnormalf32(vf32_t vf32)</a>	77
7.5.4.7	<a href="#">vec_all_iszerof32(vf32_t vf32)</a>	78
7.5.4.8	<a href="#">vec_any_isfinitef32(vf32_t vf32)</a>	78
7.5.4.9	<a href="#">vec_any_isinff32(vf32_t vf32)</a>	79
7.5.4.10	<a href="#">vec_any_isnanf32(vf32_t vf32)</a>	79
7.5.4.11	<a href="#">vec_any_isnormalf32(vf32_t vf32)</a>	80
7.5.4.12	<a href="#">vec_any_issubnormalf32(vf32_t vf32)</a>	80
7.5.4.13	<a href="#">vec_any_iszerof32(vf32_t vf32)</a>	81
7.5.4.14	<a href="#">vec_copysignf32(vf32_t vf32x, vf32_t vf32y)</a>	81
7.5.4.15	<a href="#">vec_isfinitef32(vf32_t vf32)</a>	82
7.5.4.16	<a href="#">vec_isinff32(vf32_t vf32)</a>	82
7.5.4.17	<a href="#">vec_isnanf32(vf32_t vf32)</a>	83
7.5.4.18	<a href="#">vec_isnormalf32(vf32_t vf32)</a>	83
7.5.4.19	<a href="#">vec_issubnormalf32(vf32_t vf32)</a>	84
7.5.4.20	<a href="#">vec_iszerof32(vf32_t vf32)</a>	84
7.6	<a href="#">src/vec_f64_ppc.h File Reference</a>	85
7.6.1	<a href="#">Detailed Description</a>	86
7.6.2	<a href="#">Examples</a>	87
7.6.3	<a href="#">Performance data.</a>	88
7.6.4	<a href="#">Function Documentation</a>	88
7.6.4.1	<a href="#">vec_absf64(vf64_t vf64x)</a>	88
7.6.4.2	<a href="#">vec_all_isfinitef64(vf64_t vf64)</a>	88
7.6.4.3	<a href="#">vec_all_isinff64(vf64_t vf64)</a>	89
7.6.4.4	<a href="#">vec_all_isnanf64(vf64_t vf64)</a>	89
7.6.4.5	<a href="#">vec_all_isnormalf64(vf64_t vf64)</a>	90

7.6.4.6	<code>vec_all_issubnormalf64(vf64_t vf64)</code>	90
7.6.4.7	<code>vec_all_iszerof64(vf64_t vf64)</code>	91
7.6.4.8	<code>vec_any_isfinitef64(vf64_t vf64)</code>	91
7.6.4.9	<code>vec_any_isinff64(vf64_t vf64)</code>	92
7.6.4.10	<code>vec_any_isnanf64(vf64_t vf64)</code>	92
7.6.4.11	<code>vec_any_isnormalf64(vf64_t vf64)</code>	93
7.6.4.12	<code>vec_any_issubnormalf64(vf64_t vf64)</code>	93
7.6.4.13	<code>vec_any_iszerof64(vf64_t vf64)</code>	94
7.6.4.14	<code>vec_copysignf64(vf64_t vf64x, vf64_t vf64y)</code>	94
7.6.4.15	<code>vec_isfinitef64(vf64_t vf64)</code>	95
7.6.4.16	<code>vec_isinff64(vf64_t vf64)</code>	95
7.6.4.17	<code>vec_isnanf64(vf64_t vf64)</code>	96
7.6.4.18	<code>vec_isnormalf64(vf64_t vf64)</code>	96
7.6.4.19	<code>vec_issubnormalf64(vf64_t vf64)</code>	97
7.6.4.20	<code>vec_iszerof64(vf64_t vf64)</code>	97
7.6.4.21	<code>vec_pack_longdouble(vf64_t lval)</code>	98
7.6.4.22	<code>vec_unpack_longdouble(long double lval)</code>	98
7.7	<code>src/vec_int128_ppc.h</code> File Reference	98
7.7.1	Detailed Description	101
7.7.2	Performance data.	102
7.7.3	Function Documentation	102
7.7.3.1	<code>vec_addcq(vui128_t *cout, vui128_t a, vui128_t b)</code>	102
7.7.3.2	<code>vec_addcuq(vui128_t a, vui128_t b)</code>	103
7.7.3.3	<code>vec_addecuq(vui128_t a, vui128_t b, vui128_t ci)</code>	103
7.7.3.4	<code>vec_addeq(vui128_t *cout, vui128_t a, vui128_t b, vui128_t ci)</code>	104
7.7.3.5	<code>vec_addeuqm(vui128_t a, vui128_t b, vui128_t ci)</code>	104
7.7.3.6	<code>vec_adduqm(vui128_t a, vui128_t b)</code>	105
7.7.3.7	<code>vec_clzq(vui128_t vra)</code>	105
7.7.3.8	<code>vec_cmpeqsq(vi128_t vra, vi128_t vrb)</code>	106
7.7.3.9	<code>vec_cmpequq(vui128_t vra, vui128_t vrb)</code>	106



7.7.3.10	<code>vec_cmpgesq(vi128_t vra, vi128_t vrb)</code>	106
7.7.3.11	<code>vec_cmpgeuq(vui128_t vra, vui128_t vrb)</code>	107
7.7.3.12	<code>vec_cmpgtsq(vi128_t vra, vi128_t vrb)</code>	107
7.7.3.13	<code>vec_cmpgtuq(vui128_t vra, vui128_t vrb)</code>	108
7.7.3.14	<code>vec_cmplesq(vi128_t vra, vi128_t vrb)</code>	108
7.7.3.15	<code>vec_cmpleuq(vui128_t vra, vui128_t vrb)</code>	109
7.7.3.16	<code>vec_cmpltsg(vi128_t vra, vi128_t vrb)</code>	109
7.7.3.17	<code>vec_cmpltuq(vui128_t vra, vui128_t vrb)</code>	110
7.7.3.18	<code>vec_cmpnesq(vi128_t vra, vi128_t vrb)</code>	110
7.7.3.19	<code>vec_cmpneuq(vui128_t vra, vui128_t vrb)</code>	111
7.7.3.20	<code>vec_cmpsq_all_eq(vi128_t vra, vi128_t vrb)</code>	111
7.7.3.21	<code>vec_cmpsq_all_ge(vi128_t vra, vi128_t vrb)</code>	112
7.7.3.22	<code>vec_cmpsq_all_gt(vi128_t vra, vi128_t vrb)</code>	112
7.7.3.23	<code>vec_cmpsq_all_le(vi128_t vra, vi128_t vrb)</code>	113
7.7.3.24	<code>vec_cmpsq_all_lt(vi128_t vra, vi128_t vrb)</code>	113
7.7.3.25	<code>vec_cmpsq_all_ne(vi128_t vra, vi128_t vrb)</code>	113
7.7.3.26	<code>vec_cmpuq_all_eq(vui128_t vra, vui128_t vrb)</code>	114
7.7.3.27	<code>vec_cmpuq_all_ge(vui128_t vra, vui128_t vrb)</code>	114
7.7.3.28	<code>vec_cmpuq_all_gt(vui128_t vra, vui128_t vrb)</code>	115
7.7.3.29	<code>vec_cmpuq_all_le(vui128_t vra, vui128_t vrb)</code>	115
7.7.3.30	<code>vec_cmpuq_all_lt(vui128_t vra, vui128_t vrb)</code>	115
7.7.3.31	<code>vec_cmpuq_all_ne(vui128_t vra, vui128_t vrb)</code>	116
7.7.3.32	<code>vec_cmul100cuq(vui128_t *cout, vui128_t a)</code>	116
7.7.3.33	<code>vec_cmul100ecuq(vui128_t *cout, vui128_t a, vui128_t cin)</code>	117
7.7.3.34	<code>vec_cmul10cuq(vui128_t *cout, vui128_t a)</code>	117
7.7.3.35	<code>vec_cmul10ecuq(vui128_t *cout, vui128_t a, vui128_t cin)</code>	118
7.7.3.36	<code>vec_msumudm(vui64_t a, vui64_t b, vui128_t c)</code>	118
7.7.3.37	<code>vec_mul10cuq(vui128_t a)</code>	119
7.7.3.38	<code>vec_mul10ecuq(vui128_t a, vui128_t cin)</code>	119
7.7.3.39	<code>vec_mul10euq(vui128_t a, vui128_t cin)</code>	119

7.7.3.40	<code>vec_mul10uq(vui128_t a)</code>	120
7.7.3.41	<code>vec_muleud(vui64_t a, vui64_t b)</code>	120
7.7.3.42	<code>vec_mulhud(vui64_t vra, vui64_t vrb)</code>	121
7.7.3.43	<code>vec_mulluq(vui128_t a, vui128_t b)</code>	121
7.7.3.44	<code>vec_muloud(vui64_t a, vui64_t b)</code>	122
7.7.3.45	<code>vec_muludm(vui64_t vra, vui64_t vrb)</code>	122
7.7.3.46	<code>vec_muludq(vui128_t *mulu, vui128_t a, vui128_t b)</code>	123
7.7.3.47	<code>vec_popcntq(vui128_t vra)</code>	123
7.7.3.48	<code>vec_revbq(vui128_t vra)</code>	124
7.7.3.49	<code>vec_setb_cyq(vui128_t vcy)</code>	124
7.7.3.50	<code>vec_setb_ncq(vui128_t vcy)</code>	124
7.7.3.51	<code>vec_setb_sq(vui128_t vra)</code>	125
7.7.3.52	<code>vec_sldq(vui128_t vrw, vui128_t vrx, vui128_t vrb)</code>	125
7.7.3.53	<code>vec_slq(vui128_t vra, vui128_t vrb)</code>	126
7.7.3.54	<code>vec_slq4(vui128_t vra)</code>	126
7.7.3.55	<code>vec_slq5(vui128_t vra)</code>	126
7.7.3.56	<code>vec_slqi(vui128_t vra, const unsigned int shb)</code>	127
7.7.3.57	<code>vec_srq(vui128_t vra, vui128_t vrb)</code>	127
7.7.3.58	<code>vec_srq4(vui128_t vra)</code>	128
7.7.3.59	<code>vec_srq5(vui128_t vra)</code>	128
7.7.3.60	<code>vec_srqi(vui128_t vra, const unsigned int shb)</code>	128
7.7.3.61	<code>vec_subcuq(vui128_t vra, vui128_t vrb)</code>	129
7.7.3.62	<code>vec_subecuq(vui128_t vra, vui128_t vrb, vui128_t vrc)</code>	129
7.7.3.63	<code>vec_subeuqm(vui128_t vra, vui128_t vrb, vui128_t vrc)</code>	129
7.7.3.64	<code>vec_subuqm(vui128_t vra, vui128_t vrb)</code>	131
7.7.3.65	<code>vec_vmuleud(vui64_t a, vui64_t b)</code>	131
7.7.3.66	<code>vec_vmuloud(vui64_t a, vui64_t b)</code>	132
7.8	<code>src/vec_int16_ppc.h</code> File Reference	132
7.8.1	Detailed Description	133
7.8.2	Endian problems with halfword operations	134

7.8.2.1	Multiply High Unsigned Halfword Example . . . . .	137
7.8.3	Examples, Divide by integer constant . . . . .	137
7.8.3.1	Divide by constant 10 examples . . . . .	138
7.8.3.2	Divide by constant 10000 example . . . . .	139
7.8.4	Performance data. . . . .	139
7.8.4.1	More information. . . . .	140
7.8.5	Function Documentation . . . . .	140
7.8.5.1	vec_absduh(vui16_t vra, vui16_t vrb) . . . . .	140
7.8.5.2	vec_clzh(vui16_t vra) . . . . .	141
7.8.5.3	vec_mrgahh(vui32_t vra, vui32_t vrb) . . . . .	141
7.8.5.4	vec_mrgalh(vui32_t vra, vui32_t vrb) . . . . .	142
7.8.5.5	vec_mrgeh(vui16_t vra, vui16_t vrb) . . . . .	142
7.8.5.6	vec_mrgoh(vui16_t vra, vui16_t vrb) . . . . .	143
7.8.5.7	vec_mulhsh(vi16_t vra, vi16_t vrb) . . . . .	143
7.8.5.8	vec_mulhuh(vui16_t vra, vui16_t vrb) . . . . .	143
7.8.5.9	vec_muluhm(vui16_t vra, vui16_t vrb) . . . . .	144
7.8.5.10	vec_popcnth(vui16_t vra) . . . . .	144
7.8.5.11	vec_revbh(vui16_t vra) . . . . .	145
7.8.5.12	vec_slhi(vui16_t vra, const unsigned int shb) . . . . .	145
7.8.5.13	vec_srahi(vi16_t vra, const unsigned int shb) . . . . .	146
7.8.5.14	vec_srhi(vui16_t vra, const unsigned int shb) . . . . .	146
7.8.5.15	vec_vmrgeh(vui16_t vra, vui16_t vrb) . . . . .	147
7.8.5.16	vec_vmrgoh(vui16_t vra, vui16_t vrb) . . . . .	147
7.9	src/vec_int32_ppc.h File Reference . . . . .	148
7.9.1	Detailed Description . . . . .	149
7.9.2	Endian problems with word operations . . . . .	150
7.9.2.1	Vector Merge Algebraic High Word example . . . . .	151
7.9.2.2	Vector Multiply High Unsigned Word example . . . . .	152
7.9.3	Vector Word Examples . . . . .	152
7.9.3.1	Vectorized TimeBase conversion example . . . . .	154

7.9.4	Performance data. . . . .	154
7.9.5	Function Documentation . . . . .	154
7.9.5.1	vec_absduw(vui32_t vra, vui32_t vrb) . . . . .	154
7.9.5.2	vec_clzw(vui32_t vra) . . . . .	155
7.9.5.3	vec_mrgahw(vui64_t vra, vui64_t vrb) . . . . .	155
7.9.5.4	vec_mrgalw(vui64_t vra, vui64_t vrb) . . . . .	156
7.9.5.5	vec_mrgew(vui32_t vra, vui32_t vrb) . . . . .	156
7.9.5.6	vec_mrgow(vui32_t vra, vui32_t vrb) . . . . .	157
7.9.5.7	vec_mulesw(vi32_t a, vi32_t b) . . . . .	157
7.9.5.8	vec_muleuw(vui32_t a, vui32_t b) . . . . .	158
7.9.5.9	vec_mulhsw(vi32_t vra, vi32_t vrb) . . . . .	158
7.9.5.10	vec_mulhuw(vui32_t vra, vui32_t vrb) . . . . .	159
7.9.5.11	vec_mulosw(vi32_t a, vi32_t b) . . . . .	159
7.9.5.12	vec_mulouw(vui32_t a, vui32_t b) . . . . .	160
7.9.5.13	vec_muluwm(vui32_t a, vui32_t b) . . . . .	160
7.9.5.14	vec_popcntw(vui32_t vra) . . . . .	161
7.9.5.15	vec_revbw(vui32_t vra) . . . . .	161
7.9.5.16	vec_slwi(vui32_t vra, const unsigned int shb) . . . . .	161
7.9.5.17	vec_srawi(vi32_t vra, const unsigned int shb) . . . . .	162
7.9.5.18	vec_srwi(vui32_t vra, const unsigned int shb) . . . . .	162
7.10	src/vec_int64_ppc.h File Reference . . . . .	163
7.10.1	Detailed Description . . . . .	166
7.10.2	Some missing doubleword operations . . . . .	167
7.10.2.1	Challenges and opportunities . . . . .	170
7.10.2.2	More Challenges . . . . .	171
7.10.3	Endian problems with doubleword operations . . . . .	175
7.10.4	Vector Doubleword Examples . . . . .	176
7.10.4.1	Vectorized 64-bit TimeBase conversion example . . . . .	177
7.10.5	Performance data. . . . .	178
7.10.6	Function Documentation . . . . .	178

7.10.6.1	<code>vec_absdud(vui64_t vra, vui64_t vrb)</code>	178
7.10.6.2	<code>vec_addudm(vui64_t a, vui64_t b)</code>	179
7.10.6.3	<code>vec_clzd(vui64_t vra)</code>	179
7.10.6.4	<code>vec_cmpeqsd(vi64_t a, vi64_t b)</code>	179
7.10.6.5	<code>vec_cmpequd(vui64_t a, vui64_t b)</code>	180
7.10.6.6	<code>vec_cmpgesd(vi64_t a, vi64_t b)</code>	180
7.10.6.7	<code>vec_cmpgeud(vui64_t a, vui64_t b)</code>	181
7.10.6.8	<code>vec_cmpgtsd(vi64_t a, vi64_t b)</code>	181
7.10.6.9	<code>vec_cmpgtud(vui64_t a, vui64_t b)</code>	182
7.10.6.10	<code>vec_cmplepd(vi64_t a, vi64_t b)</code>	182
7.10.6.11	<code>vec_cmpleud(vui64_t a, vui64_t b)</code>	183
7.10.6.12	<code>vec_cmpltsd(vi64_t a, vi64_t b)</code>	183
7.10.6.13	<code>vec_cmpltud(vui64_t a, vui64_t b)</code>	183
7.10.6.14	<code>vec_cmpnesd(vi64_t a, vi64_t b)</code>	184
7.10.6.15	<code>vec_cmpneud(vui64_t a, vui64_t b)</code>	184
7.10.6.16	<code>vec_cmpsd_all_eq(vi64_t a, vi64_t b)</code>	185
7.10.6.17	<code>vec_cmpsd_all_ge(vi64_t a, vi64_t b)</code>	185
7.10.6.18	<code>vec_cmpsd_all_gt(vi64_t a, vi64_t b)</code>	186
7.10.6.19	<code>vec_cmpsd_all_le(vi64_t a, vi64_t b)</code>	186
7.10.6.20	<code>vec_cmpsd_all_lt(vi64_t a, vi64_t b)</code>	186
7.10.6.21	<code>vec_cmpsd_all_ne(vi64_t a, vi64_t b)</code>	187
7.10.6.22	<code>vec_cmpsd_any_eq(vi64_t a, vi64_t b)</code>	187
7.10.6.23	<code>vec_cmpsd_any_ge(vi64_t a, vi64_t b)</code>	188
7.10.6.24	<code>vec_cmpsd_any_gt(vi64_t a, vi64_t b)</code>	188
7.10.6.25	<code>vec_cmpsd_any_le(vi64_t a, vi64_t b)</code>	189
7.10.6.26	<code>vec_cmpsd_any_lt(vi64_t a, vi64_t b)</code>	189
7.10.6.27	<code>vec_cmpsd_any_ne(vi64_t a, vi64_t b)</code>	189
7.10.6.28	<code>vec_cmpud_all_eq(vui64_t a, vui64_t b)</code>	190
7.10.6.29	<code>vec_cmpud_all_ge(vui64_t a, vui64_t b)</code>	190
7.10.6.30	<code>vec_cmpud_all_gt(vui64_t a, vui64_t b)</code>	191

7.10.6.31	<code>vec_cmpud_all_le(vui64_t a, vui64_t b)</code>	191
7.10.6.32	<code>vec_cmpud_all_lt(vui64_t a, vui64_t b)</code>	191
7.10.6.33	<code>vec_cmpud_all_ne(vui64_t a, vui64_t b)</code>	192
7.10.6.34	<code>vec_cmpud_any_eq(vui64_t a, vui64_t b)</code>	192
7.10.6.35	<code>vec_cmpud_any_ge(vui64_t a, vui64_t b)</code>	193
7.10.6.36	<code>vec_cmpud_any_gt(vui64_t a, vui64_t b)</code>	193
7.10.6.37	<code>vec_cmpud_any_le(vui64_t a, vui64_t b)</code>	194
7.10.6.38	<code>vec_cmpud_any_lt(vui64_t a, vui64_t b)</code>	194
7.10.6.39	<code>vec_cmpud_any_ne(vui64_t a, vui64_t b)</code>	194
7.10.6.40	<code>vec_maxsd(vi64_t vra, vi64_t vrb)</code>	195
7.10.6.41	<code>vec_maxud(vui64_t vra, vui64_t vrb)</code>	195
7.10.6.42	<code>vec_minsd(vi64_t vra, vi64_t vrb)</code>	196
7.10.6.43	<code>vec_minud(vui64_t vra, vui64_t vrb)</code>	196
7.10.6.44	<code>vec_mrgahd(vui128_t vra, vui128_t vrb)</code>	196
7.10.6.45	<code>vec_mrgald(vui128_t vra, vui128_t vrb)</code>	197
7.10.6.46	<code>vec_mrged(vui64_t __VA, vui64_t __VB)</code>	197
7.10.6.47	<code>vec_mrghd(vui64_t __VA, vui64_t __VB)</code>	198
7.10.6.48	<code>vec_mrgld(vui64_t __VA, vui64_t __VB)</code>	198
7.10.6.49	<code>vec_mrgod(vui64_t __VA, vui64_t __VB)</code>	199
7.10.6.50	<code>vec_pasted(vui64_t __VH, vui64_t __VL)</code>	199
7.10.6.51	<code>vec_permdi(vui64_t vra, vui64_t vrb, const int ctl)</code>	199
7.10.6.52	<code>vec_popcntd(vui64_t vra)</code>	200
7.10.6.53	<code>vec_revbd(vui64_t vra)</code>	201
7.10.6.54	<code>vec_rldi(vui64_t vra, const unsigned int shb)</code>	201
7.10.6.55	<code>vec_sldi(vui64_t vra, const unsigned int shb)</code>	201
7.10.6.56	<code>vec_splatd(vui64_t vra, const int ctl)</code>	202
7.10.6.57	<code>vec_spltd(vui64_t vra, const int ctl)</code>	202
7.10.6.58	<code>vec_sradi(vi64_t vra, const unsigned int shb)</code>	203
7.10.6.59	<code>vec_srdi(vui64_t vra, const unsigned int shb)</code>	203
7.10.6.60	<code>vec_subudm(vui64_t a, vui64_t b)</code>	204
7.10.6.61	<code>vec_swapd(vui64_t vra)</code>	204
7.10.6.62	<code>vec_vmsumuwmm(vui32_t vra, vui32_t vrb, vui64_t vrc)</code>	205
7.10.6.63	<code>vec_vpkudum(vui64_t vra, vui64_t vrb)</code>	205
7.10.6.64	<code>vec_vrld(vui64_t vra, vui64_t vrb)</code>	206
7.10.6.65	<code>vec_vslld(vui64_t vra, vui64_t vrb)</code>	206
7.10.6.66	<code>vec_vsradi(vi64_t vra, vui64_t vrb)</code>	207
7.10.6.67	<code>vec_vsrdd(vui64_t vra, vui64_t vrb)</code>	207
7.10.6.68	<code>vec_xxspltd(vui64_t vra, const int ctl)</code>	208

# Chapter 1

## POWER Vector Library (pveclib)

A library of useful vector functions for POWER. This library fills in the gap between the instructions defined in the POWER Instruction Set Architecture (**PowerISA**) and higher level library APIs. The intent is to improve the productivity of application developers who need to optimize their applications or dependent libraries for POWER.

### Authors

Steven Munroe

### Copyright

2017-2018 IBM Corporation. Licensed under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at: <http://www.apache.org/licenses/LICENSE-2.0>.

Unless required by applicable law or agreed to in writing, software and documentation distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

## 1.1 Notices

IBM, the IBM logo, and ibm.com are trademarks or registered trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the Web at "Copyright and trademark information" at <http://www.ibm.com/legal/copytrade.shtml>.

The following terms are trademarks or registered trademarks licensed by Power.org in the United States and/or other countries: Power ISA™, Power Architecture™. Information on the list of U.S. trademarks licensed by Power.org may be found at <http://www.power.org/about/brand-center/>.

The following terms are trademarks or registered trademarks of Freescale Semiconductor in the United States and/or other countries: AltiVec™. Information on the list of U.S. trademarks owned by Freescale Semiconductor may be found at [http://www.freescale.com/files/abstract/help\\_page/TERMSOFUSE.html](http://www.freescale.com/files/abstract/help_page/TERMSOFUSE.html).

### 1.1.1 Reference Documentation

- Power Instruction Set Architecture, Versions 2.06B, 2.07B and 3.0B, IBM, 2010-2017. <http://www.power.org> and more recently from <https://www-355.ibm.com/systems/power/openpower/>
- ALTIVEC PIM: AltiVec™ Technology Programming Interface Manual, Freescale Semiconductor, 1999. [http://www.freescale.com/files/32bit/doc/ref\\_manual/ALTIVECPIM.pdf](http://www.freescale.com/files/32bit/doc/ref_manual/ALTIVECPIM.pdf)
- 64-bit PowerPC ELF Application Binary Interface Supplement 1.9. <http://refspecs.linuxfoundation.org/ELF/ppc64/PPC-elf64abi.html>
- OpenPOWER ELF V2 application binary interface (ABI), OpenPOWER Foundation, 2017. [https://openpowerfoundation.org/?resource\\_lib=64-bit-elf-v2-abi-specification-power-archi](https://openpowerfoundation.org/?resource_lib=64-bit-elf-v2-abi-specification-power-archi)
- Using the GNU Compiler Collection (GCC), Free Software Foundation, 1988-2018. <https://gcc.gnu.org/onlinedocs/>
- POWER8 Processor User's Manual for the Single-Chip Module <https://ibm.ent.box.com/s/649rlau0zjcc0yrulqf4cgx5wk3pgbfbk>
- POWER9 Processor User's Manual <https://ibm.ent.box.com/s/8uj02ysel62meji4voujw29wwkhsz6a4>
- Warren, Henry S. Jr, Hacker's Delight, 2nd Edition, Upper Saddle River, NJ: Addison Wesley, 2013.

## 1.2 Rationale

The C/C++ language compilers (that support PowerISA) may implement vector intrinsic functions (compiler built-ins as embodied by `altivec.h`). These vector intrinsics offer an alternative to assembler programming, but do little to reduce the complexity of the underlying PowerISA. Higher level vector intrinsic operations are needed to improve productivity and encourage developers to optimize their applications for PowerISA. Another key goal is to smooth over the complexity of the evolving PowerISA and compiler support.

For example: the PowerISA 2.07 (POWER8) provides population count and count leading zero operations on vectors of byte, halfword, word, and doubleword elements but not on the whole vector as a `__int128` value. Before PowerISA 2.07, neither operation was supported, for any element size.

Another example: The original **Altivec** (AKA Vector Multimedia Extension (**VMX**)) provided Vector Multiply Odd / Even operations for signed / unsigned byte and halfword elements. The PowerISA 2.07 added Vector Multiply Even/Odd operations for signed / unsigned word elements. This release also added a Vector Multiply Unsigned Word Modulo operation. This was important to allow auto vectorization of C loops using 32-bit (int) multiply.

But PowerISA 2.07 did not add support for doubleword or quadword (`__int128`) multiply directly. Nor did it fill in the missing multiply modulo operations for byte and halfword. However it did add support for doubleword and quadword add / subtract modulo. This can be helpful, if you are willing to apply grade school arithmetic (add, carry the 1) to vector elements.

PowerISA 3.0 (POWER9) did add a Vector Multiply-Sum Unsigned Doubleword Modulo operation. With this instruction (and a generated vector of zeros as input) you can effectively implement the simple doubleword integer multiply modulo operation in a few instructions. Similarly for Vector Multiply-Sum Unsigned Halfword Modulo. But this may not be obvious.

This history embodies a set of trade-offs negotiated between the Software and Processor design architects at specific points in time. But most programmers would prefer to use a set of operators applied across the supported element types and sizes.



### 1.2.1 POWER Vector Library Goals

Obviously many useful operations can be constructed from existing PowerISA operations and GCC `<altivec.h>` built-ins but the implementation may not be obvious. The optimum sequence will vary across the PowerISA levels as new instructions are added. And finally the compiler's built-in support for new PowerISA instructions evolves with the compiler's release cycle.

So the goal of this project is to provide well crafted implementations of useful vector and large number operations.

- Provide equivalent functions across versions of the PowerISA. This includes some of the most useful vector instructions added to POWER9 (PowerISA 3.0B). Many of these operations can be implemented as inline function in a few vector instructions on earlier PowerISA versions.
- Provide equivalent functions across versions of the compiler. For example built-ins provided in later versions of the compiler can be implemented as inline functions with inline asm in earlier compiler versions.
- Provide complete arithmetic operations across supported C types. For example multiply modulo and even/odd for int, long, and `__int128`.
- Provide complete extended arithmetic (carry / extend / multiple high) operations across supported C types. For example add / subtract with carry and extend for int, long, and `__int128`.
- Provide higher order functions not provided directly by the PowerISA. For example vector SIMD implementation for ASCII `__isalpha`, etc. As another example full `__int128` implementations of Count Leading Zeros, Population Count, Shift left/right immediate, and integer divide.
- Such implementations should be small enough to inline and allow the compiler opportunity to apply common optimization techniques.

#### 1.2.1.1 POWER Vector Library Intrinsic headers

The POWER Vector Library will be primarily delivered as C language inline functions in headers files.

- [vec\\_common\\_ppc.h](#) Typedefs and helper macros
- [vec\\_int128\\_ppc.h](#) Operations on vector `__int128` values
- [vec\\_int64\\_ppc.h](#) Operations on vector long int (64-bit) values
- [vec\\_int32\\_ppc.h](#) Operations on vector int (32-bit) values
- [vec\\_int16\\_ppc.h](#) Operations on vector short int (16-bit) values
- [vec\\_char\\_ppc.h](#) Operations on vector char (values) values
- [vec\\_bcd\\_ppc.h](#) Operations on vectors of Binary Code Decimal and Zoned Decimal values
- [vec\\_f128\\_ppc.h](#) Operations on vector `_Float128` values
- [vec\\_f64\\_ppc.h](#) Operations on vector double values
- [vec\\_f32\\_ppc.h](#) Operations on vector float values

**Note**

The list above is complete in the current public github as a first pass. A backlog of functions remain to be implemented across these headers. Development continues while we work on the backlog listed in: [Issue #13 TODOs](#)

The goal is to provide high quality implementations that adapt to the specifics of the compile target (`-mcpu=`) and compiler (`<altivec.h>`) version you are using. Initially pveclib will focus on the GCC compiler and `-mcpu=[power7|power8|power9]` for Linux. Testing will focus on Little Endian (**powerpc64le** for power8 and power9 targets. Any testing for Big Endian (**powerpc64** will be initially restricted to power7 and power8 targets.

Expanding pveclib support beyond this list to include:

- additional compilers (ie Clang)
- additional PPC platforms (970, power6, ...)
- Larger functions that just happen to use vector registers (Checksum, Crypto, compress/decompress, lower precision neural networks, ...)

will largely depend on additional skilled practitioners joining this project and contributing (code and platform testing) on a sustained basis.

### 1.2.2 How pveclib is different from compiler vector built-ins

The PowerPC vector built-ins evolved from the original [AltiVec \(TM\) Technology Programming Interface Manual](#) (PIM). The PIM defined the minimal extensions to the application binary interface (ABI) required to support the Vector Facility. This included new keywords (`vector`, `pixel`, `bool`) for defining new vector types, and new operators (built-in functions) required for any supporting and compliant C language compiler.

The vector built-in function support included:

- generic AltiVec operations, like `vec_add()`
- specific AltiVec operations (instructions, like `vec_vaddubm()`)
- predicates computed from AltiVec operations, like `vec_all_eq()` which are also generic

See [Background on the evolution of <altivec.h>](#) for more details.

There are clear advantages with the compiler implementing the vector operations as built-ins:

- The compiler can access the C language type information and vector extensions to implement the function overloading required to process generic operations.
- Built-ins can be generated inline, which eliminates function call overhead and allows more compact code generation.
- The compiler can then apply higher order optimization across built-ins including: Local and global register allocation. Global common subexpression elimination. Loop-invariant code motion.
- The compiler can automatically select the best instructions for the *target* processor ISA level (from the `-mcpu` compiler option).

While this is an improvement over writing assembler code, it does not provide much function beyond the specific operations specified in the PowerISA.

Another issue is that generic operations were not uniformly applicable across vector types. For example:

- `vec_add` / `vec_sub` applied to float, int, short and char.
- Later compilers added support for double (with POWER7 and the Vector Scalar Extensions (VSX) facility)
- Integer long (64-bit) and `__int128` support for POWER8 (PowerISA 2.07B).

But `vec_mul` / `vec_div` did not:

- `vec_mul` applied to float (and later double, with POWER7 VSX).
- `vec_mule` / `vec_mulo` (Multiply even / odd elements) applied to [signed | unsigned] integer short and char. Later compilers added support for vector int after POWER8 added vector multiply word instructions.
- `vec_div` was not included in the original PIM as AltiVec (VMX) only included vector reciprocal estimate for float and no vector integer divide for any size. Later compilers added support for `vec_div` float / double after POWER7 (VSX) added vector divide single/double-precision instructions.

#### Note

While the processor you (plan to) use, may support the specific instructions you want to exploit, the compiler you are using may not support, the generic or specific vector operations, for the element size/types, you want to use. This is common for GCC versions installed by "Enterprise Linux" distributions. They tend to freeze the GCC version early and maintain that GCC version for long term stability. One solution is to use the [IBM Advance toolchain for Linux on Power](#) (AT). AT is free for download and new AT versions are released yearly (usually in August) with the latest stable GCC from that spring.

This all can be very frustrating, at minimum, or even a show stopper, if you are on a tight schedule for your project. Especially if you are not familiar with the evolving history of the PowerISA and supporting compilers.

#### 1.2.2.1 What can we do about this?

First the Binutils assembler is usually updated within weeks of the public release of the PowerISA document. So while your compiler may not support the latest vector operations as built-in operations, an older compiler with an updated assembler, may support the instructions as inline assembler.

Sequences of inline assembler instructions can be wrapped within C language static inline functions and placed in a header files for shared use. If you are careful with the input / output register *constraints* the GCC compiler can provide local register allocation and minimize parameter marshaling overhead. This is very close (in function) to a specific AltiVec (built-in) operation.

## Note

Using GCC's inline assembler can be challenging even for the experienced programmer. The register constraints have grown in complexity as new facilities and categories were added. The fact that some (VMX) instructions are restricted to the original 32 Vector Registers (**VRs**) (the high half of the Vector-Scalar Registers **VSRs**), while others (Binary and Decimal Floating-Point) are restricted to the original 32 Floating-Point Registers (**FPRs** (overlapping the low half of the VSRs), and the new VSX instructions can access all 64 VSRs, is just one source of complexity. So it is very important to get your input/output constraints correct if you want inline assembler code to work correctly.

In-line assembler should be reserved for the first implementation using the latest PowerISA. Where possible you should use existing vector built-ins to implement specific operations for wider element types, support older hardware, or higher order operations. Again wrapping these implementations in static inline functions for collection in header files for reuse and distribution is recommended.

The PowerISA vector facility has all the instructions you need to implement extended precision operations for add, subtract, and multiply. Add / subtract with carry-out and permute or double vector shift and grade-school arithmetic is all you need.

For example the Vector Add Unsigned Quadword Modulo introduced in POWER8 (PowerISA 2.07B) can be implemented for POWER7 and earlier machines in 10-11 instructions. This uses a combination of Vector Add Unsigned Word Modulo (`vadduwm`), Vector Add and Write Carry-Out Unsigned Word (`vaddcuw`), and Vector Shift Left Double by Octet Immediate (`vsldoi`), to propagate the word carries through the quadword.

For POWER8 and later, C vector integer (modulo) multiply can be implemented in a single Vector Unsigned Word Modulo (`vmuluwm`) instruction. This was added explicitly to address vectorizing loops using `int` multiply in C language code. And some newer compilers do support generic `vec_mul()` for vector `int`. But this is not documented. Similarly for `char` (byte) and `short` (halfword) elements.

POWER8 also introduced Vector Multiply Even Signed Word (`vmulesw`) and Vector Multiply Odd Signed Word (`vmulosw`) instructions. So you would expect the generic `vec_mule` and `vec_mulo` operations to be extended to support *vector int*, as these operations have long been supported for `char` and `short`. Sadly this is not supported as of GCC 7.3. We hope to see this implemented for GCC 8.

So what will the compiler do for vector multiply `int` (modulo, even, or odd) for targeting power7? Older compilers will reject this as a *invalid parameter combination* .... A newer compiler may implement the equivalent function in a short sequence of VMX instructions from PowerISA 2.06 or earlier. And GCC 7.3 does support `vec_mul` for element types `char`, `short`, and `int`. These sequences are in the 2-7 instruction range depending on the operation and element type. This includes some constant loads and permute control vectors that can be factored and reused across operations.

Once the pattern is understood it is not hard to write equivalent sequences using operations from the original `<altivec.h>`. With a little care these sequences will be compatible with older compilers and older PowerISA versions. These concepts can be extended to operations that PowerISA and the compiler does not support yet. For example; a processor that may not have multiply even/odd/modulo of the required width (word, doubleword, or quadword). This might take 10-12 instructions to implement the next element size bigger than the current processor. A full 128-bit by 128-bit multiply with 256-bit result only requires 32 instructions on a POWER8 (using multiple word even/odd).

Also many of the operations missing from the vector facility, exist in the Fixed-point, Floating-point, or Decimal Floating-point scalar facilities. There will be some loss of efficiency in the data transfer but compared to a complex operation like divide or decimal conversions, this can be a workable solution. On older POWER processors (before power7/8) transfers between register banks (GPR, FPR, VR) had to go through memory. But with the VSX facility (POWER7) FPRs and VRs overlap with the lower and upper halves of the 64 VSR registers. So FPR <-> VSR transfer are 0-2 cycles latency. And with power8 we have direct transfer (GPR <-> FPR | VR | VSR) instructions in the 4-5 cycle latency range.

For example POWER8 added Binary Coded Decimal (**BCD**) add/subtract for signed 31 digit vector values. The vector unit does not support BCD multiply / divide. But the Decimal Floating-Point (**DFP**) facility (introduced with PowerISA 2.05 and Power6) supports up to 34-digit (`__Decimal128`) precision and all the expected (add/subtract/multiply/divide/...) arithmetic operations. DFP also supports conversion to/from 31-digit BCD and `__Decimal128` precision. This is all supported with a hardware Decimal Floating-Point Unit (**DFU**).

So `bcd_add` / `bcd_sub` can be generated as a single instruction on POWER8 and later, and 10-11 instructions for Power6/7. This count include the VSR <-> FPRp transfers, BCD <-> DFP conversions, and DFP add/sub. Similarly `bcd_mul` / `bcd_div` are implemented in 11 instructions using register transfer and the DFU operations for Power6/7/8.

**Note**

So why does anybody care about BCD and DFP? Sometimes you get large numbers in decimal that you need converted to binary for extended computation. Sometimes you need to display the results of your extended binary computation in decimal. The multiply by 10 and BCD vector operations help simplify and speed-up these conversions.

And finally: Henry S. Warren's wonderful book *Hacker's Delight* provides inspiration for SIMD versions of; count leading zeros, population count, parity, etc.

**1.2.2.2 General Endian Issues**

For POWER8, IBM made the explicit decision to support Little Endian (**LE**) data format in the Linux ecosystem. The goal was to enhance application code portability across Linux platforms. This goal was integrated into the OpenPOWER ELF V2 Application Binary Interface **ABI** specification.

The POWER8 processor architecturally supports an *Endian Mode* and supports both BE and LE storage access in hardware. However register to register operations do not change (are not effected by endian mode). The ABI extends the LE storage format to vector register (logical) element numbering.

This has no effect for most `altivec.h` operations where the input elements and the results "stay in their lanes". For operations of the form  $(T[n] = A[n] \text{ op } B[n])$ , it does not matter if elements are numbered  $[0, 1, 2, 3]$  or  $[3, 2, 1, 0]$ .

But there are cases where element renumbering can change the results. Changing element numbering does change the even / odd relationship for merge and integer multiply. For **LE** targets, operations accessing even vector elements are implemented using the equivalent odd instruction (and visa versa) and inputs are swapped. Similarly for high and low merges. Inputs are also swapped for Pack, Unpack, and Permute operations and the permute select vector is inverted. The above is just a sampling of a larger list *LE transforms*. The OpenPOWER ABI specification provides a helpful table of [Endian-Sensitive Operations](#).

**Note**

This means that the vector built-ins provided by `altivec.h` may not generate the instructions you expect.

This does matter when doing extended precision arithmetic. Here we need to maintain most-to-least significant byte order and align "digit" columns for summing partial products. Many of these operation were defined long before Little Endian was seriously considered and are decidedly Big Endian in register format. Basically any operation where the element size changes from input to output is suspect for **LE** targets.

The coding for these higher level operations is complicated by *Little Endian* (LE) support as specified in the OpenPOWER ABI and as implemented in the compilers. Little Endian changes the effective vector element numbering and the location of even and odd elements.

This is a general problem for using vectors to implement extended precision arithmetic. The multiply even/odd operations being the primary example. The products are double-wide and in BE order in the vector register. This is reinforced by the Vector Add/Subtract Unsigned Quadword instructions. And the products from even multiply instruction are always *algebraically* higher then the odd products. The pack, unpack, and sum operations have similar issues.

This matters when you need to align (shift) the partial products or select the *algebraically* high or lower portion of the products. The (high to low) order of elements for the multiply has to match the order of the largest element size used in accumulating partial sums. This is normally a quadword (`vadduqm` instruction).

So the element order is fixed while the element numbering and the partial products (between even and odd) will change between BE and LE. This effects splatting and octet shift operations required to align partial product for summing. These are the places where careful programming is required, to nullify the compilers LE transforms, so we will get the correct numerical answer.

So what can the Power Vector Library do to help?

- Be aware of these mandated LE transforms and if required provide compliant inline assembler implementations for LE.
- Where required for correctness provide LE specific implementations that have the effect of nullifying the unwanted transforms.
- Provide higher level operations that help pveclib and applications code in an endian neutral way and get correct results.

#### See also

[Endian problems with word operations](#)  
[Vector Multiply-by-10 Unsigned Quadword example](#)

#### 1.2.2.3 So what can the Power Vector Library project do?

Clearly the PowerISA provides multiple, extensive, and powerful computational facilities that continue to evolve and grow. But the best instruction sequence for a specific computation depends on which POWER processor(s) you have or plan to support. It can also depend on the specific compiler version you use, unless you are willing to write some of your application code in assembler. Even then you need to be aware of the PowerISA versions and when specific instructions were introduced. This can be frustrating if you just want to port your application to POWER for a quick evaluation.

So you would like to start evaluating how to leverage this power for key algorithms at the heart of your application.

- But you are working with an older POWER processor (until the latest POWER box is delivered).
- Or the latest POWER machine just arrived at your site (or cloud) but you are stuck using an older/stable Linux distro version (with an older distro compiler).
- Or you need extended precision multiply for your crypto code but you are not really an assembler level programmer (or don't want to be).
- Or you would like to program with higher level operations to improve your own productivity.

Someone with the right background (knowledge of the PowerISA, assembler level programming, compilers and the vector built-ins, ...) can solve any of the issues described above. But you don't have time for this.

There should be an easier way to exploit the POWER vector hardware without getting lost in the details. And this extends beyond classical vector (Single Instruction Multiple Data (SIMD)) programming to exploiting larger data width (128-bit and beyond), and larger register space (64 x 128 Vector Scalar Registers)

##### 1.2.2.3.1 Vector Add Unsigned Quadword Modulo example

Here is an example of what can be done:

```
static inline vui128_t
vec_adduqm (vui128_t a, vui128_t b)
{
    vui32_t t;
#ifdef _ARCH_PWR8
#ifndef vec_vadduqm
__asm__(
    "vadduqm %0,%1,%2;"
    : "=v" (t)
    : "v" (a),
    "v" (b)
    : );
#else
    t = (vui32_t) vec_vadduqm (a, b);
#endif
#endif
}
```

```

#endif
#else
    vui32_t c, c2;
    vui32_t z= { 0,0,0,0};

    c = vec_vaddcuw ((vui32_t)a, (vui32_t)b);
    t = vec_vadduwm ((vui32_t)a, (vui32_t)b);
    c = vec_sld (c, z, 4);
    c2 = vec_vaddcuw (t, c);
    t = vec_vadduwm (t, c);
    c = vec_sld (c2, z, 4);
    c2 = vec_vaddcuw (t, c);
    t = vec_vadduwm (t, c);
    c = vec_sld (c2, z, 4);
    t = vec_vadduwm (t, c);
#endif
    return ((vui128_t) t);
}

```

The **\_ARCH\_PWR8** macro is defined by the compiler when it targets POWER8 (PowerISA 2.07) or later. This is the first processor and PowerISA level to support vector quadword add/subtract. Otherwise we need to use the vector word add modulo and vector word add and write carry-out word, to add 32-bit chunks and propagate the carries through the quadword.

One little detail remains. Support for `vec_vadduqm` was added to GCC in March of 2014, after GCC 4.8 was released and GCC 4.9's feature freeze. So the only guarantee is that this feature is in GCC 5.0 and later. At some point this change was backported to GCC 4.8 and 4.9 as it is included in the current GCC 4.8/4.9 documentation. When or if these backports were propagated to a specific Linux Distro version or update is difficult to determine. So support for this vector built-in depends on the specific version of the GCC compiler, or if specific Distro update includes these specific backports for the GCC 4.8/4.9 compiler they support. The:

```
#ifndef vec_vadduqm
```

C preprocessor conditional checks if the **vec\_vadduqm** is defined in `<altivec.h>`. If defined we can assume that the compiler implements **\_\_builtin\_vec\_vadduqm** and that `<altivec.h>` includes the macro definition:

```
#define vec_vadduqm __builtin_vec_vadduqm
```

For **\_ARCH\_PWR7** and earlier we need a little grade school arithmetic using Vector Add Unsigned Word Modulo (**vadduwm**) and Vector Add and Write Carry-Out Unsigned Word (**vaddcuw**). This treats the vector `__int128` as 4 32-bit binary digits. The first instruction sums each (32-bit digit) column and the second records the carry out of the high order bit. This leaves the carry bit in the original (word) column, so use a shift left to line up the carries with the next higher word.

To propagate any carries across all 4 (word) digits, repeat this (add / carry / shift) sequence three times. Then a final add modulo word to complete the 128-bit add. This sequence requires 10-11 instructions. The 11th instruction is a vector splat word 0 immediate, which is needed in the shift left (`vsldoi`) instructions. This is common in vector codes and the compiler can usually reuse this register across several blocks of code and inline functions.

For POWER7/8 these instructions are all 2 cycle latency and 2 per cycle throughput. The `vadduwm` / `vaddcuw` instruction pairs should issue in the same cycle and execute in parallel. So the expected latency for this sequence is 14 cycles. For POWER8 the `vadduqm` instruction has a 4 cycle latency.

Similarly for the carry / extend forms which can be combined to support wider (256, 512, 1024, ...) extended arithmetic.

See also

[vec\\_addcuq](#), [vec\\_addeuqm](#), and [vec\\_addecuq](#)

### 1.2.2.3.2 Vector Multiply-by-10 Unsigned Quadword example

PowerISA 3.0 (POWER9) added this instruction and it's extend / carry forms to speed up decimal to binary conversion for large numbers. But this operation is general useful and not that hard to implement for earlier processors.

```
static inline vui128_t
vec_mul10uq (vui128_t a)
{
    vui32_t t;
#ifdef _ARCH_PWR9
    __asm__(
        "vmul10uq %0,%1;\n"
        : "=v" (t)
        : "v" (a)
        : );
#else
    vui16_t ts = (vui16_t) a;
    vui16_t t10;
    vui32_t t_odd, t_even;
    vui32_t z = { 0, 0, 0, 0 };
    t10 = vec_splat_u16(10);
    if __BYTE_ORDER__ == __ORDER_LITTLE_ENDIAN__
        t_even = vec_vmulouh (ts, t10);
        t_odd = vec_vmuleuh (ts, t10);
    else
        t_even = vec_vmuleuh (ts, t10);
        t_odd = vec_vmulouh (ts, t10);
    #endif
    t_even = vec_sld (t_even, z, 2);
#ifdef _ARCH_PWR8
    t = (vui32_t) vec_vadduqm ((vui128_t) t_even, (vui128_t) t_odd);
#else
    t = (vui32_t) vec_adduqm ((vui128_t) t_even, (vui128_t) t_odd);
#endif
    #endif
    return ((vui128_t) t);
}
```

Notice that under the `_ARCH_PWR9` conditional, there is no check for the specific `vec_vmul10uq` built-in. As of this writing `vec_vmul10uq` is not included in the *OpenPOWER ELF2 ABI* documentation nor in the latest GCC trunk source code.

#### Note

The *OpenPOWER ELF2 ABI* does define `bcd_mul10` which (from the description) will actually generate Decimal Shift (**bc**ds). This instruction shifts 4-bit nibbles (BCD digits) left or right while preserving the BCD sign nibble in bits 124-127. While this is a handy instruction to have, it is not the same operation as `vec_vmul10uq`, which is a true 128-bit binary multiply by 10. As of this writing `bcd_mul10` support is not included in the latest GCC trunk source code.

For `_ARCH_PWR8` and earlier we need a little grade school arithmetic using **Vector Multiply Even/Odd Unsigned Halfword**. This treats the vector `__int128` as 8 16-bit binary digits. We multiply each of these 16-bit digits by 10, which is done in two (even and odd) parts. The result is 4 32-bit (2 16-bit digits) partial products for the even digits and 4 32-bit products for the odd digits. The vector register (independent of endian); the even product elements are higher order and odd product elements are lower order.

The even digit partial products are offset right by 16-bits in the register. If we shift the even products left 1 (16-bit) digit, the even digits are lined up in columns with the odd digits. Now we can sum across partial products to get the final 128 bit product.

Notice also the conditional code for endian around the `vec_vmulouh` and `vec_vmuleuh` built-ins:

```
#if __BYTE_ORDER__ == __ORDER_LITTLE_ENDIAN__
```



Little endian (**LE**) changes the element numbering. This also changes the meaning of even / odd and this effects the code generated by compilers. But the relationship of high and low order bytes, within multiplication products, is defined by the hardware and does not change. (See: [General Endian Issues](#)) So the pveclib implementation needs to pre-swap the even/odd partial product multiplies for LE. This in effect nullifies the even / odd swap hidden in the compilers **LE** code generation and the resulting code gives the correct results.

Now we are ready to sum the partial product *digits* while propagating the digit carries across the 128-bit product. For **\_ARCH\_PWR8** we can use **Vector Add Unsigned Quadword Modulo** which handles all the internal carries in hardware. Before **\_ARCH\_PWR8** we only have **Vector Add Unsigned Word Modulo** and **Vector Add and Write Carry-Out Unsigned Word**.

We see these instructions used in the **else** leg of the pveclib **vec\_adduqm** implementation above. We can assume that this implementation is correct and tested for supported platforms. So here we use another pveclib function to complete the implementation of **Vector Multiply-by-10 Unsigned Quadword**.

Again similarly for the carry / extend forms which can be combined to support wider (256, 512, 1024, ...) extended decimal to binary conversions.

See also

[vec\\_mul10cuq](#), [vec\\_mul10euq](#), and [vec\\_mul10ecuq](#)

And similarly for full 128-bit x 128-bit multiply which combined with the add quadword carry / extended forms above can be used to implement wider (256, 512, 1024, ...) multiply operations.

See also

[vec\\_mulluq](#) and [vec\\_muluq](#)  
[Vector Merge Algebraic High Word example](#)  
[Vector Multiply High Unsigned Word example](#)

#### 1.2.2.4 Returning extended quadword results.

Extended quadword add, subtract and multiply results can exceed the width of a single 128-bit vector. A 128-bit add can produce 129-bit results. A unsigned 128-bit by 128-bit multiply result can produce 256-bit results. This is simplified for the *modulo* case where any result bits above the low order 128 can be discarded. But extended arithmetic requires returning the full precision result. Returning double wide quadword results are a complication for both RISC processor and C language library design.

##### 1.2.2.4.1 PowerISA and Implementation.

For a RISC processor, encoding multiple return registers forces hard trade-offs in a fixed sized instruction format. Also building a vector register file that can support at least one (or more) double wide register writes per cycle is challenging. For a super-scalar machine with multiple vector execution pipelines, the processor can issue and complete multiple instructions per cycle. As most operations return single vector results, this is a higher priority than optimizing for double wide results.

The PowerISA addresses this by splitting these operations into two instructions that execute independently. Here independent means that given the same inputs, one instruction does not depend on the result of the other. Independent instructions can execute out-of-order, or if the processor has multiple vector execution pipelines, can execute (issue and complete) concurrently.

The original VMX implementation had Vector Add/Subtract Unsigned Word Modulo (**vadduwm** / **vsubuwm**), paired with Vector Add/Subtract and Write Carry-out Unsigned Word (**vaddcuw** / **vsubcuw**). Most usage ignores the carry-out and only uses the add/sub modulo instructions. Applications requiring extended precision, pair the add/sub modulo with add/sub write carry-out, to capture the carry and propagate it to higher order bits.

The (four word) carries are generated into the same *word lane* as the source addends and modulo result. Propagating the carries require a separate shift (to align the carry-out with the low order (carry-in) bit of the next higher word) and another add word modulo.

POWER8 (PowerISA 2.07B) added full Vector Add/Subtract Unsigned Quadword Modulo (**vadduqm** / **vsubuqm**) instructions, paired with corresponding Write Carry-out instructions. (**vaddcuq** / **vsubcuq**). A further improvement over the word instructions was the addition of three operand *Extend* forms which combine add/subtract with carry-in (**vaddeuqm**, **vsubeuqm**, **vaddecuq** and **vsubecuq**). This simplifies propagating the carry-out into higher quadword operations.

See also

[vec\\_adduqm](#), [vec\\_addcuq](#), [vec\\_addeuqm](#), [vec\\_addecuq](#)

POWER9 (PowerISA 3.0B) added Vector Multiply-by-10 Unsigned Quadword (Modulo is implied), paired with Vector Multiply-by-10 and Write Carry-out Unsigned Quadword (**vmul10uq** / **vmul10cuq**). And the *Extend* forms (**vmul10euq** / **vmul10ecuq**) simplifies the digit (0-9) carry-in for extended precision decimal to binary conversions.

See also

[vec\\_mul10uq](#), [vec\\_mul10cuq](#), [vec\\_mul10euq](#), [vec\\_mul10ecuq](#)

The VMX integer multiply operations are split into multiply even/odd instructions by element size. The product requires the next larger element size (twice as many bits). So a vector multiply byte would generate 16 halfword products (256-bits in total). Requiring separate even and odd multiply instructions cuts the total generated product bits (per instruction) in half. It also simplifies the hardware design by keeping the generated product in adjacent element lanes. So each vector multiply even or odd byte operation generates 8 halfword products (128-bits) per instruction.

This multiply even/odd technique applies to most element sizes from byte up to doubleword. The original V←MX supports multiply even/odd byte and halfword operations. In the original VMX, arithmetic operations were restricted to byte, halfword, and word elements. Multiply halfword products fit within the integer word element. No multiply byte/halfword modulo instructions were provided, but could be implemented via a vmule, vmulo, vperm sequence.

POWER8 (PowerISA 2.07B) added multiply even/odd word and multiply modulo word instructions.

See also

[vec\\_muleuw](#), [vec\\_mulouw](#), [vec\\_muluwm](#)

The latest PowerISA (3.0B for POWER9) does add a doubleword integer multiply via **Vector Multiply-Sum unsigned Doubleword Modulo**. This is a departure from the Multiply even/odd byte/halfword/word instructions available in earlier Power processors. But careful conditioning of the inputs can generate the equivalent of multiply even/odd unsigned doubleword.

See also

[vec\\_msumudm](#), [vec\\_muleud](#), [vec\\_muloud](#)

This (multiply even/odd) technique breaks down when the input element size is quadword or larger. A quadword integer multiply forces a different split. The easiest next step would be a high/low split (like the Fixed-point integer multiply). A multiply low (modulo) quadword would be a useful function. Paired with multiply high quadword provides the double quadword product. This would provide the basis for higher (multi-quadword) precision multiplies.

See also

[vec\\_mulluq](#), [vec\\_muludq](#)

#### 1.2.2.4.2 C Language restrictions.

The Power Vector Library is implemented using C language (inline) functions and this imposes its own restrictions. Standard C language allows an arbitrary number of formal parameters and one return value per function. Parameters and return values with simple C types are normally transferred (passed / returned) efficiently in local (high performance) hardware registers. Aggregate types (struct, union, and arrays of arbitrary size) are normally handled by pointer indirection. The details are defined in the appropriate Application Binary Interface (ABI) documentation.

The POWER processor provides lots (96) of registers so we want to use registers wherever possible. Especially when our application is composed of collections of small functions. And more especially when these functions are small enough to inline and we want the compiler to perform local register allocation and common subexpression elimination optimizations across these functions. The PowerISA defines 3 kinds of registers;

- General Purpose Registers (GPRs),
- Floating-point Registers (FPRs),
- Vector registers (VRs),

with 32 of each kind. We will ignore the various special registers for now.

The PowerPC64 64-bit ELF (and OpenPOWER ELF V2) ABIs normally pass simple arguments and return values in a single register (of the appropriate kind) per value. Arguments of aggregate types are passed as storage pointers in General Purpose Registers (GPRs).

The language specification, the language implementation, and the ABI provide some exceptions. The C99 language adds `_Complex` floating types which are composed of real and imaginary parts. GCC adds `_Complex` integer types. For PowerPC ABIs complex values are held in a pair of registers of the appropriate kind. C99 also adds double word integers as the *long long int* type. This only matters for PowerPC 32-bit ABIs. For PowerPC64 ABIs *long long* and *long* are both 64-bit integers and are held in 64-bit GPRs.

GCC also adds the `__int128` type for some targets including the PowerPC64 ABIs. Values of `__int128` type are held (for operations, parameter passing and function return) in 64-bit GPR pairs. GCC adds `__ibm128` and `_Decimal128` floating point types which are held in Floating-point Registers pairs. GCC recently added the `__float128` floating point type which are held in single vector register. Similarly for vector `__int128`.

GCC defines Generic Vector Extensions that allow typedefs for vectors of various element sizes/types and generic SIMD (arithmetic, logical, and element indexing) operations. For PowerPC64 ABIs this is currently restricted to 16-byte vectors as defined in `<altivec.h>`. For currently available compilers attempts to define vector types with larger (32 or 64 byte) *vector\_size* values are treated as arrays of scalar elements. Only *vector\_size(16)* variables are passed and returned in vector registers.

The OpenPOWER 64-Bit ELF V2 ABI Specification makes specific provisions for passing/returning *homogeneous aggregates* of multiple like data types. Such aggregates can be passed as up to eight floating-point or vector registers. This is defined for the Little Endian ELF V2 ABI and is not applicable to Big Endian ELF V1 targets. Also current GCC versions, in common use, do not fully implement this ABI feature.

So we have shown that there are mechanisms for functions to return multiple vector register values. But none are really practical at this time as they not yet available (function or optimal code generation) in current GCC compilers, that are in common use.

Returning pairs of vector `__int128` values as `_Complex __float128` would be awkward at best. And it is not clear when or if `_Complex vector __int128` will be supported. GCC's Generic Vector Extensions are only implemented for *vector\_size(16)*. And current GCC compilers can generate some sub-optimal code for passing/returning *homogeneous aggregates* as suggested in the OpenPOWER ABI.

#### 1.2.2.4.3 Subsetting the problem.

We can simplify this problem by remembering that:

- Only a subset of the pveclib functions need to return more than one 128-bit vector.
- The PowerISA normally splits these cases into multiple instructions anyway.
- So far these functions are small and fully inlined.

So we have two (or three) options given the current state of GCC compilers in common use:

- Mimic the PowerISA and split the operation into two functions, where each function only returns (up to) 128-bits of the result.
- Use pointer parameters to return a second vector value in addition to the function return.
- Support both and let the user decide which works best.

The add/subtract quadword operations provide good examples. For example adding two 256-bit unsigned integer values and returning the 257-bit (the high / low sum and the carry) result looks like this:

```
s1 = vec_vadduqm (a1, b1); // sum low 128-bits a1+b1
c1 = vec_vaddcuq (a1, b1); // write-carry from low a1+b1
s0 = vec_vaddeuqm (a0, b0, c1); // Add-extend high 128-bits a0+b0+c1
c0 = vec_vaddecuq (a0, b0, c1); // write-carry from high a0+b0+c1
```

This sequence uses the built-ins from `<altivec.h>` and generates instructions that will execute on POWER8 and POWER9. The compiler must target POWER8 (`-mcpu=power8`) or higher. In fact the compile will fail if the target is POWER7.

Now let's look at the pveclib version of these operations from `<vec_int128_ppc.h>`:

```
s1 = vec_adduqm (a1, b1); // sum low 128-bits a1+b1
c1 = vec_addcuq (a1, b1); // write-carry from low a1+b1
s0 = vec_addeuqm (a0, b0, c1); // Add-extend high 128-bits a0+b0+c1
c0 = vec_addecuq (a0, b0, c1); // write-carry from high a0+b0+c1
```

Looks almost the same but the operations do not use the 'v' prefix on the operation name. This sequence generates the same instructions for (`-mcpu=power8`) as the `<altivec.h>` version above. It will also generate a different (slightly longer) instruction sequence for (`-mcpu=power7`) which is functionally equivalent.

The pveclib `<vec_int128_ppc.h>` header also provides a coding style alternative:

```
s1 = vec_addcq (&c1, a1, b1);
s0 = vec_addeq (&c0, a0, b0, c1);
```

Here `vec_addcq` combines the `adduqm/addcuq` operations into a *add and carry quadword* operation. The first parameter is a pointer to vector to receive the carry-out while the 128-bit modulo sum is the function return value. Similarly `vec_addeq` combines the `addeuqm/addecuq` operations into a *add with extend and carry quadword* operation.

As these functions are inlined by the compiler the implied store / reload of the carry can be converted into a simple register assignment. For (`-mcpu=power8`) the compiler should generate the same instruction sequence as the two previous examples.

For (`-mcpu=power7`) these functions will expand into a different (slightly longer) instruction sequence which is functionally equivalent to the instruction sequence generated for (`-mcpu=power8`).

For older processors (power7 and earlier) and under some circumstances instructions generated for this "combined form" may perform better than the "split form" equivalent from the second example. Here the compiler may not recognize all the common subexpressions, as the "split forms" are expanded before optimization.

### 1.2.3 Background on the evolution of <altivec.h>

The original **AltiVec (TM) Technology Programming Interface Manual** defined the minimal vector extensions to the application binary interface (ABI), new keywords (vector, pixel, bool) for defining new vector types, and new operators (built-in functions).

- generic AltiVec operations, like `vec_add()`
- specific AltiVec operations (instructions, like `vec_addubm()`)
- predicates computed from a AltiVec operation like `vec_all_eq()`

A generic operation generates specific instructions based on the types of the actual parameters. So a generic `vec_add` operation, with vector char parameters, will generate the (specific) vector add unsigned byte modulo (`vadubm`) instruction. Predicates are used within if statement conditional clauses to access the condition code from vector operations that set Condition Register 6 (vector SIMD compares and Decimal Integer arithmetic and format conversions).

The PIM defined a set of compiler built-ins for vector instructions (see section "4.4 Generic and Specific AltiVec Operations") that compilers should support. The document suggests that any required typedefs and supporting macro definitions be collected into an include file named <altivec.h>.

The built-ins defined by the PIM closely match the vector instructions of the underlying PowerISA. For example: `vec_mul`, `vec_mule` / `vec_mulo`, and `vec_muleub` / `vec_muloub`.

- `vec_mul` is defined for float and double and will (usually) generate a single instruction for the type. This is a simpler case as floating point operations usually stay in their lanes (result elements are the same size as the input operand elements).
- `vec_mule` / `vec_mulo` (multiply even / odd) are defined for integer multiply as integer products require twice as many bits as the inputs (the results don't stay in their lane).

The RISC philosophy resists and POWER Architecture avoids instructions that write to more than one register. So the hardware and PowerISA vector integer multiply generate even and odd product results (from even and odd input elements) from two instructions executing separately. The PIM defines and compiler supports these operations as overloaded built-ins and selects the specific instructions based on the operand (char or short) type.

This is complicated as the PowerISA evolves. The original AltiVec (VMX) provided vector multiply (even / odd) operations for byte (char) and halfword (short) integers. Multiple even / odd word (int) instructions were not introduced until PowerISA V2.07 (POWER8). PowerISA 2.07 also introduced vector multiply word modulo which is included under the generic `vec_mul`.

As the PowerISA evolved adding new vector (VMX) instructions, new facilities (Vector Scalar Extended (VSX)), and specialized vector categories (little endian, AES, SHA2, RAID), these new operators were added to <altivec.h>. This included new specific and generic operations and additional vector element types (long (64-bit) int, `__int128`, double and quad precision (`__Float128`) float).

However the PIM documents were primarily focused on embedded processors and were not updated to include the vector extensions implemented by the server processors. So any documentation for new vector operations were relegated to the various compilers. This was a haphazard process and some divergence in operation naming did occur between compilers.

In the run up to the POWER8 launch and the OpenPOWER initiative it was recognized that switching to Little Endian would require a new and well documented Application Binary Interface (**ABI**). It was also recognized that new

<altivec.h> extensions needed to be documented in a common place so the various compilers could implement a common vector built-in API. So ...

The [OpenPOWER ELF V2 application binary interface \(ABI\)](#): Chapter 6. Vector Programming Interfaces and Appendix A. Predefined Functions for Vector Programming document the current and proposed vector built-ins we expect all C/C++ compilers to implement for the PowerISA.

The ABI also defines many overloaded built-in functions as generic operations. Here the compiler selects a specific PowerISA implementation based on the operand (vector element) types. The ABI also defines the (big/little) endian behavior and the compiler may select different instructions based on the endianness of the target.

Also note that the vector element numbering changes between big and little endian, and so does the meaning of even and odd. Both affect what the compiler supports and the instruction sequence generated.

- **vec\_muleub** and **vec\_muloub** (multiply even / odd unsigned byte) are examples of non-overloaded built-ins provided by the GCC compiler but not defined in the ABI. One would assume these built-ins will generate the matching instruction, however the GCC compiler will adjust the generated instruction based on the target endianness (even / odd is reversed for little endian).

The ABI also defines `vec_mul` as an overloaded operation on integer types, where only the low order half (modulo element size) of the product is returned. The PowerISA does not provide direct multiply modulo instructions for all the integer sizes / types. So this requires a multiple-instruction sequence to implement. Also integer `vec_mul` is defined in the ABI as "phased in" and is only implemented in the latest GCC versions.

This is a small sample of the complexity we encounter programming at this low level (vector intrinsic) API. Partially this is due to RISC design philosophy where there is a trade-off of software complexity for simpler (hopefully faster) hardware design.

#### 1.2.4 pveclib is not a vector math library

The pveclib does not implement general purpose vector math operations. These should continue to be developed and improved within existing projects (ie LAPACK, OpenBLAS, ATLAS and libmvec).

We believe that pveclib will be helpful to implementors of vector math libraries by providing a higher level, more portable, and more consistent vector interface for the PowerISA. Similarly for implementors of extended arithmetic, cryptographic, compression/decompression, and pattern matching / search libraries.

### 1.3 Performance data.

It is useful to provide basic performance data for each pveclib function. This is challenging as these functions are small and intended to be in-lined within larger functions (algorithms). As such they are subject to both the compiler's instruction scheduling and common subexpression optimizations plus the processors super-scalar and out-of-order execution design features.

As pveclib functions are normally only a few instructions, the actual timing will depend on the context it is in (the instructions that it depends on for data and instructions that proceed them in the pipelines).

The simplest approach is to use the same performance metrics as the Power Processor Users Manuals Performance Profile. This is normally per instruction latency in cycles and throughput in instructions issued per cycle. There may also be additional information for special conditions that may apply.

For example the vector float absolute value function. For recent PowerISA implementations this a single (VSX **xvabssp**) instruction which we can look up in the POWER8 / POWER9 Processor User's Manuals (**UM**).

processor	Latency	Throughput
power8	6-7	2/cycle
power9	2	2/cycle

The POWER8 UM specifies a latency of "6 cycles to FPU (+1 cycle to other VSU ops" for this class of VSX single precision FPU instructions. So the minimum latency is 6 cycles if the register result is input to another VSX single precision FPU instruction. Otherwise if the result is input to a VSU logical or integer instruction then the latency is 7 cycles. The POWER9 UM shows the pipeline improvement of 2 cycles latency for simple FPU instructions like this. Both processors support dual pipelines for a 2/cycle throughput capability.

A more complicated example:

```
static inline vb32_t
vec_isnanf32 (vf32_t vf32)
{
    vui32_t tmp2;
    const vui32_t expmask = CONST_VINT128_W(0x7f800000, 0x7f800000, 0x7f800000,
                                              0x7f800000);

    #if _ARCH_PWR9
    // P9 has a 2 cycle xvbsssp and eliminates a const load.
    tmp2 = (vui32_t) vec_abs (vf32);
    #else
    const vui32_t signmask = CONST_VINT128_W(0x80000000, 0x80000000, 0x80000000,
                                              0x80000000);
    tmp2 = vec_andc ((vui32_t)vf32, signmask);
    #endif
    return vec_cmpgt (tmp2, expmask);
}
```

Here we want to test for *Not A Number* without triggering any of the associate floating-point exceptions (VXSNAN or VXVC). For this test the sign bit does not effect the result so we need to zero the sign bit before the actual test. The vector abs would work for this, but we know from the example above that this instruction has a high latency as we are definitely passing the result to a non-FPU instruction (vector compare greater than unsigned word).

So the code needs to load two constant vectors masks, then vector and-compliment to clear the sign-bit, before comparing each word for greater then infinity. The generated code should look something like this:

```
addis    r9,r2,.rodata.cst16+0x10@ha
addis    r10,r2,.rodata.cst16+0x20@ha
addi     r9,r9,.rodata.cst16+0x10@l
addi     r10,r10,.rodata.cst16+0x20@l
lvx      v0,0,r10 # load vector const signmask
lvx      v12,0,r9 # load vector const expmask
xxlandc  vs34,vs34,vs32
vcmpgtuw v2,v2,v12
```

So six instructions to load the const masks and two instructions for the actual vec\_isnanf32 function. The first six instructions are only needed once for each containing function, can be hoisted out of loops and into the function prologue, can be *commoned* with the same constant for other pveclib functions, or executed out-of-order and early by the processor.

Most of the time, constant setup does not contribute measurably to the over all performance of vec\_isnanf32. When it does it is limited by the longest (in cycles latency) of the various independent paths that load constants. In this case the const load sequence is composed of three pairs of instructions that can issue and execute in parallel. The addis/addi FXU instructions supports throughput of 6/cycle and the lvx load supports 2/cycle. So the two vector constant load sequences can execute in parallel and the latency is same as a single const load.

For POWER8 it appears to be (2+2+5=) 9 cycles latency for the const load. While the core vec\_isnanf32 function (xxlandc/vcmpgtuw) is a dependent sequence and runs (2+2) 4 cycles latency. Similar analysis for POWER9 where the addis/addi/lvx sequence is still listed as (2+2+5) 9 cycles latency. While the xxlandc/vcmpgtuw sequence increases to (2+3) 5 cycles.

The next interesting question is what can we say about throughput (if anything) for this example. The thought experiment is "what would happen if?";

- two or more instances of `vec_isnanf32` are used within a single function,
- in close proximity in the code,
- with independent data as input,

could the generated instructions execute in parallel and to what extent. This illustrated by the following (contrived) example:

```
int
test512_all_f32_nan (vf32_t val0, vf32_t val1, vf32_t val2,
                    vf32_t val3)
{
    const vb32_t alltrue = { -1, -1, -1, -1 };
    vb32_t nan0, nan1, nan2, nan3;

    nan0 = vec_isnanf32 (val0);
    nan1 = vec_isnanf32 (val1);
    nan2 = vec_isnanf32 (val2);
    nan3 = vec_isnanf32 (val3);

    nan0 = vec_and (nan0, nan1);
    nan2 = vec_and (nan2, nan3);
    nan0 = vec_and (nan2, nan0);

    return vec_all_eq(nan0, alltrue);
}
```

which tests 4 X vector float (16 X float) values and returns true if all 16 floats are NaN. Recent compilers will generates something like the following PowerISA code:

```
addis    r9,r2,-2
addis    r10,r2,-2
vspltisw v13,-1      # load vector const alltrue
addi     r9,r9,21184
addi     r10,r10,-13760
lvx      v0,0,r9      # load vector const signmask
lvx      v1,0,r10     # load vector const expmask
xxlandc  vs35,vs35,vs32
xxlandc  vs34,vs34,vs32
xxlandc  vs37,vs37,vs32
xxlandc  vs36,vs36,vs32
vcmpgtuw v3,v3,v1      # nan1 = vec_isnanf32 (val1);
vcmpgtuw v2,v2,v1      # nan0 = vec_isnanf32 (val0);
vcmpgtuw v5,v5,v1      # nan3 = vec_isnanf32 (val3);
vcmpgtuw v4,v4,v1      # nan2 = vec_isnanf32 (val2);
xxland  vs35,vs35,vs34  # nan0 = vec_and (nan0, nan1);
xxland  vs36,vs37,vs36  # nan2 = vec_and (nan2, nan3);
xxland  vs36,vs36,vs36  # nan0 = vec_and (nan2, nan0);
vcmpequw v4,v4,v13     # vec_all_eq(nan0, alltrue);
...
```

first the generated code loading the vector constants for signmask, expmask, and alltrue. We see that the code is generated only once for each constant. Then the compiler generate the core `vec_isnanf32` function four times and interleaves the instructions. This enables parallel pipeline execution where conditions allow. Finally the 16X isnan results are reduced to 8X, then 4X, then to a single condition code.

For this exercise we will ignore the constant load as in any realistic usage it will be *commoned* across several pveclib functions and hoisted out of any loops. The reduction code is not part of the `vec_isnanf32` implementation and also ignored. The sequence of 4X `xxlandc` and 4X `vcmpgtuw` in the middle it the interesting part.

For POWER8 both `xxlandc` and `vcmpgtuw` are listed as 2 cycles latency and throughput of 2 per cycle. So we can assume that (only) the first two `xxlandc` will issue in the same cycle (assuming the input vectors are ready). The issue of the next two `xxlandc` instructions will be delay by 1 cycle. The following `vcmpgtuw` instruction are dependent on the `xxlandc` results and will not execute until their input vectors are ready. The first two `vcmpgtuw` instruction will execute 2 cycles (latency) after the first two `xxlandc` instructions execute. Execution of the second two `vcmpgtuw` instructions will be delayed 1 cycle due to the issue delay in the second pair of `xxlandc` instructions.

So at least for this example and this set of simplifying assumptions we suggest that the throughput metric for `vec_isnanf32` is 2/cycle. For latency metric we offer the range with the latency for the core function (without and constant load overhead) first. Followed by the total latency (the sum of the constant load and core function latency). For the `vec_isnanf32` example the metrics are:



processor	Latency	Throughput
power8	4-13	2/cycle
power9	5-14	2/cycle

Looking at a slightly more complicated example where core functions implementation can execute more than one instruction per cycle. Consider:

```
static inline vb32_t
vec_isnormalf32 (vf32_t vf32)
{
    vui32_t tmp, tmp2;
    const vui32_t expmask = CONST_VINT128_W(0x7f800000, 0x7f800000, 0x7f800000,
                                              0x7f800000);
    const vui32_t minnorm = CONST_VINT128_W(0x00800000, 0x00800000, 0x00800000,
                                              0x00800000);

    #if _ARCH_PWR9
    // P9 has a 2 cycle xwabssp and eliminates a const load.
    tmp2 = (vui32_t) vec_abs (vf32);
    #else
    const vui32_t signmask = CONST_VINT128_W(0x80000000, 0x80000000, 0x80000000,
                                              0x80000000);
    tmp2 = vec_andc ((vui32_t)vf32, signmask);
    #endif
    tmp = vec_and ((vui32_t) vf32, expmask);
    tmp2 = (vui32_t) vec_cmplt (tmp2, minnorm);
    tmp = (vui32_t) vec_cmpeq (tmp, expmask);

    return (vb32_t)vec_nor (tmp, tmp2);
}
```

which requires two (independent) masking operations (sign and exponent), two (independent) compares that are dependent on the masking operations, and a final *not OR* operation dependent on the compare results.

The generated POWER8 code looks like this:

```
addis    r10,r2,-2
addis    r8,r2,-2
addi     r10,r10,21184
addi     r8,r8,-13760
addis    r9,r2,-2
lvx      v13,0,r8
addi     r9,r9,21200
lvx      v1,0,r10
lvx      v0,0,r9
xxland   vs33,vs33,vs34
xxlandc  vs34,vs45,vs34
vcmpgtuw v0,v0,v1
vcmpewuw v2,v2,v13
xxlnor   vs34,vs32,vs34
```

Note this this sequence needs to load 3 vector constants. In previous examples we have noted that POWER8 *lvx* supports 2/cycle throughput. But with good scheduling, the 3rd vector constant load, will only add 1 additional cycle to the timing (10 cycles).

Once the constant masks are loaded the *xxland/xxlandc* instructions can execute in parallel. The *vcmpgtuw/vcmpequw* can also execute in parallel but are delayed waiting for the results of masking operations. Finally the *xxnor* is dependent on the data from both compare instructions.

For POWER8 the latencies are 2 cycles each, and assuming parallel execution of *xxland/xxlandc* and *vcmpgtuw/vcmpequw* we can assume (2+2+2=) 6 cycles minimum latency and another 10 cycles for the constant loads (if needed).

While the POWER8 core has ample resources (10 issue ports across 16 execution units), this specific sequence is restricted to the two *issue ports and VMX execution units* for this class of (simple vector integer and logical) instructions. For *vec\_isnormalf32* this allows for a lower latency (6 cycles vs the expected 10, over 5 instructions), it also implies that both of the POWER8 cores *VMX execution units* are busy for 2 out of the 6 cycles.

So while the individual instructions have can have a throughput of 2/cycle, `vec_isnormalf32` can not. It is plausible for two executions of `vec_isnormalf32` to interleave with a delay of 1 cycle for the second sequence. To keep the table information simple for now, just say the throughput of `vec_isnormalf32` is 1/cycle.

After that it gets complicated. For example after the first two instances of `vec_isnormalf32` are issued, both *VMX execution units* are busy for 4 cycles. So either the first instructions of the third `vec_isnormalf32` will be delayed until the fifth cycle. Or the compiler scheduler will interleave instructions across the instances of `vec_isnormalf32` and the latencies of individual `vec_isnormalf32` results will increase. This is too complicated to put in a simple table.

For POWER9 the sequence is slightly different

```
addis    r10,r2,-2
addis    r9,r2,-2
xvabssp  vs45,vs34
addi     r10,r10,-14016
addi     r9,r9,-13920
lvx      v1,0,r10
lvx      v0,0,r9
xxland   vs34,vs34,vs33
vcmpgtuw v0,v0,v13
vcmpgequ v2,v2,v1
xxlnor   vs34,vs32,vs34
```

We use `vec_abs` (`xvabssp`) to replace the `sigmask` and `vec_andc` and so only need to load two vector constants. So the constant load overhead is reduced to 9 cycles. However the the vector compares are now 3 cycles for (2+3+2=) 7 cycles for the core sequence. The final table for `vec_isnormalf32`:

processor	Latency	Throughput
power8	6-16	1/cycle
power9	7-16	1/cycle

### 1.3.1 Additional analysis and tools.

The overview above is simplified analysis based on the instruction latency and throughput numbers published in the Processor User's Manuals (see [Reference Documentation](#)). These values are *best case* (input data is ready, SMT1 mode, no cache misses, mispredicted branches, or other hazards) for each instruction in isolation.

#### Note

This information is intended as a guide for compiler and application developers wishing to optimize for the platform. Any performance tables provided for pveclib functions are in this spirit.

Of course the actual performance is complicated by the overall environment and how the pveclib functions are used. It would be unusual for pveclib functions to be used in isolation. The compiler will in-line pveclib functions and look for sub-expressions it can hoist out of loops or share across pveclib function instances. The The compiler will also model the processor and schedule instructions across the larger containing function. So in actual use the instruction sequences for the examples above are likely to be interleaved with instructions from other pveclib functions and user written code.

Larger functions that use pveclib and even some of the more complicated pveclib functions (like `vec_muludq`) defy simple analysis. For these cases it is better to use POWER specific analysis tools. To understand the overall pipeline flows and identify hazards the instruction trace driven performance simulator is recommended.

The **IBM Advance Toolchain** includes an updated (POWER enabled) Valgrind tool and instruction trace plug-in (`itrace`). The `itrace` tool (`-tool=itrace`) collects instruction traces for the whole program or specific functions (via `-fname=` option).

**Note**

The Valgrind package provided by the Linux Distro may not be enabled for the latest POWER processor. Nor will it include the itrace plug-in or the associated vgi2qt conversion tool.

Instruction trace files are processed by the **Performance Simulator** (sim\_ppc) models. Performance simulators are specific to each processor generation (POWER7-9) and provides a cycle accurate modeling for instruction trace streams. The results of the model (a pipe file) can viewed via one the interactive display tools (scrollpv, jviewer) or passed to an analysis tool like **pipestat**.



## Chapter 2

## Todo List

Member [vec\\_bcdmul](#) (vui32\_t a, vui32\_t b)

This is tricky as the product can be up to 62 digits and \_Decimal128 can only hold 34 digits. So products that are more than 34 digits will lose precision on the right during the conversion back to BCD.



## Chapter 3

# Deprecated List

**Member `vec_slq4` (`vui128_t vra`)**

Vector Shift Left 4-bits Quadword. Replaced by `vec_slqi` with `shb param = 4`.

**Member `vec_slq5` (`vui128_t vra`)**

Vector Shift Left 5-bits Quadword. Replaced by `vec_slqi` with `shb param = 5`.

**Member `vec_spltd` (`vui64_t vra, const int ctl`)**

Vector splat doubleword. Duplicate the selected doubleword element across the doubleword elements of the result.

**Member `vec_srq4` (`vui128_t vra`)**

Vector Shift right 4-bits Quadword. Replaced by `vec_srq` with `shb param = 4`.

**Member `vec_srq5` (`vui128_t vra`)**

Vector Shift right 5-bits Quadword. Replaced by `vec_srq` with `shb param = 5`.





## Chapter 4

# Class Index

### 4.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

<a href="#">__VEC_U_128</a>	Union used to transfer 128-bit data between vector and non-vector types . . . . .	<a href="#">31</a>
<a href="#">__VF_128</a>	Union used to transfer 128-bit data between vector and __float128 types . . . . .	<a href="#">32</a>



## Chapter 5

# File Index

### 5.1 File List

Here is a list of all documented files with brief descriptions:

doc/ <a href="#">pveclibmaindox.h</a> . . . . .	??
src/ <a href="#">vec_bcd_ppc.h</a> Header package containing a collection of Binary Coded Decimal ( <b>BCD</b> ) computation functions implemented with PowerISA VMX, VSX, and DFP instructions . . . . .	33
src/ <a href="#">vec_char_ppc.h</a> Header package containing a collection of 128-bit SIMD operations over 8-bit integer (char) ele- ments . . . . .	36
src/ <a href="#">vec_common_ppc.h</a> Common definitions and typedef used by the collection of Power Vector Library (pveclib) headers	50
src/ <a href="#">vec_f128_ppc.h</a> Header package containing a collection of 128-bit SIMD operations over Quad-Precision floating point elements . . . . .	55
src/ <a href="#">vec_f32_ppc.h</a> Header package containing a collection of 128-bit SIMD operations over 4x32-bit floating point elements . . . . .	72
src/ <a href="#">vec_f64_ppc.h</a> Header package containing a collection of 128-bit SIMD operations over 64-bit double-precision floating point elements . . . . .	85
src/ <a href="#">vec_int128_ppc.h</a> Header package containing a collection of 128-bit computation functions implemented with PowerISA VMX and VSX instructions . . . . .	98
src/ <a href="#">vec_int16_ppc.h</a> Header package containing a collection of 128-bit SIMD operations over 16-bit integer elements	132
src/ <a href="#">vec_int32_ppc.h</a> Header package containing a collection of 128-bit SIMD operations over 32-bit integer elements	148
src/ <a href="#">vec_int64_ppc.h</a> Header package containing a collection of 128-bit SIMD operations over 64-bit integer elements	163



## Chapter 6

# Class Documentation

### 6.1 \_\_VEC\_U\_128 Union Reference

Union used to transfer 128-bit data between vector and non-vector types.

```
#include <vec_common_ppc.h>
```

#### Public Attributes

- unsigned \_\_int128 [i128](#)  
*Signed 128-bit integer from pair of 64-bit GPRs.*
- unsigned \_\_int128 [ui128](#)  
*Unsigned 128-bit integer from pair of 64-bit GPRs.*
- \_Decimal128 [dpc128](#)  
*128 bit Decimal Float from pair of double float registers.*
- long double [ldbl128](#)  
*IBM long double float from pair of double float registers.*
- [vui8\\_t vx16](#)  
*128 bit Vector of 16 unsigned char elements.*
- [vui16\\_t vx8](#)  
*128 bit Vector of 8 unsigned short int elements.*
- [vui32\\_t vx4](#)  
*128 bit Vector of 4 unsigned int elements.*
- [vui64\\_t vx2](#)  
*128 bit Vector of 2 unsigned long int (64-bit) elements.*
- [vui128\\_t vx1](#)  
*128 bit Vector of 1 unsigned \_\_int128 element.*
- [vf64\\_t vf2](#)  
*128 bit Vector of 2 double float elements.*
- struct {  
    uint64\_t [lower](#)  
    uint64\_t [upper](#)  
} [ulong](#)  
  
*Struct of two unsigned long int (64-bit GPR) fields.*

### 6.1.1 Detailed Description

Union used to transfer 128-bit data between vector and non-vector types.

The documentation for this union was generated from the following file:

- [src/vec\\_common\\_ppc.h](#)

## 6.2 \_\_VF\_128 Union Reference

Union used to transfer 128-bit data between vector and \_\_float128 types.

```
#include <vec_f128_ppc.h>
```

### Public Attributes

- [vui8\\_t vx16](#)  
*union field of vector unsigned char elements.*
- [vui16\\_t vx8](#)  
*union field of vector unsigned short elements.*
- [vui32\\_t vx4](#)  
*union field of vector unsigned int elements.*
- [vui64\\_t vx2](#)  
*union field of vector unsigned long long elements.*
- [vui128\\_t vx1](#)  
*union field of vector unsigned \_\_int128 elements.*
- [vb128\\_t vbool1](#)  
*union field of vector \_\_bool \_\_int128 elements.*
- [\\_\\_binary128 vf1](#)  
*union field of \_\_float128 elements.*

### 6.2.1 Detailed Description

Union used to transfer 128-bit data between vector and \_\_float128 types.

The documentation for this union was generated from the following file:

- [src/vec\\_f128\\_ppc.h](#)

# Chapter 7

## File Documentation

### 7.1 src/vec\_bcd\_ppc.h File Reference

Header package containing a collection of Binary Coded Decimal (**BCD**) computation functions implemented with PowerISA VMX, VSX, and DFP instructions.

```
#include <vec_common_ppc.h>
#include <vec_char_ppc.h>
```

#### Functions

- static `_Decimal128 bcd_qauntize0` (`_Decimal128 val`)  
*Quantize (truncate) a `_Decimal128` value before convert to BCD.*
- static `_Decimal128 vec_BCD2DFP` (`vui32_t val`)  
*Convert a Vector Signed BCD value to `__Decimal128`.*
- static `vui32_t vec_DFP2BCD` (`_Decimal128 val`)  
*Convert a `__Decimal128` value to Vector BCD.*
- static `vui32_t vec_bcdadd` (`vui32_t a`, `vui32_t b`)  
*Add two Vector Signed BCD 31 digit values.*
- static `vui32_t vec_bcddiv` (`vui32_t a`, `vui32_t b`)  
*Divide a Vector Signed BCD 31 digit value by another BCD value.*
- static `vui32_t vec_bcdmul` (`vui32_t a`, `vui32_t b`)  
*Multiply two Vector Signed BCD 31 digit values.*
- static `vui32_t vec_bcdsub` (`vui32_t a`, `vui32_t b`)  
*Subtract two Vector Signed BCD 31 digit values.*
- static `vui8_t vec_bcdctb100s` (`vui8_t vra`)  
*Vector Decimal Convert Binary Coded Decimal (BCD) digit pairs to binary integer bytes.*

#### 7.1.1 Detailed Description

Header package containing a collection of Binary Coded Decimal (**BCD**) computation functions implemented with PowerISA VMX, VSX, and DFP instructions.

#### Note

BCD and DFP support requires at least PowerISA 2.05 (POWER6) or later server level processor support. Plan to extend this header to cover the POWER9 Zoned decimal operations and conversions to / from BCD and DFP.

## 7.1.2 Function Documentation

### 7.1.2.1 `static __Decimal128 bcd_qauntize0 ( __Decimal128 val ) [inline],[static]`

Quantize (truncate) a `__Decimal128` value before convert to BCD.

The truncate (round toward 0) and justify right the input `__Decimal128` value so that the unit digit is in the right most position. This supports BCD multiply and divide using DFP instructions by truncating fractional digits before conversion back to BCD.

#### Parameters

<i>val</i>	a <code>__Decimal128</code> value.
------------	------------------------------------

#### Returns

The quantized `__Decimal128` value in a double float pair.

### 7.1.2.2 `static __Decimal128 vec_BCD2DFP ( vui32_t val ) [inline],[static]`

Convert a Vector Signed BCD value to `__Decimal128`.

The BCD vector is permuted into a double float pair before conversion to DPD format via the DFP Encode BCD To DPD Quad instruction.

#### Parameters

<i>val</i>	a 128-bit vector treated a signed BCD 31 digit value.
------------	---

#### Returns

a `__Decimal128` in a double float pair.

### 7.1.2.3 `static vui32_t vec_bcdadd ( vui32_t a, vui32_t b ) [inline],[static]`

Add two Vector Signed BCD 31 digit values.

Two Signed 31 digit values are added and lower 31 digits of the sum are returned. Overflow (carry-out) is ignored.

#### Parameters

<i>a</i>	a 128-bit vector treated a signed BCD 31 digit value.
<i>b</i>	a 128-bit vector treated a signed BCD 31 digit value.

#### Returns

a 128-bit vector which is the lower 31 digits of (a + b).



#### 7.1.2.4 static vui8\_t vec\_bcdctb100s ( vui8\_t vra ) [inline],[static]

Vector Decimal Convert Binary Coded Decimal (BCD) digit pairs to binary integer bytes.

Vector Convert 2 adjacent BCD digits, in each byte, to the corresponding binary char int value. Input values should be valid BCD nibbles in the range 0-9. The result will be binary char int values in the range 0-99.

This can be used as the first stage operation in a wider BCD to binary conversions. Basically the result of this stage are binary coded 100s "digits".

##### Parameters

<i>vra</i>	a 128-bit vector treated as a vector unsigned char of BCD pairs.
------------	--

##### Returns

128-bit vector unsigned char, For each byte, BCD digit pairs converted to the equivalent binary representation in the range 0-99.

#### 7.1.2.5 static vui32\_t vec\_bcddiv ( vui32\_t a, vui32\_t b ) [inline],[static]

Divide a Vector Signed BCD 31 digit value by another BCD value.

One Signed 31 digit value is divided by a second 31 digit value and the quotient is returned.

##### Parameters

<i>a</i>	a 128-bit vector treated a signed BCD 31 digit value.
<i>b</i>	a 128-bit vector treated a signed BCD 31 digit value.

##### Returns

a 128-bit vector which is the lower 31 digits of (a / b).

#### 7.1.2.6 static vui32\_t vec\_bcdmul ( vui32\_t a, vui32\_t b ) [inline],[static]

Multiply two Vector Signed BCD 31 digit values.

Two Signed 31 digit values are multiplied and the lower 31 digits of the product are returned. Overflow is ignored.

**Todo** This is tricky as the product can be up to 62 digits and \_Decimal128 can only hold 34 digits. So products that are more than 34 digits will loose precision on the right during the conversion back to BCD.

##### Parameters

<i>a</i>	a 128-bit vector treated a signed BCD 31 digit value.
<i>b</i>	a 128-bit vector treated a signed BCD 31 digit value.

**Returns**

a 128-bit vector which is the lower 31 digits of (a \* b).

**7.1.2.7** `static vui32_t vec_bcdsub ( vui32_t a, vui32_t b )` `[inline],[static]`

Subtract two Vector Signed BCD 31 digit values.

Subtract Signed 31 digit values and return the lower 31 digits of of the result. Overflow (carry-out/barrow) is ignored.

**Parameters**

<i>a</i>	a 128-bit vector treated a signed BCD 31 digit value.
<i>b</i>	a 128-bit vector treated a signed BCD 31 digit value.

**Returns**

a 128-bit vector which is the lower 31 digits of (a - b).

**7.1.2.8** `static vui32_t vec_DFP2BCD ( __Decimal128 val )` `[inline],[static]`

Convert a \_\_Decimal128 value to Vector BCD.

The \_\_Decimal128 value is converted to a signed BCD 31 digit value via "DFP Decode DPD To BCD Quad". The conversion result is still in a double float register pair and so is permuted into single vector register for use.

**Parameters**

<i>val</i>	a __Decimal128 in a double float pair.
------------	--

**Returns**

a 128-bit vector treated a signed BCD 31 digit value.

## 7.2 src/vec\_char\_ppc.h File Reference

Header package containing a collection of 128-bit SIMD operations over 8-bit integer (char) elements.

```
#include <vec_common_ppc.h>
```

**Functions**

- static [vui8\\_t vec\\_absdub](#) ([vui8\\_t](#) vra, [vui8\\_t](#) vrb)  
*Vector Absolute Difference Unsigned byte.*
- static [vui8\\_t vec\\_clzb](#) ([vui8\\_t](#) vra)

- Count leading zeros for a vector unsigned char (byte) elements.*
  - static `vui8_t vec_isalnum (vui8_t vec_str)`  
*Vector isalpha.*
  - static `vui8_t vec_isalpha (vui8_t vec_str)`  
*Vector isalnum.*
  - static `vui8_t vec_isdigit (vui8_t vec_str)`  
*Vector isdigit.*
  - static `vui8_t vec_mrgahb (vui16_t vra, vui16_t vrb)`  
*Vector Merge Algebraic High Byte operation.*
  - static `vui8_t vec_mrgalb (vui16_t vra, vui16_t vrb)`  
*Vector Merge Algebraic Low Byte operation.*
  - static `vui8_t vec_mrgeb (vui8_t vra, vui8_t vrb)`  
*Vector Merge Even Bytes operation.*
  - static `vui8_t vec_mrgob (vui8_t vra, vui8_t vrb)`  
*Vector Merge Odd Halfwords operation.*
  - static `vi8_t vec_mulhsb (vi8_t vra, vi8_t vrb)`  
*Vector Multiply High Signed Bytes.*
  - static `vui8_t vec_mulhub (vui8_t vra, vui8_t vrb)`  
*Vector Multiply High Unsigned Bytes.*
  - static `vui8_t vec_mulubm (vui8_t vra, vui8_t vrb)`  
*Vector Multiply Unsigned Byte Modulo.*
  - static `vui8_t vec_popcntb (vui8_t vra)`  
*Vector Population Count byte.*
  - static `vui8_t vec_slbi (vui8_t vra, const unsigned int shb)`  
*Vector Shift left Byte Immediate.*
  - static `vi8_t vec_srabi (vi8_t vra, const unsigned int shb)`  
*Vector Shift Right Algebraic Byte Immediate.*
  - static `vui8_t vec_srbi (vui8_t vra, const unsigned int shb)`  
*Vector Shift Right Byte Immediate.*
  - static `vui8_t vec_shift_leftdo (vui8_t vrw, vui8_t vrx, vui8_t vrb)`  
*Shift left double quadword by octet. Return a vector unsigned char that is the left most 16 chars after shifting left 0-15 octets (chars) of the 32 char double vector (vrw||vrx). The octet shift amount is from bits 121:124 of vrb.*
  - static `vui8_t vec_toupper (vui8_t vec_str)`  
*Vector toupper.*
  - static `vui8_t vec_tolower (vui8_t vec_str)`  
*Vector tolower.*
  - static `vui8_t vec_vmrgeb (vui8_t vra, vui8_t vrb)`  
*Vector Merge Even Bytes.*
  - static `vui8_t vec_vmrgob (vui8_t vra, vui8_t vrb)`  
*Vector Merge Odd Byte.*

### 7.2.1 Detailed Description

Header package containing a collection of 128-bit SIMD operations over 8-bit integer (char) elements.

Most of these operations are implemented in a single VMX or VSX instruction on newer (POWER6/POWER7/POWER8/POWER9) processors. This header serves to fill in functional gaps for older (POWER7, POWER8) processors and provides in-line assembler implementations for older compilers that do not provide the build-ins.

Most vector char (8-bit integer) operations are already covered by the original VMX (AKA AltiVec) instructions. VMX intrinsic (compiler built-ins) operations are defined in `<altivec.h>` and described in the compiler documentation. PowerISA 2.07B (POWER8) added several useful byte operations (count leading zeros, population count) not included in the original VMX. PowerISA 3.0B (POWER9) adds several more (absolute difference, compare not equal, count trailing zeros, extend sign, extract/insert, and reverse bytes). Most of these intrinsic (compiler built-ins) operations are defined in `<altivec.h>` and described in the compiler documentation.

**Note**

The compiler disables associated `<altivec.h>` built-ins if the **mcpu** target does not enable the specific instruction. For example if you compile with **-mcpu=power7**, `vec_vclz` and `vec_vclzb` will not be defined. But `vec_clzb` is always defined in this header, will generate the minimum code, appropriate for the target, and produce correct results.

This header covers operations that are either:

- Implemented in later processors and useful to programmers if the same operations are available on slightly older processors. This is required even if the operation is defined in the OpenPOWER ABI or `<altivec.h>`, as the compiler disables the associated built-ins if the **mcpu** target does not enable the instruction.
- Defined in the OpenPOWER ABI but *not* yet defined in `<altivec.h>` provided by available compilers in common use. Examples include Count Leading Zeros and Population Count.
- Commonly used operations, not covered by the ABI or `<altivec.h>`, and require multiple instructions or are not obvious. Examples include the multiply high, ASCII character tests, and shift immediate operations.

### 7.2.2 Endian problems with byte operations

It would be useful to provide a vector multiply high byte (return the high order 8-bits of the 16-bit product) operation. This can be used for multiplicative inverse (effectively integer divide) operations. Neither integer multiply high nor divide are available as vector instructions. However the multiply high byte operation can be composed from the existing multiply even/odd byte operations followed by the vector merge even byte operation. Similarly a multiply low (modulo) byte operation can be composed from the existing multiply even/odd byte operations followed by the vector merge odd byte operation.

As a prerequisite we need to provide the merge even/odd byte operations. While PowerISA has added these operations for word and doubleword, instructions are not defined for byte and halfword. Fortunately vector merge operations are just a special case of vector permute. So the `vec_vmrgeb()` and `vec_vmrgeb()` implementation can use `vec_perm` and appropriate selection vectors to provide these merge operations.

As described for other element sizes this is complicated by *little-endian* (LE) support as specified in the OpenPOWER ABI and as implemented in the compilers. Little-endian changes the effective vector element numbering and the location of even and odd elements. This means that the vector built-ins provided by `altivec.h` may not generate the instructions you would expect.

**See also**

[Endian problems with halfword operations](#)  
[General Endian Issues](#)

So this header defines endian independent byte operations `vec_vmrgeb()` and `vec_vmrgeb()`. These operations are used in the implementation of the endian sensitive `vec_mrgbeb()` and `vec_mrgbeb()`. These support the OpenPOWER ABI mandated merge even/odd semantic.

We also provide the merge algebraic high/low operations `vec_mrgahb()` and `vec_mrgalb()` to simplify extended precision arithmetic. These implementations use `vec_vmrgeb()` and `vec_vmrgeb()` as extended precision byte order does not change with endian. These operations are used in turn to implement multiply byte high/low/modulo (`vec_↵_mulhsb()`, `vec_mulhub()`, `vec_mulubm()`).

These operations provide a basis for using the multiplicative inverse as a alternative to integer divide.

**See also**

[Examples, Divide by integer constant](#)

### 7.2.3 Performance data.

The performance characteristics of the merge and multiply byte operations are very similar to the halfword implementations. (see [Performance data](#)).

#### 7.2.3.1 More information.

High level performance estimates are provided as an aid to function selection when evaluating algorithms. For background on how *Latency* and *Throughput* are derived see: [Performance data](#).

### 7.2.4 Function Documentation

#### 7.2.4.1 static vui8\_t vec\_absdub ( vui8\_t vra, vui8\_t vrb ) [inline],[static]

Vector Absolute Difference Unsigned byte.

Compute the absolute difference for each byte. For each unsigned byte, subtract B[i] from A[i] and return the absolute value of the difference.

processor	Latency	Throughput
power8	4	1/cycle
power9	3	2/cycle

#### Parameters

<i>vra</i>	vector of 16 unsigned bytes
<i>vr</i> <i>b</i>	vector of 16 unsigned bytes

#### Returns

vector of the absolute difference.

#### 7.2.4.2 static vui8\_t vec\_clzb ( vui8\_t vra ) [inline],[static]

Count leading zeros for a vector unsigned char (byte) elements.

Count the number of leading '0' bits (0-7) within each byte element of a 128-bit vector.

For POWER8 (PowerISA 2.07B) or later use the Vector Count Leading Zeros byte instruction **vclzb**. Otherwise use sequence of pre 2.07 VMX instructions. SIMDized count leading zeros inspired by:

Warren, Henry S. Jr and *Hacker's Delight*, 2nd Edition, Addison Wesley, 2013. Chapter 5 Counting Bits, Figure 5-12.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

## Parameters

<i>vra</i>	128-bit vector treated as 16 x 8-bit integer (byte) elements.
------------	---

## Returns

128-bit vector with the Leading Zeros count for each byte element.

**7.2.4.3** `static vui8_t vec_isalnum ( vui8_t vec_str )` `[inline],[static]`

Vector isalpha.

Return a vector boolean char with a true indicator for any character that is either Lower Case Alpha ASCII or Upper Case ASCII. False otherwise.

processor	Latency	Throughput
power8	10-20	1/cycle
power9	11-21	1/cycle

## Parameters

<i>vec_str</i>	vector of 16 ASCII characters
----------------	-------------------------------

## Returns

vector bool char of the isalpha operation applied to each character of *vec\_str*. For each byte 0xff indicates true (isalpha), 0x00 indicates false.

**7.2.4.4** `static vui8_t vec_isalpha ( vui8_t vec_str )` `[inline],[static]`

Vector isalnum.

Return a vector boolean char with a true indicator for any character that is either Lower Case Alpha ASCII, Upper Case ASCII, or numeric ASCII. False otherwise.

processor	Latency	Throughput
power8	9-18	1/cycle
power9	10-19	1/cycle

## Parameters

<i>vec_str</i>	vector of 16 ASCII characters
----------------	-------------------------------

**Returns**

vector bool char of the isalnum operation applied to each character of vec\_str. For each byte 0xff indicates true (isalpha), 0x00 indicates false.

**7.2.4.5** static vui8\_t vec\_isdigit ( vui8\_t vec\_str ) [inline],[static]

Vector isdigit.

Return a vector boolean char with a true indicator for any character that is ASCII decimal digit. False otherwise.

processor	Latency	Throughput
power8	4-13	1/cycle
power9	5-14	1/cycle

**Parameters**

vec_str	vector of 16 ASCII characters
---------	-------------------------------

**Returns**

vector bool char of the isdigit operation applied to each character of vec\_str. For each byte 0xff indicates true (isdigit), 0x00 indicates false.

**7.2.4.6** static vui8\_t vec\_mrgahb ( vui16\_t vra, vui16\_t vrb ) [inline],[static]

Vector Merge Algebraic High Byte operation.

Merge only the high byte from 16 x Algebraic halfwords across vectors vra and vrb. This is effectively the Vector Merge Even Byte operation that is not modified for Endian.

For example merge the high 8-bits from each of 16 x 16-bit products as generated by vec\_muleub/vec\_muloub. This result is effectively a vector multiply high unsigned byte.

processor	Latency	Throughput
power8	2-13	2/cycle
power9	3-14	2/cycle

**Parameters**

vra	128-bit vector unsigned short.
vrh	128-bit vector unsigned short.

**Returns**

A vector merge from only the high bytes of the 16 x Algebraic halfwords across vra and vrb.

**7.2.4.7** `static vui8_t vec_mrgalb ( vui16_t vra, vui16_t vrb ) [inline],[static]`

Vector Merge Algebraic Low Byte operation.

Merge only the low bytes from 16 x Algebraic halfwords across vectors *vra* and *vr**b*. This is effectively the Vector Merge Odd Bytes operation that is not modified for Endian.

For example merge the low 8-bits from each of 16 x 16-bit products as generated by `vec_muleub/vec_muloub`. This result is effectively a vector multiply low unsigned byte.

processor	Latency	Throughput
power8	2-13	2/cycle
power9	3-14	2/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned int.
<i>vr</i> <i>b</i>	128-bit vector unsigned int.

#### Returns

A vector merge from only the high halfwords of the 8 x Algebraic words across *vra* and *vr**b*.

**7.2.4.8** `static vui8_t vec_mrgeb ( vui8_t vra, vui8_t vrb ) [inline],[static]`

Vector Merge Even Bytes operation.

Merge the even byte elements from the concatenation of 2 x vectors (*vra* and *vr**b*).

#### Note

The element numbering changes between Big and Little Endian. So the compiler and this implementation adjusts the generated code to reflect this.

processor	Latency	Throughput
power8	2-13	2/cycle
power9	3-14	2/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned char.
<i>vr</i> <i>b</i>	128-bit vector unsigned char.

#### Returns

A vector merge from only the even bytes of *vra* and *vr**b*.



7.2.4.9 `static vui8_t vec_mrgob ( vui8_t vra, vui8_t vrb ) [inline],[static]`

Vector Merge Odd Halfwords operation.

Merge the odd halfword elements from the concatenation of 2 x vectors (vra and vrb).

#### Note

The element numbering changes between Big and Little Endian. So the compiler and this implementation adjusts the generated code to reflect this.

processor	Latency	Throughput
power8	2-13	2/cycle
power9	3-14	2/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned char.
<i>vrb</i>	128-bit vector unsigned char.

#### Returns

A vector merge from only the odd bytes of vra and vrb.

7.2.4.10 `static vi8_t vec_mulhsb ( vi8_t vra, vi8_t vrb ) [inline],[static]`

Vector Multiply High Signed Bytes.

Multiple the corresponding byte elements of two vector signed char values and return the high order 8-bits, for each 16-bit product element.

processor	Latency	Throughput
power8	9-13	1/cycle
power9	10-14	1/cycle

#### Parameters

<i>vra</i>	128-bit vector signed char.
<i>vrb</i>	128-bit vector signed char.

#### Returns

vector of the high order 8-bits of the product of the byte elements from vra and vrb.

7.2.4.11 `static vui8_t vec_mulhub ( vui8_t vra, vui8_t vrb ) [inline],[static]`

Vector Multiply High Unsigned Bytes.

Multiple the corresponding byte elements of two vector unsigned char values and return the high order 8-bits, for each 16-bit product element.

processor	Latency	Throughput
power8	9-13	1/cycle
power9	10-14	1/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned char.
<i>vrh</i>	128-bit vector unsigned char.

#### Returns

vector of the high order 8-bits of the product of the byte elements from *vra* and *vrh*.

**7.2.4.12** `static vui8_t vec_mulubm ( vui8_t vra, vui8_t vrh ) [inline],[static]`

Vector Multiply Unsigned Byte Modulo.

Multiple the corresponding byte elements of two vector unsigned char values and return the low order 8-bits of the 16-bit product for each element.

#### Note

`vec_mulubm` can be used for unsigned or signed char integers. It is the vector equivalent of Multiply Low Byte.

processor	Latency	Throughput
power8	9-13	1/cycle
power9	10-14	1/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned char.
<i>vrh</i>	128-bit vector unsigned char.

#### Returns

vector of the low order 8-bits of the unsigned product of the byte elements from *vra* and *vrh*.

**7.2.4.13** `static vui8_t vec_popcntb ( vui8_t vra ) [inline],[static]`

Vector Population Count byte.

Count the number of '1' bits (0-8) within each byte element of a 128-bit vector.

For POWER8 (PowerISA 2.07B) or later use the Vector Population Count Byte instruction. Otherwise use simple Vector (VMX) instructions to count bits in bytes in parallel. SIMDized population count inspired by:

Warren, Henry S. Jr and *Hacker's Delight*, 2nd Edition, Addison Wesley, 2013. Chapter 5 Counting Bits, Figure 5-2.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

#### Parameters

<i>vra</i>	128-bit vector treated as 16 x 8-bit integers (byte) elements.
------------	--

#### Returns

128-bit vector with the population count for each byte element.

**7.2.4.14** `static vui8_t vec_shift_leftdo ( vui8_t vrw, vui8_t vrx, vui8_t vrb )` `[inline], [static]`

Shift left double quadword by octet. Return a vector unsigned char that is the left most 16 chars after shifting left 0-15 octets (chars) of the 32 char double vector (*vrw*||*vrx*). The octet shift amount is from bits 121:124 of *vrb*.

This sequence can be used to align a unaligned 16 char substring based on the result of a vector count leading zero of of the compare boolean.

processor	Latency	Throughput
power8	6-8	1/cycle
power9	8-9	1/cycle

#### Parameters

<i>vrw</i>	upper 16-bytes of the 32-byte double vector.
<i>vrx</i>	lower 16-bytes of the 32-byte double vector.
<i>vrb</i>	Shift amount in bits 121:124.

#### Returns

upper 16-bytes of left shifted double vector.

**7.2.4.15** `static vui8_t vec_slbi ( vui8_t vra, const unsigned int shb )` `[inline], [static]`

Vector Shift left Byte Immediate.

Shift left each byte element [0-15], 0-7 bits, as specified by an immediate value. The shift amount is a const unsigned int in the range 0-7. A shift count of 0 returns the original value of *vra*. Shift counts greater than 7 bits return zero.

processor	Latency	Throughput
power8	4-11	2/cycle
power9	5-11	2/cycle

## Parameters

<i>vra</i>	a 128-bit vector treated as a vector unsigned char.
<i>shb</i>	Shift amount in the range 0-7.

## Returns

128-bit vector unsigned char, shifted left shb bits.

**7.2.4.16** `static vi8_t vec_srabi ( vi8_t vra, const unsigned int shb )` `[inline], [static]`

Vector Shift Right Algebraic Byte Immediate.

Shift right each byte element [0-15], 0-7 bits, as specified by an immediate value. The shift amount is a const unsigned int in the range 0-7. A shift count of 0 returns the original value of vra. Shift counts greater than 7 bits return the sign bit propagated to each bit of each element.

processor	Latency	Throughput
power8	4-11	2/cycle
power9	5-11	2/cycle

## Parameters

<i>vra</i>	a 128-bit vector treated as a vector signed char.
<i>shb</i>	Shift amount in the range 0-7.

## Returns

128-bit vector signed char, shifted right shb bits.

**7.2.4.17** `static vui8_t vec_srbi ( vui8_t vra, const unsigned int shb )` `[inline], [static]`

Vector Shift Right Byte Immediate.

Shift right each byte element [0-15], 0-7 bits, as specified by an immediate value. The shift amount is a const unsigned int in the range 0-7. A shift count of 0 returns the original value of vra. Shift counts greater than 7 bits return zero.

processor	Latency	Throughput
power8	4-11	2/cycle
power9	5-11	2/cycle

## Parameters

<i>vra</i>	a 128-bit vector treated as a vector unsigned char.
<i>shb</i>	Shift amount in the range 0-7.

**Returns**

128-bit vector unsigned char, shifted right shb bits.

**7.2.4.18** `static vui8_t vec_tolower ( vui8_t vec_str )` `[inline],[static]`

Vector tolower.

Convert any Upper Case Alpha ASCII characters within a vector unsigned char into the equivalent Lower Case character. Return the result as a vector unsigned char.

processor	Latency	Throughput
power8	8-17	1/cycle
power9	9-18	1/cycle

**Parameters**

<code>vec_str</code>	vector of 16 ASCII characters
----------------------	-------------------------------

**Returns**

vector char converted to lower case.

**7.2.4.19** `static vui8_t vec_toupper ( vui8_t vec_str )` `[inline],[static]`

Vector toupper.

Convert any Lower Case Alpha ASCII characters within a vector unsigned char into the equivalent Upper Case character. Return the result as a vector unsigned char.

processor	Latency	Throughput
power8	8-17	1/cycle
power9	9-18	1/cycle

**Parameters**

<code>vec_str</code>	vector of 16 ASCII characters
----------------------	-------------------------------

**Returns**

vector char converted to upper case.

**7.2.4.20** `static vui8_t vec_vmrgeb ( vui8_t vra, vui8_t vrb )` `[inline],[static]`

Vector Merge Even Bytes.

Merge the even byte elements from the concatenation of 2 x vectors (vra and vrb).

**Note**

This function implements the operation of a Vector Merge Even Bytes instruction, if the PowerISA included such an instruction. This implementation is NOT Endian sensitive and the function is stable across BE/LE implementations. Using Big Endian element numbering:

- res[0] = vra[0];
- res[1] = vrb[0];
- res[2] = vra[2];
- res[3] = vrb[2];
- res[4] = vra[4];
- res[5] = vrb[4];
- res[6] = vra[6];
- res[7] = vrb[6];
- res[8] = vra[8];
- res[9] = vrb[8];
- res[10] = vra[10];
- res[11] = vrb[10];
- res[12] = vra[12];
- res[13] = vrb[12];
- res[14] = vra[14];
- res[15] = vrb[14];

processor	Latency	Throughput
power8	2-13	2/cycle
power9	3-14	2/cycle

**Parameters**

<i>vra</i>	128-bit vector unsigned char.
<i>vrb</i>	128-bit vector unsigned char.

**Returns**

A vector merge from only the even bytes of *vra* and *vrb*.

**7.2.4.21** `static vui8_t vec_vmrgob ( vui8_t vra, vui8_t vrb ) [inline], [static]`

Vector Merge Odd Byte.

Merge the odd byte elements from the concatenation of 2 x vectors (*vra* and *vrb*).

**Note**

This function implements the operation of a Vector Merge Odd Bytes instruction, if the PowerISA included such an instruction. This implementation is NOT Endian sensitive and the function is stable across BE/LE implementations. Using Big Endian element numbering:

- res[0] = vra[1];

- `res[1] = vrb[1];`
- `res[2] = vra[3];`
- `res[3] = vrb[3];`
- `res[4] = vra[5];`
- `res[5] = vrb[5];`
- `res[6] = vra[7];`
- `res[7] = vrb[7];`
- `res[8] = vra[9];`
- `res[9] = vrb[9];`
- `res[10] = vra[11];`
- `res[11] = vrb[11];`
- `res[12] = vra[13];`
- `res[13] = vrb[13];`
- `res[14] = vra[15];`
- `res[15] = vrb[15];`

processor	Latency	Throughput
power8	2-13	2/cycle
power9	3-14	2/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned char.
<i>vrb</i>	128-bit vector unsigned char.

#### Returns

A vector merge from only the odd bytes of *vra* and *vrb*.

## 7.3 src/vec\_common\_ppc.h File Reference

Common definitions and typedef used by the collection of Power Vector Library (pveclib) headers.

```
#include <stdint.h>
#include <altivec.h>
```

#### Classes

- union [\\_\\_VEC\\_U\\_128](#)

*Union used to transfer 128-bit data between vector and non-vector types.*



## Macros

- #define `CONST_VINT64_DW`(\_\_dw0, \_\_dw1) {\_\_dw1, \_\_dw0}  
*Arrange elements of dword initializer in high->low order.*
- #define `CONST_VINT128_DW`(\_\_dw0, \_\_dw1) (`vui64_t`){\_\_dw1, \_\_dw0}  
*Initializer for 128-bits vector, as two unsigned long long elements in high->low order. May require an explicit cast.*
- #define `CONST_VINT128_DW128`(\_\_dw0, \_\_dw1) (`vui128_t`)((`vui64_t`){\_\_dw1, \_\_dw0})  
*A vector unsigned \_\_int128 initializer, as two unsigned long long elements in high->low order.*
- #define `CONST_VINT128_W`(\_\_w0, \_\_w1, \_\_w2, \_\_w3) (`vui32_t`){\_\_w3, \_\_w2, \_\_w1, \_\_w0}  
*Arrange word elements of a unsigned int initializer in high->low order. May require an explicit cast.*
- #define `CONST_VINT32_W`(\_\_w0, \_\_w1, \_\_w2, \_\_w3) {\_\_w3, \_\_w2, \_\_w1, \_\_w0}  
*Arrange elements of word initializer in high->low order.*
- #define `CONST_VINT128_H`(\_\_hw0, \_\_hw1, \_\_hw2, \_\_hw3, \_\_hw4, \_\_hw5, \_\_hw6, \_\_hw7) (`vui16_t`){\_\_hw7, \_\_hw6, \_\_hw5, \_\_hw4, \_\_hw3, \_\_hw2, \_\_hw1, \_\_hw0}  
*Arrange halfword elements of a unsigned int initializer in high->low order. May require an explicit cast.*
- #define `CONST_VINT16_H`(\_\_hw0, \_\_hw1, \_\_hw2, \_\_hw3, \_\_hw4, \_\_hw5, \_\_hw6, \_\_hw7) {\_\_hw7, \_\_hw6, \_\_hw5, \_\_hw4, \_\_hw3, \_\_hw2, \_\_hw1, \_\_hw0}  
*Arrange elements of halfword initializer in high->low order.*
- #define `CONST_VINT128_B`(\_b0, \_b1, \_b2, \_b3, \_b4, \_b5, \_b6, \_b7, \_b8, \_b9, \_b10, \_b11, \_b12, \_b13, \_b14, \_b15) (`vui8_t`){\_b15, \_b14, \_b13, \_b12, \_b11, \_b10, \_b9, \_b8, \_b7, \_b6, \_b5, \_b4, \_b3, \_b2, \_b1, \_b0}  
*Arrange byte elements of a unsigned int initializer in high->low order. May require an explicit cast.*
- #define `CONST_VINT8_B`(\_b0, \_b1, \_b2, \_b3, \_b4, \_b5, \_b6, \_b7, \_b8, \_b9, \_b10, \_b11, \_b12, \_b13, \_b14, \_b15) {\_b15, \_b14, \_b13, \_b12, \_b11, \_b10, \_b9, \_b8, \_b7, \_b6, \_b5, \_b4, \_b3, \_b2, \_b1, \_b0}  
*Arrange elements of byte initializer in high->low order.*
- #define `VEC_DW_H` 1  
*Element index for high order dword.*
- #define `VEC_DW_L` 0  
*Element index for low order dword.*
- #define `VEC_W_H` 3  
*Element index for highest order word.*
- #define `VEC_W_L` 0  
*Element index for lowest order word.*
- #define `VEC_WE_0` 3  
*Element index for vector splat word 0.*
- #define `VEC_WE_1` 2  
*Element index for vector splat word 1.*
- #define `VEC_WE_2` 1  
*Element index for vector splat word 2.*
- #define `VEC_WE_3` 0  
*Element index for vector splat word 3.*
- #define `VEC_HW_H` 7  
*Element index for highest order hword.*
- #define `VEC_HW_L` 0  
*Element index for lowest order hword.*
- #define `VEC_BYTE_L` 0  
*Element index for lowest order byte.*
- #define `VEC_BYTE_L_DWH` 8  
*Element index for lowest order byte of the high dword.*
- #define `VEC_BYTE_L_DWL` 0  
*Element index for lowest order byte of the low dword.*
- #define `VEC_BYTE_H` 15  
*Element index for highest order byte.*
- #define `VEC_BYTE_HHW` 14  
*Element index for second lowest order byte.*

## Typedefs

- typedef \_\_vector unsigned char [vui8\\_t](#)  
*vector of 8-bit unsigned char elements.*
- typedef \_\_vector unsigned short [vui16\\_t](#)  
*vector of 16-bit unsigned short elements.*
- typedef \_\_vector unsigned int [vui32\\_t](#)  
*vector of 32-bit unsigned int elements.*
- typedef \_\_vector unsigned long long [vui64\\_t](#)  
*vector of 64-bit unsigned long long elements.*
- typedef \_\_vector signed char [vi8\\_t](#)  
*vector of 8-bit signed char elements.*
- typedef \_\_vector short [vi16\\_t](#)  
*vector of 16-bit signed short elements.*
- typedef \_\_vector int [vi32\\_t](#)  
*vector of 32-bit signed int elements.*
- typedef \_\_vector long long [vi64\\_t](#)  
*vector of 64-bit signed long long elements.*
- typedef \_\_vector float [vf32\\_t](#)  
*vector of 32-bit float elements.*
- typedef \_\_vector double [vf64\\_t](#)  
*vector of 64-bit double elements.*
- typedef \_\_vector \_\_bool char [vb8\\_t](#)  
*vector of 8-bit bool char elements.*
- typedef \_\_vector \_\_bool short [vb16\\_t](#)  
*vector of 16-bit bool short elements.*
- typedef \_\_vector \_\_bool int [vb32\\_t](#)  
*vector of 32-bit bool int elements.*
- typedef \_\_vector \_\_bool long long [vb64\\_t](#)  
*vector of 64-bit bool long long elements.*
- typedef \_\_vector int [vi128\\_t](#)  
*vector of one 128-bit signed \_\_int128 element.*
- typedef \_\_vector unsigned int [vui128\\_t](#)  
*vector of one 128-bit unsigned \_\_int128 element.*
- typedef \_\_vector \_\_bool unsigned int [vb128\\_t](#)  
*vector of 128-bit bool \_\_int128 elements.*

### 7.3.1 Detailed Description

Common definitions and typedef used by the collection of Power Vector Library (pveclib) headers.

This includes:

- Typedefs as short names of common vector types.
- Union used to transfer 128-bit data between vector and non-vector types.
- Helper macros that make declaring constants and accessing elements a little easier.

### 7.3.2 Consistent vector type naming

Type names should be short, concise, and consistent. The ABI defines the vector types as extensions of the existing C Language types. So while *vector unsigned long long* is consistent it is neither short or concise. Pveclib uses the following naming convention for typedefs used in its operations, function prototypes, and internal variables.

- Starting with the **v** prefix for vector.
- followed by one of the element classes:
  - **i** for signed integer.
  - **ui** for unsigned integer.
  - **f** for floating-point.
  - **b** for bool.
- followed by the element size in bits:
  - 8, 16, 32, 64, 128
- Ending with the **\_t** suffix signifying a typedef.

For example: `vi32_t` is a vector int, `vui32_t` is a vector unsigned int, `vb32_t` is a vector bool int, and `vf32_t` is vector float.

### 7.3.3 Transferring 128-bit types

The OpenPOWER ABI and the GCC compiler define a number of 128-bit scalar types that are not vector types:

- `__int128` (a general purpose register pair)
- `_Decimal128` (a floating-point even/odd register pair)
- `__ibm128` (a floating-point register pair)
- `__float128` (a vector register)

These are not cast nor assignment compatible with any vector type. However it may be useful to transfer to/from vector types for conversion or manipulation within an operation. For example:

- Conversions between `__float128` and `__int128`, `__ibm128`, and `_Decimal128` types.
- Conversions between vector BCD integers and `__int128` and `_Decimal128` types.
- Conversions between vector `__int128` and `__float128`, `__ibm128`, and `_Decimal128` types.

Here we use the `__VEC_U_128` union to affect the transfer between the various types. We assume (fervently hope) that the compiler will recognize and optimize these as registers to registers transfers using the hardware instructions provided.

The vector to/from `__float128` transfer should be the simplest as `__float128` operations are defined over the vector register set. However `__float128` types are defined in the PowerISA and OpenPOWER ABI, as scalars that just happens to use vector registers for parameter passing and operations. This distinction between scalars and vector prevents a direct cast between types. The `__VEC_U_128` union is the simplest work around but in most cases no code should generated for this transfer. For example: `vec_xfer_bin128_2_vui128t()` and `vec_xfer_vui128t_2_bin128()`.

Any vector to/from `__int128` transfer requires a transfer between vector and general purpose registers. POWER8 (PowerISA 2.07B) added Move to/from Vector Scalar Register (`mfvsr`, `mtvsr`) instructions. Again the `__VEC_U_128` union is used to effect the transfer and the compiler should leverage the move instructions in the generated code.

Any vector to/from `__ibm128` or `_Decimal128` requires a transfer between a pair of FPRs and a Vector Scalar Register (VSR). Technically this is transfer between the upper doubleword of two VSRs in the lower bank (VS $\leftarrow$ R0-31) and another VSR. POWER7 (PowerISA 2.06B) provides the Permute Doubleword Immediate (`xxpermdi`) instruction. Again the `__VEC_U_128` union is used to effect the transfer and the compiler should leverage the Permute Doubleword Immediate instructions in the generate code. For example: `vec_BCD2DFP()` and `vec_DF $\leftarrow$ P2BCD()`.

### 7.3.4 Endian and vector constants

Vector constants are often needed for: masking operations, range checks, permute selection, and radix conversion. Also compiler support for large integer and floating-point constants may be limited by the compiler. For example the GCC compilers support the (vector) `__int128` type but do not directly support `__int128` (39 digit) decimal constants. Another example is `__float128` where the type and Q suffix constants are recent additions. In both cases we need to construct: large numeric constants, special values (infinity and NaN), masks for manipulating the sign bit and exponent bits. Often these values will be constructed from vectors of word or doubleword constants.

#### Note

GCC does not support expressing an integer constant of type `__int128` for targets where long long integer is less than 128 bits wide. This applies to the PowerPC target as the long long type is reserved for 64-bit integers. This was verified in GCC 8.2, GCC `__float128` support for the PowerPC target began with GCC 6. In GCC 6 `__float128` support is off by default and has to be explicitly enabled via the `-mfloat128` option. Starting with GCC 7, `__float128` is enabled by default with VSX support.

Defining large constants for vectors is complicated by *little-endian* (LE) support as specified in the OpenPOWER ABI and as implemented in the compilers. Little-endian changes the effective vector element numbering and the order of constant elements in initializers. But the `__int128` numerical order of magnitude or floating-point format does not change in registers. The high order bits are on the left and the low order bits are on the right.

So for example:

```
const vui32_t signmask = { 0x80000000, 0, 0, 0 };
const vui32_t expmask = { 0x7fff0000, 0, 0, 0 };
```

are correct sign and exponent masks for `__float128` in big endian (BE) but would be incorrect for little endian (LE). To get correct results for both endians, one could code something like this:

```
#if __BYTE_ORDER__ == __ORDER_LITTLE_ENDIAN__
    const vui32_t signmask = { 0, 0, 0, 0x80000000 };
    const vui32_t expmask = { 0, 0, 0, 0x7fff0000 };
#else
    const vui32_t signmask = { 0x80000000, 0, 0, 0 };
    const vui32_t expmask = { 0x7fff0000, 0, 0, 0 };
#endif
```

But this gets tedious after the first dozen times. Also this can be confusing because it does not appear to match the floating-point format diagrams in the PowerISA. The sign-bit and the exponent are always on the left.

So this header provides endian sensitive macros that maintain consistent "magnitude" order. For example:

```
const vui32_t signmask = CONST_VINT128_W (0x80000000, 0, 0, 0);
const vui32_t expmask = CONST_VINT128_W (0x7fff0000, 0, 0, 0);
```

This is always correct in either endian.

Another example; the multiplicative inverse for `__int128 10**32` is 211857340822306639531405861550393824741. The GCC compiler will not accept this constant in a vector `__int128` initializer. The next best thing would be

```
// The multiplicative inverse for 1 / 10**32 is
// 211857340822306639531405861550393824741
// or 0x9f623d5a8a732974cfbc31db4b0295e5
const vui128_t mulinv_10to32 =
    (vui128_t) CONST_VINT128_DW128 ( 0x9f623d5a8a732974UL,
                                      0xcfb31db4b0295e5UL );
```

Here we use the `CONST_VINT128_DW128` macro to maintain magnitude order across endian. Again the high order bits are on the left and the low order bits are on the right.

#### See also

[Endian problems with word operations](#)  
[General Endian Issues](#)

## 7.4 src/vec\_f128\_ppc.h File Reference

Header package containing a collection of 128-bit SIMD operations over Quad-Precision floating point elements.

```
#include <vec_common_ppc.h>
#include <vec_int128_ppc.h>
#include <vec_f64_ppc.h>
```

### Classes

- union [\\_\\_VF\\_128](#)

*Union used to transfer 128-bit data between vector and `__float128` types.*

### Typedefs

- typedef [vui128\\_t](#) [vf128\\_t](#)  
*vector of 128-bit binary128 element. Same as `__float128` for PPC.*
- typedef [vf128\\_t](#) [\\_\\_Float128](#)  
*Define `__Float128` if not defined by the compiler. Same as `__float128` for PPC.*
- typedef [vf128\\_t](#) [\\_\\_binary128](#)  
*Define `__binary128` if not defined by the compiler. Same as `__float128` for PPC.*
- typedef [vf128\\_t](#) [\\_\\_float128](#)  
*Define `__float128` if not defined by the compiler. Same as `__float128` for PPC.*
- typedef long double [\\_\\_IBM128](#)  
*Define `__IBM128` if not defined by the compiler. Same as old long double for PPC.*

### Functions

- static [vui8\\_t](#) [vec\\_xfer\\_bin128\\_2\\_vui8t](#) ([\\_\\_binary128](#) f128)  
*Transfer function from a `__binary128` scalar to a vector char.*
- static [vui16\\_t](#) [vec\\_xfer\\_bin128\\_2\\_vui16t](#) ([\\_\\_binary128](#) f128)  
*Transfer function from a `__binary128` scalar to a vector short int.*
- static [vui32\\_t](#) [vec\\_xfer\\_bin128\\_2\\_vui32t](#) ([\\_\\_binary128](#) f128)  
*Transfer function from a `__binary128` scalar to a vector int.*
- static [vui64\\_t](#) [vec\\_xfer\\_bin128\\_2\\_vui64t](#) ([\\_\\_binary128](#) f128)  
*Transfer function from a `__binary128` scalar to a vector long long int.*
- static [vui128\\_t](#) [vec\\_xfer\\_bin128\\_2\\_vui128t](#) ([\\_\\_binary128](#) f128)  
*Transfer function from a `__binary128` scalar to a vector `__int128`.*
- static [\\_\\_binary128](#) [vec\\_xfer\\_vui8t\\_2\\_bin128](#) ([vui8\\_t](#) f128)  
*Transfer a vector unsigned char to `__binary128` scalar.*
- static [\\_\\_binary128](#) [vec\\_xfer\\_vui16t\\_2\\_bin128](#) ([vui16\\_t](#) f128)  
*Transfer a vector unsigned short to `__binary128` scalar.*
- static [\\_\\_binary128](#) [vec\\_xfer\\_vui32t\\_2\\_bin128](#) ([vui32\\_t](#) f128)  
*Transfer a vector unsigned int to `__binary128` scalar.*
- static [\\_\\_binary128](#) [vec\\_xfer\\_vui64t\\_2\\_bin128](#) ([vui64\\_t](#) f128)  
*Transfer a vector unsigned long long to `__binary128` scalar.*
- static [\\_\\_binary128](#) [vec\\_xfer\\_vui128t\\_2\\_bin128](#) ([vui128\\_t](#) f128)  
*Transfer a vector unsigned `__int128` to `__binary128` scalar.*

- static `__binary128 vec_absf128 (__binary128 f128)`  
*Clear the sign bit of \_\_float128 input and return the resulting positive \_\_float128 value.*
- static int `vec_all_isfinitef128 (__binary128 f128)`  
*Return true if the \_\_float128 value is Finite (Not NaN nor Inf).*
- static int `vec_all_isinff128 (__binary128 f128)`  
*Return true if the \_\_float128 value is infinity.*
- static int `vec_all_isnanf128 (__binary128 f128)`  
*Return true if the \_\_float128 value is Not a Number (NaN).*
- static int `vec_all_isnormalf128 (__binary128 f128)`  
*Return true if the \_\_float128 value is normal (Not NaN, Inf, denormal, or zero).*
- static int `vec_all_issubnormalf128 (__binary128 f128)`  
*Return true if the \_\_float128 value is subnormal (denormal).*
- static int `vec_all_iszerof128 (__binary128 f128)`  
*Return true if the \_\_float128 value is +/-0.0.*
- static `__binary128 vec_copysignf128 (__binary128 f128x, __binary128 f128y)`  
*Copy the sign bit from f128y and merge with the magnitude from f128x. The merged result is returned as a \_\_float128 value.*
- static `__binary128 vec_const_huge_valf128 ()`  
*return a positive infinity.*
- static `__binary128 vec_const_inff128 ()`  
*return a positive infinity.*
- static `__binary128 vec_const_nanf128 ()`  
*return a quiet NaN.*
- static `__binary128 vec_const_nansf128 ()`  
*return a signaling NaN.*
- static `vb128_t vec_isfinitef128 (__binary128 f128)`  
*Return 128-bit vector boolean true if the \_\_float128 value is Finite (Not NaN nor Inf).*
- static int `vec_isinf_signf128 (__binary128 f128)`  
*Return true (nonzero) value if the \_\_float128 value is infinity. For infinity indicate the sign as +1 for positive infinity and -1 for negative infinity.*
- static `vb128_t vec_isinff128 (__binary128 f128)`  
*Return a 128-bit vector boolean true if the \_\_float128 value is infinity.*
- static `vb128_t vec_isnanf128 (__binary128 f128)`  
*Return 128-bit vector boolean true if the \_\_float128 value is Not a Number (NaN).*
- static `vb128_t vec_isnormalf128 (__binary128 f128)`  
*Return 128-bit vector boolean true if the \_\_float128 value is normal (Not NaN, Inf, denormal, or zero).*
- static `vb128_t vec_issubnormalf128 (__binary128 f128)`  
*Return 128-bit vector boolean true value, if the \_\_float128 value is subnormal (denormal).*
- static `vb128_t vec_iszerof128 (__binary128 f128)`  
*Return 128-bit vector boolean true value, if the value that is +/-0.0.*
- static `vb128_t vec_setb_qp (__binary128 f128)`  
*Vector Set Bool from Quadword Floating-point.*
- static int `vec_signbitf128 (__binary128 f128)`  
*Return int boolean true if the \_\_float128 value is negative (sign bit is '1').*

### 7.4.1 Detailed Description

Header package containing a collection of 128-bit SIMD operations over Quad-Precision floating point elements.

PowerISA 3.0 added Quad-Precision floating point type and operations to the Vector-Scalar Extension (VSX) facility. The first hardware implementation is available in POWER9.

While all Quad-Precision operations are on 128-bit vector registers, they are defined as scalars in the PowerISA. The OpenPOWER ABI also treats the `__float128` type as scalar that just happens to use vector registers for parameter passing and operations. As such no operations using `__float128` (`_Float128`, or `__ieee128`) as parameter or return value are defined as vector built-ins in the ABI or `<altivec.h>`.

#### Note

GCC 8.2 does document some built-ins, using the *scalar* prefix (`scalar_extract_exp`, `scalar_extract_sig`, `scalar_test_data_class`), that do accept the `__ieee128` type. This work seems to be incomplete as `scalar_↔_exp_cmp_*` for the `__ieee128` type are not present. GCC 7.3 defines vector and scalar forms of the `extract/insert_exp` for float and double but not for `__ieee128`. These built-ins are not defined in GCC 6.4. See [compiler documentation](#). These are useful operations and can be implement in a few vector logical instruction for earlier machines. So it seems reasonable to add these to pveclib for both vector and scalar forms.

Quad-Precision is not supported in hardware until POWER9. However the compiler and runtime supports the `__↔float128` type and arithmetic operations via soft-float emulation for earlier processors. The soft-float implementation follows the ABI and passes `__float128` parameters and return values in vector registers.

So it is not unreasonable for this header to provide vector forms of the `__float128` classification functions (`isnormal/subnormal/finite/inf/nan/zero`, `copysign`, and `abs`). These functions can be implemented directly using (one or more) POWER9 instructions, or a few vector logical and integer compare instructions for POWER7/8. Each is comfortably small enough to be in-lined and inherently faster than the equivalent POSIX or compiler built-in runtime functions. Performing these operations in-line and directly in vector registers (VRs) avoids call/return and VR `<->` GPR transfer overhead.

#### Note

The compiler disables associated `<altivec.h>` built-ins if the `mcpu` target does not enable the specific instruction. For example if you compile with `-mcpu=power8`, Quad-Precision floating-point operations useful for floating point classification are not defined. This header provides the appropriate substitutions, will generate the minimum code, appropriate for the target, and produce correct results.

Most ppc64le compilers will default to `-mcpu=power8` if `-mcpu` is not specified.

This header covers operations that are any of the following:

- Implemented in hardware instructions in newer processors, but useful to programmers on slightly older processors (even if the equivalent function requires more instructions).
- Defined in the OpenPOWER ABI but *not* yet defined in `<altivec.h>` provided by available compilers in common use. Examples include `scalar_test_neg`, `scalar_test_data_class`, etc.
- Providing special vector float tests for special conditions without generating extraneous floating-point exceptions. This is important for implementing `__float128` forms of ISO C99 Math functions. Examples include `vector_isnan`, `isinf`, etc.
- Commonly used operations, not covered by the ABI or `<altivec.h>`, and require multiple instructions or are not obvious.

### 7.4.2 Examples

For example: using the the classification functions for implementing the math library function sine and cosine. The Posix specification requires that special input values are processed without raising extraneous floating point exceptions and return specific floating point values in response. For example the `sin()` function.

- If the input *value* is NaN then return a NaN.
- If the input *value* is  $\pm 0.0$  then return *value*.
- If the input *value* is subnormal then return *value*.
- If the input *value* is  $\pm \text{Inf}$  then return a NaN.
- Otherwise compute and return `sin(value)`.

The following code example uses functions from this header to address the POSIX requirements for special values input to `sinf128()`:

```
__binary128
test_sinf128 (__binary128 value)
{
    __binary128 result;

    if (vec_all_isnormalf128 (value))
    {
        // body of taylor series.
        ...
    }
    else
    {
        if (vec_all_isinff128 (value))
            result = vec_const_nanf128 ();
        else
            result = value;
    }
    return result;
}
```

For another example the `cos()` function.

- If the input *value* is NaN then return a NaN.
- If the input *value* is  $\pm 0.0$  then return `1.0`.
- If the input *value* is  $\pm \text{Inf}$  then return a NaN.
- Otherwise compute and return `cos(value)`.

The following code example uses functions from this header to address the Posix requirements for special values input to `cosf128()`:

```
__binary128
test_cosf128 (__binary128 value)
{
    __binary128 result;

    if (vec_all_isfinitef128 (value))
    {
        if (vec_all_iszerof128 (value))
            result = 1.0Q;
        else
        {
            // body of taylor series ...
        }
    }
    else
    {
        if (vec_all_isinff128 (value))
            result = vec_const_nanf128 ();
        else
            result = value;
    }
    return result;
}
```

Neither example raises floating point exceptions or sets **errno**, as appropriate for a vector math library.



### 7.4.3 Performance data

High level performance estimates are provided as an aid to function selection when evaluating algorithms. For background on how *Latency* and *Throughput* are derived see: [Performance data](#).

### 7.4.4 Function Documentation

**7.4.4.1** `static __binary128 vec_absf128 ( __binary128 f128 ) [inline],[static]`

Clear the sign bit of \_\_float128 input and return the resulting positive \_\_float128 value.

processor	Latency	Throughput
power8	2-11	2/cycle
power9	2	4/cycle

#### Parameters

<i>f128</i>	a __float128 value containing a signed value.
-------------	---

#### Returns

a \_\_float128 value with magnitude from f128 and a positive sign of f128.

**7.4.4.2** `static int vec_all_isfinitef128 ( __binary128 f128 ) [inline],[static]`

Return true if the \_\_float128 value is Finite (Not NaN nor Inf).

A IEEE Binary128 finite value has an exponent between 0x0000 and 0x7ffe (a 0x7fff indicates NaN or Inf). The significand can be any value. Using the !vec\_all\_eq compare conditional verify this condition and avoids a vector -> GPR transfer for platforms before PowerISA-2.07. The sign bit is ignored.

processor	Latency	Throughput
power8	4-20	2/cycle
power9	3	2/cycle

#### Note

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal \_\_float128 compare can.

#### Parameters

<i>f128</i>	a __float128 value in vector.
-------------	-------------------------------

**Returns**

an int containing 0 or 1.

**7.4.4.3** `static int vec_all_isinff128 ( __binary128 f128 ) [inline],[static]`

Return true if the `__float128` value is infinity.

A IEEE Binary128 infinity has a exponent of 0x7fff and significand of all zeros. Using the `vec_all_eq` compare conditional verifies both conditions and avoids a vector -> GPR transfer for platforms before PowerISA-2.07.

processor	Latency	Throughput
power8	4-20	2/cycle
power9	3	2/cycle

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal `__float128` compare can.

**Parameters**

<i>f128</i>	a <code>__float128</code> value in vector.
-------------	--

**Returns**

an int containing 0 or 1.

**7.4.4.4** `static int vec_all_isnanf128 ( __binary128 f128 ) [inline],[static]`

Return true if the `__float128` value is Not a Number (NaN).

A IEEE Binary128 NaN has a exponent of 0x7fff and nonzero significand. Using the combined `vec_all_eq` / `vec_all_neq` / `any_gt` compare conditional verify both conditions and avoids a vector -> GPR transfer for platforms before PowerISA-2.07. The sign bit is ignored.

processor	Latency	Throughput
power8	6-29	1/cycle
power9	3	2/cycle

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal `__float128` compare can.

**Parameters**

<i>f128</i>	a <code>__float128</code> value in vector.
-------------	--

**Returns**

an int containing 0 or 1.

**7.4.4.5 static int vec\_all\_isnormalf128 ( \_\_binary128 f128 ) [inline],[static]**

Return true if the \_\_float128 value is normal (Not NaN, Inf, denormal, or zero).

A IEEE Binary128 normal value has an exponent between 0x0001 and 0x7ffe (a 0x7fff indicates NaN or Inf). The significand can be any value (except 0 if the exponent is zero). Using the combined vec\_all\_ne compares conditional verify both conditions and avoids a vector -> GPR transfer for platforms before PowerISA-2.07. The sign bit is ignored.

processor	Latency	Throughput
power8	4-29	1/cycle
power9	3	2/cycle

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal \_\_float128 compare can.

**Parameters**

<i>f128</i>	a __float128 value in vector.
-------------	-------------------------------

**Returns**

an int containing 0 or 1.

**7.4.4.6 static int vec\_all\_issubnormalf128 ( \_\_binary128 f128 ) [inline],[static]**

Return true if the \_\_float128 value is subnormal (denormal).

A IEEE Binary128 subnormal has an exponent of 0x0000 and a nonzero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal \_\_float128 compare can.

processor	Latency	Throughput
power8	8-29	1/cycle
power9	3	2/cycle

**Parameters**

<i>f128</i>	a vector of __binary128 values.
-------------	---------------------------------

**Returns**

a boolean int, true if the `__float128` value is subnormal.

**7.4.4.7** `static int vec_all_iszerof128 ( __binary128 f128 ) [inline],[static]`

Return true if the `__float128` value is `+0.0`.

A IEEE Binary128 zero has an exponent of `0x0000` and a zero significand. The sign bit is ignored.

**Note**

This function will not raise `VXSNAN` or `VXVC` (`FE_INVALID`) exceptions. A normal `__float128` compare can.

processor	Latency	Throughput
power8	4-20	1/cycle
power9	3	2/cycle

**Parameters**

<i>f128</i>	a vector of <code>__binary64</code> values.
-------------	---

**Returns**

a boolean int, true if the `__float128` value is `+/- zero`.

**7.4.4.8** `static __binary128 vec_const_huge_valf128 ( ) [inline],[static]`

return a positive infinity.

**Returns**

`const __float128` positive infinity.

**7.4.4.9** `static __binary128 vec_const_inff128 ( ) [inline],[static]`

return a positive infinity.

**Returns**

a `const __float128` positive infinity.

7.4.4.10 `static __binary128 vec_const_nanf128 ( ) [inline],[static]`

return a quiet NaN.

#### Returns

a const \_\_float128 quiet NaN.

7.4.4.11 `static __binary128 vec_const_nansf128 ( ) [inline],[static]`

return a signaling NaN.

#### Returns

a const \_\_float128 signaling NaN.

7.4.4.12 `static __binary128 vec_copysignf128 ( __binary128 f128x, __binary128 f128y ) [inline],[static]`

Copy the sign bit from f128y and merge with the magnitude from f128x. The merged result is returned as a \_\_float128 value.

processor	Latency	Throughput
power8	2-11	2/cycle
power9	2	4/cycle

#### Parameters

<i>f128x</i>	a __float128 value containing the magnitude.
<i>f128y</i>	a __float128 value containing the sign bit.

#### Returns

a \_\_float128 value with magnitude from f128x and the sign of f128y.

7.4.4.13 `static vb128_t vec_isfinitef128 ( __binary128 f128 ) [inline],[static]`

Return 128-bit vector boolean true if the \_\_float128 value is Finite (Not NaN nor Inf).

A IEEE Binary128 finite value has an exponent between 0x0000 and 0x7ffe (a 0x7fff indicates NaN or Inf). The significand can be any value. Using the vec\_cmpeq conditional to generate the predicate mask for NaN / Inf and then invert this for the finite condition. The sign bit is ignored.

processor	Latency	Throughput
power8	8-17	2/cycle
power9	6	2/cycle

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal `__float128` compare can.

**Parameters**

<i>f128</i>	a <code>__float128</code> value in vector.
-------------	--

**Returns**

a vector boolean containing all 0s or 1s.

**7.4.4.14** `static int vec_isinf_signf128 ( __binary128 f128 )` `[inline],[static]`

Return true (nonzero) value if the `__float128` value is infinity. For infinity indicate the sign as +1 for positive infinity and -1 for negative infinity.

A IEEE Binary128 infinity has a exponent of 0x7fff and significand of all zeros. Using the `vec_all_eq` compare conditional verifies both conditions. A subsequent `vec_any_gt` checks the sign bit and set the result appropriately. The sign bit is ignored.

This sequence avoids a vector -> GPR transfer for platforms before PowerISA-2.07.

processor	Latency	Throughput
power8	12-32	1/cycle
power9	3-12	2/cycle

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal `__float128` compare can.

**Parameters**

<i>f128</i>	a <code>__float128</code> value in vector.
-------------	--

**Returns**

an int containing 0 if not infinity and +1/-1 otherwise.

**7.4.4.15** `static vb128_t vec_isinff128 ( __binary128 f128 )` `[inline],[static]`

Return a 128-bit vector boolean true if the `__float128` value is infinity.

A IEEE Binary128 infinity has a exponent of 0x7fff and significand of all zeros. The sign bit is ignored.

processor	Latency	Throughput
power8	8-17	2/cycle
power9	6	2/cycle

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal \_\_float128 compare can.

**Parameters**

<i>f128</i>	a __float128 value in vector.
-------------	-------------------------------

**Returns**

a vector boolean containing all 0s or 1s..

**7.4.4.16** static vb128\_t vec\_isnanf128 ( \_\_binary128 *f128* ) [inline],[static]

Return 128-bit vector boolean true if the \_\_float128 value is Not a Number (NaN).

A IEEE Binary128 NaN has a exponent of 0x7fff and nonzero significand. This requires a combination of verifying the exponent and that any bit of the significand is nonzero. Using the combined vec\_all\_eq / vec\_any\_gt compare conditional verify both conditions before negating the result from zero to all ones.. The sign bit is ignored.

processor	Latency	Throughput
power8	10-19	1/cycle
power9	6	2/cycle

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal \_\_float128 compare can.

**Parameters**

<i>f128</i>	a __float128 value in vector.
-------------	-------------------------------

**Returns**

a vector boolean containing all 0s or 1s.

**7.4.4.17** static vb128\_t vec\_isnormalf128 ( \_\_binary128 *f128* ) [inline],[static]

Return 128-bit vector boolean true if the \_\_float128 value is normal (Not NaN, Inf, denormal, or zero).

A IEEE Binary128 normal value has an exponent between 0x0001 and 0x7ffe (a 0x7fff indicates NaN or Inf). The significand can be any value (except 0 if the exponent is zero). The sign bit is ignored.

processor	Latency	Throughput
power8	10-19	2/cycle
power9	6	2/cycle

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal \_\_float128 compare can.

**Parameters**

<i>f128</i>	a __float128 value in vector.
-------------	-------------------------------

**Returns**

a vector boolean containing all 0s or 1s.

**7.4.4.18** `static vb128_t vec_issubnormalf128 ( __binary128 f128 ) [inline],[static]`

Return 128-bit vector boolean true value, if the \_\_float128 value is subnormal (denormal).

A IEEE Binary128 subnormal has an exponent of 0x0000 and a nonzero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal \_\_float128 compare can.

processor	Latency	Throughput
power8	16-25	1/cycle
power9	6	1/cycle

**Parameters**

<i>f128</i>	a vector of __binary64 values.
-------------	--------------------------------

**Returns**

a vector boolean long long, each containing all 0s(false) or 1s(true).

**7.4.4.19** `static vb128_t vec_iszerof128 ( __binary128 f128 ) [inline],[static]`

Return 128-bit vector boolean true value, if the value that is +-0.0.

A IEEE Binary64 zero has an exponent of 0x000 and a zero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal \_\_float128 compare can.

processor	Latency	Throughput
power8	8-17	2/cycle
power9	6	2/cycle



## Parameters

<i>f128</i>	a vector of <code>__binary32</code> values.
-------------	---

## Returns

a vector boolean int, each containing all 0s(false) or 1s(true).

**7.4.4.20** `static vb128_t vec_setb_qp ( __binary128 f128 ) [inline],[static]`

Vector Set Bool from Quadword Floating-point.

If the quadword's sign bit is '1' then return a vector bool `__int128` that is all '1's. Otherwise return all '0's.

processor	Latency	Throughput
power8	4 - 6	2/cycle
power9	6	2/cycle

## Parameters

<i>f128</i>	a 128-bit vector treated a signed <code>__int128</code> .
-------------	---

## Returns

a 128-bit vector bool of all '1's if the sign bit is '1'. Otherwise all '0's.

**7.4.4.21** `static int vec_signbitf128 ( __binary128 f128 ) [inline],[static]`

Return int boolean true if the `__float128` value is negative (sign bit is '1').

Anding with a signmask and then `vec_all_eq` compare with that mask generates the boolean of the sign bit.

processor	Latency	Throughput
power8	4-20	2/cycle
power9	3	2/cycle

## Parameters

<i>f128</i>	a <code>__float128</code> value in vector.
-------------	--

## Returns

a int boolean indicating the sign bit.

**7.4.4.22** `static vui128_t vec_xfer_bin128_2_vui128t ( __binary128 f128 ) [inline],[static]`

Transfer function from a `__binary128` scalar to a vector `__int128`.

The compiler does not allow direct transfer (assignment or type cast) between `__binary128` (`__float128`) scalars and vector types. This despite the fact the the ABI and ISA require `__binary128` in vector registers (VRs).

#### Note

this function uses a union to effect the (logical) transfer. The compiler should not generate any code for this.

#### Parameters

<i>f128</i>	a <code>__binary128</code> floating point scalar value.
-------------	---

#### Returns

The original value as a 128-bit vector `__int128`.

**7.4.4.23** `static vui16_t vec_xfer_bin128_2_vui16t ( __binary128 f128 ) [inline],[static]`

Transfer function from a `__binary128` scalar to a vector short int.

The compiler does not allow direct transfer (assignment or type cast) between `__binary128` (`__float128`) scalars and vector types. This despite the fact the the ABI and ISA require `__binary128` in vector registers (VRs).

#### Note

this function uses a union to effect the (logical) transfer. The compiler should not generate any code for this.

#### Parameters

<i>f128</i>	a <code>__binary128</code> floating point scalar value.
-------------	---

#### Returns

The original value as a 128-bit vector short int.

**7.4.4.24** `static vui32_t vec_xfer_bin128_2_vui32t ( __binary128 f128 ) [inline],[static]`

Transfer function from a `__binary128` scalar to a vector int.

The compiler does not allow direct transfer (assignment or type cast) between `__binary128` (`__float128`) scalars and vector types. This despite the fact the the ABI and ISA require `__binary128` in vector registers (VRs).

#### Note

this function uses a union to effect the (logical) transfer. The compiler should not generate any code for this.

## Parameters

<i>f128</i>	a <code>__binary128</code> floating point scalar value.
-------------	---

## Returns

The original value as a 128-bit vector int.

**7.4.4.25** `static vui64_t vec_xfer_bin128_2_vui64t ( __binary128 f128 ) [inline],[static]`

Transfer function from a `__binary128` scalar to a vector long long int.

The compiler does not allow direct transfer (assignment or type cast) between `__binary128` (`__float128`) scalars and vector types. This despite the fact the the ABI and ISA require `__binary128` in vector registers (VRs).

## Note

this function uses a union to effect the (logical) transfer. The compiler should not generate any code for this.

## Parameters

<i>f128</i>	a <code>__binary128</code> floating point scalar value.
-------------	---

## Returns

The original value as a 128-bit vector long long int.

**7.4.4.26** `static vui8_t vec_xfer_bin128_2_vui8t ( __binary128 f128 ) [inline],[static]`

Transfer function from a `__binary128` scalar to a vector char.

The compiler does not allow direct transfer (assignment or type cast) between `__binary128` (`__float128`) scalars and vector types. This despite the fact the the ABI and ISA require `__binary128` in vector registers (VRs).

## Note

this function uses a union to effect the (logical) transfer. The compiler should not generate any code for this.

## Parameters

<i>f128</i>	a <code>__binary128</code> floating point scalar value.
-------------	---

## Returns

The original value as a 128-bit vector char.

**7.4.4.27** `static __binary128 vec_xfer_vui128t_2_bin128 ( vui128_t f128 ) [inline],[static]`

Transfer a vector unsigned \_\_int128 to \_\_binary128 scalar.

The compiler does not allow direct transfer (assignment or type cast) between \_\_binary128 (\_\_float128) scalars and vector types. This despite the fact the the ABI and ISA require \_\_binary128 in vector registers (VRs).

#### Note

this function uses a union to effect the (logical) transfer. The compiler should not generate any code for this.

#### Parameters

<i>f128</i>	a vector unsigned __int128 value.
-------------	-----------------------------------

#### Returns

The original value returned as a \_\_binary128 scalar.

**7.4.4.28** `static __binary128 vec_xfer_vui16t_2_bin128 ( vui16_t f128 ) [inline],[static]`

Transfer a vector unsigned short to \_\_binary128 scalar.

The compiler does not allow direct transfer (assignment or type cast) between \_\_binary128 (\_\_float128) scalars and vector types. This despite the fact the the ABI and ISA require \_\_binary128 in vector registers (VRs).

#### Note

this function uses a union to effect the (logical) transfer. The compiler should not generate any code for this.

#### Parameters

<i>f128</i>	a vector unsigned short value.
-------------	--------------------------------

#### Returns

The original value returned as a \_\_binary128 scalar.

**7.4.4.29** `static __binary128 vec_xfer_vui32t_2_bin128 ( vui32_t f128 ) [inline],[static]`

Transfer a vector unsigned int to \_\_binary128 scalar.

The compiler does not allow direct transfer (assignment or type cast) between \_\_binary128 (\_\_float128) scalars and vector types. This despite the fact the the ABI and ISA require \_\_binary128 in vector registers (VRs).

#### Note

this function uses a union to effect the (logical) transfer. The compiler should not generate any code for this.

**Parameters**

<i>f128</i>	a vector unsigned int value.
-------------	------------------------------

**Returns**

The original value returned as a `__binary128` scalar.

**7.4.4.30** `static __binary128 vec_xfer_vui64t_2_bin128 ( vui64_t f128 ) [inline],[static]`

Transfer a vector unsigned long long to `__binary128` scalar.

The compiler does not allow direct transfer (assignment or type cast) between `__binary128` (`__float128`) scalars and vector types. This despite the fact the the ABI and ISA require `__binary128` in vector registers (VRs).

**Note**

this function uses a union to effect the (logical) transfer. The compiler should not generate any code for this.

**Parameters**

<i>f128</i>	a vector unsigned long long value.
-------------	------------------------------------

**Returns**

The original value returned as a `__binary128` scalar.

**7.4.4.31** `static __binary128 vec_xfer_vui8t_2_bin128 ( vui8_t f128 ) [inline],[static]`

Transfer a vector unsigned char to `__binary128` scalar.

The compiler does not allow direct transfer (assignment or type cast) between `__binary128` (`__float128`) scalars and vector types. This despite the fact the the ABI and ISA require `__binary128` in vector registers (VRs).

**Note**

this function uses a union to effect the (logical) transfer. The compiler should not generate any code for this.

**Parameters**

<i>f128</i>	a vector unsigned char value.
-------------	-------------------------------

**Returns**

The original value returned as a `__binary128` scalar.

## 7.5 src/vec\_f32\_ppc.h File Reference

Header package containing a collection of 128-bit SIMD operations over 4x32-bit floating point elements.

```
#include <vec_common_ppc.h>
```

### Typedefs

- typedef [vf32\\_t](#) [\\_\\_vbinary32](#)  
*typedef \_\_vbinary32 to vector of 4 xfloat elements.*

### Functions

- static [vf32\\_t](#) [vec\\_absf32](#) ([vf32\\_t](#) vf32x)  
*Vector float absolute value.*
- static int [vec\\_all\\_isfinitef32](#) ([vf32\\_t](#) vf32)  
*Return true if all 4x32-bit vector float values are Finite (Not NaN nor Inf).*
- static int [vec\\_all\\_isinff32](#) ([vf32\\_t](#) vf32)  
*Return true if all 4x32-bit vector float values are infinity.*
- static int [vec\\_all\\_isnanf32](#) ([vf32\\_t](#) vf32)  
*Return true if all of 4x32-bit vector float values are NaN.*
- static int [vec\\_all\\_isnormalf32](#) ([vf32\\_t](#) vf32)  
*Return true if all of 4x32-bit vector float values are normal (Not NaN, Inf, denormal, or zero).*
- static int [vec\\_all\\_issubnormalf32](#) ([vf32\\_t](#) vf32)  
*Return true if all of 4x32-bit vector float values is subnormal (denormal).*
- static int [vec\\_all\\_iszerof32](#) ([vf32\\_t](#) vf32)  
*Return true if all of 4x32-bit vector float values are +-0.0.*
- static int [vec\\_any\\_isfinitef32](#) ([vf32\\_t](#) vf32)  
*Return true if any 4x32-bit vector float values are Finite (Not NaN nor Inf).*
- static int [vec\\_any\\_isinff32](#) ([vf32\\_t](#) vf32)  
*Return true if any 4x32-bit vector float values are infinity.*
- static int [vec\\_any\\_isnanf32](#) ([vf32\\_t](#) vf32)  
*Return true if any of 4x32-bit vector float values are NaN.*
- static int [vec\\_any\\_isnormalf32](#) ([vf32\\_t](#) vf32)  
*Return true if any of 4x32-bit vector float values are normal (Not NaN, Inf, denormal, or zero).*
- static int [vec\\_any\\_issubnormalf32](#) ([vf32\\_t](#) vf32)  
*Return true if any of 4x32-bit vector float values is subnormal (denormal).*
- static int [vec\\_any\\_iszerof32](#) ([vf32\\_t](#) vf32)  
*Return true if any of 4x32-bit vector float values are +-0.0.*
- static [vf32\\_t](#) [vec\\_copysignf32](#) ([vf32\\_t](#) vf32x, [vf32\\_t](#) vf32y)  
*Copy the sign bit from vf32y merged with magnitude from vf32x and return the resulting vector float values.*
- static [vb32\\_t](#) [vec\\_isfinitef32](#) ([vf32\\_t](#) vf32)  
*Return 4x32-bit vector boolean true values for each float element that is Finite (Not NaN nor Inf).*
- static [vb32\\_t](#) [vec\\_isinff32](#) ([vf32\\_t](#) vf32)  
*Return 4x32-bit vector boolean true values for each float, if infinity.*
- static [vb32\\_t](#) [vec\\_isnanf32](#) ([vf32\\_t](#) vf32)  
*Return 4x32-bit vector boolean true values, for each float NaN value.*
- static [vb32\\_t](#) [vec\\_isnormalf32](#) ([vf32\\_t](#) vf32)  
*Return 4x32-bit vector boolean true values, for each float value, if normal (Not NaN, Inf, denormal, or zero).*
- static [vb32\\_t](#) [vec\\_issubnormalf32](#) ([vf32\\_t](#) vf32)  
*Return 4x32-bit vector boolean true values, for each float value that is subnormal (denormal).*
- static [vb32\\_t](#) [vec\\_iszerof32](#) ([vf32\\_t](#) vf32)  
*Return 4x32-bit vector boolean true values, for each float value that is +-0.0.*

### 7.5.1 Detailed Description

Header package containing a collection of 128-bit SIMD operations over 4x32-bit floating point elements.

Most vector float (32-bit float) operations are implemented with PowerISA VMX instructions either defined by the original VMX (a.k.a. Altivec) or added to later versions of the PowerISA. POWER8 added the Vector Scalar Extended (VSX) with access to additional vector registers (64 total) and operations. Most of these operations (compiler built-ins, or intrinsics) are defined in `<altivec.h>` and described in the [compiler documentation](#).

#### Note

The compiler disables associated `<altivec.h>` built-ins if the `mcpu` target does not enable the specific instruction. For example if you compile with `-mcpu=power7`, some of the wordwise pack, unpack and merge operations useful for conversions are not defined and the equivalent `vec_perm` and `permute` control must be used instead. This header will provide the appropriate substitutions, will generate the minimum code, appropriate for the target, and produce correct results.

Most ppc64le compilers will default to `-mcpu=power8` if not specified.

Most of these operations are implemented in a single instruction on newer (POWER8/POWER9) processors. This header serves to fill in functional gaps for older (POWER7, POWER8) processors and provides an inline assembler implementation for older compilers that do not provide the built-ins.

POWER9 adds useful vector float operations, including: test data class, extract exponent, extract significand, and insert exponent. These operations are common in math library implementations.

#### Note

GCC 7.3 defines vector forms of the test data class, extract significand, and extract/insert\_exp for float and double. These built-ins are not defined in GCC 6.4. See [compiler documentation](#). These are useful operations and can be implement in a few vector logical instruction for earlier machines.

So it is reasonable for this header to provide vector forms of the floating point classification functions (isnormal/subnormal/finite/inf/nan/zero, etc.). These functions can be implemented directly using (one or more) POWER9 instructions, or a few vector logical and integer compare instructions for POWER7/8. Each is comfortably small enough to be in-lined and inherently faster than the equivalent POSIX or compiler built-in runtime scalar functions.

This header covers operations that are any of the following:

- Implemented in hardware instructions in newer processors, but useful to programmers on slightly older processors (even if the equivalent function requires more instructions). Examples include the floating point test data class, extract exponent, extract significand, and insert exponent operations.
- Defined in the OpenPOWER ABI but *not* yet defined in `<altivec.h>` provided by available compilers in common use. Examples include vector float even/odd.
- Providing special vector float tests for special conditions without generating extraneous floating-point exceptions. This is important for implementing vectorized forms of ISO C99 Math functions.
- Commonly used operations, not covered by the ABI or `<altivec.h>`, and require multiple instructions or are not obvious.

## 7.5.2 Examples

For example: using the the classification functions for implementing the math library function sine and cosine. The POSIX specification requires that special input values are processed without raising extraneous floating point exceptions and return specific floating point values in response. For example the `sin()` function.

- If the input *value* is NaN then return a NaN.
- If the input *value* is  $\pm 0.0$  then return *value*.
- If the input *value* is subnormal then return *value*.
- If the input *value* is  $\pm \text{Inf}$  then return a NaN.
- Otherwise compute and return `sin(value)`.

The following code example uses functions from this header to address the POSIX requirements for special values input to for a vectorized `sinf()`:

```
vf32_t
test_vec_sinf32 (vf32_t value)
{
    const vf32_t vec_f0 = { 0.0, 0.0, 0.0, 0.0 };
    const vui32_t vec_f32_qnan =
        { 0x7f800001, 0x7fc00000, 0x7fc00000, 0x7fc00000 };
    vf32_t result;
    vb32_t normmask, infmask;

    normmask = vec_isnormalf32 (value);
    if (vec_any_isnormalf32 (value))
    {
        // replace non-normal input values with safe values.
        vf32_t safeval = vec_sel (vec_f0, value, normmask);
        // body of vec_sin(safeval) computation elided for this example.
    }
    else
        result = value;

    // merge non-normal input values back into result
    result = vec_sel (value, result, normmask);
    // Inf input value elements return quiet-nan
    infmask = vec_isinff32 (value);
    result = vec_sel (result, (vf32_t) vec_f32_qnan, infmask);

    return result;
}
```

The code generated for this fragment runs between 24 (`-mcpu=power9`) and 40 (`-mcpu=power8`) instructions. The normal execution path is 14 to 25 instructions respectively.

Another example the `cos()` function.

- If the input *value* is NaN then return a NaN.
- If the input *value* is  $\pm 0.0$  then return `1.0`.
- If the input *value* is  $\pm \text{Inf}$  then return a NaN.
- Otherwise compute and return `cos(value)`.

The following code example uses functions from this header to address the POSIX requirements for special values input to vectorized `cosf()`:



```

vf32_t
test_vec_cosf32 (vf32_t value)
{
    vf32_t result;
    const vf32_t vec_f0 = { 0.0, 0.0, 0.0, 0.0 };
    const vf32_t vec_f1 = { 1.0, 1.0, 1.0, 1.0 };
    const vui32_t vec_f32_qnan =
        { 0x7f800001, 0x7fc00000, 0x7fc00000, 0x7fc00000 };
    vb32_t finitemask, infmask, zeromask;

    finitemask = vec_isfinitef32 (value);
    if (vec_any_isfinitef32 (value))
    {
        // replace non-finite input values with safe values
        vf32_t safeval = vec_sel (vec_f0, value, finitemask);
        // body of vec_sin(safeval) computation elided for this example
    }
    else
        result = value;

    // merge non-finite input values back into result
    result = vec_sel (value, result, finitemask);
    // Set +-0.0 input elements to exactly 1.0 in result
    zeromask = vec_iszerof32 (value);
    result = vec_sel (result, vec_f1, zeromask);
    // Set Inf input elements to quiet-nan in result
    infmask = vec_isinfff32 (value);
    result = vec_sel (result, (vf32_t) vec_f32_qnan, infmask);

    return result;
}

```

Neither example raises floating point exceptions or sets **errno**, as appropriate for a vector math library.

### 7.5.3 Performance data.

High level performance estimates are provided as an aid to function selection when evaluating algorithms. For background on how *Latency* and *Throughput* are derived see: [Performance data](#).

### 7.5.4 Function Documentation

#### 7.5.4.1 static vf32\_t vec\_absf32 ( vf32\_t vf32x ) [inline],[static]

Vector float absolute value.

processor	Latency	Throughput
power8	6-7	2/cycle
power9	2	2/cycle

#### Parameters

vf32x	vector float values containing the magnitudes.
-------	--

#### Returns

vector absolute values of 4x float elements of vf32x.

#### 7.5.4.2 static int vec\_all\_isfinitef32 ( vf32\_t vf32 ) [inline],[static]

Return true if all 4x32-bit vector float values are Finite (Not NaN nor Inf).

A IEEE Binary32 finite value has an exponent between 0x000 and 0x7f0 (a 0x7f8 indicates NaN or Inf). The significand can be any value. The sign bit is ignored.

#### Note

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	4-20	2/cycle
power9	6	1/cycle

#### Parameters

<code>vf32</code>	a vector of <code>__binary32</code> values.
-------------------	---

#### Returns

an int containing 0 or 1.

**7.5.4.3** `static int vec_all_isinff32 ( vf32_t vf32 )` `[inline]`, `[static]`

Return true if all 4x32-bit vector float values are infinity.

A IEEE Binary32 infinity has a exponent of 0x7f8 and significand of all zeros. The sign bit is ignored.

#### Note

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	6-20	2/cycle
power9	6	1/cycle

#### Parameters

<code>vf32</code>	a vector of <code>__binary32</code> values.
-------------------	---

#### Returns

boolean int, true if all 4 float values are infinity

**7.5.4.4** `static int vec_all_isnanf32 ( vf32_t vf32 )` `[inline]`, `[static]`

Return true if all of 4x32-bit vector float values are NaN.

A IEEE Binary32 NaN value has an exponent between 0x7f8 and the significand is nonzero. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	6-20	2/cycle
power9	6	1/cycle

**Parameters**

<b>vf32</b>	a vector of __binary32 values.
-------------	--------------------------------

**Returns**

a boolean int, true if all of 4 vector float values are NaN.

**7.5.4.5** static int vec\_all\_isnormalf32 ( vf32\_t vf32 ) [inline], [static]

Return true if all of 4x32-bit vector float values are normal (Not NaN, Inf, denormal, or zero).

A IEEE Binary32 normal value has an exponent between 0x008 and 0x7f (a 0x7f8 indicates NaN or Inf). The significand can be any value (except 0 if the exponent is zero). The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	6-20	1/cycle
power9	6	1/cycle

**Parameters**

<b>vf32</b>	a vector of __binary32 values.
-------------	--------------------------------

**Returns**

a boolean int, true if all of 4 vector float values are normal.

**7.5.4.6** static int vec\_all\_issubnormalf32 ( vf32\_t vf32 ) [inline], [static]

Return true if all of 4x32-bit vector float values is subnormal (denormal).

A IEEE Binary32 subnormal has an exponent of 0x000 and a nonzero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	10-30	1/cycle
power9	6	1/cycle

**Parameters**

<code>vf32</code>	a vector of <code>__binary32</code> values.
-------------------	---

**Returns**

a boolean int, true if all of 4 vector float values are subnormal.

**7.5.4.7** `static int vec_all_iszerof32 ( vf32_t vf32 )` `[inline],[static]`

Return true if all of 4x32-bit vector float values are  $\pm 0.0$ .

A IEEE Binary32 zero has an exponent of 0x000 and a zero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	6-20	2/cycle
power9	6	1/cycle

**Parameters**

<code>vf32</code>	a vector of <code>__binary32</code> values.
-------------------	---

**Returns**

a boolean int, true if all of 4 vector float values are  $\pm$  zero.

**7.5.4.8** `static int vec_any_isfinitef32 ( vf32_t vf32 )` `[inline],[static]`

Return true if any 4x32-bit vector float values are Finite (Not NaN nor Inf).

A IEEE Binary32 finite value has an exponent between 0x000 and 0x7f0 (a 0x7f8 indicates NaN or Inf). The significand can be any value. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	4-20	2/cycle
power9	6	1/cycle

**Parameters**

<b>vf32</b>	a vector of __binary32 values.
-------------	--------------------------------

**Returns**

an int containing 0 or 1.

**7.5.4.9** `static int vec_any_isinff32 ( vf32_t vf32 ) [inline],[static]`

Return true if any 4x32-bit vector float values are infinity.

A IEEE Binary32 infinity has a exponent of 0x7f8 and significand of all zeros.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	6-20	2/cycle
power9	6	2/cycle

**Parameters**

<b>vf32</b>	a vector of __binary32 values.
-------------	--------------------------------

**Returns**

boolean int, true if any of 4 float values are infinity

**7.5.4.10** `static int vec_any_isnanf32 ( vf32_t vf32 ) [inline],[static]`

Return true if any of 4x32-bit vector float values are NaN.

A IEEE Binary32 NaN value has an exponent between 0x7f8 and the significand is nonzero. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	6-20	2/cycle
power9	6	2/cycle

**Parameters**

<code>vf32</code>	a vector of <code>__binary32</code> values.
-------------------	---

**Returns**

a boolean int, true if any of 4 vector float values are NaN.

**7.5.4.11** `static int vec_any_isnormalf32 ( vf32_t vf32 )` `[inline]`, `[static]`

Return true if any of 4x32-bit vector float values are normal (Not NaN, Inf, denormal, or zero).

A IEEE Binary32 normal value has an exponent between 0x008 and 0x7f (a 0x7f8 indicates NaN or Inf). The significand can be any value (expect 0 if the exponent is zero). The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	10-24	1/cycle
power9	6	1/cycle

**Parameters**

<code>vf32</code>	a vector of <code>__binary32</code> values.
-------------------	---

**Returns**

a boolean int, true if any of 4 vector float values are normal.

**7.5.4.12** `static int vec_any_issubnormalf32 ( vf32_t vf32 )` `[inline]`, `[static]`

Return true if any of 4x32-bit vector float values is subnormal (denormal).

A IEEE Binary32 subnormal has an exponent of 0x000 and a nonzero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	10-18	1/cycle
power9	6	1/cycle

**Parameters**

<b>vf32</b>	a vector of __binary32 values.
-------------	--------------------------------

**Returns**

if any of 4 vector float values are subnormal.

**7.5.4.13** `static int vec_any_iszerof32 ( vf32_t vf32 ) [inline],[static]`

Return true if any of 4x32-bit vector float values are +/-0.0.

A IEEE Binary32 zero has an exponent of 0x000 and a zero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	6-20	2/cycle
power9	6	1/cycle

**Parameters**

<b>vf32</b>	a vector of __binary32 values.
-------------	--------------------------------

**Returns**

a boolean int, true if any of 4 vector float values are +/- zero.

**7.5.4.14** `static vf32_t vec_copysignf32 ( vf32_t vf32x, vf32_t vf32y ) [inline],[static]`

Copy the sign bit from vf32y merged with magnitude from vf32x and return the resulting vector float values.

processor	Latency	Throughput
power8	6-7	2/cycle
power9	2	2/cycle

**Parameters**

<b>vf32x</b>	vector float values containing the magnitudes.
--------------	--

## Parameters

<code>vf32y</code>	vector float values containing the sign bits.
--------------------	---

## Returns

vector float values with magnitude from `vf32x` and the sign of `vf32y`.

#### 7.5.4.15 `static vb32_t vec_isfinitef32 ( vf32_t vf32 ) [inline],[static]`

Return 4x32-bit vector boolean true values for each float element that is Finite (Not NaN nor Inf).

A IEEE Binary32 finite value has an exponent between 0x000 and 0x7f0 (a 0x7f8 indicates NaN or Inf). The significand can be any value. Using the `vec_cmpeq` conditional to generate the predicate mask for NaN / Inf and then invert this for the finite condition. The sign bit is ignored.

## Note

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	6-15	2/cycle
power9	5	2/cycle

## Parameters

<code>vf32</code>	a vector of <code>__binary32</code> values.
-------------------	---

## Returns

a vector boolean int, each containing all 0s(false) or 1s(true).

#### 7.5.4.16 `static vb32_t vec_isinff32 ( vf32_t vf32 ) [inline],[static]`

Return 4x32-bit vector boolean true values for each float, if infinity.

A IEEE Binary32 infinity has a exponent of 0x7f8 and significand of all zeros.

## Note

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	4-13	2/cycle
power9	3	2/cycle



## Parameters

<code>vf32</code>	a vector of <code>__binary32</code> values.
-------------------	---

## Returns

a vector boolean int, each containing all 0s(false) or 1s(true).

**7.5.4.17** `static vb32_t vec_isnanf32 ( vf32_t vf32 )` `[inline]`, `[static]`

Return 4x32-bit vector boolean true values, for each float NaN value.

A IEEE Binary32 NaN value has an exponent between 0x7f8 and the significand is nonzero. The sign bit is ignored.

processor	Latency	Throughput
power8	4-13	2/cycle
power9	3	2/cycle

## Parameters

<code>vf32</code>	a vector of <code>__binary32</code> values.
-------------------	---

## Returns

a vector boolean int, each containing all 0s(false) or 1s(true).

**7.5.4.18** `static vb32_t vec_isnormalf32 ( vf32_t vf32 )` `[inline]`, `[static]`

Return 4x32-bit vector boolean true values, for each float value, if normal (Not NaN, Inf, denormal, or zero).

A IEEE Binary32 normal value has an exponent between 0x008 and 0x7f (a 0x7f8 indicates NaN or Inf). The significand can be any value (except 0 if the exponent is zero). The sign bit is ignored.

## Note

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	6-15	1/cycle
power9	5	1/cycle

## Parameters

<code>vf32</code>	a vector of <code>__binary32</code> values.
-------------------	---

**Returns**

a vector boolean int, each containing all 0s(false) or 1s(true).

**7.5.4.19** `static vb32_t vec_issubnormalf32 ( vf32_t vf32 )` `[inline],[static]`

Return 4x32-bit vector boolean true values, for each float value that is subnormal (denormal).

A IEEE Binary32 subnormal has an exponent of 0x000 and a nonzero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	6-16	1/cycle
power9	3	1/cycle

**Parameters**

<code>vf32</code>	a vector of <code>__binary32</code> values.
-------------------	---

**Returns**

a vector boolean int, each containing all 0s(false) or 1s(true).

**7.5.4.20** `static vb32_t vec_iszerof32 ( vf32_t vf32 )` `[inline],[static]`

Return 4x32-bit vector boolean true values, for each float value that is +/-0.0.

A IEEE Binary32 zero has an exponent of 0x000 and a zero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal float compare can.

processor	Latency	Throughput
power8	4-13	2/cycle
power9	5	2/cycle

**Parameters**

<code>vf32</code>	a vector of <code>__binary32</code> values.
-------------------	---

## Returns

a vector boolean int, each containing all 0s(false) or 1s(true).

## 7.6 src/vec\_f64\_ppc.h File Reference

Header package containing a collection of 128-bit SIMD operations over 64-bit double-precision floating point elements.

```
#include <vec_common_ppc.h>
#include <vec_int64_ppc.h>
```

### Functions

- static [vf64\\_t](#) [vec\\_absf64](#) ([vf64\\_t](#) vf64x)  
*Vector double absolute value.*
- static int [vec\\_all\\_isfinitef64](#) ([vf64\\_t](#) vf64)  
*Return true if all 2x64-bit vector double values are Finite (Not NaN nor Inf).*
- static int [vec\\_all\\_isinff64](#) ([vf64\\_t](#) vf64)  
*Return true if all 2x64-bit vector double values are infinity.*
- static int [vec\\_all\\_isnanf64](#) ([vf64\\_t](#) vf64)  
*Return true if all 2x64-bit vector double values are NaN.*
- static int [vec\\_all\\_isnormalf64](#) ([vf64\\_t](#) vf64)  
*Return true if all 2x64-bit vector double values are normal (Not NaN, Inf, denormal, or zero).*
- static int [vec\\_all\\_issubnormalf64](#) ([vf64\\_t](#) vf64)  
*Return true if all 2x64-bit vector double values are subnormal (denormal).*
- static int [vec\\_all\\_iszerof64](#) ([vf64\\_t](#) vf64)  
*Return true if all 2x64-bit vector double values are +-0.0.*
- static int [vec\\_any\\_isfinitef64](#) ([vf64\\_t](#) vf64)  
*Return true if any of 2x64-bit vector double values are Finite (Not NaN nor Inf).*
- static int [vec\\_any\\_isinff64](#) ([vf64\\_t](#) vf64)  
*Return true if any of 2x64-bit vector double values are infinity.*
- static int [vec\\_any\\_isnanf64](#) ([vf64\\_t](#) vf64)  
*Return true if any of 2x64-bit vector double values are NaN.*
- static int [vec\\_any\\_isnormalf64](#) ([vf64\\_t](#) vf64)  
*Return true if any of 2x64-bit vector double values are normal (Not NaN, Inf, denormal, or zero).*
- static int [vec\\_any\\_issubnormalf64](#) ([vf64\\_t](#) vf64)  
*Return true if any of 2x64-bit vector double values is subnormal (denormal).*
- static int [vec\\_any\\_iszerof64](#) ([vf64\\_t](#) vf64)  
*Return true if any of 2x64-bit vector double values are +-0.0.*
- static [vf64\\_t](#) [vec\\_copysignf64](#) ([vf64\\_t](#) vf64x, [vf64\\_t](#) vf64y)  
*Copy the sign bit from vf64y merged with magnitude from vf64x and return the resulting vector double values.*
- static [vb64\\_t](#) [vec\\_isfinitef64](#) ([vf64\\_t](#) vf64)  
*Return 2x64-bit vector boolean true values for each double element that is Finite (Not NaN nor Inf).*
- static [vb64\\_t](#) [vec\\_isinff64](#) ([vf64\\_t](#) vf64)  
*Return 2x64-bit vector boolean true values for each double, if infinity.*
- static [vb64\\_t](#) [vec\\_isnanf64](#) ([vf64\\_t](#) vf64)  
*Return 2x64-bit vector boolean true values, for each double NaN value.*
- static [vb64\\_t](#) [vec\\_isnormalf64](#) ([vf64\\_t](#) vf64)

- Return 2x64-bit vector boolean true values, for each double value, if normal (Not NaN, Inf, denormal, or zero).*
- static `vb64_t vec_issubnormalf64 (vf64_t vf64)`  
*Return 2x64-bit vector boolean true values, for each double value that is subnormal (denormal).*
- static `vb64_t vec_iszerof64 (vf64_t vf64)`  
*Return 2x64-bit vector boolean true values, for each double value that is +/-0.0.*
- static long double `vec_pack_longdouble (vf64_t lval)`  
*Copy the pair of doubles from a vector to IBM long double.*
- static `vf64_t vec_unpack_longdouble (long double lval)`  
*Copy the pair of doubles from a IBM long double to a vector double.*

### 7.6.1 Detailed Description

Header package containing a collection of 128-bit SIMD operations over 64-bit double-precision floating point elements.

Many vector double-precision (64-bit float) operations are implemented with PowerISA-2.06 Vector Scalar Extended (VSX) (POWER7 and later) instructions. Most VSX instructions provide access to 64 combined scalar/vector registers. PowerISA-3.0 (POWER9) provides additional vector double operations: convert with round, convert to/from integer, insert/extract exponent and significand, and test data class. Most of these operations (compiler built-ins, or intrinsics) are defined in `<altivec.h>` and described in the [compiler documentation](#).

#### Note

The compiler disables associated `<altivec.h>` built-ins if the `mcpu` target does not enable the specific instruction. For example if you compile with `-mcpu=power8`, the double-precision vector converts, insert/extract and test data class built-ins are not defined. This header provides the appropriate substitutions, will generate the minimum code, appropriate for the target, and produce correct results.

Most ppc64le compilers will default to `-mcpu=power8` if not specified.

GCC 7.3 defines vector forms of the test data class, extract significand, and extract/insert\_exp for float and double. These built-ins are not defined in GCC 6.4. See [compiler documentation](#). These are useful operations and can be implemented in a few vector logical instructions for earlier machines.

So it is reasonable for this header to provide vector forms of the double-precision floating point classification functions (isnormal/subnormal/finite/inf/nan/zero, etc.). These functions can be implemented directly using (one or more) POWER9 instructions, or a few vector logical and integer compare instructions for POWER7/8. Each is comfortably small enough to be in-lined and inherently faster than the equivalent POSIX or compiler built-in runtime scalar functions.

Most of these operations are implemented in a few instructions on newer (POWER7/POWER8/POWER9) processors. This header serves to fill in functional gaps for older (POWER7, POWER8) processors and provides an inline assembler implementation for older compilers that do not provide the built-ins.

This header covers operations that are any of the following:

- Implemented in hardware instructions in newer processors, but useful to programmers on slightly older processors (even if the equivalent function requires more instructions).
- Defined in the OpenPOWER ABI but *not* yet defined in `<altivec.h>` provided by available compilers in common use. Examples include vector double even/odd conversions.
- Providing special vector double tests for special conditions without generating extraneous floating-point exceptions. This is important for implementing vectorized forms of ISO C99 Math functions. Examples include vector double isnan, isinf, etc.
- Commonly used operations, not covered by the ABI or `<altivec.h>`, and require multiple instructions or are not obvious. For example, converts that change element size and imply converting two vectors into one vector of smaller elements, or one vector into two vectors of larger elements. Another example is the special case of packing/unpacking an IBM long double between a pair of floating-point registers (FPRs) and a single vector register (VR).

## 7.6.2 Examples

For example: using the the classification functions for implementing the math library function sine and cosine. The POSIX specification requires that special input values are processed without raising extraneous floating point exceptions and return specific floating point values in response. For example, the `sin()` function.

- If the input *value* is NaN then return a NaN.
- If the input *value* is  $\pm 0.0$  then return *value*.
- If the input *value* is subnormal then return *value*.
- If the input *value* is  $\pm \text{Inf}$  then return a quiet-NaN.
- Otherwise compute and return `sin(value)`.

The following code example uses functions from this header to address the POSIX requirements for special values input to for a vectorized `sinf()`:

```
vf64_t
test_vec_sinf64 (vf64_t value)
{
    const vf64_t vec_f0 = { 0.0, 0.0 };
    const vui64_t vec_f64_qnan =
        { 0x7ff8000000000000, 0x7ff8000000000000 };
    vf64_t result;
    vb64_t normmask, infmask;

    normmask = vec_isnormalf64 (value);
    if (vec_any_isnormalf64 (value))
    {
        // replace non-normal input values with safe values.
        vf64_t safeval = vec_sel (vec_f0, value, normmask);
        // body of vec_sin(safeval) computation elided for this example.
    }
    else
        result = value;

    // merge non-normal input values back into result
    result = vec_sel (value, result, normmask);
    // Inf input value elements return quiet-nan.
    infmask = vec_isinff64 (value);
    result = vec_sel (result, (vf64_t) vec_f64_qnan, infmask);

    return result;
}
```

The code generated for this fragment runs between 24 (`-mcpu=power9`) and 40 (`-mcpu=power8`) instructions. The normal execution path is 14 to 25 instructions respectively.

Another example the `cos()` function.

- If the input *value* is NaN then return a NaN.
- If the input *value* is  $\pm 0.0$  then return *1.0*.
- If the input *value* is  $\pm \text{Inf}$  then return a quiet-NaN.
- Otherwise compute and return `cos(value)`.

The following code example uses functions from this header to address the POSIX requirements for special values input to vectorized `cosf()`:

```

vf64_t
test_vec_cosf64 (vf64_t value)
{
    vf64_t result;
    const vf64_t vec_f0 = { 0.0, 0.0 };
    const vf64_t vec_f1 = { 1.0, 1.0 };
    const vui64_t vec_f64_qnan =
        { 0x7ff8000000000000, 0x7ff8000000000000 };
    vb64_t finitemask, infmask, zeromask;

    finitemask = vec_isfinitef64 (value);
    if (vec_any_isfinitef64 (value))
    {
        // replace non-finite input values with safe values.
        vf64_t safeval = vec_sel (vec_f0, value, finitemask);
        // body of vec_sin(safeval) computation elided for this example.
    }
    else
        result = value;

    // merge non-finite input values back into result
    result = vec_sel (value, result, finitemask);
    // Set +-0.0 input elements to exactly 1.0 in result.
    zeromask = vec_iszerof64 (value);
    result = vec_sel (result, vec_f1, zeromask);
    // Set Inf input elements to quiet-nan in result.
    infmask = vec_isinff64 (value);
    result = vec_sel (result, (vf64_t) vec_f64_qnan, infmask);

    return result;
}

```

Neither example raises floating point exceptions or sets **errno**, as appropriate for a vector math library.

### 7.6.3 Performance data.

High level performance estimates are provided as an aid to function selection when evaluating algorithms. For background on how *Latency* and *Throughput* are derived see: [Performance data](#).

### 7.6.4 Function Documentation

#### 7.6.4.1 static vf64\_t vec\_absf64 ( vf64\_t vf64x ) [inline],[static]

Vector double absolute value.

processor	Latency	Throughput
power8	6-7	2/cycle
power9	2	2/cycle

#### Parameters

vf64x	vector double values containing the magnitudes.
-------	---

#### Returns

vector double absolute values of vf64x.

#### 7.6.4.2 static int vec\_all\_isfinitef64 ( vf64\_t vf64 ) [inline],[static]

Return true if all 2x64-bit vector double values are Finite (Not NaN nor Inf).

A IEEE Binary64 finite value has an exponent between 0x000 and 0x7fe (a 0x7ff indicates NaN or Inf). The significand can be any value. The sign bit is ignored.

processor	Latency	Throughput
power8	4-20	2/cycle
power9	6	1/cycle

#### Note

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal \_\_binary64 compare can.

#### Parameters

<i>vf64</i>	a vector of __binary64 values.
-------------	--------------------------------

#### Returns

an int containing 0 or 1.

**7.6.4.3** static int vec\_all\_isinff64 ( vf64\_t vf64 ) [inline], [static]

Return true if all 2x64-bit vector double values are infinity.

A IEEE Binary64 infinity has a exponent of 0x7ff and significand of all zeros. The sign bit is ignored.

#### Note

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	6-20	2/cycle
power9	6	1/cycle

#### Parameters

<i>vf64</i>	a vector of __binary64 values.
-------------	--------------------------------

#### Returns

boolean int, true if all 2 double values are infinity

**7.6.4.4** static int vec\_all\_isnanf64 ( vf64\_t vf64 ) [inline], [static]

Return true if all 2x64-bit vector double values are NaN.

A IEEE Binary64 NaN value has an exponent between 0x7ff and the significand is nonzero. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	6-20	2/cycle
power9	6	1/cycle

**Parameters**

<b>vf64</b>	a vector of __binary64 values.
-------------	--------------------------------

**Returns**

a boolean int, true if all 2 vector double values are NaN.

**7.6.4.5** static int vec\_all\_isnormalf64 ( vf64\_t vf64 ) [inline], [static]

Return true if all 2x64-bit vector double values are normal (Not NaN, Inf, denormal, or zero).

A IEEE Binary64 normal value has an exponent between 0x001 and 0x7fe (a 0x7ff indicates NaN or Inf). The significand can be any value (except 0 if the exponent is zero). The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	10-28	1/cycle
power9	6	1/cycle

**Parameters**

<b>vf64</b>	a vector of __binary64 values.
-------------	--------------------------------

**Returns**

a boolean int, true if all 2 vector double values are normal.

**7.6.4.6** static int vec\_all\_issubnormalf64 ( vf64\_t vf64 ) [inline], [static]

Return true if all 2x64-bit vector double values are subnormal (denormal).

A IEEE Binary64 subnormal has an exponent of 0x000 and a nonzero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.



processor	Latency	Throughput
power8	10-30	1/cycle
power9	6	1/cycle

**Parameters**

<i>vf64</i>	a vector of <code>__binary64</code> values.
-------------	---

**Returns**

a boolean int, true if all of 2 vector double values are subnormal.

**7.6.4.7** `static int vec_all_iszerof64 ( vf64_t vf64 )` `[inline], [static]`

Return true if all 2x64-bit vector double values are +/-0.0.

A IEEE Binary64 zero has an exponent of 0x000 and a zero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	6-20	2/cycle
power9	6	1/cycle

**Parameters**

<i>vf64</i>	a vector of <code>__binary64</code> values.
-------------	---

**Returns**

a boolean int, true if all 2 vector double values are +/- zero.

**7.6.4.8** `static int vec_any_isfinitef64 ( vf64_t vf64 )` `[inline], [static]`

Return true if any of 2x64-bit vector double values are Finite (Not NaN nor Inf).

A IEEE Binary64 finite value has an exponent between 0x000 and 0x7fe (a 0x7ff indicates NaN or Inf). The significand can be any value. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	4-20	2/cycle
power9	6	1/cycle

**Parameters**

<i>vf64</i>	a vector of <code>__binary64</code> values.
-------------	---

**Returns**

an int containing 0 or 1.

**7.6.4.9** `static int vec_any_isinff64 ( vf64_t vf64 )` `[inline]`, `[static]`

Return true if any of 2x64-bit vector double values are infinity.

A IEEE Binary64 infinity has a exponent of 0x7ff and significand of all zeros.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	6-20	2/cycle
power9	6	1/cycle

**Parameters**

<i>vf64</i>	a vector of <code>__binary32</code> values.
-------------	---

**Returns**

boolean int, true if any of 2 double values are infinity

**7.6.4.10** `static int vec_any_isnanf64 ( vf64_t vf64 )` `[inline]`, `[static]`

Return true if any of 2x64-bit vector double values are NaN.

A IEEE Binary64 NaN value has an exponent between 0x7ff and the significand is nonzero. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	6-20	2/cycle
power9	6	1/cycle

**Parameters**

<b>vf64</b>	a vector of __binary64 values.
-------------	--------------------------------

**Returns**

a boolean int, true if any of 2 vector double values are NaN.

**7.6.4.11** `static int vec_any_isnormalf64 ( vf64_t vf64 )` `[inline]`, `[static]`

Return true if any of 2x64-bit vector double values are normal (Not NaN, Inf, denormal, or zero).

A IEEE Binary64 normal value has an exponent between 0x001 and 0x7fe (a 0x7ff indicates NaN or Inf). The significand can be any value (except 0 if the exponent is zero). The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	6-20	1/cycle
power9	6	1/cycle

**Parameters**

<b>vf64</b>	a vector of __binary64 values.
-------------	--------------------------------

**Returns**

a boolean int, true if any of 2 vector double values are normal.

**7.6.4.12** `static int vec_any_issubnormalf64 ( vf64_t vf64 )` `[inline]`, `[static]`

Return true if any of 2x64-bit vector double values is subnormal (denormal).

A IEEE Binary64 subnormal has an exponent of 0x000 and a nonzero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	10-18	1/cycle
power9	6	1/cycle

**Parameters**

<b>vf64</b>	a vector of __binary64 values.
-------------	--------------------------------

**Returns**

true if any of 2 vector double values are subnormal.

**7.6.4.13** `static int vec_any_iszerof64 ( vf64_t vf64 ) [inline],[static]`

Return true if any of 2x64-bit vector double values are +/-0.0.

A IEEE Binary64 zero has an exponent of 0x000 and a zero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	6-20	2/cycle
power9	6	1/cycle

**Parameters**

<b>vf64</b>	a vector of __binary64 values.
-------------	--------------------------------

**Returns**

a boolean int, true if any of 2 vector double values are +/- zero.

**7.6.4.14** `static vf64_t vec_copysignf64 ( vf64_t vf64x, vf64_t vf64y ) [inline],[static]`

Copy the sign bit from vf64y merged with magnitude from vf64x and return the resulting vector double values.

processor	Latency	Throughput
power8	6-7	2/cycle
power9	2	2/cycle

**Parameters**

<b>vf64x</b>	vector double values containing the magnitudes.
--------------	---

## Parameters

<i>vf64y</i>	vector double values containing the sign bits.
--------------	--

## Returns

vector double values with magnitude from *vf64x* and the sign of *vf64y*.

## 7.6.4.15 static vb64\_t vec\_isfinitef64 ( vf64\_t vf64 ) [inline],[static]

Return 2x64-bit vector boolean true values for each double element that is Finite (Not NaN nor Inf).

A IEEE Binary64 finite value has an exponent between 0x000 and 0x7fe (a 0x7ff indicates NaN or Inf). The significand can be any value.

Using the `vec_cmpeq` conditional to generate the predicate mask for NaN / Inf and then invert this for the finite condition. The sign bit is ignored.

## Note

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	6-15	2/cycle
power9	5	2/cycle

## Parameters

<i>vf64</i>	a vector of <code>__binary64</code> values.
-------------	---

## Returns

a vector boolean long, each containing all 0s(false) or 1s(true).

## 7.6.4.16 static vb64\_t vec\_isinff64 ( vf64\_t vf64 ) [inline],[static]

Return 2x64-bit vector boolean true values for each double, if infinity.

A IEEE Binary64 infinity has a exponent of 0x7ff and significand of all zeros.

## Note

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	4-13	2/cycle
power9	3	2/cycle

## Parameters

<code>vf64</code>	a vector of <code>__binary64</code> values.
-------------------	---

## Returns

a vector boolean long long, each containing all 0s(false) or 1s(true).

**7.6.4.17** `static vb64_t vec_isnanf64 ( vf64_t vf64 )` `[inline]`, `[static]`

Return 2x64-bit vector boolean true values, for each double NaN value.

A IEEE Binary64 NaN value has an exponent between 0x7ff and the significand is nonzero. The sign bit is ignored.

processor	Latency	Throughput
power8	4-13	2/cycle
power9	3	2/cycle

## Parameters

<code>vf64</code>	a vector of <code>__binary64</code> values.
-------------------	---

## Returns

a vector boolean long long, each containing all 0s(false) or 1s(true).

**7.6.4.18** `static vb64_t vec_isnormalf64 ( vf64_t vf64 )` `[inline]`, `[static]`

Return 2x64-bit vector boolean true values, for each double value, if normal (Not NaN, Inf, denormal, or zero).

A IEEE Binary64 normal value has an exponent between 0x001 and 0x7ffe (a 0x7ff indicates NaN or Inf). The significand can be any value (except 0 if the exponent is zero).

## Note

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	6-15	1/cycle
power9	5	1/cycle

## Parameters

<code>vf64</code>	a vector of <code>__binary64</code> values.
-------------------	---

**Returns**

a vector boolean long long, each containing all 0s(false) or 1s(true).

**7.6.4.19 static vb64\_t vec\_issubnormalf64 ( vf64\_t vf64 ) [inline],[static]**

Return 2x64-bit vector boolean true values, for each double value that is subnormal (denormal).

A IEEE Binary64 subnormal has an exponent of 0x000 and a nonzero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	6-16	1/cycle
power9	3	1/cycle

**Parameters**

<b>vf64</b>	a vector of __binary64 values.
-------------	--------------------------------

**Returns**

a vector boolean long long, each containing all 0s(false) or 1s(true).

**7.6.4.20 static vb64\_t vec\_iszerof64 ( vf64\_t vf64 ) [inline],[static]**

Return 2x64-bit vector boolean true values, for each double value that is +-0.0.

A IEEE Binary64 zero has an exponent of 0x000 and a zero significand. The sign bit is ignored.

**Note**

This function will not raise VXSNaN or VXVC (FE\_INVALID) exceptions. A normal double compare can.

processor	Latency	Throughput
power8	4-13	2/cycle
power9	3	2/cycle

**Parameters**

<b>vf64</b>	a vector of __binary32 values.
-------------	--------------------------------

**Returns**

a vector boolean int, each containing all 0s(false) or 1s(true).

**7.6.4.21** `static long double vec_pack_longdouble ( vf64_t lval ) [inline],[static]`

Copy the pair of doubles from a vector to IBM long double.

**Parameters**

<i>lval</i>	vector double values containing the IBM long double.
-------------	--

**Returns**

IBM long double as FPR pair.

**7.6.4.22** `static vf64_t vec_unpack_longdouble ( long double lval ) [inline],[static]`

Copy the pair of doubles from a IBM long double to a vector double.

**Parameters**

<i>lval</i>	IBM long double as FPR pair.
-------------	------------------------------

**Returns**

vector double values containing the IBM long double.

## 7.7 src/vec\_int128\_ppc.h File Reference

Header package containing a collection of 128-bit computation functions implemented with PowerISA VMX and VSX instructions.

```
#include <vec_common_ppc.h>
#include <vec_int64_ppc.h>
```

**Functions**

- static [vui128\\_t vec\\_addcuq](#) ([vui128\\_t](#) a, [vui128\\_t](#) b)  
*Vector Add & write Carry Unsigned Quadword.*
- static [vui128\\_t vec\\_addecuq](#) ([vui128\\_t](#) a, [vui128\\_t](#) b, [vui128\\_t](#) ci)  
*Vector Add Extended & write Carry Unsigned Quadword.*
- static [vui128\\_t vec\\_addeuqm](#) ([vui128\\_t](#) a, [vui128\\_t](#) b, [vui128\\_t](#) ci)  
*Vector Add Extended Unsigned Quadword Modulo.*
- static [vui128\\_t vec\\_adduqm](#) ([vui128\\_t](#) a, [vui128\\_t](#) b)



- Vector Add Unsigned Quadword Modulo.*
- static [vui128\\_t](#) [vec\\_addcq](#) ([vui128\\_t](#) \*cout, [vui128\\_t](#) a, [vui128\\_t](#) b)
- Vector Add with carry Unsigned Quadword.*
- static [vui128\\_t](#) [vec\\_addeq](#) ([vui128\\_t](#) \*cout, [vui128\\_t](#) a, [vui128\\_t](#) b, [vui128\\_t](#) ci)
- Vector Add Extend with carry Unsigned Quadword.*
- static [vui128\\_t](#) [vec\\_clzq](#) ([vui128\\_t](#) vra)
- Count leading zeros for a vector \_\_int128.*
- static [vb128\\_t](#) [vec\\_cmpeqsq](#) ([vi128\\_t](#) vra, [vi128\\_t](#) vrb)
- Vector Compare Equal Signed Quadword.*
- static [vb128\\_t](#) [vec\\_cmpequq](#) ([vui128\\_t](#) vra, [vui128\\_t](#) vrb)
- Vector Compare Equal Unsigned Quadword.*
- static [vb128\\_t](#) [vec\\_cmpgesq](#) ([vi128\\_t](#) vra, [vi128\\_t](#) vrb)
- Vector Compare Greater Than or Equal Signed Quadword.*
- static [vb128\\_t](#) [vec\\_cmpgeuq](#) ([vui128\\_t](#) vra, [vui128\\_t](#) vrb)
- Vector Compare Greater Than or Equal Unsigned Quadword.*
- static [vb128\\_t](#) [vec\\_cmpgtsq](#) ([vi128\\_t](#) vra, [vi128\\_t](#) vrb)
- Vector Compare Greater Than Signed Quadword.*
- static [vb128\\_t](#) [vec\\_cmpgtuq](#) ([vui128\\_t](#) vra, [vui128\\_t](#) vrb)
- Vector Compare Greater Than Unsigned Quadword.*
- static [vb128\\_t](#) [vec\\_cmplesq](#) ([vi128\\_t](#) vra, [vi128\\_t](#) vrb)
- Vector Compare Less Than or Equal Signed Quadword.*
- static [vb128\\_t](#) [vec\\_cmpleuq](#) ([vui128\\_t](#) vra, [vui128\\_t](#) vrb)
- Vector Compare Less Than or Equal Unsigned Quadword.*
- static [vb128\\_t](#) [vec\\_cmpltseq](#) ([vi128\\_t](#) vra, [vi128\\_t](#) vrb)
- Vector Compare Less Than Signed Quadword.*
- static [vb128\\_t](#) [vec\\_cmpltuq](#) ([vui128\\_t](#) vra, [vui128\\_t](#) vrb)
- Vector Compare Less Than Unsigned Quadword.*
- static [vb128\\_t](#) [vec\\_cmpnesq](#) ([vi128\\_t](#) vra, [vi128\\_t](#) vrb)
- Vector Compare Equal Signed Quadword.*
- static [vb128\\_t](#) [vec\\_cmpneuq](#) ([vui128\\_t](#) vra, [vui128\\_t](#) vrb)
- Vector Compare Not Equal Unsigned Quadword.*
- static int [vec\\_cmpsq\\_all\\_eq](#) ([vi128\\_t](#) vra, [vi128\\_t](#) vrb)
- Vector Compare all Equal Signed Quadword.*
- static int [vec\\_cmpsq\\_all\\_ge](#) ([vi128\\_t](#) vra, [vi128\\_t](#) vrb)
- Vector Compare any Greater Than or Equal Signed Quadword.*
- static int [vec\\_cmpsq\\_all\\_gt](#) ([vi128\\_t](#) vra, [vi128\\_t](#) vrb)
- Vector Compare any Greater Than Signed Quadword.*
- static int [vec\\_cmpsq\\_all\\_le](#) ([vi128\\_t](#) vra, [vi128\\_t](#) vrb)
- Vector Compare any Less Than or Equal Signed Quadword.*
- static int [vec\\_cmpsq\\_all\\_lt](#) ([vi128\\_t](#) vra, [vi128\\_t](#) vrb)
- Vector Compare any Less Than Signed Quadword.*
- static int [vec\\_cmpsq\\_all\\_ne](#) ([vi128\\_t](#) vra, [vi128\\_t](#) vrb)
- Vector Compare all Not Equal Signed Quadword.*
- static int [vec\\_cmpuq\\_all\\_eq](#) ([vui128\\_t](#) vra, [vui128\\_t](#) vrb)
- Vector Compare all Equal Unsigned Quadword.*
- static int [vec\\_cmpuq\\_all\\_ge](#) ([vui128\\_t](#) vra, [vui128\\_t](#) vrb)
- Vector Compare any Greater Than or Equal Unsigned Quadword.*
- static int [vec\\_cmpuq\\_all\\_gt](#) ([vui128\\_t](#) vra, [vui128\\_t](#) vrb)
- Vector Compare any Greater Than Unsigned Quadword.*
- static int [vec\\_cmpuq\\_all\\_le](#) ([vui128\\_t](#) vra, [vui128\\_t](#) vrb)
- Vector Compare any Less Than or Equal Unsigned Quadword.*

- static int [vec\\_cmpuq\\_all\\_lt](#) ([vui128\\_t](#) vra, [vui128\\_t](#) vrb)  
*Vector Compare any Less Than Unsigned Quadword.*
- static int [vec\\_cmpuq\\_all\\_ne](#) ([vui128\\_t](#) vra, [vui128\\_t](#) vrb)  
*Vector Compare all Not Equal Unsigned Quadword.*
- static [vui128\\_t](#) [vec\\_cmul10ecuq](#) ([vui128\\_t](#) \*cout, [vui128\\_t](#) a, [vui128\\_t](#) cin)  
*Vector combined Multiply by 10 Extended & write Carry Unsigned Quadword.*
- static [vui128\\_t](#) [vec\\_cmul10cuq](#) ([vui128\\_t](#) \*cout, [vui128\\_t](#) a)  
*Vector combined Multiply by 10 & write Carry Unsigned Quadword.*
- static [vui128\\_t](#) [vec\\_mul10cuq](#) ([vui128\\_t](#) a)  
*Vector Multiply by 10 & write Carry Unsigned Quadword.*
- static [vui128\\_t](#) [vec\\_mul10ecuq](#) ([vui128\\_t](#) a, [vui128\\_t](#) cin)  
*Vector Multiply by 10 Extended & write Carry Unsigned Quadword.*
- static [vui128\\_t](#) [vec\\_mul10euq](#) ([vui128\\_t](#) a, [vui128\\_t](#) cin)  
*Vector Multiply by 10 extended Unsigned Quadword.*
- static [vui128\\_t](#) [vec\\_mul10uq](#) ([vui128\\_t](#) a)  
*Vector Multiply by 10 Unsigned Quadword.*
- static [vui128\\_t](#) [vec\\_cmul100cuq](#) ([vui128\\_t](#) \*cout, [vui128\\_t](#) a)  
*Vector combined Multiply by 100 & write Carry Unsigned Quadword.*
- static [vui128\\_t](#) [vec\\_cmul100ecuq](#) ([vui128\\_t](#) \*cout, [vui128\\_t](#) a, [vui128\\_t](#) cin)  
*Vector combined Multiply by 100 Extended & write Carry Unsigned Quadword.*
- static [vui128\\_t](#) [vec\\_msumudm](#) ([vui64\\_t](#) a, [vui64\\_t](#) b, [vui128\\_t](#) c)  
*Vector Multiply-Sum Unsigned Doubleword Modulo.*
- static [vui128\\_t](#) [vec\\_muleud](#) ([vui64\\_t](#) a, [vui64\\_t](#) b)  
*Vector multiply even unsigned doublewords.*
- static [vui64\\_t](#) [vec\\_mulhud](#) ([vui64\\_t](#) vra, [vui64\\_t](#) vrb)  
*Vector Multiply High Unsigned Doubleword.*
- static [vui128\\_t](#) [vec\\_muloud](#) ([vui64\\_t](#) a, [vui64\\_t](#) b)  
*Vector multiply odd unsigned doublewords.*
- static [vui64\\_t](#) [vec\\_muludm](#) ([vui64\\_t](#) vra, [vui64\\_t](#) vrb)  
*Vector Multiply Unsigned Doubleword Modulo.*
- static [vui128\\_t](#) [vec\\_muludq](#) ([vui128\\_t](#) \*mulu, [vui128\\_t](#) a, [vui128\\_t](#) b)  
*Vector Multiply Unsigned double Quadword.*
- static [vui128\\_t](#) [vec\\_mulluq](#) ([vui128\\_t](#) a, [vui128\\_t](#) b)  
*Vector Multiply Low Unsigned Quadword.*
- static [vui128\\_t](#) [vec\\_popcntq](#) ([vui128\\_t](#) vra)  
*Population Count vector \_\_int128.*
- static [vui128\\_t](#) [vec\\_revbq](#) ([vui128\\_t](#) vra)  
*byte reverse quadword for a vector \_\_int128.*
- static [vb128\\_t](#) [vec\\_setb\\_cyq](#) ([vui128\\_t](#) vcy)  
*Vector Set Bool from Quadword Carry.*
- static [vb128\\_t](#) [vec\\_setb\\_ncq](#) ([vui128\\_t](#) vcy)  
*Vector Set Bool from Quadword not Carry.*
- static [vb128\\_t](#) [vec\\_setb\\_sq](#) ([vui128\\_t](#) vra)  
*Vector Set Bool from Signed Quadword.*
- static [vui128\\_t](#) [vec\\_sldq](#) ([vui128\\_t](#) vrw, [vui128\\_t](#) vrx, [vui128\\_t](#) vrb)  
*Vector Shift Left double Quadword.*
- static [vui128\\_t](#) [vec\\_slqi](#) ([vui128\\_t](#) vra, const unsigned int shb)  
*Vector Shift left Quadword Immediate.*
- static [vui128\\_t](#) [vec\\_slq](#) ([vui128\\_t](#) vra, [vui128\\_t](#) vrb)  
*Vector Shift Left Quadword.*
- static [vui128\\_t](#) [vec\\_srqi](#) ([vui128\\_t](#) vra, const unsigned int shb)

*Vector Shift right Quadword Immediate.*

- static `vui128_t vec_srq (vui128_t vra, vui128_t vrb)`

*Vector Shift right Quadword.*

- static `vui128_t vec_slq4 (vui128_t vra)`
- static `vui128_t vec_slq5 (vui128_t vra)`
- static `vui128_t vec_srq4 (vui128_t vra)`
- static `vui128_t vec_srq5 (vui128_t vra)`
- static `vui128_t vec_subcuq (vui128_t vra, vui128_t vrb)`

*Vector Subtract and Write Carry Unsigned Quadword.*

- static `vui128_t vec_subecuq (vui128_t vra, vui128_t vrb, vui128_t vrc)`

*Vector Subtract Extended and Write Carry Unsigned Quadword.*

- static `vui128_t vec_subeuqm (vui128_t vra, vui128_t vrb, vui128_t vrc)`

*Vector Subtract Extended Unsigned Quadword Modulo.*

- static `vui128_t vec_subuqm (vui128_t vra, vui128_t vrb)`

*Vector subtract Unsigned Quadword Modulo.*

- static `vui128_t vec_vmuleud (vui64_t a, vui64_t b)`

*Vector multiply even unsigned doublewords.*

- static `vui128_t vec_vmuloud (vui64_t a, vui64_t b)`

*Vector multiply odd unsigned doublewords.*

### 7.7.1 Detailed Description

Header package containing a collection of 128-bit computation functions implemented with PowerISA VMX and VSX instructions.

The original VMX (AKA AltiVec) only defined a few instructions that operated on the 128-bit vector as a whole. This included the vector shift left/right (bit), vector shift left/right by octet (byte), vector shift left double by octet (select a contiguous 16-bytes from 2 concatenated vectors) 256-bit), and generalized vector permute (select any 16-bytes from 2 concatenated vectors). Use of these instructions can be complicated when;

- the shift amount is more than 8 bits,
- the shift amount is not a multiple of 8-bits (octet),
- the shift amount is a constant and needs to be generated/loaded before use.

These operations require pre-conditions to avoid multiple instructions or require a combination of (bit and octet shift) instructions to get the quadword result. The compiler <altivec.h> built-ins only supports individual instructions. So using these operations quickly inspires a need for a header (like this) to contain implementations of the common operations.

The VSX facility (introduced with POWER7) did not add any integer doubleword (64-bit) or quadword (128-bit) operations. However it did add a useful doubleword permute immediate and word wise; merge, shift, and splat immediate operations. Otherwise vector `__int128` (128-bit elements) operations have to be implemented using VMX word and halfword element integer operations for POWER7.

POWER8 added multiply word operations that produce the full doubleword product and full quadword add / subtract (with carry extend). The add quadword is useful to sum the partial products for a full 128 x 128-bit multiply. The add quadword write carry and extend forms, simplify extending arithmetic to 256-bits and beyond.

While POWER8 provided quadword integer add and subtract operations, it did not provide quadword Signed/↔ Unsigned integer compare operations. It is possible to implement quadword compare operations using existing word / doubleword compares and the the new quadword subtract, but this requires some study. Pveclib provides easy to use quadword compare operations.

POWER9 (PowerISA 3.0B) adds the **Vector Multiply-Sum unsigned Doubleword Modulo** instruction. Aspects of this instruction mean it needs to be used carefully as part of larger quadword multiply. It performs only two of the four required doubleword multiplies. The final quadword modulo sum will discard any overflow/carry from the potential 130-bit result. With careful pre-conditioning of doubleword inputs the results are can not overflow from 128-bits. Then separate add quadword add/write carry operations can be used to complete the sum of partial products. These techniques are used in the POWER9 specific implementations of `vec_muleud`, `vec_muloud`, `vec_mulluq`, and `vec_muludq`.

PowerISA 3.0B also defined additional: Binary Coded Decimal (BCD) and Zoned character format conversions. String processing operations. Vector Parity operations. Integer Extend Sign Operations. Integer Absolute Difference Operations. All of these seem to be useful additions to pveclib for older (POWER7/8) processors and across element sizes (including quadword elements).

Most of these intrinsic (compiler built-in) operations are defined in `<altivec.h>` and described in the compiler documentation. However it took several compiler releases for all the new POWER8 64-bit and 128-bit integer vector intrinsics to be added to **altivec.h**. This support started with the GCC 4.9 but was not complete across function/type and bug free until GCC 6.0.

#### Note

The compiler disables associated `<altivec.h>` built-ins if the **mcpu** target does not enable the specific instruction. For example if you compile with **-mcpu=power7**, `vec_vadduqm` and `vec_vsubudm` will not be defined. But `vec_adduqm()` and `vec_subudm()` and always be defined in this header, will generate the minimum code, appropriate for the target, and produce correct results.

Most of these operations are implemented in a single instruction on newer (POWER8/POWER9) processors. So this header serves to fill in functional gaps for older (POWER7, POWER8) processors and provides a in-line assembler implementation for older compilers that do not provide the build-ins.

This header covers operations that are either:

- Operations implemented in hardware instructions for later processors and useful to programmers, on slightly older processors, even if the equivalent function requires more instructions. Examples include quadword byte reverse, add and subtract.
- Defined in the OpenPOWER ABI but *not* yet defined in `<altivec.h>` provided by available compilers in common use. Examples include quadword byte reverse, add and subtract.
- Are commonly used operations, not covered by the ABI or `<altivec.h>`, and require multiple instructions or are not obvious. Examples include quadword; Signed and Unsigned compare, shift immediate, multiply, multiply by 10 immediate, count leading zeros and population count.

#### Note

The Multiply sum/even/odd doubleword operations are currently implemented here (in `<vec_int128_ppc.h>`) which resolves a dependency on Add Quadword. These functions (`vec_msumudm`, `vec_muleud`, `vec_muloud`) all produce a quadword results and may use the `vec_adduqm` implementation to sum partial products.

See [Returning extended quadword results](#). for more background on extended quadword computation.

### 7.7.2 Performance data.

High level performance estimates are provided as an aid to function selection when evaluating algorithms. For background on how *Latency* and *Throughput* are derived see: [Performance data](#).

### 7.7.3 Function Documentation

**7.7.3.1** `static vui128_t vec_addcq ( vui128_t * cout, vui128_t a, vui128_t b )` `[inline]`, `[static]`

Vector Add with carry Unsigned Quadword.

Add two vector `__int128` values and return sum and the carry out.

processor	Latency	Throughput
power8	8	1/2 cycles
power9	6	2/cycle

**Parameters**

<i>*cout</i>	carry out from the sum of a and b.
<i>a</i>	128-bit vector treated a __int128.
<i>b</i>	128-bit vector treated a __int128.

**Returns**

\_\_int128 (lower 128-bits) sum of a and b.

**7.7.3.2** `static vui128_t vec_addcuq ( vui128_t a, vui128_t b )` `[inline], [static]`

Vector Add & write Carry Unsigned Quadword.

Add two vector \_\_int128 values and return the carry out.

processor	Latency	Throughput
power8	4	2/2 cycles
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated a __int128.
<i>b</i>	128-bit vector treated a __int128.

**Returns**

\_\_int128 carry of the sum of a and b.

**7.7.3.3** `static vui128_t vec_addecuq ( vui128_t a, vui128_t b, vui128_t ci )` `[inline], [static]`

Vector Add Extended & write Carry Unsigned Quadword.

Add two vector \_\_int128 values plus a carry-in (0|1) and return the carry out bit.

processor	Latency	Throughput
power8	4	2/2 cycles
power9	3	2/cycle

## Parameters

<i>a</i>	128-bit vector treated a __int128.
<i>b</i>	128-bit vector treated a __int128.
<i>ci</i>	Carry-in from vector bit[127].

## Returns

carry-out in bit[127] of the sum of  $a + b + c$ .

```
7.7.3.4 static vui128_t vec_addeq ( vui128_t * cout, vui128_t a, vui128_t b, vui128_t ci ) [inline],
[static]
```

Vector Add Extend with carry Unsigned Quadword.

Add two vector \_\_int128 values plus a carry-in (0|1) and return sum and the carry out.

processor	Latency	Throughput
power8	8	1/2 cycles
power9	6	2/cycle

## Parameters

<i>*cout</i>	carry out from the sum of a and b.
<i>a</i>	128-bit vector treated a __int128.
<i>b</i>	128-bit vector treated a __int128.
<i>ci</i>	Carry-in from vector bit[127].

## Returns

\_\_int128 (lower 128-bits) sum of  $a + b + c$ .

```
7.7.3.5 static vui128_t vec_addeuqm ( vui128_t a, vui128_t b, vui128_t ci ) [inline],[static]
```

Vector Add Extended Unsigned Quadword Modulo.

Add two vector \_\_int128 values plus a carry (0|1) and return the modulo 128-bit result.

processor	Latency	Throughput
power8	4	2/2 cycles
power9	3	2/cycle

## Parameters

<i>a</i>	128-bit vector treated a __int128.
<i>b</i>	128-bit vector treated a __int128.

## Parameters

<i>ci</i>	Carry-in from vector bit[127].
-----------	--------------------------------

## Returns

\_\_int128 sum of a + b + c, modulo 128-bits.

**7.7.3.6** `static vui128_t vec_adduqm ( vui128_t a, vui128_t b )` `[inline], [static]`

Vector Add Unsigned Quadword Modulo.

Add two vector \_\_int128 values and return result modulo 128-bits.

processor	Latency	Throughput
power8	4	2/2 cycles
power9	3	2/cycle

## Parameters

<i>a</i>	128-bit vector treated as a __int128.
<i>b</i>	128-bit vector treated as a __int128.

## Returns

\_\_int128 sum of a and b.

**7.7.3.7** `static vui128_t vec_clzq ( vui128_t vra )` `[inline], [static]`

Count leading zeros for a vector \_\_int128.

Count leading zeros for a vector \_\_int128 and return the count in a vector suitable for use with vector shift (left|right) and vector shift (left|right) by octet instructions.

processor	Latency	Throughput
power8	19-28	1/cycle
power9	25-36	1/cycle

## Parameters

<i>vra</i>	a 128-bit vector treated a __int128.
------------	--------------------------------------

## Returns

a 128-bit vector with bits 121:127 containing the count of leading zeros.

**7.7.3.8** `static vb128_t vec_cmpeqsq ( vi128_t vra, vi128_t vrb ) [inline],[static]`

Vector Compare Equal Signed Quadword.

Compare signed `__int128` (128-bit) integers and return all '1's, if `vra == vrb`, otherwise all '0's. We use `vec_cmpequq` as it works for both signed and unsigned compares.

processor	Latency	Throughput
power8	6	2/cycle
power9	7	2/cycle

#### Parameters

<i>vra</i>	128-bit vector treated as an signed <code>__int128</code> .
<i>vrb</i>	128-bit vector treated as an signed <code>__int128</code> .

#### Returns

128-bit vector boolean reflecting vector signed `__int128` compare equal.

**7.7.3.9** `static vb128_t vec_cmpequq ( vui128_t vra, vui128_t vrb ) [inline],[static]`

Vector Compare Equal Unsigned Quadword.

Compare unsigned `__int128` (128-bit) integers and return all '1's, if `vra == vrb`, otherwise all '0's.

For POWER8 (PowerISA 2.07B) or later, use the Vector Compare Equal Unsigned DoubleWord (**vcmpqud**) instruction. To get the correct quadword result, the doubleword element equal truth values are swapped, then *anded* with the original compare results. Otherwise use vector word compare and additional boolean logic to insure all word elements are equal.

processor	Latency	Throughput
power8	6	2/cycle
power9	7	2/cycle

#### Parameters

<i>vra</i>	128-bit vector treated as an unsigned <code>__int128s</code> .
<i>vrb</i>	128-bit vector treated as an unsigned <code>__int128</code> .

#### Returns

128-bit vector boolean reflecting vector unsigned `__int128` compare equal.

**7.7.3.10** `static vb128_t vec_cmpgesq ( vi128_t vra, vi128_t vrb ) [inline],[static]`

Vector Compare Greater Than or Equal Signed Quadword.



Compare signed `__int128` (128-bit) integers and return all '1's, if `vra >= vrb`, otherwise all '0's.

Flip the operand sign bits and use `vec_cmpgeuq` for signed compare.

processor	Latency	Throughput
power8	10-16	1/ 2cycles
power9	8-14	1/cycle

#### Parameters

<i>vra</i>	128-bit vector treated as an signed <code>__int128</code> .
<i>vrb</i>	128-bit vector treated as an signed <code>__int128</code> .

#### Returns

128-bit vector boolean reflecting vector signed `__int128` compare greater than.

```
7.7.3.11 static vb128_t vec_cmpgeuq ( vui128_t vra, vui128_t vrb ) [inline],[static]
```

Vector Compare Greater Than or Equal Unsigned Quadword.

Compare unsigned `__int128` (128-bit) integers and return all '1's, if `vra >= vrb`, otherwise all '0's.

For POWER8 (PowerISA 2.07B) or later, use the Vector Subtract & write Carry QuadWord (**vsubcuq**) instruction. This generates a carry for greater than or equal and NOT carry for less than. Then use `vec_setb_cyq` to convert the carry into a vector bool. Here we use the pveclib implementations (`vec_subcuq()` and `vec_setb_cyq()`), instead of `<altivec.h>` intrinsics, to address older compilers and POWER7.

processor	Latency	Throughput
power8	8	2/ 2cycles
power9	6	2/cycle

#### Parameters

<i>vra</i>	128-bit vector treated as an unsigned <code>__int128</code> .
<i>vrb</i>	128-bit vector treated as an unsigned <code>__int128</code> .

#### Returns

128-bit vector boolean reflecting vector unsigned `__int128` compare greater than.

```
7.7.3.12 static vb128_t vec_cmpgtsq ( vi128_t vra, vi128_t vrb ) [inline],[static]
```

Vector Compare Greater Than Signed Quadword.

Compare signed `__int128` (128-bit) integers and return all '1's, if `vra > vrb`, otherwise all '0's.

Flip the operand sign bits and use `vec_cmpgtuq` for signed compare.

processor	Latency	Throughput
power8	10-16	1/ 2cycles
power9	8-14	1/cycle

#### Parameters

<i>vra</i>	128-bit vector treated as an signed <code>__int128</code> .
<i>vrb</i>	128-bit vector treated as an signed <code>__int128</code> .

#### Returns

128-bit vector boolean reflecting vector signed `__int128` compare greater than.

**7.7.3.13** `static vb128_t vec_cmpgtuq ( vui128_t vra, vui128_t vrb )` `[inline]`, `[static]`

Vector Compare Greater Than Unsigned Quadword.

Compare unsigned `__int128` (128-bit) integers and return all '1's, if `vra > vrb`, otherwise all '0's.

For POWER8 (PowerISA 2.07B) or later, use the Vector Subtract & write Carry QuadWord (**vsubcuq**) instruction with the parameters reversed. This generates a carry for less than or equal and NOT carry for greater than. Then use `vec_setb_ncq` to convert the carry into a vector bool. Here we use the pveclib implementations ([vec\\_subcuq\(\)](#) and [vec\\_setb\\_ncq\(\)](#)), instead of `<altivec.h>` intrinsics, to address older compilers and POWER7.

processor	Latency	Throughput
power8	8	2/ 2cycles
power9	6	2/cycle

#### Parameters

<i>vra</i>	128-bit vector treated as an unsigned <code>__int128</code> .
<i>vrb</i>	128-bit vector treated as an unsigned <code>__int128</code> .

#### Returns

128-bit vector boolean reflecting vector unsigned `__int128` compare greater than.

**7.7.3.14** `static vb128_t vec_cmpleuq ( vi128_t vra, vi128_t vrb )` `[inline]`, `[static]`

Vector Compare Less Than or Equal Signed Quadword.

Compare signed `__int128` (128-bit) integers and return all '1's, if `vra <= vrb`, otherwise all '0's.

Flip the operand sign bits and use `vec_cmpleuq` for signed compare.

processor	Latency	Throughput
power8	10-16	1/ 2cycles
power9	8-14	1/cycle

**Parameters**

<i>vra</i>	128-bit vector treated as an signed __int128.
<i>vr</i> <i>b</i>	128-bit vector treated as an signed __int128.

**Returns**

128-bit vector boolean reflecting vector signed \_\_int128 compare less than or equal.

**7.7.3.15** `static vb128_t vec_cmpleuq ( vui128_t vra, vui128_t vr`*b*`)` [*inline*], [*static*]

Vector Compare Less Than or Equal Unsigned Quadword.

Compare unsigned \_\_int128 (128-bit) integers and return all '1's, if *vra* <= *vr**b*, otherwise all '0's.

For POWER8 (PowerISA 2.07B) or later, use the Vector Subtract & write Carry QuadWord (**vsubcuq**) instruction. This generates a carry for greater than or equal and NOT carry for less than. Then use `vec_setb_ncq` to convert the carry into a vector bool. Here we use the pveclib implementations ([vec\\_subcuq\(\)](#) and [vec\\_setb\\_cyq\(\)](#)), instead of <altivec.h> intrinsics, to address older compilers and POWER7.

processor	Latency	Throughput
power8	8	2/ 2cycles
power9	6	2/cycle

**Parameters**

<i>vra</i>	128-bit vector treated as an unsigned __int128.
<i>vr</i> <i>b</i>	128-bit vector treated as an unsigned __int128.

**Returns**

128-bit vector boolean reflecting vector unsigned \_\_int128 compare less than or equal.

**7.7.3.16** `static vb128_t vec_cmpltuq ( vi128_t vra, vi128_t vr`*b*`)` [*inline*], [*static*]

Vector Compare Less Than Signed Quadword.

Compare signed \_\_int128 (128-bit) integers and return all '1's, if *vra* < *vr**b*, otherwise all '0's.

Flip the operand sign bits and use `vec_cmpltuq` for signed compare.

processor	Latency	Throughput
power8	10-16	1/ 2cycles
power9	8-14	1/cycle

**Parameters**

<i>vra</i>	128-bit vector treated as an signed __int128.
<i>vrb</i>	128-bit vector treated as an signed __int128.

**Returns**

128-bit vector boolean reflecting vector unsigned \_\_int128 compare less than.

**7.7.3.17** `static vb128_t vec_cmpltuq ( vui128_t vra, vui128_t vrb ) [inline],[static]`

Vector Compare Less Than Unsigned Quadword.

Compare unsigned \_\_int128 (128-bit) integers and return all '1's, if *vra* < *vrb*, otherwise all '0's.

For POWER8 (PowerISA 2.07B) or later, use the Vector Subtract & write Carry QuadWord (**vsubcuq**) instruction. This generates a carry for greater than or equal and NOT carry for less than. Then use `vec_setb_ncq` to convert the carry into a vector bool. Here we use the pveclib implementations (`vec_subcuq()` and `vec_setb_ncq()`), instead of `<altivec.h>` intrinsics, to address older compilers and POWER7.

processor	Latency	Throughput
power8	8	2/ 2cycles
power9	6	2/cycle

**Parameters**

<i>vra</i>	128-bit vector treated as an unsigned __int128.
<i>vrb</i>	128-bit vector treated as an unsigned __int128.

**Returns**

128-bit vector boolean reflecting vector unsigned \_\_int128 compare less than.

**7.7.3.18** `static vb128_t vec_cmpnesq ( vi128_t vra, vi128_t vrb ) [inline],[static]`

Vector Compare Equal Signed Quadword.

Compare signed \_\_int128 (128-bit) integers and return all '1's, if *vra* != *vrb*, otherwise all '0's. We use `vec_cmpequq` as it works for both signed and unsigned compares.

processor	Latency	Throughput
power8	6	2/cycle
power9	7	2/cycle

## Parameters

<i>vra</i>	128-bit vector treated as an signed __int128.
<i>vrb</i>	128-bit vector treated as an signed __int128.

## Returns

128-bit vector boolean reflecting vector signed \_\_int128 compare not equal.

**7.7.3.19** `static vb128_t vec_cmpneuq ( vui128_t vra, vui128_t vrb ) [inline], [static]`

Vector Compare Not Equal Unsigned Quadword.

Compare unsigned \_\_int128 (128-bit) integers and return all '1's, if *vra* != *vr*b, otherwise all '0's.

For POWER8 (PowerISA 2.07B) or later, use the Vector Compare Equal Unsigned DoubleWord (**vc<sub>mpequd</sub>**) instruction. To get the correct quadword result, the doubleword element equal truth values are swapped, then *not anded* with the original compare results. Otherwise use vector word compare and additional boolean logic to insure all word elements are equal.

processor	Latency	Throughput
power8	6	2/cycle
power9	7	2/cycle

## Parameters

<i>vra</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>vr</i> b	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

## Returns

128-bit vector boolean reflecting vector unsigned \_\_int128 compare equal.

**7.7.3.20** `static int vec_cmpsq_all_eq ( vi128_t vra, vi128_t vrb ) [inline], [static]`

Vector Compare all Equal Signed Quadword.

Compare vector signed \_\_int128 values and return true if *vra* and *vr*b are equal.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

## Parameters

<i>vra</i>	128-bit vector treated as an vector signed __int128 (qword) element.
------------	--

**Parameters**

<i>vr<b>b</b></i>	128-bit vector treated as an vector signed <code>__int128</code> (qword) element.
-------------------	---

**Returns**

boolean int for all 128-bits, true if equal, false otherwise.

**7.7.3.21** `static int vec_cmpsq_all_ge ( vi128_t vra, vi128_t vrb )` `[inline],[static]`

Vector Compare any Greater Than or Equal Signed Quadword.

Compare vector unsigned `__int128` values and return true if `vra >= vrb`.

processor	Latency	Throughput
power8	10-15	1/ 2cycles
power9	8	1/cycle

**Parameters**

<i>vra</i>	128-bit vector treated as an vector signed <code>__int128</code> (qword) element.
<i>vr<b>b</b></i>	128-bit vector treated as an vector signed <code>__int128</code> (qword) element.

**Returns**

boolean int for all 128-bits, true if Greater Than or Equal, false otherwise.

**7.7.3.22** `static int vec_cmpsq_all_gt ( vi128_t vra, vi128_t vrb )` `[inline],[static]`

Vector Compare any Greater Than Signed Quadword.

Compare vector signed `__int128` values and return true if `vra > vrb`.

processor	Latency	Throughput
power8	10-15	1/ 2cycles
power9	8	1/cycle

**Parameters**

<i>vra</i>	128-bit vector treated as an vector signed <code>__int128</code> (qword) element.
<i>vr<b>b</b></i>	128-bit vector treated as an vector signed <code>__int128</code> (qword) element.

**Returns**

boolean int for all 128-bits, true if Greater Than, false otherwise.

**7.7.3.23** `static int vec_cmpsq_all_le ( vi128_t vra, vi128_t vrb ) [inline],[static]`

Vector Compare any Less Than or Equal Signed Quadword.

Compare vector signed \_\_int128 values and return true if  $vra \leq vrb$ .

processor	Latency	Throughput
power8	10-15	1/ 2cycles
power9	8	1/cycle

#### Parameters

<i>vra</i>	128-bit vector treated as an vector signed __int128 (qword) element.
<i>vrb</i>	128-bit vector treated as an vector signed __int128 (qword) element.

#### Returns

boolean int for all 128-bits, true if Less Than or Equal, false otherwise.

**7.7.3.24** `static int vec_cmpsq_all_lt ( vi128_t vra, vi128_t vrb ) [inline],[static]`

Vector Compare any Less Than Signed Quadword.

Compare vector signed \_\_int128 values and return true if  $vra < vrb$ .

processor	Latency	Throughput
power8	10-15	1/ 2cycles
power9	8	1/cycle

#### Parameters

<i>vra</i>	128-bit vector treated as an vector signed __int128 (qword) element.
<i>vrb</i>	128-bit vector treated as an vector signed __int128 (qword) element.

#### Returns

boolean int for all 128-bits, true if Less Than, false otherwise.

**7.7.3.25** `static int vec_cmpsq_all_ne ( vi128_t vra, vi128_t vrb ) [inline],[static]`

Vector Compare all Not Equal Signed Quadword.

Compare vector signed \_\_int128 values and return true if *vra* and *vrb* are not equal.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>vra</i>	128-bit vector treated as an vector signed <code>__int128</code> (qword) element.
<i>vrb</i>	128-bit vector treated as an vector signed <code>__int128</code> (qword) element.

**Returns**

boolean `__int128` for all 128-bits, true if equal, false otherwise.

**7.7.3.26** `static int vec_cmpuq_all_eq( vui128_t vra, vui128_t vrb )` `[inline],[static]`

Vector Compare all Equal Unsigned Quadword.

Compare vector unsigned `__int128` values and return true if *vra* and *vrb* are equal.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>vra</i>	128-bit vector treated as an vector unsigned <code>__int128</code> (qword) element.
<i>vrb</i>	128-bit vector treated as an vector unsigned <code>__int128</code> (qword) element.

**Returns**

boolean int for all 128-bits, true if equal, false otherwise.

**7.7.3.27** `static int vec_cmpuq_all_ge( vui128_t vra, vui128_t vrb )` `[inline],[static]`

Vector Compare any Greater Than or Equal Unsigned Quadword.

Compare vector unsigned `__int128` values and return true if *vra*  $\geq$  *vrb*.

processor	Latency	Throughput
power8	8-13	2/ 2cycles
power9	6	2/cycle

**Parameters**

<i>vra</i>	128-bit vector treated as an vector unsigned <code>__int128</code> (qword) element.
<i>vrb</i>	128-bit vector treated as an vector unsigned <code>__int128</code> (qword) element.



**Returns**

boolean int for all 128-bits, true if Greater Than or Equal, false otherwise.

**7.7.3.28** `static int vec_cmpuq_all_gt ( vui128_t vra, vui128_t vrb )` `[inline], [static]`

Vector Compare any Greater Than Unsigned Quadword.

Compare vector unsigned \_\_int128 values and return true if  $vra > vrb$ .

processor	Latency	Throughput
power8	8-13	2/ 2cycles
power9	6	2/cycle

**Parameters**

<i>vra</i>	128-bit vector treated as an vector unsigned __int128 (qword) element.
<i>vrb</i>	128-bit vector treated as an vector unsigned __int128 (qword) element.

**Returns**

boolean int for all 128-bits, true if Greater Than, false otherwise.

**7.7.3.29** `static int vec_cmpuq_all_le ( vui128_t vra, vui128_t vrb )` `[inline], [static]`

Vector Compare any Less Than or Equal Unsigned Quadword.

Compare vector unsigned \_\_int128 values and return true if  $vra \leq vrb$ .

processor	Latency	Throughput
power8	8-13	2/ 2cycles
power9	6	2/cycle

**Parameters**

<i>vra</i>	128-bit vector treated as an vector unsigned __int128 (qword) element.
<i>vrb</i>	128-bit vector treated as an vector unsigned __int128 (qword) element.

**Returns**

boolean int for all 128-bits, true if Less Than or Equal, false otherwise.

**7.7.3.30** `static int vec_cmpuq_all_lt ( vui128_t vra, vui128_t vrb )` `[inline], [static]`

Vector Compare any Less Than Unsigned Quadword.

Compare vector unsigned \_\_int128 values and return true if  $vra < vrb$ .

processor	Latency	Throughput
power8	8-13	2/ 2cycles
power9	6	2/cycle

## Parameters

<i>vra</i>	128-bit vector treated as an vector unsigned __int128 (qword) element.
<i>vrh</i>	128-bit vector treated as an vector unsigned __int128 (qword) element.

## Returns

boolean int for all 128-bits, true if Less Than, false otherwise.

**7.7.3.31** `static int vec_cmpuq_all_ne ( vui128_t vra, vui128_t vrh )` `[inline],[static]`

Vector Compare all Not Equal Unsigned Quadword.

Compare vector unsigned \_\_int128 values and return true if vra and vrh are not equal.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

## Parameters

<i>vra</i>	128-bit vector treated as an vector unsigned __int128 (qword) element.
<i>vrh</i>	128-bit vector treated as an vector unsigned __int128 (qword) element.

## Returns

boolean \_\_int128 for all 128-bits, true if equal, false otherwise.

**7.7.3.32** `static vui128_t vec_cmulo100cuq ( vui128_t * cout, vui128_t a )` `[inline],[static]`

Vector combined Multiply by 100 & write Carry Unsigned Quadword.

compute the product of a 128 bit values a \* 100. Only the low order 128 bits of the product are returned.

processor	Latency	Throughput
power8	13-15	1/cycle
power9	6	1/cycle

## Parameters

<i>*cout</i>	pointer to upper 128-bits of the product.
--------------	---

## Parameters

<i>a</i>	128-bit vector treated as a <code>__int128</code> .
----------	---

## Returns

vector `__int128` (lower 128-bits of the 256-bit product)  $a * 100$ .

**7.7.3.33** `static vui128_t vec_cmul100ecuq ( vui128_t * cout, vui128_t a, vui128_t cin )` `[inline], [static]`

Vector combined Multiply by 100 Extended & write Carry Unsigned Quadword.

Compute the product of a 128 bit value  $a * 100 + \text{digit}(\text{cin})$ . The function return its low order 128 bits of the extended product. The first parameter (*\*cout*) it the address of the vector to receive the generated carry out in the range 0-99.

processor	Latency	Throughput
power8	15-17	1/cycle
power9	9	1/cycle

## Parameters

<i>*cout</i>	pointer to upper 128-bits of the product.
<i>a</i>	128-bit vector treated as a unsigned <code>__int128</code> .
<i>cin</i>	values 0-99 in bits 120:127 of a vector.

## Returns

vector `__int128` (lower 128-bits of the 256-bit product)  $a * 100$ .

**7.7.3.34** `static vui128_t vec_cmul10cuq ( vui128_t * cout, vui128_t a )` `[inline], [static]`

Vector combined Multiply by 10 & write Carry Unsigned Quadword.

compute the product of a 128 bit values  $a * 10$ . Only the low order 128 bits of the product are returned.

processor	Latency	Throughput
power8	13-15	1/cycle
power9	3	1/ 2cycles

## Parameters

<i>*cout</i>	pointer to upper 128-bits of the product.
<i>a</i>	128-bit vector treated as a <code>__int128</code> .

**Returns**

vector \_\_int128 (lower 128-bits of the 256-bit product)  $a * 10$ .

**7.7.3.35** `static vui128_t vec_cmul10ecuq ( vui128_t * cout, vui128_t a, vui128_t cin )` `[inline], [static]`

Vector combined Multiply by 10 Extended & write Carry Unsigned Quadword.

Compute the product of a 128 bit value  $a * 10 + \text{digit}(\text{cin})$ . Only the low order 128 bits of the extended product are returned.

processor	Latency	Throughput
power8	13-15	1/cycle
power9	3	1/ 2cycles

**Parameters**

<i>*cout</i>	pointer to upper 128-bits of the product.
<i>a</i>	128-bit vector treated as a unsigned __int128.
<i>cin</i>	values 0-9 in bits 124:127 of a vector.

**Returns**

vector \_\_int128 (upper 128-bits of the 256-bit product)  $a * 10$ .

**7.7.3.36** `static vui128_t vec_msumudm ( vui64_t a, vui64_t b, vui128_t c )` `[inline], [static]`

Vector Multiply-Sum Unsigned Doubleword Modulo.

compute the even and odd produ256 bit product of two 128 bit values a, b. Only the low order 128 bits of the product are returned.

processor	Latency	Throughput
power8	30-32	1/cycle
power9	5-7	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as __vector unsigned long int.
<i>b</i>	128-bit vector treated as __vector unsigned long int.
<i>c</i>	128-bit vector treated as __vector unsigned __int128.

**Returns**

\_\_vector unsigned Modulo Sum of the 128-bit even / odd products of operands a and b plus the unsigned \_\_int128 operand c.

**7.7.3.37** `static vui128_t vec_mul10cuq( vui128_t a ) [inline],[static]`

Vector Multiply by 10 & write Carry Unsigned Quadword.

compute the product of a 128 bit value  $a * 10$ . Only the high order 128 bits of the product are returned. This will be binary coded decimal value 0-9 in bits 124-127, Bits 0-123 will be '0'.

processor	Latency	Throughput
power8	13-15	1/cycle
power9	3	1/cycle

#### Parameters

<i>a</i>	128-bit vector treated as a <code>__int128</code> .
----------	---

#### Returns

`__int128` (upper 128-bits of the 256-bit product)  $a * 10 >> 128$ .

**7.7.3.38** `static vui128_t vec_mul10ecuq( vui128_t a, vui128_t cin ) [inline],[static]`

Vector Multiply by 10 Extended & write Carry Unsigned Quadword.

Compute the product of a 128 bit value  $a * 10 + \text{digit}(\text{cin})$ . Only the low order 128 bits of the extended product are returned.

processor	Latency	Throughput
power8	15-17	1/cycle
power9	3	1/cycle

#### Parameters

<i>a</i>	128-bit vector treated as a unsigned <code>__int128</code> .
<i>cin</i>	values 0-9 in bits 124:127 of a vector.

#### Returns

`__int128` (upper 128-bits of the 256-bit product)  $a * 10 >> 128$ .

**7.7.3.39** `static vui128_t vec_mul10euq( vui128_t a, vui128_t cin ) [inline],[static]`

Vector Multiply by 10 extended Unsigned Quadword.

compute the product of a 128 bit value  $a * 10 + \text{digit}(\text{cin})$ . Only the low order 128 bits of the extended product are returned.

processor	Latency	Throughput
power8	13-15	1/cycle
power9	3	1/cycle

#### Parameters

<i>a</i>	128-bit vector treated as a unsigned <code>__int128</code> .
<i>cin</i>	values 0-9 in bits 124:127 of a vector.

#### Returns

`__int128` (lower 128-bits)  $a * 10$ .

**7.7.3.40** `static vui128_t vec_mul10uq ( vui128_t a )` `[inline]`, `[static]`

Vector Multiply by 10 Unsigned Quadword.

compute the product of a 128 bit value  $a * 10$ . Only the low order 128 bits of the product are returned.

processor	Latency	Throughput
power8	13-15	1/cycle
power9	3	1/cycle

#### Parameters

<i>a</i>	128-bit vector treated as a <code>__int128</code> .
----------	---

#### Returns

`__int128` (lower 128-bits)  $a * 10$ .

**7.7.3.41** `static vui128_t vec_muleud ( vui64_t a, vui64_t b )` `[inline]`, `[static]`

Vector multiply even unsigned doublewords.

Multiple the even 64-bit doublewords of two vector unsigned long values and return the unsigned `__int128` product of the even doublewords.

#### Note

The element numbering changes between big and little-endian. So the compiler and this implementation adjusts the generated code to reflect this.

processor	Latency	Throughput
power8	21-23	1/cycle
power9	8-13	2/cycle

## Parameters

<i>a</i>	128-bit vector unsigned long.
<i>b</i>	128-bit vector unsigned long.

## Returns

vector unsigned \_\_int128 product of the even double words of a and b.

**7.7.3.42** `static vui64_t vec_mulhud ( vui64_t vra, vui64_t vrb )` `[inline], [static]`

Vector Multiply High Unsigned Doubleword.

Multiple the corresponding doubleword elements of two vector unsigned long values and return the high order 64-bits, from each 128-bit product.

processor	Latency	Throughput
power8	28-32	1/cycle
power9	11-16	1/cycle

## Note

This operation can be used to effectively perform a divide by multiplying by the scaled multiplicative inverse (reciprocal).

Warren, Henry S. Jr and *Hacker's Delight*, 2nd Edition, Addison Wesley, 2013. Chapter 10, Integer Division by Constants.

## Parameters

<i>vra</i>	128-bit vector unsigned long.
<i>vrb</i>	128-bit vector unsigned long.

## Returns

vector of the high order 64-bits of the signed product of the doubleword elements from vra and vrb.

**7.7.3.43** `static vui128_t vec_mulluq ( vui128_t a, vui128_t b )` `[inline], [static]`

Vector Multiply Low Unsigned Quadword.

compute the 256 bit product of two 128 bit values a, b. Only the low order 128 bits of the product are returned.

processor	Latency	Throughput
power8	42-48	1/cycle
power9	16-20	2/cycle

## Parameters

<i>a</i>	128-bit vector treated a <code>__int128</code> .
<i>b</i>	128-bit vector treated a <code>__int128</code> .

## Returns

`__int128` (lower 128-bits)  $a * b$ .

**7.7.3.44** `static vui128_t vec_muloud ( vui64_t a, vui64_t b )` `[inline],[static]`

Vector multiply odd unsigned doublewords.

Multiple the odd 64-bit doublewords of two vector unsigned long values and return the unsigned `__int128` product of the odd doublewords.

## Note

The element numbering changes between big and little-endian. So the compiler and this implementation adjusts the generated code to reflect this.

processor	Latency	Throughput
power8	21-23	1/cycle
power9	8-13	2/cycle

## Parameters

<i>a</i>	128-bit vector unsigned long.
<i>b</i>	128-bit vector unsigned long. @ return vector unsigned <code>__int128</code> product of the odd double words of <i>a</i> and <i>b</i> .

**7.7.3.45** `static vui64_t vec_muludm ( vui64_t vra, vui64_t vrb )` `[inline],[static]`

Vector Multiply Unsigned Doubleword Modulo.

Multiple the corresponding doubleword elements of two vector unsigned long values and return the low order 64-bits of the 128-bit product for each element.

## Note

`vec_muludm` can be used for unsigned or signed integers. It is the vector equivalent of Multiply Low Doubleword.

processor	Latency	Throughput
power8	19-28	1/cycle
power9	11-16	1/cycle



## Parameters

<i>vra</i>	128-bit vector unsigned long.
<i>vrb</i>	128-bit vector unsigned long.

## Returns

vector of the low order 64-bits of the unsigned product of the doubleword elements from *vra* and *vrb*.

**7.7.3.46** `static vui128_t vec_muludq ( vui128_t * mulu, vui128_t a, vui128_t b )` `[inline], [static]`

Vector Multiply Unsigned double Quadword.

compute the 256 bit product of two 128 bit values *a*, *b*. The low order 128 bits of the product are returned, while the high order 128-bits are "stored" via the *mulu* pointer.

processor	Latency	Throughput
power8	56-64	1/cycle
power9	33-39	1/cycle

## Parameters

<i>*mulu</i>	pointer to upper 128-bits of the product.
<i>a</i>	128-bit vector treated a <code>__int128</code> .
<i>b</i>	128-bit vector treated a <code>__int128</code> .

## Returns

vector unsigned `__int128` (lower 128-bits) of *a* \* *b*.

**7.7.3.47** `static vui128_t vec_popcntq ( vui128_t vra )` `[inline], [static]`

Population Count vector `__int128`.

Count the number of '1' bits within a vector `__int128` and return the count (0-128) in a vector `__int128`.

processor	Latency	Throughput
power8	15	2/2 cycles
power9	16	2/cycle

## Parameters

<i>vra</i>	a 128-bit vector treated a <code>__int128</code> .
------------	--

**Returns**

a 128-bit vector with bits 121:127 containing the population count.

**7.7.3.48** `static vui128_t vec_revbq ( vui128_t vra )` `[inline],[static]`

byte reverse quadword for a vector `__int128`.

Return the bytes / octets of a 128-bit vector in reverse order.

processor	Latency	Throughput
power8	2-13	2 cycle
power9	3	2/cycle

**Parameters**

<code>vra</code>	a 128-bit vector treated a <code>__int128</code> .
------------------	--

**Returns**

a 128-bit vector with the bytes in reserve order.

**7.7.3.49** `static vb128_t vec_setb_cyq ( vui128_t vcy )` `[inline],[static]`

Vector Set Bool from Quadword Carry.

If the vector quadword carry bit (`vcy.bit[127]`) is '1' then return a vector bool `__int128` that is all '1's. Otherwise return all '0's.

processor	Latency	Throughput
power8	4 - 6	2/2 cycles
power9	3 - 5	2/cycle

Vector quadword carries are normally the result of a *write-Carry* operation. For example; `vec_addcuq()`, `vec_↵addecuq()`, `vec_subcuq()`, `vec_subecuq()`, `vec_addcq()`, `vec_addecq()`.

**Parameters**

<code>vcy</code>	a 128-bit vector generated from a <i>write-Carry</i> operation.
------------------	---

**Returns**

a 128-bit vector bool of all '1's if the carry bit is '1'. Otherwise all '0's.

**7.7.3.50** `static vb128_t vec_setb_ncq ( vui128_t vcy )` `[inline],[static]`

Vector Set Bool from Quadword not Carry.

If the vector quadword carry bit (`vcy.bit[127]`) is '1' then return a vector `bool __int128` that is all '0's. Otherwise return all '1's.

processor	Latency	Throughput
power8	4 - 6	2/2 cycles
power9	3 - 5	2/cycle

Vector quadword carries are normally the result of a *write-Carry* operation. For example; `vec_addcuq()`, `vec_addecuq()`, `vec_subcuq()`, `vec_subecuq()`, `vec_addcq()`, `vec_addeq()`.

#### Parameters

<code>vcy</code>	a 128-bit vector generated from a <i>write-Carry</i> operation.
------------------	---

#### Returns

a 128-bit vector `bool` of all '1's if the carry bit is '0'. Otherwise all '0's.

**7.7.3.51** `static vb128_t vec_setb_sq ( vi128_t vra ) [inline],[static]`

Vector Set Bool from Signed Quadword.

If the quadword's sign bit is '1' then return a vector `bool __int128` that is all '1's. Otherwise return all '0's.

processor	Latency	Throughput
power8	4 - 6	2/cycle
power9	5 - 8	2/cycle

#### Parameters

<code>vra</code>	a 128-bit vector treated a signed <code>__int128</code> .
------------------	---

#### Returns

a 128-bit vector `bool` of all '1's if the sign bit is '1'. Otherwise all '0's.

**7.7.3.52** `static vui128_t vec_sldq ( vui128_t vrw, vui128_t vrx, vui128_t vrb ) [inline],[static]`

Vector Shift Left double Quadword.

Vector Shift Left double Quadword 0-127 bits. Return a vector `__int128` that is the left most 128-bits after shifting left 0-127-bits of the 32-byte double vector (`vrw||vrx`). The shift amount is from bits 121:127 of `vrb`.

processor	Latency	Throughput
power8	10	1 cycle
power9	14	1/cycle

## Parameters

<i>vrw</i>	upper 128-bits of the 256-bit double vector.
<i>vrx</i>	lower 128-bits of the 256-bit double vector.
<i>vrb</i>	Shift amount in bits 121:127.

## Returns

high 128-bits of left shifted double vector.

**7.7.3.53** `static vui128_t vec_slq( vui128_t vra, vui128_t vrb )` `[inline]`, `[static]`

Vector Shift Left Quadword.

Vector Shift Left Quadword 0-127 bits. The shift amount is from bits 121-127 of vrb.

processor	Latency	Throughput
power8	4	1/cycle
power9	6	1/cycle

## Parameters

<i>vra</i>	a 128-bit vector treated as a <code>__int128</code> .
<i>vrb</i>	Shift amount in bits 121:127.

## Returns

Left shifted vector.

**7.7.3.54** `static vui128_t vec_slq4( vui128_t vra )` `[inline]`, `[static]`

**Deprecated** Vector Shift Left 4-bits Quadword. Replaced by `vec_slqi` with `shb` param = 4.

Vector Shift Left Quadword 0-127 bits. The shift amount is from bits 121-127 of vrb.

## Parameters

<i>vra</i>	a 128-bit vector treated a <code>__int128</code> .
------------	--

## Returns

Left shifted vector.

**7.7.3.55** `static vui128_t vec_slq5( vui128_t vra )` `[inline]`, `[static]`

**Deprecated** Vector Shift Left 5-bits Quadword. Replaced by `vec_slqi` with `shb` param = 5.

Vector Shift Left Quadword 0-127 bits. The shift amount is from bits 121-127 of vrb.

```
@param vra a 128-bit vector treated a __int128.
@return Left shifted vector.
```

**7.7.3.56** `static vui128_t vec_slqi ( vui128_t vra, const unsigned int shb )` `[inline], [static]`

Vector Shift left Quadword Immediate.

Shift left Quadword 0-127 bits. The shift amount is a const unsigned int in the range 0-127. A shift count of 0 returns the original value of vra. Shift counts greater then 127 bits return zero.

processor	Latency	Throughput
power8	2-13	2 cycle
power9	3-15	2/cycle

#### Parameters

<i>vra</i>	a 128-bit vector treated as a __int128.
<i>shb</i>	Shift amount in the range 0-127.

#### Returns

128-bit vector shifted left shb bits.

**7.7.3.57** `static vui128_t vec_srqi ( vui128_t vra, vui128_t vrb )` `[inline], [static]`

Vector Shift right Quadword.

Vector Shift Right Quadword 0-127 bits. The shift amount is from bits 121-127 of vrb.

processor	Latency	Throughput
power8	4	1/cycle
power9	6	1/cycle

#### Parameters

<i>vra</i>	a 128-bit vector treated as a __int128.
<i>vrb</i>	Shift amount in bits 121:127.

#### Returns

Right shifted vector.

**7.7.3.58** `static vui128_t vec_srq4( vui128_t vra ) [inline],[static]`

**Deprecated** Vector Shift right 4-bits Quadword. Replaced by `vec_srq` with `shb` param = 4.

Vector Shift Right Quadword 0-127 bits. The shift amount is from bits 121-127 of `vrb`.

#### Parameters

<code>vra</code>	a 128-bit vector treated as a <code>__int128</code> .
------------------	---

#### Returns

Right shifted vector.

**7.7.3.59** `static vui128_t vec_srq5( vui128_t vra ) [inline],[static]`

**Deprecated** Vector Shift right 5-bits Quadword. Replaced by `vec_srq` with `shb` param = 5.

Vector Shift Right Quadword 0-127 bits. The shift amount is from bits 121-127 of `vrb`.

#### Parameters

<code>vra</code>	a 128-bit vector treated as a <code>__int128</code> .
------------------	---

#### Returns

Right shifted vector.

**7.7.3.60** `static vui128_t vec_srq( vui128_t vra, const unsigned int shb ) [inline],[static]`

Vector Shift right Quadword Immediate.

Shift right Quadword 0-127 bits. The shift amount is a const unsigned int in the range 0-127. A shift count of 0 returns the original value of `vra`. Shift counts greater than 127 bits return zero.

processor	Latency	Throughput
power8	2-13	2 cycle
power9	3-15	2/cycle

#### Parameters

<code>vra</code>	a 128-bit vector treated as a <code>__int128</code> .
<code>shb</code>	Shift amount in the range 0-127.

**Returns**

128-bit vector shifted right shb bits.

**7.7.3.61** `static vui128_t vec_subcuq ( vui128_t vra, vui128_t vrb ) [inline],[static]`

Vector Subtract and Write Carry Unsigned Quadword.

Generate the carry-out of the sum ( $vra + \text{NOT}(vrb) + 1$ ).

processor	Latency	Throughput
power8	4	2/2 cycles
power9	3	2/cycle

**Parameters**

<i>vra</i>	128-bit vector treated a unsigned __int128.
<i>vrb</i>	128-bit vector treated a unsigned __int128.

**Returns**

\_\_int128 carry from the unsigned difference  $vra - vrb$ .

**7.7.3.62** `static vui128_t vec_subecuq ( vui128_t vra, vui128_t vrb, vui128_t vrc ) [inline],[static]`

Vector Subtract Extended and Write Carry Unsigned Quadword.

Generate the carry-out of the sum ( $vra + \text{NOT}(vrb) + vrc.\text{bit}[127]$ ).

processor	Latency	Throughput
power8	4	2/2 cycles
power9	3	2/cycle

**Parameters**

<i>vra</i>	128-bit vector treated a unsigned __int128.
<i>vrb</i>	128-bit vector treated a unsigned __int128.
<i>vrc</i>	128-bit vector carry-in from bit 127.

**Returns**

\_\_int128 carry from the extended \_\_int128 difference.

**7.7.3.63** `static vui128_t vec_subeuqm ( vui128_t vra, vui128_t vrb, vui128_t vrc ) [inline],[static]`

Vector Subtract Extended Unsigned Quadword Modulo.

Subtract two vector `__int128` values and return result modulo 128-bits.



processor	Latency	Throughput
power8	4	2/2 cycles
power9	3	2/cycle

## Parameters

<i>vra</i>	128-bit vector treated an unsigned __int128.
<i>vrb</i>	128-bit vector treated an unsigned __int128.
<i>vrc</i>	128-bit vector carry-in from bit 127.

## Returns

\_\_int128 unsigned difference of vra minus vrb.

**7.7.3.64** `static vui128_t vec_subuqm ( vui128_t vra, vui128_t vrb )` `[inline], [static]`

Vector subtract Unsigned Quadword Modulo.

Subtract two vector \_\_int128 values and return result modulo 128-bits.

processor	Latency	Throughput
power8	4	2/2 cycles
power9	3	2/cycle

## Parameters

<i>vra</i>	128-bit vector treated an unsigned __int128.
<i>vrb</i>	128-bit vector treated an unsigned __int128.

## Returns

\_\_int128 unsigned difference of vra minus vrb.

**7.7.3.65** `static vui128_t vec_vmuleud ( vui64_t a, vui64_t b )` `[inline], [static]`

Vector multiply even unsigned doublewords.

Multiple the even 64-bit doublewords of two vector unsigned long values and return the unsigned \_\_int128 product of the even doublewords.

## Note

This function implements the operation of a Vector Multiply Even Doubleword instruction, if the PowerISA included such an instruction. This implementation is NOT endian sensitive and the function is stable across BE/LE implementations.

processor	Latency	Throughput
power8	21-23	1/cycle
power9	8-13	2/cycle

#### Parameters

<i>a</i>	128-bit vector unsigned long.
<i>b</i>	128-bit vector unsigned long.

#### Returns

vector unsigned \_\_int128 product of the even double words of a and b.

**7.7.3.66** `static vui128_t vec_vmuloud ( vui64_t a, vui64_t b )` `[inline], [static]`

Vector multiply odd unsigned doublewords.

Multiple the odd 64-bit doublewords of two vector unsigned long values and return the unsigned \_\_int128 product of the odd doublewords.

#### Note

This function implements the operation of a Vector Multiply Odd Doubleword instruction, if the PowerISA included such an instruction. This implementation is NOT endian sensitive and the function is stable across BE/LE implementations.

processor	Latency	Throughput
power8	21-23	1/cycle
power9	8-13	2/cycle

#### Parameters

<i>a</i>	128-bit vector unsigned long.
<i>b</i>	128-bit vector unsigned long. @ return vector unsigned __int128 product of the odd double words of a and b.

## 7.8 src/vec\_int16\_ppc.h File Reference

Header package containing a collection of 128-bit SIMD operations over 16-bit integer elements.

```
#include <vec_char_ppc.h>
```

#### Functions

- static `vui16_t vec_absduh (vui16_t vra, vui16_t vrb)`

- Vector Absolute Difference Unsigned halfword.*
- static `vui16_t vec_clzh` (`vui16_t` vra)
- Count Leading Zeros for a vector unsigned short (halfword) elements.*
- static `vui16_t vec_mrgahh` (`vui32_t` vra, `vui32_t` vrb)
- Vector Merge Algebraic High Halfword operation.*
- static `vui16_t vec_mrgalh` (`vui32_t` vra, `vui32_t` vrb)
- Vector Merge Algebraic Low Halfword operation.*
- static `vui16_t vec_mrgeh` (`vui16_t` vra, `vui16_t` vrb)
- Vector Merge Even Halfwords operation.*
- static `vui16_t vec_mrgoh` (`vui16_t` vra, `vui16_t` vrb)
- Vector Merge Odd Halfwords operation.*
- static `vi16_t vec_mulhsh` (`vi16_t` vra, `vi16_t` vrb)
- Vector Multiply High Signed halfword.*
- static `vui16_t vec_mulhuh` (`vui16_t` vra, `vui16_t` vrb)
- Vector Multiply High Unsigned halfword.*
- static `vui16_t vec_muluhm` (`vui16_t` vra, `vui16_t` vrb)
- Vector Multiply Unsigned halfword Modulo.*
- static `vui16_t vec_popcnth` (`vui16_t` vra)
- Vector Population Count halfword.*
- static `vui16_t vec_revbh` (`vui16_t` vra)
- byte reverse each halfword of a vector unsigned short.*
- static `vui16_t vec_slhi` (`vui16_t` vra, const unsigned int shb)
- Vector Shift left Halfword Immediate.*
- static `vui16_t vec_srhi` (`vui16_t` vra, const unsigned int shb)
- Vector Shift Right Halfword Immediate.*
- static `vi16_t vec_srahi` (`vi16_t` vra, const unsigned int shb)
- Vector Shift Right Algebraic Halfword Immediate.*
- static `vui16_t vec_vmrgeh` (`vui16_t` vra, `vui16_t` vrb)
- Vector Merge Even Halfwords.*
- static `vui16_t vec_vmrgoh` (`vui16_t` vra, `vui16_t` vrb)
- Vector Merge Odd Halfwords.*

### 7.8.1 Detailed Description

Header package containing a collection of 128-bit SIMD operations over 16-bit integer elements.

Most of these operations are implemented in a single instruction on newer (POWER6/POWER7POWER8/POWER9) processors. This header serves to fill in functional gaps for older (POWER7, POWER8) processors and provides a in-line assembler implementation for older compilers that do not provide the build-ins.

Most vector short (16-bit integer halfword) operations are implemented with PowerISA VMX instructions either defined by the original VMX (AKA AltiVec) or added to later versions of the PowerISA. PowerISA 2.07B (POWER8) added several useful halfword operations (count leading zeros, population count) not included in the original VMX. PowerISA 3.0B (POWER9) adds several more (absolute difference, compare not equal, count trailing zeros, extend sign, extract/insert, and reverse bytes). Most of these intrinsic (compiler built-ins) operations are defined in `<altivec.h>` and described in the compiler documentation.

**Note**

The compiler disables associated `<altivec.h>` built-ins if the **mcpu** target does not enable the specific instruction. For example if you compile with **-mcpu=power7**, `vec_vclz` and `vec_vclzh` will not be defined. Another example if you compile with **-mcpu=power8**, `vec_revb` will not be defined. But `vec_vclzh` and `vec_revbh` is always defined in this header. This header provides the appropriate substitutions, will generate the minimum code, appropriate for the target, and produce correct results.

Most ppc64le compilers will default to **-mcpu=power8** if not specified.

This header covers operations that are either:

- Implemented in hardware instructions for later processors and useful to programmers, on slightly older processors, even if the equivalent function requires more instructions. Examples include Count Leading Zeros, Population Count and Byte Reverse.
- Defined in the OpenPOWER ABI but *not* yet defined in `<altivec.h>` provided by available compilers in common use. Examples include Count Leading Zeros, Population Count and Byte Reverse.
- Commonly used operations, not covered by the ABI or `<altivec.h>`, and require multiple instructions or are not obvious. Examples include the multiply high and shift immediate operations.

### 7.8.2 Endian problems with halfword operations

It would be useful to provide a vector multiply high halfword (return the high order 16-bits of the 32-bit product) operation. This can be used for multiplicative inverse (effectively integer divide) operations. Neither integer multiply high nor divide are available as vector instructions. However the multiply high halfword operation can be composed from the existing multiply even/odd halfword operations followed by the vector merge even halfword operation. Similarly a multiply low (modulo) halfword operation can be composed from the existing multiply even/odd halfword operations followed by the vector merge odd halfword operation.

As a prerequisite we need to provide the merge even/odd halfword operations. While PowerISA has added these operations for word and doubleword, instructions are not defined for byte and halfword. Fortunately vector merge operations are just a special case of vector permute. So the [vec\\_vmrgeh\(\)](#) and [vec\\_vmrgeh\(\)](#) implementation can use `vec_perm` and appropriate selection vectors to provide these merge operations.

But this is complicated by *little-endian* (LE) support as specified in the OpenPOWER ABI and as implemented in the compilers. Little-endian changes the effective vector element numbering and the location of even and odd elements. This means that the vector built-ins provided by `altivec.h` may not generate the instructions you would expect.

**See also**

[General Endian Issues](#)  
[Endian problems with word operations](#)

The OpenPOWER ABI provides a helpful table of [Endian Sensitive Operations](#). For for `vec_mule` (`vmuleuh`, `vmulesh`):

Replace with `vmulouh` and so on, for LE.

For for `vec_mulo` (`vmulouh`, `vmulosh`):

Replace with `vmuleuh` and so on, for LE.

Also for `vec_perm` (`vperm`) it specifies:

For LE, Swap input arguments and complement the selection vector.

The above is just a sampling of a larger list of Endian Sensitive Operations.

The obvious coding for Vector Multiply High Halfword would be:

```
vui16_t
test_mulhw (vui16_t vra, vui16_t vrb)
{
    return vec_mergee ((vui16_t)vec_mule (vra, vrb),
                       (vui16_t)vec_mulo (vra, vrb));
}
```

A couple problems with this:

- vec\_mergee is only defined for vector int/long and float/double (word/doubleword) types.
- vec\_mergee is endian sensitive and would produce the wrong results in LE mode.
- vec\_mule/vec\_mulo are endian sensitive and produce the wrong results in LE mode.

The first step is to implement Vector Merge Even Halfword operation:

```
static inline vui16_t
vec_vmrgheh (vui16_t vra, vui16_t vrb)
{
    #if __BYTE_ORDER__ == __ORDER_LITTLE_ENDIAN__
        vui16_t permute =
            { 0x0302, 0x1312, 0x0706, 0x1716, 0x0B0A, 0x1B1A, 0x0F0E, 0x1F1E };
        return vec_perm (vrb, vra, (vui8_t)permute);
    #else
        vui16_t permute =
            { 0x0001, 0x1011, 0x0405, 0x1415, 0x0809, 0x1819, 0x0C0D, 0x1C1D };
        return vec_perm (vra, vrb, (vui8_t)permute);
    #endif
}
```

For big-endian we have a straight forward vec\_perm with a permute select vector interleaving even halfwords from vectors vra and vrb.

For little-endian we need to nullify the LE transform applied by the compiler. So the select vector looks like it interleaves odd halfwords from vectors vrb and vra. It also reverses byte numbering within halfwords. The compiler transforms this back into the operation we wanted in the first place. The result is *not* endian sensitive and is stable across BE/LE implementations. Similarly for the Vector Merge Odd Halfword operation.

As good OpenPOWER ABI citizens we should also provide endian sensitive operations [vec\\_mrgeh\(\)](#) [vec\\_mrgoh\(\)](#). For example:

```
static inline vui16_t
vec_mrgeh (vui16_t vra, vui16_t vrb)
{
    #if __BYTE_ORDER__ == __ORDER_LITTLE_ENDIAN__
        return vec_vmrgoh ((vui16_t) vrb, (vui16_t) vra);
    #else
        return vec_vmrgheh ((vui16_t) vra, (vui16_t) vrb);
    #endif
}
```

**Note**

This is essentially what the compiler would do for `vec_mergee`.

Also to follow that pattern established for [vec\\_int32\\_ppc.h](#) we should provide implementations for Vector Merge Algebraic High/Low Halfword. For example:

```
static inline vuil6_t
vec_mrgahh (vui32_t vra, vui32_t vrb)
{
    return vec_vmrgch ((vui16_t) vra, (vui16_t) vrb);
}
```

This is simpler as we can use the endian invariant [vec\\_vmrgch\(\)](#) operation. Similarly for Vector Merge Algebraic Low Halfword using [vec\\_vmrgoh\(\)](#).

**Note**

The inputs are defined as 32-bit to match the results from multiply even/odd halfword.

Now we have all the parts we need to implement multiply high/low halfword. For example Multiply High Unsigned Halfword:

```
static inline vuil6_t
vec_mulhuh (vui16_t vra, vui16_t vrb)
{
    #if __BYTE_ORDER__ == __ORDER_LITTLE_ENDIAN__
        return vec_mrgahh (vec_mulo (vra, vrb), vec_mule (vra, vrb));
    #else
        return vec_mrgahh (vec_mule (vra, vrb), vec_mulo (vra, vrb));
    #endif
}
```

Similarly for Multiply High Signed Halfword.

**Note**

For LE we need to nullify the compiler transform by reversing of the order of `vec_mulo/vec_mule`. This is required to get the algebraically correct (multiply high) result.

Finally we can implement the Multiply Low Halfword which by PowerISA conventions is called Multiply Unsigned Halfword Modulo:

```
static inline vuil6_t
vec_muluhm (vui16_t vra, vui16_t vrb)
{
    #if __BYTE_ORDER__ == __ORDER_LITTLE_ENDIAN__
        return vec_mrgalh (vec_mulo (vra, vrb), vec_mule (vra, vrb));
    #else
        return vec_mrgalh (vec_mule (vra, vrb), vec_mulo (vra, vrb));
    #endif
}
```

**Note**

We use the endian stable [vec\\_mrgalh\(\)](#) for multiply low. Again for LE we have to nullify the compiler transform by reversing of the order of `vec_mulo/vec_mule`. This is required to get the algebraically correct (multiply high) result.

[vec\\_muluhm\(\)](#) works for signed and unsigned multiply low (modulo).

### 7.8.2.1 Multiply High Unsigned Halfword Example

So what does the compiler generate after unwinding three levels of inline functions. For this test case:

```
vuil6_t
__test_mulhuh (vuil6_t a, vuil6_t b)
{
    return vec_mulhuh (a, b);
}
```

The GCC 8 compiler targeting powerpc64le and -mcpu=power8 generates:

```
addis    r9,r2,.rodata.cst16@ha
vmulouh  v1,v2,v3
vmuleuh  v2,v2,v3
addi     r9,r9,.rodata.cst16@l
lvx      v0,0,r9
xxlnor   vs32,vs32,vs32
vperm    v2,v2,v1,v0
```

The addis, addi, lvx instruction sequence loads the permute selection constant vector. The xxlnor instruction complements the selection vector for LE. These instructions are only needed once per function and can be hoisted out of loops and shared across instances of `vec_mulhuh()`. Which might look like this:

```
addis    r9,r2,.rodata.cst16@ha
addi     r9,r9,.rodata.cst16@l
lvx      v0,0,r9
xxlnor   vs32,vs32,vs32
...
Loop:
vmulouh  v1,v2,v3
vmuleuh  v2,v2,v3
vperm    v2,v2,v1,v0
...
```

The vmulouh, vmuleuh, vperm instruction sequence is the core of the function. They multiply the elements and selects/merges the high order 16-bits of each product into the result vector.

## 7.8.3 Examples, Divide by integer constant

Suppose we have a requirement to convert an array of 16-bit unsigned short values to decimal. The classic *itoa* implementation performs a sequence of divide / modulo by 10 operations that produce one (decimal) value per iteration, until the divide returns 0.

For this example we want to vectorize the operation and the PowerISA (and most other platforms) does not provide a vector integer divide instruction. But we do have vector integer multiply. As we will see the multiply high defined above is very handy for applying the multiplicative inverse. Also, the conversion divide is a constant value applied across the vector which simplifies the coding.

Here we can use the multiplicative inverse which is a scaled fixed point fraction calculated from the original divisor. This works nicely if the fixed radix point is just before the 16-bit fraction and we have a multiply high (`vec_mulhuh()`) operation. Multiplying a 16-bit unsigned integer by a 16-bit unsigned fraction generates a 32-bit product with 16-bits above (integer) and below (fraction) the radix point. The high 16-bits of the product is a good approximation of the integer quotient.

It turns out that generating the multiplicative inverse can be tricky. To produce correct results over the full range, requires possible pre-scaling and post-shifting, and sometimes a corrective addition. Fortunately, the mathematics are well understood and are commonly used in optimizing compilers. Even better, Henry Warren's book has a whole chapter on this topic.

### See also

"Hacker's Delight, 2nd Edition," Henry S. Warren, Jr, Addison Wesley, 2013. Chapter 10, Integer Division by Constants.

### 7.8.3.1 Divide by constant 10 examples

In the chapter above;

Figure 10-2 Computing the magic number for unsigned division.

provides a sample C function for generating the magic number (actually a struct containing; the magic multiplicative inverse, "add" indicator, and the shift amount). For the 16-bit unsigned divisor 10, this is { 52429, 0, 3 }:

- the multiplier is 52429.
- no corrective add of the dividend is required.
- the final shift is 3-bits right.

Which could look like this:

```
vuil6_t
__test_div10 (vuil6_t n)
{
    vuil6_t q;
    // M= 52429, a=0, s=3
    vuil6_t magic = vec_splats ((unsigned short) 52429);
    const int s = 3;

    q = vec_mulhuh (magic, n);
    return vec_srhi (q, s);
}
```

But we also need the modulo to extract each digit. The simplest and oldest technique is to multiply the quotient by the divisor (constant 10) and subtract that from the original dividend. Here we can use the `vec_muluhm()` operation we defined above. Which could look like this:

```
vuil6_t
__test_mod10 (vuil6_t n)
{
    vuil6_t q;
    // M= 52429, a=0, s=3
    vuil6_t magic = vec_splats ((unsigned short) 52429);
    vuil6_t c_10 = vec_splats ((unsigned short) 10);
    const int s = 3;
    vuil6_t tmp, rem, q_10;

    q = vec_mulhuh (magic, n);
    q_10 = vec_srhi (q, s);
    tmp = vec_muluhm (q_10, c_10);
    rem = vec_sub (n, tmp);
    return rem;
}
```

#### Note

`vec_sub()` and `vec_splats()` are an existing `altivec.h` generic built-ins.



### 7.8.3.2 Divide by constant 10000 example

As we mentioned above, some divisors require an add before the shift as a correction. For the 16-bit unsigned divisor 10000 this is { 41839, 1, 14 }:

- the multiplier is 41839.
- corrective add of the dividend is required.
- the final shift is 14-bits right.

In this case the perfect multiplier is too large ( $\geq 2^{16}$ ). So the magic multiplier is reduced by  $2^{16}$  and to correct for this we need to add the dividend to the product. This add may generate a carry that must be included in the shift. Here `vec_avg` handles the 17-bit sum internally before shifting right 1. But `vec_avg` adds an extra +1 (for rounding) that we don't want. So we use (n-1) for the product correction then complete the operation with shift right (s-1). Which could look like this:

```
vuil6_t
__test_div10000 (vuil6_t n)
{
    vuil6_t result, q;
    // M= 41839, a=1, s=14
    vuil6_t magic = vec_splats ((unsigned short) 41839);
    const int s = 14;
    vuil6_t tmp, rem;

    q = vec_mulhuh (magic, n);
    {
        const vuil6_t vec_ones = vec_splat_u16 ( 1 );
        vuil6_t n_1 = vec_sub (n, vec_ones);
        // avg = (q + (n-1) + 1) >> 1
        q = vec_avg (q, n_1);
        result = vec_srhi (q, (s - 1));
    }
    return result;
}
```

#### Note

`vec_avg()`, `vec_sub()`, `vec_splats()` and `vec_splat_u16()` are existing `altivec.h` generic built-ins.

The modulo computation remains the same as [Divide by constant 10 examples](#).

## 7.8.4 Performance data.

We can use the example above (see [Multiply High Unsigned Halfword Example](#)) to illustrate the performance metrics `pveclib` provides. For `vec_mulhuh()` the core operation is the sequence `vmulouh/vmuleuh/vperm`. This represents the best case latency, when it is used multiple times in a single larger function.

The compiler notes that `vmulouh/vmuleuh` are independent instructions that can execute concurrently (in separate vector pipelines). The compiler schedules them to issue in same cycle. The latency for `vmulouh/vmuleuh` is listed as 7 cycle and the throughput of 2 per cycle (there are 2 vector pipes for multiply). As we assume this function will use both vector pipelines, the throughput for this function is reduced to 1 per cycle.

We still need to select/merge the results. The `vperm` instruction is dependent on the execution of both `vmulouh/vmuleuh` and load of the select vector complete. For this case we assume that the load of the permute select vector has already executed. The processor can not issue the `vperm` until both `vmulouh/vmuleuh` instructions execute. The latency for `vperm` is 2 cycles (3 on POWER9). So the best case latency for this operation is is (7 + 2 = 9) cycles (10 on POWER9).

Looking at the first or only execution of `vec_mulhuh()` in a function defines the worse case latency. Here we have to include the permute select vector load and (for LE) the select vector complement. However this case provides additional multiple pipe parallelism that needs to be accounted for in the latencies.

The compiler notes that `addis/vmulouh/vmuleuh` are independent instructions that can execute concurrently in separate pipelines. So the compiler schedules them to issue in same cycle. The latency for `vmulouh/vmuleuh` is 7 cycles while the `addis` latency is only 2 cycles. The dependent `addi` instruction can issue in the 3rd cycle, while `vmulouh/vmuleuh` are still executing. The `addi` also has a 2 cycle latency, so the dependent `lvx` can issue in the 5th cycle, while `vmulouh/vmuleuh` are still executing. The `lvx` has a latency of 5 cycles and will not complete execution until 2 cycles after `vmulouh/vmuleuh`. The dependent `xxlnor` is waiting of the load (`lvx`) and has a latency of 2 cycles.

So there are two independent instruction sequences; `vmulouh/vmuleuh` and `addis/addi/lvx/xxlnor`. Both must complete execution before the `vperm` can issue and complete the operation. The later sequence has the longer ( $2+2+5+2=11$ ) latency and dominates the timing. So the worst latency for the full sequence is ( $2+2+5+2+2 = 13$ ) cycles (14 on POWER9).

processor	Latency	Throughput
power8	9-13	1/cycle
power9	10-14	1/cycle

#### 7.8.4.1 More information.

High level performance estimates are provided as an aid to function selection when evaluating algorithms. For background on how *Latency* and *Throughput* are derived see: [Performance data](#).

## 7.8.5 Function Documentation

**7.8.5.1** `static vui16_t vec_absduh ( vui16_t vra, vui16_t vrb ) [inline],[static]`

Vector Absolute Difference Unsigned halfword.

Compute the absolute difference for each halfword. For each unsigned halfword, subtract `VRB[i]` from `VRA[i]` and return the absolute value of the difference.

processor	Latency	Throughput
power8	4	1/cycle
power9	3	2/cycle

#### Parameters

<code>vra</code>	vector of 8 x unsigned halfword
<code>vrb</code>	vector of 8 x unsigned halfword

#### Returns

vector of the absolute differences.

### 7.8.5.2 static vui16\_t vec\_clzh ( vui16\_t vra ) [inline],[static]

Count Leading Zeros for a vector unsigned short (halfword) elements.

Count the number of leading '0' bits (0-16) within each halfword element of a 128-bit vector.

For POWER8 (PowerISA 2.07B) or later use the Vector Count Leading Zeros Halfword instruction **vcclzh**. Otherwise use sequence of pre 2.07 VMX instructions.

#### Note

SIMDized count leading zeros inspired by: Warren, Henry S. Jr and *Hacker's Delight*, 2nd Edition, Addison Wesley, 2013. Chapter 5 Counting Bits, Figure 5-12.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

#### Parameters

<i>vra</i>	128-bit vector treated as 8 x 16-bit integer (halfword) elements.
------------	---

#### Returns

128-bit vector with the Leading Zeros count for each halfword element.

### 7.8.5.3 static vui16\_t vec\_mrgahh ( vui32\_t vra, vui32\_t vrb ) [inline],[static]

Vector Merge Algebraic High Halfword operation.

Merge only the high halfwords from 8 x Algebraic words across vectors vra and vrb. This is effectively the Vector Merge Even Halfword operation that is not modified for endian.

For example merge the high 16-bits from each of 8 x 32-bit products as generated by vec\_muleuh/vec\_mulouh. This result is effectively a vector multiply high unsigned halfword.

processor	Latency	Throughput
power8	2-13	2/cycle
power9	3-14	2/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned int.
<i>vrh</i>	128-bit vector unsigned int.

#### Returns

A vector merge from only the high halfwords of the 8 x Algebraic words across vra and vrb.

**7.8.5.4** `static vui16_t vec_mrgalh ( vui32_t vra, vui32_t vrb ) [inline],[static]`

Vector Merge Algebraic Low Halfword operation.

Merge only the low halfwords from 8 x Algebraic words across vectors *vra* and *vr**b*. This is effectively the Vector Merge Odd Halfword operation that is not modified for endian.

For example merge the low 16-bits from each of 8 x 32-bit products as generated by `vec_muleuh/vec_mulouh`. This result is effectively a vector multiply low unsigned halfword.

processor	Latency	Throughput
power8	2-13	2/cycle
power9	3-14	2/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned int.
<i>vr</i> <i>b</i>	128-bit vector unsigned int.

#### Returns

A vector merge from only the high halfwords of the 8 x Algebraic words across *vra* and *vr**b*.

**7.8.5.5** `static vui16_t vec_mrgeh ( vui16_t vra, vui16_t vrb ) [inline],[static]`

Vector Merge Even Halfwords operation.

Merge the even halfword elements from the concatenation of 2 x vectors (*vra* and *vr**b*).

#### Note

The element numbering changes between big and little-endian. So the compiler and this implementation adjusts the generated code to reflect this.

processor	Latency	Throughput
power8	2-13	2/cycle
power9	3-14	2/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned short.
<i>vr</i> <i>b</i>	128-bit vector unsigned short.

#### Returns

A vector merge from only the even halfwords of *vra* and *vr**b*.

7.8.5.6 `static vui16_t vec_mrgoh ( vui16_t vra, vui16_t vrb ) [inline],[static]`

Vector Merge Odd Halfwords operation.

Merge the odd halfword elements from the concatenation of 2 x vectors (vra and vrb).

#### Note

The element numbering changes between big and little-endian. So the compiler and this implementation adjusts the generated code to reflect this.

processor	Latency	Throughput
power8	2-13	2/cycle
power9	3-14	2/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned short.
<i>vrb</i>	128-bit vector unsigned short.

#### Returns

A vector merge from only the odd halfwords of vra and vrb.

7.8.5.7 `static vi16_t vec_mulhsh ( vi16_t vra, vi16_t vrb ) [inline],[static]`

Vector Multiply High Signed halfword.

Multiple the corresponding halfword elements of two vector signed short values and return the high order 16-bits, for each 32-bit product element.

processor	Latency	Throughput
power8	9-13	1/cycle
power9	10-14	1/cycle

#### Parameters

<i>vra</i>	128-bit vector signed short.
<i>vrb</i>	128-bit vector signed short.

#### Returns

vector of the high order 16-bits of the product of the halfword elements from vra and vrb.

7.8.5.8 `static vui16_t vec_mulhuh ( vui16_t vra, vui16_t vrb ) [inline],[static]`

Vector Multiply High Unsigned halfword.

Multiply the corresponding halfword elements of two vector unsigned short values and return the high order 16-bits, for each 32-bit product element.

processor	Latency	Throughput
power8	9-13	1/cycle
power9	10-14	1/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned short.
<i>vrh</i>	128-bit vector unsigned short.

#### Returns

vector of the high order 16-bits of the product of the halfword elements from *vra* and *vrh*.

**7.8.5.9** `static vui16_t vec_muluhm ( vui16_t vra, vui16_t vrh ) [inline],[static]`

Vector Multiply Unsigned halfword Modulo.

Multiply the corresponding halfword elements of two vector unsigned short values and return the low order 16-bits of the 32-bit product for each element.

#### Note

`vec_muluhm` can be used for unsigned or signed short integers. It is the vector equivalent of Multiply Low Halfword.

processor	Latency	Throughput
power8	9-13	1/cycle
power9	10-14	1/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned short.
<i>vrh</i>	128-bit vector unsigned short.

#### Returns

vector of the low order 16-bits of the unsigned product of the halfword elements from *vra* and *vrh*.

**7.8.5.10** `static vui16_t vec_popcnth ( vui16_t vra ) [inline],[static]`

Vector Population Count halfword.

Count the number of '1' bits (0-16) within each byte element of a 128-bit vector.

For POWER8 (PowerISA 2.07B) or later use the Vector Population Count Halfword instruction. Otherwise use simple Vector (VMX) instructions to count bits in bytes in parallel.

**Note**

SIMDized population count inspired by: Warren, Henry S. Jr and *Hacker's Delight*, 2nd Edition, Addison Wesley, 2013. Chapter 5 Counting Bits, Figure 5-2.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

**Parameters**

<i>vra</i>	128-bit vector treated as 8 x 16-bit integers (halfword) elements.
------------	--

**Returns**

128-bit vector with the population count for each halfword element.

#### 7.8.5.11 `static vui16_t vec_revbh ( vui16_t vra ) [inline],[static]`

byte reverse each halfword of a vector unsigned short.

For each halfword of the input vector, reverse the order of bytes / octets within the halfword.

processor	Latency	Throughput
power8	2-11	2/cycle
power9	3	2/cycle

**Parameters**

<i>vra</i>	a 128-bit vector unsigned short.
------------	----------------------------------

**Returns**

a 128-bit vector with the bytes of each halfword reversed.

#### 7.8.5.12 `static vui16_t vec_slhi ( vui16_t vra, const unsigned int shb ) [inline],[static]`

Vector Shift left Halfword Immediate.

Shift left each halfword element [0-7], 0-15 bits, as specified by an immediate value. The shift amount is a const unsigned int in the range 0-15. A shift count of 0 returns the original value of *vra*. Shift counts greater than 15 bits return zero.

processor	Latency	Throughput
power8	4-11	2/cycle
power9	5-11	2/cycle

## Parameters

<i>vra</i>	a 128-bit vector treated as a vector unsigned short.
<i>shb</i>	Shift amount in the range 0-15.

## Returns

128-bit vector unsigned short, shifted left *shb* bits.

**7.8.5.13** `static vi16_t vec_srahi ( vi16_t vra, const unsigned int shb )` `[inline]`, `[static]`

Vector Shift Right Algebraic Halfword Immediate.

Shift right algebraic each halfword element [0-7], 0-15 bits, as specified by an immediate value. The shift amount is a const unsigned int in the range 0-15. A shift count of 0 returns the original value of *vra*. Shift counts greater than 7 bits return the sign bit propagated to each bit of each element.

processor	Latency	Throughput
power8	4-11	2/cycle
power9	5-11	2/cycle

## Parameters

<i>vra</i>	a 128-bit vector treated as a vector signed char.
<i>shb</i>	Shift amount in the range 0-7.

## Returns

128-bit vector signed short, shifted right *shb* bits.

**7.8.5.14** `static vui16_t vec_srhi ( vui16_t vra, const unsigned int shb )` `[inline]`, `[static]`

Vector Shift Right Halfword Immediate.

Shift right each halfword element [0-7], 0-15 bits, as specified by an immediate value. The shift amount is a const unsigned int in the range 0-15. A shift count of 0 returns the original value of *vra*. Shift counts greater than 15 bits return zero.

processor	Latency	Throughput
power8	4-11	2/cycle
power9	5-11	2/cycle

## Parameters

<i>vra</i>	a 128-bit vector treated as a vector unsigned short.
<i>shb</i>	Shift amount in the range 0-15.



**Returns**

128-bit vector unsigned short, shifted right shb bits.

**7.8.5.15** `static vui16_t vec_vmrgeh ( vui16_t vra, vui16_t vrb )` `[inline],[static]`

Vector Merge Even Halfwords.

Merge the even halfword elements from the concatenation of 2 x vectors (vra and vrb).

**Note**

This function implements the operation of a Vector Merge Even Halfword instruction, if the PowerISA included such an instruction. This implementation is NOT endian sensitive and the function is stable across BE/LE implementations. Using big-endian element numbering:

- res[0] = vra[0];
- res[1] = vrb[0];
- res[2] = vra[2];
- res[3] = vrb[2];
- res[4] = vra[4];
- res[5] = vrb[4];
- res[6] = vra[6];
- res[7] = vrb[6];

processor	Latency	Throughput
power8	2-13	2/cycle
power9	3-14	2/cycle

**Parameters**

<i>vra</i>	128-bit vector unsigned short.
<i>vrb</i>	128-bit vector unsigned short.

**Returns**

A vector merge from only the even halfwords of vra and vrb.

**7.8.5.16** `static vui16_t vec_vmrgeh ( vui16_t vra, vui16_t vrb )` `[inline],[static]`

Vector Merge Odd Halfwords.

Merge the odd halfword elements from the concatenation of 2 x vectors (vra and vrb).

**Note**

This function implements the operation of a Vector Merge Odd Halfword instruction, if the PowerISA included such an instruction. This implementation is NOT endian sensitive and the function is stable across BE/LE implementations. Using big-endian element numbering:

- `res[0] = vra[1];`
- `res[1] = vrb[1];`
- `res[2] = vra[3];`
- `res[3] = vrb[3];`
- `res[4] = vra[5];`
- `res[5] = vrb[5];`
- `res[6] = vra[7];`
- `res[7] = vrb[7];`

processor	Latency	Throughput
power8	2-13	2/cycle
power9	3-14	2/cycle

**Parameters**

<i>vra</i>	128-bit vector unsigned short.
<i>vrb</i>	128-bit vector unsigned short.

**Returns**

A vector merge from only the odd halfwords of *vra* and *vrb*.

## 7.9 `src/vec_int32_ppc.h` File Reference

Header package containing a collection of 128-bit SIMD operations over 32-bit integer elements.

```
#include <vec_int16_ppc.h>
```

**Functions**

- static [vui32\\_t vec\\_absduw](#) ([vui32\\_t](#) vra, [vui32\\_t](#) vrb)  
*Vector Absolute Difference Unsigned Word.*
- static [vui32\\_t vec\\_clzw](#) ([vui32\\_t](#) vra)  
*Vector Count Leading Zeros word.*
- static [vui32\\_t vec\\_mrgahw](#) ([vui64\\_t](#) vra, [vui64\\_t](#) vrb)  
*Vector Merge Algebraic High Words.*
- static [vui32\\_t vec\\_mrgalw](#) ([vui64\\_t](#) vra, [vui64\\_t](#) vrb)  
*Vector merge Algebraic low words.*
- static [vui32\\_t vec\\_mrgew](#) ([vui32\\_t](#) vra, [vui32\\_t](#) vrb)  
*Vector Merge Even Words.*
- static [vui32\\_t vec\\_mrgow](#) ([vui32\\_t](#) vra, [vui32\\_t](#) vrb)

*Vector Merge Odd Words.*

- static `vi64_t vec_mulesw (vi32_t a, vi32_t b)`

*Vector multiply even signed words.*

- static `vi64_t vec_mulosw (vi32_t a, vi32_t b)`

*Vector multiply odd signed words.*

- static `vui64_t vec_muleuw (vui32_t a, vui32_t b)`

*Vector multiply even unsigned words.*

- static `vui64_t vec_mulouw (vui32_t a, vui32_t b)`

*Vector multiply odd unsigned words.*

- static `vi32_t vec_mulhsw (vi32_t vra, vi32_t vrb)`

*Vector Multiply High Signed Word.*

- static `vui32_t vec_mulhuw (vui32_t vra, vui32_t vrb)`

*Vector Multiply High Unsigned Word.*

- static `vui32_t vec_muluwm (vui32_t a, vui32_t b)`

*Vector Multiply Unsigned Word Modulo.*

- static `vui32_t vec_popcntw (vui32_t vra)`

*Vector Population Count word.*

- static `vui32_t vec_revbw (vui32_t vra)`

*byte reverse each word of a vector unsigned int.*

- static `vui32_t vec_slwi (vui32_t vra, const unsigned int shb)`

*Vector Shift left Word Immediate.*

- static `vi32_t vec_srawi (vi32_t vra, const unsigned int shb)`

*Vector Shift Right Algebraic Word Immediate.*

- static `vui32_t vec_srwi (vui32_t vra, const unsigned int shb)`

*Vector Shift Right Word Immediate.*

### 7.9.1 Detailed Description

Header package containing a collection of 128-bit SIMD operations over 32-bit integer elements.

Most of these operations are implemented in a single instruction on newer (POWER8/POWER9) processors. This header serves to fill in functional gaps for older (POWER7, POWER8) processors and provides a in-line assembler implementation for older compilers that do not provide the build-ins.

Most vector int (32-bit integer word) operations are implemented with PowerISA VMX instructions either defined by the original VMX (AKA AltiVec) or added to later versions of the PowerISA. Vector word-wise merge, shift, and splat operations were added with VSX in PowerISA 2.06B (POWER7). PowerISA 2.07B (POWER8) added several useful word wise operations (multiply, merge even/odd, count leading zeros, population count) not included in the original VMX. PowerISA 3.0B (POWER9) adds several more (compare not equal, count trailing zeros, extend sign, extract/insert, and parity). Most of these intrinsic (compiler built-ins) operations are defined in `<altivec.h>` and described in the compiler documentation.

#### Note

The compiler disables associated `<altivec.h>` built-ins if the `mcpu` target does not enable the specific instruction. For example if you compile with `-mcpu=power7`, `vec_vclz` and `vec_vclzw` will not be defined. Another example if you compile with `-mcpu=power8`, `vec_revb` will not be defined. This header provides the appropriate substitutions, will generate the minimum code, appropriate for the target, and produce correct results. Most ppc64le compilers will default to `-mcpu=power8` if not specified.

The newly introduced vector operations imply some useful composite operations. For example, we can make the vector multiply even/odd/modulo word operations available for older compilers. And provide implementations for older (POWER7 and earlier) processors using the original VMX operations.

This header covers operations that are either:

- Implemented in hardware instructions for later processors and useful to programmers, on slightly older processors, even if the equivalent function requires more instructions. Examples include the multiply even/odd/modulo word operations.
- Defined in the OpenPOWER ABI but *not* yet defined in `<altivec.h>` provided by available compilers in common use. Examples include Count Leading Zeros, Population Count and Byte Reverse.
- Commonly used operations, not covered by the ABI or `<altivec.h>`, and require multiple instructions or are not obvious. Examples include the shift immediate, merge algebraic high/low, and multiply high operations.

### 7.9.2 Endian problems with word operations

It would be useful to provide a vector multiply high word (return the high order 32-bits of the 64-bit product) operation. This can be used for multiplicative inverse (effectively integer divide) operations. Neither integer multiply high nor divide are available as vector instructions. However the multiply high word operation can be composed from the existing multiply even/odd word operations followed by the vector merge even word instruction.

As a prerequisite we need to provide the merge even/odd word operations for older compilers and an implementation for older (POWER7) processors. Fortunately vector merge operations are just a special case of vector permute. So the POWER7 (and earlier) implementation can use `vec_perm` and appropriate selection vectors to provide these merge operations.

But this is complicated by *little-endian* (LE) support as specified in the OpenPOWER ABI and as implemented in the compilers. Little-endian changes the effective vector element numbering and the location of even and odd elements. This means that the vector built-ins provided by `altivec.h` may not generate the instructions you would expect.

See also

[General Endian Issues](#)

The OpenPOWER ABI provides a helpful table of [Endian Sensitive Operations](#). For `vec_mergee` (`vmrgew`) it specifies:

Swap inputs and use `vmrgow`, for LE.

Also for `vec_mule` (`vmuleuw`, `vmulesw`):

Replace with `vmulouw` and so on, for LE.

Also for `vec_perm` (`vperm`) it specifies:

For LE, Swap input arguments and complement the selection vector.

The above is just a sampling of a larger list of Endian Sensitive Operations.

So the obvious coding for Vector Multiply High Word:

```

vui32_t
test_mulhw (vui32_t vra, vui32_t vrb)
{
    return vec_mergee ((vui32_t)vec_mule (vra, vrb),
                       (vui32_t)vec_mulo (vra, vrb));
}

```

Would produce the expected code and correct results when compiled for BE:

```

<test_mulhw>:
    vmuleuw  v0,v2,v3
    vmuluuw  v2,v2,v3
    vmrgew   v2,v0,v2
    blr

```

But the following and wrong code for LE:

```

<test_mulhw>:
    vmulouw  v0,v2,v3
    vmuleuw  v2,v2,v3
    vmrgow   v2,v2,v0
    blr

```

The compiler swapped the multiplies even for odd and odd of even. That is somewhat mitigated by swapping the input arguments in the merge. But changing the merge from even to odd actually returns the low order 32-bits of the product. This is not the correct result for multiply high.

This header provides implementations of vector merge even/odd word ([vec\\_mrgew\(\)](#) and [vec\\_mrgow\(\)](#)) that support older compilers and older (POWER7) processor. Similarly for the multiply Even/odd unsigned/signed word instructions ([vec\\_mulesw\(\)](#), [vec\\_mulosw\(\)](#), [vec\\_muleuw\(\)](#) and [vec\\_mulouw\(\)](#)). These implementations include the mandated LE transforms.

#### 7.9.2.1 Vector Merge Algebraic High Word example

This header also provides the higher level operations Vector Merge Algebraic High/low Word ([vec\\_mrgahw\(\)](#) and [vec\\_mrgalw\(\)](#)). These implementations generate the correct merge even/odd word instruction for the operation independent of endian.

##### Note

The parameters are vector unsigned long ([vui64\\_t](#)) to match results from [vec\\_muleuw\(\)](#) and [vec\\_mulouw\(\)](#).

```

static inline vui32_t
vec_mrgahw (vui64_t vra, vui64_t vrb)
{
    vui32_t res;
#ifdef _ARCH_PWR8
#ifdef vec_vmrgew // Use altivec.h builtins
#if __BYTE_ORDER__ == __ORDER_LITTLE_ENDIAN__
    // really want vmrgew here! So do the opposite.
    res = vec_vmrgow ((vui32_t)vrb, (vui32_t)vra);
#else
    res = vec_vmrgew ((vui32_t)vra, (vui32_t)vrb);
#endif
#else // Generate vmrgew directly in assembler
    __asm__(
        "vmrgew %0,%1,%2;\n"
        : "=v" (res)
        : "v" (vra),
        "v" (vrb)
        : );
#endif
#else // POWER7 and earlier, Assume BE only
    const vui32_t vconstp =
        CONST_VINT32_W(0x00010203, 0x10111213, 0x08090a0b, 0x18191a1b);
    res = (vui32_t) vec_perm ((vui8_t) vra, (vui8_t) vrb, (
        vui8_t) vconstp);
#endif
    return (res);
}

```

The implementation is a bit complicated so that it can nullify the unwanted LE transformation of [vec\\_vmrgew\(\)](#), in addition to handling older compilers and processors.

### 7.9.2.2 Vector Multiply High Unsigned Word example

Now we can implement Vector Multiply High Unsigned Word (`vec_mulhuw()`):

```
static inline vui32_t
vec_mulhuw (vui32_t vra, vui32_t vrb)
{
    #if __BYTE_ORDER__ == __ORDER_LITTLE_ENDIAN__
        return vec_mrgahw (vec_mulouw (vra, vrb), vec_muleuw (vra, vrb));
    #else
        return vec_mrgahw (vec_muleuw (vra, vrb), vec_mulouw (vra, vrb));
    #endif
}
```

Again the implementation is more complicated than expected as we still have to nullify the LE transformation associated with multiply even/odd.

The good news is all this complexity is contained within pveclib and the generated code is still just 3 instructions.

```
vmulouw v0,v2,v3
vmuleuw v2,v2,v3
vmrgew v2,v2,v0
```

## 7.9.3 Vector Word Examples

Suppose we have a requirement to convert an array of 32-bit time-interval values that need to convert to timespec format. For simplicity we will also assume that the array is nicely (Quadword) aligned and an integer multiple of 4 words.

The PowerISA provides a 64-bit TimeBase register that clocks at a constant 512MHz. The TimeBase can be read directly as either the full 64-bit value or as 32-bit upper and lower halves. For this example we assume that the lower 32-bits of the TimeBase is sufficient to compute intervals (~8.38 seconds). TimeBase values of adjacent events are subtracted to generate the intervals stored in the array.

The timespec format is a struct of unsigned int fields for seconds and microseconds. So the task is to convert the 512MHz TimeBase intervals to microseconds and then split the integer seconds and microseconds for the timespec.

First the TimeBase to microseconds conversion is simply  $(1000000 / 512000000)$  which reduces to  $(1 / 512)$  or divide by 512. The vector unit does not provide integer divide but luckily, 512 is a power of 2 and we can shift right. If we don't care for the niceties of rounding we can simply shift right 9 bits:

```
tb_usec = vec_srwi (*tb++, 9);
```

But if we decide that rounding is important we can leverage the Vector Average Unsigned Word (`vavguw`) instruction. Here we need to add 256 ( $512 / 2 = 256$ ) to the timeBase interval before we shift right.

But we need to reverse engineer the `vavguw` operation to get the results we want. For each word, `vavguw` computes the sum of A and B plus 1, then shifts the 33-bit sum right 1 bit. We can effectively round by passing the rounding factor as the B operand to the `vec_avg()` built-in. But we get a +1 and 1 bit right shift for free. So in this case the rounding constant is  $256-1 = 255$ . And we only need to shift an additional 8 bits to complete the conversion:

```
const vui32_t rnd_512 =
{ (256-1), (256-1), (256-1), (256-1) };
// Convert 512MHz timebase to microseconds with rounding.
tmp = vec_avg (*tb++, rnd_512);
tb_usec = vec_srwi (tmp, 8);
```

**Note**

`vec_avg()` is an existing `altivec.h` generic built-in.

Next we need to separate TimeBase microseconds into the integer seconds and microseconds. Normally scalar codes would use integer divide/modulo by 1000000. Did I mention that the PowerISA vector unit does not have a integer divide operation?

Instead we can use the multiplicative inverse which is a scaled fixed point fraction calculated from the original divisor. This works nicely if the fixed radix point is just before the 32-bit fraction and we have a multiply high (`vec_mulhuw()`) operation. Multiplying a 32-bit unsigned integer by a 32-bit unsigned fraction generates a 64-bit product with 32-bits above (integer) and below (fraction) the radix point. The high 32-bits of the product is the integer quotient.

It turns out that generating the multiplicative inverse can be tricky. To produce correct results over the full analysis, possible pre-scaling and post-shifting, and sometimes a corrective addition is necessary. Fortunately the mathematics are well understood and are commonly used in optimizing compilers. Even better, Henry Warren's book has a whole chapter on this topic.

**See also**

"Hacker's Delight, 2nd Edition," Henry S. Warren, Jr, Addison Wesley, 2013. Chapter 10, Integer Division by Constants.

In the chapter above;

Figure 10-2 Computing the magic number for unsigned division.

provides a sample C function for generating the magic number (actually a struct containing; the magic multiplicative inverse, "add" indicator, and the shift amount.). For the divisor 1000000 this is { 1125899907, 0 , 18 }:

- the multiplier is 1125899907.
- no corrective add of the dividend is required.
- the final shift is 18-bits right.

```
const vui32_t mul_invs_lm =
{ 1125899907, 1125899907, 1125899907 };
const int shift_lm = 18;

tmp = vec_mulhuw (tb_usec, mul_invs_lm);
seconds = vec_srwi (tmp, shift_lm);
```

Now we need to compute the remainder to get microseconds.

```
const vui32_t usec_sec =
{ 1000000, 1000000, 1000000, 1000000 };

tmp = vec_muluwm (seconds, usec_sec);
useconds = vec_sub (tb_usec, tmp);
```

Finally we need to merge the vectors of seconds and useconds into vectors of timespec.

```
timespec1 = vec_mergeh (seconds, useconds);
timespec2 = vec_mergel (seconds, useconds);
```

**Note**

`vec_sub()`, `vec_mergeh()`, and `vec_mergel()` are an existing `altivec.h` generic built-ins.

### 7.9.3.1 Vectorized TimeBase conversion example

Here is the complete vectorized TimeBase to timespec conversion example:

```
void
example_convert_timebase (vui32_t *tb, vui32_t *timespec, int n)
{
    const vui32_t rnd_512 =
        { (256-1), (256-1), (256-1), (256-1) };
    // Magic numbers for multiplicative inverse to divide by 1,000,000
    // are 1125899907 and shift right 18 bits.
    const vui32_t mul_invs_lm =
        { 1125899907, 1125899907, 1125899907, 1125899907 };
    const int shift_lm = 18;
    // Need const for microseconds/second to extract remainder.
    const vui32_t usec_sec =
        { 1000000, 1000000, 1000000, 1000000 };
    vui32_t tmp, tb_usec, seconds, useconds;
    vui32_t timespec1, timespec2;
    int i;

    for (i = 0; i < n; i++)
    {
        // Convert 512MHz timebase to microseconds with rounding.
        tmp = vec_avg (*tb++, rnd_512);
        tb_usec = vec_srwi (tmp, 8);
        // extract integer seconds from tb_usec.
        tmp = vec_mulhuw (tb_usec, mul_invs_lm);
        seconds = vec_srwi (tmp, shift_lm);
        // Extract remainder microseconds.
        tmp = vec_muluwm (seconds, usec_sec);
        useconds = vec_sub (tb_usec, tmp);
        // Use merge high/low to interleave seconds and useconds in timespec.
        timespec1 = vec_mergeh (seconds, useconds);
        timespec2 = vec_mergel (seconds, useconds);
        // Store timespec.
        *timespec++ = timespec1;
        *timespec++ = timespec2;
    }
}
```

## 7.9.4 Performance data.

High level performance estimates are provided as an aid to function selection when evaluating algorithms. For background on how *Latency* and *Throughput* are derived see: [Performance data](#).

## 7.9.5 Function Documentation

### 7.9.5.1 static vui32\_t vec\_absduw ( vui32\_t vra, vui32\_t vrb ) [inline],[static]

Vector Absolute Difference Unsigned Word.

Compute the absolute difference for each word. For each unsigned word, subtract VRB[i] from VRA[i] and return the absolute value of the difference.

processor	Latency	Throughput
power8	4	1/cycle
power9	3	2/cycle

#### Parameters

<i>vra</i>	vector of 4 x unsigned words
<i>vrb</i>	vector of 4 x unsigned words



**Returns**

vector of the absolute differences.

**7.9.5.2 static vui32\_t vec\_clzw ( vui32\_t vra ) [inline],[static]**

Vector Count Leading Zeros word.

Count the number of leading '0' bits (0-32) within each word element of a 128-bit vector.

For POWER8 (PowerISA 2.07B) or later use the Vector Count Leading Zeros Word instruction **vclzw**. Otherwise use sequence of pre 2.07 VMX instructions. SIMDized count leading zeros inspired by:

Warren, Henry S. Jr and *Hacker's Delight*, 2nd Edition, Addison Wesley, 2013. Chapter 5 Counting Bits, Figure 5-12.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

**Parameters**

<i>vra</i>	128-bit vector treated as 4 x 32-bit integer (words) elements.
------------	--

**Returns**

128-bit vector with the Leading Zeros count for each word element.

**7.9.5.3 static vui32\_t vec\_mrgahw ( vui64\_t vra, vui64\_t vrb ) [inline],[static]**

Vector Merge Algebraic High Words.

Merge only the high words from 4 x Algebraic doublewords across vectors vra and vrb. This effectively the Vector Merge Even Word operation that is not modified for endian.

For example merge the high 32-bits from 4 x 64-bit products as generated by vec\_muleuw/vec\_mulouw. This result is effectively a vector multiply high unsigned word.

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

**Parameters**

<i>vra</i>	128-bit vector unsigned long.
<i>vrh</i>	128-bit vector unsigned long.

**Returns**

A vector merge from only the high words of the 4 x Algebraic doublewords across *vra* and *vr*b.

**7.9.5.4** `static vui32_t vec_mrgalw ( vui64_t vra, vui64_t vrb ) [inline],[static]`

Vector merge Algebraic low words.

Merge the arithmetic low words 4 x Algebraic doublewords across vectors *vra* and *vr*b. This is effectively the Vector Merge Odd Word operation that is not modified for endian.

For example merge the low 32-bits from 4 x 64-bit products as generated by `vec_muleuw/vec_mulouw`. This result is effectively a vector multiply low unsigned word (multiply unsigned word modulo).

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

**Parameters**

<i>vra</i>	128-bit vector unsigned long.
<i>vr</i> b	128-bit vector unsigned long.

**Returns**

A vector merge from only the low words of the 4 x Algebraic doublewords across *vra* and *vr*b.

**7.9.5.5** `static vui32_t vec_mrgew ( vui32_t vra, vui32_t vrb ) [inline],[static]`

Vector Merge Even Words.

Merge the even word elements from the concatenation of 2 x vectors (*vra* and *vr*b).

- `res[0] = vra[0];`
- `res[1] = vrb[0];`
- `res[2] = vra[2];`
- `res[3] = vrb[2];`

The element numbering changes between big and little-endian environements. So the compiler and this implementation adjusts the generated code to reflect this.

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

## Parameters

<i>vra</i>	128-bit vector unsigned int.
<i>vr<b>b</b></i>	128-bit vector unsigned int.

## Returns

A vector merge from only the even words of *vra* and *vr**b***.

7.9.5.6 `static vui32_t vec_mrgow ( vui32_t vra, vui32_t vrb )` `[inline]`, `[static]`

Vector Merge Odd Words.

Merge the odd word elements from the concatenation of 2 x vectors (*vra* and *vr**b***).

- `res[0] = vra[1];`
- `res[1] = vrb[1];`
- `res[2] = vra[3];`
- `res[3] = vrb[3];`

The element numbering changes between big and little-endian environments. So the compiler and this implementation adjusts the generated code to reflect this.

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

## Parameters

<i>vra</i>	128-bit vector unsigned int.
<i>vr<b>b</b></i>	128-bit vector unsigned int.

## Returns

A vector merge from only the even words of *vra* and *vr**b***.

7.9.5.7 `static vi64_t vec_mulesw ( vi32_t a, vi32_t b )` `[inline]`, `[static]`

Vector multiply even signed words.

Multiple the even words of two vector signed int values and return the signed long product of the even words.

processor	Latency	Throughput
power8	7	2/cycle
power9	7	2/cycle

**Parameters**

<i>a</i>	128-bit vector signed int.
<i>b</i>	128-bit vector signed int.

**Returns**

vector signed long product of the even words of *a* and *b*.

**7.9.5.8** `static vui64_t vec_muleuw ( vui32_t a, vui32_t b )` `[inline],[static]`

Vector multiply even unsigned words.

Multiple the even words of two vector unsigned int values and return the unsigned long product of the even words.

processor	Latency	Throughput
power8	7	2/cycle
power9	7	2/cycle

**Parameters**

<i>a</i>	128-bit vector unsigned int.
<i>b</i>	128-bit vector unsigned int.

**Returns**

vector unsigned long product of the even words of *a* and *b*.

**7.9.5.9** `static vi32_t vec_mulhsw ( vi32_t vra, vi32_t vrb )` `[inline],[static]`

Vector Multiply High Signed Word.

Multiple the corresponding word elements of two vector signed int values and return the high order 32-bits, for each 64-bit product element.

processor	Latency	Throughput
power8	9	1/cycle
power9	9	1/cycle

**Parameters**

<i>vra</i>	128-bit vector signed int.
<i>vrb</i>	128-bit vector signed int.

**Returns**

vector of the high order 32-bits of the product of the word elements from *vra* and *vrh*.

**7.9.5.10** `static vui32_t vec_mulhuw ( vui32_t vra, vui32_t vrh )` `[inline], [static]`

Vector Multiply High Unsigned Word.

Multiple the corresponding word elements of two vector unsigned int values and return the high order 32-bits, from each 64-bit product.

processor	Latency	Throughput
power8	9	1/cycle
power9	9	1/cycle

**Note**

This operation can be used to effectively perform a divide by multiplying by the scaled multiplicative inverse (reciprocal).

Warren, Henry S. Jr and *Hacker's Delight*, 2nd Edition, Addison Wesley, 2013. Chapter 10, Integer Division by Constants.

**Parameters**

<i>vra</i>	128-bit vector unsigned int.
<i>vrh</i>	128-bit vector unsigned int.

**Returns**

vector of the high order 32-bits of the signed product of the word elements from *vra* and *vrh*.

**7.9.5.11** `static vi64_t vec_mulosw ( vi32_t a, vi32_t b )` `[inline], [static]`

Vector multiply odd signed words.

Multiple the odd words of two vector signed int values and return the signed long product of the odd words.

processor	Latency	Throughput
power8	7	2/cycle
power9	7	2/cycle

**Parameters**

<i>a</i>	128-bit vector signed int.
<i>b</i>	128-bit vector signed int.

**Returns**

vector signed long product of the odd words of a and b.

**7.9.5.12** `static vui64_t vec_mulouw ( vui32_t a, vui32_t b )` `[inline],[static]`

Vector multiply odd unsigned words.

Multiple the odd words of two vector unsigned int values and return the unsigned long product of the odd words.

processor	Latency	Throughput
power8	7	2/cycle
power9	7	2/cycle

**Parameters**

<i>a</i>	128-bit vector unsigned int.
<i>b</i>	128-bit vector unsigned int.

**Returns**

vector unsigned long product of the odd words of a and b.

**7.9.5.13** `static vui32_t vec_muluwm ( vui32_t a, vui32_t b )` `[inline],[static]`

Vector Multiply Unsigned Word Modulo.

Multiple the corresponding word elements of two vector unsigned int values and return the low order 32-bits of the 64-bit product for each element.

**Note**

`vec_muluwm` can be used for unsigned or signed integers. It is the vector equivalent of Multiply Low Word.

processor	Latency	Throughput
power8	7	2/cycle
power9	7	2/cycle

**Parameters**

<i>a</i>	128-bit vector signed int.
<i>b</i>	128-bit vector signed int.

**Returns**

vector of the low order 32-bits of the unsigned product of the word elements from `vra` and `vrb`.

**7.9.5.14** static vui32\_t vec\_popcntw ( vui32\_t vra ) [inline],[static]

Vector Population Count word.

Count the number of '1' bits (0-32) within each word element of a 128-bit vector.

For POWER8 (PowerISA 2.07B) or later use the Vector Population Count Word instruction. Otherwise use the pveclib vec\_popcntb to count each byte then sum across with Vector Sum across Quarter Unsigned Byte Saturate.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

**Parameters**

vra	128-bit vector treated as 4 x 32-bit integer (words) elements.
-----	--

**Returns**

128-bit vector with the population count for each word element.

**7.9.5.15** static vui32\_t vec\_revbw ( vui32\_t vra ) [inline],[static]

byte reverse each word of a vector unsigned int.

For each word of the input vector, reverse the order of bytes / octets within the word.

processor	Latency	Throughput
power8	2-11	2/cycle
power9	3	2/cycle

**Parameters**

vra	a 128-bit vector unsigned int.
-----	--------------------------------

**Returns**

a 128-bit vector with the bytes of each word reversed.

**7.9.5.16** static vui32\_t vec\_slwi ( vui32\_t vra, const unsigned int shb ) [inline],[static]

Vector Shift left Word Immediate.

Shift left each word element [0-3], 0-31 bits, as specified by an immediate value. The shift amount is a const unsigned int in the range 0-31. A shift count of 0 returns the original value of vra. Shift counts greater than 31 bits return zero.

processor	Latency	Throughput
power8	4-11	2/cycle
power9	5-11	2/cycle

#### Parameters

<i>vra</i>	a 128-bit vector treated as a vector unsigned int.
<i>shb</i>	shift amount in the range 0-31.

#### Returns

128-bit vector unsigned int, shifted left *shb* bits.

**7.9.5.17** `static vi32_t vec_srawi ( vi32_t vra, const unsigned int shb )` `[inline]`, `[static]`

Vector Shift Right Algebraic Word Immediate.

Shift Right Algebraic each word element [0-3], 0-31 bits, as specified by an immediate value. The shift amount is a const unsigned int in the range 0-31. A shift count of 0 returns the original value of *vra*. Shift counts greater than 31 bits return the sign bit propagated to each bit of each element.

processor	Latency	Throughput
power8	4-11	2/cycle
power9	5-11	2/cycle

#### Parameters

<i>vra</i>	a 128-bit vector treated as a vector signed int.
<i>shb</i>	shift amount in the range 0-31.

#### Returns

128-bit vector signed int, shifted right *shb* bits.

**7.9.5.18** `static vui32_t vec_srwi ( vui32_t vra, const unsigned int shb )` `[inline]`, `[static]`

Vector Shift Right Word Immediate.

Shift right each word element [0-3], 0-31 bits, as specified by an immediate value. The shift amount is a const unsigned int in the range 0-31. A shift count of 0 returns the original value of *vra*. Shift counts greater than 31 bits return zero.

processor	Latency	Throughput
power8	4-11	2/cycle
power9	5-11	2/cycle



## Parameters

<i>vra</i>	a 128-bit vector treated as a vector unsigned char.
<i>shb</i>	shift amount in the range 0-31.

## Returns

128-bit vector unsigned int, shifted right shb bits.

## 7.10 src/vec\_int64\_ppc.h File Reference

Header package containing a collection of 128-bit SIMD operations over 64-bit integer elements.

```
#include <vec_int32_ppc.h>
```

## Functions

- static [vui64\\_t](#) [vec\\_absdud](#) ([vui64\\_t](#) vra, [vui64\\_t](#) vrb)  
*Vector Absolute Difference Unsigned doubleword.*
- static [vui64\\_t](#) [vec\\_addudm](#) ([vui64\\_t](#) a, [vui64\\_t](#) b)  
*Vector Add Unsigned Doubleword Modulo.*
- static [vui64\\_t](#) [vec\\_clzd](#) ([vui64\\_t](#) vra)  
*Count leading zeros for a vector unsigned long int.*
- static [vb64\\_t](#) [vec\\_cmpeqsd](#) ([vi64\\_t](#) a, [vi64\\_t](#) b)  
*Vector Compare Equal Signed Doubleword.*
- static [vb64\\_t](#) [vec\\_cmpequd](#) ([vui64\\_t](#) a, [vui64\\_t](#) b)  
*Vector Compare Equal Unsigned Doubleword.*
- static [vb64\\_t](#) [vec\\_cmpgesd](#) ([vi64\\_t](#) a, [vi64\\_t](#) b)  
*Vector Compare Greater Than or Equal Signed Doubleword.*
- static [vb64\\_t](#) [vec\\_cmpgeud](#) ([vui64\\_t](#) a, [vui64\\_t](#) b)  
*Vector Compare Greater Than or Equal Unsigned Doubleword.*
- static [vb64\\_t](#) [vec\\_cmpgtsd](#) ([vi64\\_t](#) a, [vi64\\_t](#) b)  
*Vector Compare Greater Than Signed Doubleword.*
- static [vb64\\_t](#) [vec\\_cmpgtud](#) ([vui64\\_t](#) a, [vui64\\_t](#) b)  
*Vector Compare Greater Than Unsigned Doubleword.*
- static [vb64\\_t](#) [vec\\_cmpleud](#) ([vui64\\_t](#) a, [vui64\\_t](#) b)  
*Vector Compare Less Than Equal Unsigned Doubleword.*
- static [vb64\\_t](#) [vec\\_cmpleud](#) ([vui64\\_t](#) a, [vui64\\_t](#) b)  
*Vector Compare Less Than Equal Unsigned Doubleword.*
- static [vb64\\_t](#) [vec\\_cmpltud](#) ([vui64\\_t](#) a, [vui64\\_t](#) b)  
*Vector Compare less Than Signed Doubleword.*
- static [vb64\\_t](#) [vec\\_cmpltud](#) ([vui64\\_t](#) a, [vui64\\_t](#) b)  
*Vector Compare less Than Unsigned Doubleword.*
- static [vb64\\_t](#) [vec\\_cmpnesd](#) ([vi64\\_t](#) a, [vi64\\_t](#) b)  
*Vector Compare Not Equal Signed Doubleword.*
- static [vb64\\_t](#) [vec\\_cmpneud](#) ([vui64\\_t](#) a, [vui64\\_t](#) b)  
*Vector Compare Not Equal Unsigned Doubleword.*

- static int [vec\\_cmpsd\\_all\\_eq](#) (vi64\_t a, vi64\_t b)  
*Vector Compare all Equal Signed Doubleword.*
- static int [vec\\_cmpsd\\_all\\_ge](#) (vi64\_t a, vi64\_t b)  
*Vector Compare all Greater Than or Equal Signed Doubleword.*
- static int [vec\\_cmpsd\\_all\\_gt](#) (vi64\_t a, vi64\_t b)  
*Vector Compare all Greater Than Signed Doubleword.*
- static int [vec\\_cmpsd\\_all\\_le](#) (vi64\_t a, vi64\_t b)  
*Vector Compare all Less than equal Signed Doubleword.*
- static int [vec\\_cmpsd\\_all\\_lt](#) (vi64\_t a, vi64\_t b)  
*Vector Compare all Less than Signed Doubleword.*
- static int [vec\\_cmpsd\\_all\\_ne](#) (vi64\_t a, vi64\_t b)  
*Vector Compare all Not Equal Signed Doubleword.*
- static int [vec\\_cmpsd\\_any\\_eq](#) (vi64\_t a, vi64\_t b)  
*Vector Compare any Equal Signed Doubleword.*
- static int [vec\\_cmpsd\\_any\\_ge](#) (vi64\_t a, vi64\_t b)  
*Vector Compare any Greater Than or Equal Signed Doubleword.*
- static int [vec\\_cmpsd\\_any\\_gt](#) (vi64\_t a, vi64\_t b)  
*Vector Compare any Greater Than Signed Doubleword.*
- static int [vec\\_cmpsd\\_any\\_le](#) (vi64\_t a, vi64\_t b)  
*Vector Compare any Less than equal Signed Doubleword.*
- static int [vec\\_cmpsd\\_any\\_lt](#) (vi64\_t a, vi64\_t b)  
*Vector Compare any Less than Signed Doubleword.*
- static int [vec\\_cmpsd\\_any\\_ne](#) (vi64\_t a, vi64\_t b)  
*Vector Compare any Not Equal Signed Doubleword.*
- static int [vec\\_cmpud\\_all\\_eq](#) (vui64\_t a, vui64\_t b)  
*Vector Compare all Equal Unsigned Doubleword.*
- static int [vec\\_cmpud\\_all\\_ge](#) (vui64\_t a, vui64\_t b)  
*Vector Compare all Greater Than or Equal Unsigned Doubleword.*
- static int [vec\\_cmpud\\_all\\_gt](#) (vui64\_t a, vui64\_t b)  
*Vector Compare all Greater Than Unsigned Doubleword.*
- static int [vec\\_cmpud\\_all\\_le](#) (vui64\_t a, vui64\_t b)  
*Vector Compare all Less than equal Unsigned Doubleword.*
- static int [vec\\_cmpud\\_all\\_lt](#) (vui64\_t a, vui64\_t b)  
*Vector Compare all Less than Unsigned Doubleword.*
- static int [vec\\_cmpud\\_all\\_ne](#) (vui64\_t a, vui64\_t b)  
*Vector Compare all Not Equal Unsigned Doubleword.*
- static int [vec\\_cmpud\\_any\\_eq](#) (vui64\_t a, vui64\_t b)  
*Vector Compare any Equal Unsigned Doubleword.*
- static int [vec\\_cmpud\\_any\\_ge](#) (vui64\_t a, vui64\_t b)  
*Vector Compare any Greater Than or Equal Unsigned Doubleword.*
- static int [vec\\_cmpud\\_any\\_gt](#) (vui64\_t a, vui64\_t b)  
*Vector Compare any Greater Than Unsigned Doubleword.*
- static int [vec\\_cmpud\\_any\\_le](#) (vui64\_t a, vui64\_t b)  
*Vector Compare any Less than equal Unsigned Doubleword.*
- static int [vec\\_cmpud\\_any\\_lt](#) (vui64\_t a, vui64\_t b)  
*Vector Compare any Less than Unsigned Doubleword.*
- static int [vec\\_cmpud\\_any\\_ne](#) (vui64\_t a, vui64\_t b)  
*Vector Compare any Not Equal Unsigned Doubleword.*
- static [vi64\\_t vec\\_maxsd](#) (vi64\_t vra, vi64\_t vrb)  
*Vector Maximum Signed Doubleword.*
- static [vui64\\_t vec\\_maxud](#) (vui64\_t vra, vui64\_t vrb)

- Vector Maximum Unsigned Doubleword.*

  - static `vi64_t vec_minsd` (`vi64_t` vra, `vi64_t` vrb)
- Vector Minimum Signed Doubleword.*

  - static `vui64_t vec_minud` (`vui64_t` vra, `vui64_t` vrb)
- Vector Minimum Unsigned Doubleword.*

  - static `vui64_t vec_mrgahd` (`vui128_t` vra, `vui128_t` vrb)
- Vector Merge Algebraic High Doublewords.*

  - static `vui64_t vec_mrgald` (`vui128_t` vra, `vui128_t` vrb)
- Vector Merge Algebraic Low Doublewords.*

  - static `vui64_t vec_mrged` (`vui64_t` \_\_VA, `vui64_t` \_\_VB)
- Vector Merge Even Doubleword. Merge the even doubleword elements from two vectors into the high and low doubleword elements of the result. This is effectively the VSX Permute doubleword Immediate operation modified for endian.*

  - static `vui64_t vec_mrghd` (`vui64_t` \_\_VA, `vui64_t` \_\_VB)
- Vector Merge High Doubleword. Merge the high doubleword elements from two vectors into the high and low doubleword elements of the result. This is effectively the VSX Permute doubleword Immediate operation modified for endian.*

  - static `vui64_t vec_mrgld` (`vui64_t` \_\_VA, `vui64_t` \_\_VB)
- Vector Merge Low Doubleword. Merge the low doubleword elements from two vectors into the high and low doubleword elements of the result. This is effectively the VSX Permute doubleword Immediate operation modified for endian.*

  - static `vui64_t vec_mrgod` (`vui64_t` \_\_VA, `vui64_t` \_\_VB)
- Vector Merge Odd Doubleword. Merge the odd doubleword elements from two vectors into the high and low doubleword elements of the result. This is effectively the VSX Permute doubleword Immediate operation modified for endian.*

  - static `vui64_t vec_pasted` (`vui64_t` \_\_VH, `vui64_t` \_\_VL)
- Vector doubleword paste. Concatenate the high doubleword of the 1st vector with the low double word of the 2nd vector.*

  - static `vui64_t vec_permdi` (`vui64_t` vra, `vui64_t` vrb, const int ctl)
- Vector Permute Doubleword Immediate. Combine a doubleword selected from the 1st (vra) vector with a doubleword selected from the 2nd (vrb) vector.*

  - static `vui64_t vec_popcntd` (`vui64_t` vra)
- Vector Population Count doubleword.*

  - static `vui64_t vec_revbd` (`vui64_t` vra)
- byte reverse each doubleword for a vector unsigned long int.*

  - static `vui64_t vec_vrld` (`vui64_t` vra, `vui64_t` vrb)
- Vector Rotate Left Doubleword.*

  - static `vui64_t vec_vsld` (`vui64_t` vra, `vui64_t` vrb)
- Vector Shift Left Doubleword.*

  - static `vui64_t vec_vsrd` (`vui64_t` vra, `vui64_t` vrb)
- Vector Shift Right Doubleword.*

  - static `vi64_t vec_vsrld` (`vi64_t` vra, `vui64_t` vrb)
- Vector Shift Right Algebraic Doubleword.*

  - static `vui64_t vec_rldi` (`vui64_t` vra, const unsigned int shb)
- Vector Rotate left Doubleword Immediate.*

  - static `vui64_t vec_sldi` (`vui64_t` vra, const unsigned int shb)
- Vector Shift left Doubleword Immediate.*

  - static `vui64_t vec_splatd` (`vui64_t` vra, const int ctl)
- Vector splat doubleword. Duplicate the selected doubleword element across the doubleword elements of the result. This is effectively the VSX Merge doubleword operation modified for endian.*

  - static `vui64_t vec_spltd` (`vui64_t` vra, const int ctl)
- Vector Shift Right Doubleword Immediate.*

  - static `vui64_t vec_srdi` (`vui64_t` vra, const unsigned int shb)

- static `vi64_t vec_sradi` (`vi64_t` vra, const unsigned int shb)  
*Vector Shift Right Algebraic Doubleword Immediate.*
- static `vui64_t vec_subudm` (`vui64_t` a, `vui64_t` b)  
*Vector Subtract Unsigned Doubleword Modulo.*
- static `vui64_t vec_swapd` (`vui64_t` vra)  
*Vector doubleword swap. Exchange the high and low doubleword elements of a vector.*
- static `vui32_t vec_vpudum` (`vui64_t` vra, `vui64_t` vrb)  
*Vector Pack Unsigned Doubleword Unsigned Modulo.*
- static `vui64_t vec_xspltd` (`vui64_t` vra, const int ctl)  
*Vector splat doubleword. Duplicate the selected doubleword element across the doubleword elements of the result.*
- static `vui64_t vec_vmsumuwmm` (`vui32_t` vra, `vui32_t` vrb, `vui64_t` vrc)  
*Vector Multiply-Sum Unsigned Word Modulo.*

### 7.10.1 Detailed Description

Header package containing a collection of 128-bit SIMD operations over 64-bit integer elements.

Most of these operations are implemented in a single instruction on newer (POWER8/POWER9) processors. This header serves to fill in functional gaps for older (POWER7, POWER8) processors and provides a in-line assembler implementation for older compilers that do not provide the build-ins.

The original VMX (AKA AltiVec) did not define any doubleword element (long long integer or double float) operations. The VSX facility (introduced with POWER7) added vector double float but did not add any integer doubleword (64-bit) operations. However it did add a useful doubleword permute immediate and word wise; merge, shift, and splat immediate operations. Otherwise vector long int (64-bit elements) operations have to be implemented using VMX word and halfword element integer operations for POWER7.

POWER8 (PowerISA 2.07B) adds important doubleword integer (add, subtract, compare, shift, rotate, ...)  $V \leftrightarrow$  MX operations. POWER8 also added multiply word operations that produce the full doubleword product and full quadword add / subtract (with carry extend).

POWER9 (PowerISA 3.0B) adds the **Vector Multiply-Sum Unsigned Doubleword Modulo** instruction. This is not the expected multiply even/odd/modulo doubleword nor a full multiply modulo quadword. But with a few extra (permutes and splat zero) instructions you can get equivalent function.

#### Note

The doubleword integer multiply implementations are included in `vec_int128_ppc.h`. This resolves a circular dependency as 64-bit by 64-bit integer multiplies require 128-bit integer addition (`vec_adduqm()`) to produce the full product.

#### See also

`vec_msumudm`, `vec_muleud`, `vec_mulhud`, `vec_muloud`, `vec_muludm`, `vec_vmuleud`, and `vec_vmuloud`

Most of these intrinsic (compiler built-in) operations are defined in `<altivec.h>` and described in the compiler documentation. However it took several compiler releases for all the new POWER8 64-bit integer vector intrinsics to be added to **altivec.h**. This support started with the GCC 4.9 but was not complete across function/type and bug free until GCC 6.0.

**Note**

The compiler disables associated `<altivec.h>` built-ins if the **mcpu** target does not enable the specific instruction. For example if you compile with **-mcpu=power7**, `vec_vclz` and `vec_vclzd` will not be defined. But `vec_clzd` is always defined in this header, will generate the minimum code, appropriate for the target, and produce correct results.

64-bit integer operations are commonly used in the implementation of optimized double float math library functions and this applies to the vector equivalents of math functions. So missing, incomplete or buggy support for vector long integer intrinsics can be a impediment to the implementation of optimized and portable vector double math libraries. This header is a prerequisite for [vec\\_f64\\_ppc.h](#) which together are intended to support the implementation of vector math libraries.

Most of these operations are implemented in a single instruction on newer (POWER8/POWER9) processors. So this header serves to fill in functional gaps for older (POWER7, POWER8) processors and provides a in-line assembler implementation for older compilers that do not provide the build-ins.

This header covers operations that are any of the following:

- Implemented in hardware instructions for later processors and useful to programmers, on slightly older processors, even if the equivalent function requires more instructions. Examples include the doubleword operations: Add, Compare, Maximum, Minimum and subtract.
- Defined in the OpenPOWER ABI but *not* yet defined in `<altivec.h>` provided by available compilers in common use. Examples include doubleword forms of: Multiply even/odd/modulo, Count Leading Zeros, Population Count, and Byte Reverse operations.
- Commonly used operations, not covered by the ABI or `<altivec.h>`, and require multiple instructions or are not obvious. Examples include doubleword forms of: Merge Algebraic High/Low, Paste, and Rotate/Shift Immediate operations.
- Commonly used operations that are useful for doubleword, but are missing from the PowerISA and OpenPOWER ABI. Examples include: Absolute Difference Doubleword and Multiply-Sum Unsigned Word Modulo.

### 7.10.2 Some missing doubleword operations

The original VMX instruction set extension was limited to byte, halfword, and word size element operations. This limited vector arithmetic operations to char, short, int and float elements. This limitation persisted until PowerISA 2.06 (POWER7) added the Vector Scalar Extensions (VSX) facility. VSX combined/extended the FPRs and VRs into 64 by 128-bit Vector/Scalar Registers (VSRs).

VSX added a large number of scalar double-precision and vector single / double-precision floating-point operations. The double-precision scalar (**xs** prefix) instructions were largely duplicates of the existing Floating-Point Facility operations, extended to access the whole (64) VSX register set. Similarly the VSX vector single precision floating-point (**xv** prefix, **sp** suffix) instructions were added to give vectorized float code access to 64 VSX registers.

The addition of VSX vector double-precision (**xv** prefix) instructions was the most significant addition. This added vector doubleword floating-point operations and provided access to all 64 VSX registers. Alas there are no doubleword (64-bit long) integer operations in the initial VSX. A few logical and permute class (**xx** prefix) operations on word/doubleword elements were tacked on. These apply equally to float and integer elements. But nothing for 64-bit integer arithmetic.

**Note**

The full title in PowerISA 2.06 is **Vector-Scalar Floating-Point Operations [Category: VSX]**.

PowerISA 2.07 (POWER8) did add a significant number of doubleword (64-bit) integer operations. Including;

- Add and subtract modulo
- Signed and unsigned compare, maximum, minimum,
- Shift and rotate
- Count leading zeros and population count

Also a number of new word (32-bit) integer operations;

- Multiply even/odd/modulo.
- Pack signed/unsigned/saturate and Unpack signed.
- Merge even/odd words

And some new quadword (128-bit) integer operations;

- Add and Subtract modulo/extend/write-carry
- Decimal Add and Subtract modulo

And some specialize operations;

- Crypto, Raid, Polynomial multiply-sum

**Note**

The operations above are all Vector Category and can only access the 32 original vector registers (VSRs 32-63).

The new VSX operations (with access to all 64 VSRs) where not directly applicable to 64-bit integer arithmetic:

- Scalar single precision floating-point
- Direct move between GPRs and VSRs
- Logical operations; equivalence, not and, or compliment

PowerISA 3.0 (POWER9) adds a few more doubleword (64-bit) integer operations. Including;

- Compare not equal
- Count Trailing zeros and parity
- Extract and Insert
- Multiply-sum modulo
- Negate

- Rotate Left under mask

Also a number of new word (32-bit) integer operations;

- Absolute Difference word
- Extend Sign word to doubleword

And some new quadword (128-bit) integer operations;

- Multiply-by-10 extend/write-carry
- Decimal convert from/to signed (binary) quadword
- Decimal convert from/to zoned (ASCII char)
- Decimal shift/round/truncate

The new VSX operations (with access to all 64 VSRs) where not directly applicable to 64-bit integer arithmetic:

- Scalar quad-precision floating-point
- Scalar and Vector convert with rounding
- Scalar and Vector extract/insert exponent/significand
- Scalar and Vector test data class
- Permute and Permute right index

An impressive list of operation that can be used for;

- Vectorizing long integer loops
- Implementing useful quadword integer operations which do not have corresponding PowerISA instructions
- implementing extended precision multiply and multiplicative inverse operations

The challenge is that useful operations available for POWER9 will need equivalent implementations for POWER8 and POWER7. Similarly for operations introduced for POWER8 will need POWER7 implementations. Also there are some obvious missing operations;

- Absolute Difference doubleword (we have byte, halfword, and word)
- Average doubleword (we have byte, halfword, and word)
- Extend Sign doubleword to quadword (we have byte, halfword, and word)
- Multiply-sum word (we have byte, halfword, and doubleword)
- Multiply even/odd doublewords (we have byte, halfword, and word)

### 7.10.2.1 Challenges and opportunities

The stated goals for pveclib are:

- Provide equivalent functions across versions of the compiler.
- Provide equivalent functions across versions of the PowerISA.
- Provide complete arithmetic operations across supported C types.

So the first step is to provide implementations for the key POWER8 doubleword integer operations for older compilers. For example some of the generic doubleword integer operations were not defined until GCC 6.0. Here we define the specific Compare Equal Unsigned Doubleword implementation:

```
static inline
vb64_t
vec_cmpequd (vui64_t a, vui64_t b)
{
    vb64_t result;
#ifdef _ARCH_PWR8
    if __GNUC__ >= 6
        result = vec_cmpeq(a, b);
    else
        __asm__(
            "vcpequd %0,%1,%2;\n"
            : "=v" (result)
            : "v" (a),
              "v" (b)
            : );
#endif
    // _ARCH_PWR7 implementation ...
    return (result);
}
```

The implementation checks if the compile target is POWER8 then checks if the compiler is new enough to use the generic vector compare built-in. If the generic built-in is not defined in `<altivec.h>` then we provide the equivalent inline assembler.

For POWER7 targets we don't have any vector compare doubleword operations and we need to define the equivalent operation using PowerISA 2.06B (and earlier) instructions. For example:

```
#else
// _ARCH_PWR7 implementation ...
vui8_t permute =
{ 0x04,0x05,0x06,0x07, 0x00,0x01,0x02,0x03,
  0x0C,0x0D,0x0E,0x0F, 0x08,0x09,0x0A,0x0B};
vui32_t r, rr;
r = (vui32_t)vec_cmpeq((vui32_t)a, (vui32_t)b);
if (vec_any_ne((vui32_t)a, (vui32_t)b))
{
    rr = vec_perm(r, r, permute);
    r = vec_and(r, rr);
}
result = (vb64_t)r;
#endif
```

Here we use compare equal unsigned word. If all words are equal use the result as is. Otherwise, if any word elements are not equal we do some extra work. For each doubleword, rotate the word compare result by 32-bits (here we use permute as we don't have rotate doubleword either). Then logical and the original word compare and rotated results to get the final doubleword compare results.

Similarly for all the doubleword compare variants. Similarly for doubleword; add, subtract, maximum, minimum, shift, rotate, count leading zeros, population count, and Byte reverse.



## 7.10.2.2 More Challenges

Now we can look at the case where vector doubleword operations, of interest, don't have an equivalent instruction. Here interesting operations include those that are supported for other element sizes and types.

The simplest example is absolute difference which was introduced in PowerISA 3.0 for byte, halfword and word elements. From the implementation of `vec_absduw()` we see how to implement the operation for POWER8 using subtract, maximum, and minimum. For example:

```
static inline vui64_t
vec_absdud (vui64_t vra, vui64_t vrb)
{
    return vec_subudm (vec_maxud (vra, vrb), vec_minud (vra, vrb));
}
```

This works because pveclib provides implementations for min, max, and sub operations that work across GCC versions and provide processor specific implementations for POWER8/9 and POWER7.

Now we need to look at the multiply doubleword situation. We need implementations for `vec_msumudm()`, `vec_muleud()`, `vec_mulhud()`, `vec_muloud()`, and `vec_muludm()`. As we saw in the implementations of `vec_int32_ppc.h` that multiply high and low/modulo can be implemented using multiply and merge even/odd of that element size. Multiply low can also be implemented using the multiply-sum and multiply odd of the next smaller element size. Also multiply-sum can be implemented using multiply even/odd and a couple of adds. And multiply even/odd can be implemented using multiply-sum by supplying zeros to appropriate inputs/elements.

The above discussion has many circular dependencies. Eventually we need to get down to an implementation on each processor using actual hardware instructions. So what multiply doubleword operations, does the PowerISA actually have from the list above:

- POWER9 added multiply-sum unsigned doubleword modulo but no multiply doubleword even/odd/modulo instructions.
- POWER8 added multiply even/odd/modulo word but no multiply-sum word instructions
- POWER7 and earlier we have the original VMX multiply even/odd halfword, and multiply-sum unsigned halfword modulo, but no multiply modulo halfword.

It seems the best implementation strategy uses;

- Multiply-sum doubleword for POWER9
- Multiply even/odd word for POWER8
- Multiply even/odd halfword for POWER7

We really care about performance and latency for POWER9/8. We need POWER7 to work correctly so we can test on and support *legacy* hardware. The rest is grade school math.

First we need to make sure we have implementations across the GCC versions 6, 7, and 8 for the instructions we need. For example:

```
static inline vuil28_t
vec_msumudm (vui64_t a, vui64_t b, vuil28_t c)
{
    vuil28_t res;
#ifdef _ARCH_PWR9
    __asm__ (
        "vmsumudm %0,%1,%2,%3;\n"
        : "=v" (res)
        : "v" (a), "v" (b), "v" (c)
        : );
#else
    vuil28_t p_even, p_odd, p_sum;

    p_even = vec_muleud (a, b);
    p_odd = vec_muloud (a, b);
    p_sum = vec_adduqm (p_even, p_odd);
    res = vec_adduqm (p_sum, c);
#endif
    return (res);
}
```

**Note**

The `_ARCH_PWR8` implementation above depends on `vec_muleud()` and `vec_muloud()` for which there are no hardware instructions. Hold that thought.

While we are it we can implement multiply-sum unsigned word modulo.

```
static inline vui64_t
vec_vmsumuwmm (vui32_t vra, vui32_t vrb, vui64_t vrc)
{
    vui64_t peven, podd, psum;

    peven = vec_muleuw (vra, vrb);
    podd = vec_mulouw (vra, vrb);
    psum = vec_addudm (peven, podd);

    return vec_addudm (psum, vrc);
}
```

We will need this later.

Now we need to provide implementations of `vec_muleud()` and `vec_muloud()`. For example:

```
static inline vui128_t
vec_muleud (vui64_t a, vui64_t b)
{
    #if __BYTE_ORDER__ == __ORDER_LITTLE_ENDIAN__
        return vec_vmuloud (a, b);
    #else
        return vec_vmuleud (a, b);
    #endif
}
```

The implementation above is just handling the pesky little endian transforms. The real implementations are in `vec_vmuleud()` and `vec_vmuloud()` which implement the operation as if the PowerISA included such an instruction. These implementation is NOT endian sensitive and the function is stable across BE/LE implementations. For example:

```
static inline vui128_t
vec_vmuleud (vui64_t a, vui64_t b)
{
    vui64_t res;
    #ifdef _ARCH_PWR9
        const vui64_t zero = { 0, 0 };
        vui64_t b_eud = vec_mrgahd ((vui128_t) b, (vui128_t) zero);
        __asm__(
            "vmsumudm %0,%1,%2,%3;\n"
            : "=v" (res)
            : "v" (a), "v" (b_eud), "v" (zero)
            : );
    #else
        #ifdef _ARCH_PWR8
            const vui64_t zero = { 0, 0 };
            vui64_t p0, p1, pp10, pp01;
            vui32_t m0, m1;

            // Need the endian invariant merge word high here
            #if __BYTE_ORDER__ == __ORDER_LITTLE_ENDIAN__
                // Nullify the little endian transform
                m0 = vec_mergel ((vui32_t) b, (vui32_t) b);
            #else
                m0 = vec_mergeh ((vui32_t) b, (vui32_t) b);
            #endif
            m1 = (vui32_t) vec_xxspltd ((vui64_t) a, 0);

            // Need the endian invariant multiply even/odd word here
            #if __BYTE_ORDER__ == __ORDER_LITTLE_ENDIAN__
                // Nullify the little endian transform
                p1 = vec_muleuw (m1, m0);
                p0 = vec_mulouw (m1, m0);
            #else
                p0 = vec_muleuw (m1, m0);
                p1 = vec_mulouw (m1, m0);
            #endif
            // res[1] = p1[1]; res[0] = p0[0];
        #endif
    #endif
}
```

```

res = vec_pasted (p0, p1);

// pp10[1] = p1[0]; pp10[0] = 0;
// pp01[1] = p0[1]; pp01[0] = 0;
// Need the endian invariant merge algebraic high/low here
pp10 = (vui64_t) vec_mrgahd ((vui128_t) zero, (
    vui128_t) p1);
pp01 = (vui64_t) vec_mrgald ((vui128_t) zero, (
    vui128_t) p0);
// pp01 = pp01 + pp10.
pp01 = (vui64_t) vec_adduqm ((vui128_t) pp01, (
    vui128_t) pp10);

// res = res + (pp01 << 32)
pp01 = (vui64_t) vec_sld ((vi32_t) pp01, (vi32_t) pp01, 4);
res = (vui64_t) vec_adduqm ((vui128_t) pp01, (
    vui128_t) res);
#else
// _ARCH_PWR7 implementation ...
#endif
#endif
return ((vui128_t) res);
}

```

The `_ARCH_PWR9` implementation uses the multiply-sum doubleword operation but implements the multiply even behavior by forcing the contents of doubleword element 1 of [VRB] and the contents of [VRC] to 0.

The `_ARCH_PWR8` implementation looks ugly but it works. It starts with some merges and splats to get inputs columns lined up for the multiply. Then we use (POWER8 instructions) Multiply even/odd unsigned word to generate doubleword partial products. Then more merges and a rotate to line up the partial products for summation as the final quadword product.

Individually `vec_vmuleud()` and `vec_vmuloud()` execute with a latency of 21-23 cycles on POWER8. Normally these operations are used and scheduled together as in the POWER8 implementation of `vec_msumudm()` or `vec_mulhud()`. Good scheduling by the compiler and pipelining keeps the POWER8 latency in the 28-32 cycle range. For example the `vec_mulhud()` implementation:

```

static inline vui64_t
vec_mulhud (vui64_t vra, vui64_t vrb)
{
    return vec_mrgahd (vec_vmuleud (vra, vrb), vec_vmuloud (vra, vrb));
}

```

Generates the following code for POWER8:

```

vspltisw v0,0
xxmrghw vs33,vs35,vs35
xxspltd vs45,vs34,0
xxmrghw vs35,vs35,vs35
vmulouw v11,v13,v1
xxspltd vs34,vs34,1
xxmrghd vs41,vs32,vs43
vmulouw v12,v2,v3
vmuleuw v13,v13,v1
vmuleuw v2,v2,v3
xxmrghd vs42,vs32,vs44
xxmrghd vs33,vs32,vs45
xxmrghd vs32,vs32,vs34
xxpermdi vs44,vs34,vs44,1
vadduqm v1,v1,v9
xxpermdi vs45,vs45,vs43,1
vadduqm v0,v0,v10
vsldoi v1,v1,v1,4
vsldoi v0,v0,v0,4
vadduqm v2,v1,v13
vadduqm v0,v0,v12
xxmrghd vs34,vs34,vs32

```

The POWER9 latencies for this operations range from 5-7 (for `vmsumudm` itself) to 11-16 (for `vec_mulhud()`). The additional latency reflects generate of zero constant vector and merges requires to condition the inputs and output. For these operations the `vec_msumudm()`, `vrc` operand is always zero. Selecting the even/odd doubleword for input requires a merge low/high. And selecting the high doubleword for multiply high require a final merge high.

`vec_mulhud()` generates the following code for POWER9:

```

xxspltib vs32,0
xxmrghd vs33,vs35,vs32
xxmrghd vs35,vs32,vs35
vmsumudm v1,v2,v1,v0
vmsumudm v2,v2,v3,v0
xxmrghd vs34,vs33,vs34

```

Wrapping up the doubleword multiplies we should look at the multiply low (AKA Multiply Unsigned Doubleword Modulo). The POWER9 implementation is similar to `vec_mulhud()` and the generated code is similar to the example above.

Multiply low doubleword is a special case, as we are discarding the highest partial doubleword product. For POWER8 we can optimize for that case using multiply odd and multiply-sum word operations. Also as we are only generating doubleword partial products we only need add doubleword modulo operations to sum the results. This avoids the more expensive add quadword operation required for the general case. The fact that `vec_vmsumuwm()` is only a software construct is not an issue. It expands into hardware multiple even/odd word and add doubleword instructions that the compiler can schedule and optimize.

Here `vec_mulouw()` generates low order partial product. Then `vec_vrld()` and `vec_vmsumuwm()` generate doubleword sums of the two middle order partial products. Then `vec_vslld()` shifts the middle order partial sum left 32-bits (discarding the unneeded high order 32-bits). Finally sum the low and middle order partial doubleword products to produce the multiply-low doubleword result. For example this POWER8 only implementation:

```

static inline vui64_t
vec_muludm (vui64_t vra, vui64_t vrb)
{
    vui64_t s32 = { 32, 32 }; // shift / rotate amount.
    vui64_t z = { 0, 0 };
    vui64_t t2, t3, t4;
    vui32_t t1;

    t1 = (vui32_t)vec_vrld (vrb, s32);
    #if __BYTE_ORDER__ == __ORDER_LITTLE_ENDIAN__
        // Nullify the little endian transform, really want mulouw here.
        t2 = vec_muleuw ((vui32_t)vra, (vui32_t)vrb);
    #else
        t2 = vec_mulouw ((vui32_t)vra, (vui32_t)vrb);
    #endif
    t3 = vec_vmsumuwm ((vui32_t)vra, t1, z);
    t4 = vec_vslld (t3, s32);
    return (vui64_t)vec_vaddudm (t4, t2);
}

```

Which generates the following for POWER8:

```

addis    r9,r2,.rodata.cst16+0x60@ha
addi     r9,r9,.rodata.cst16+0x60@l
lxv      vs33,0,r9
vmulouw  v13,v2,v3
vrld     v0,v3,v1
vmulouw  v3,v2,v0
vmuleuw  v2,v2,v0
vaddudm  v2,v3,v2
vslld    v2,v2,v1
vaddudm  v2,v13,v2

```

#### Note

The addition of zeros to the final sum of `vec_vmsumuwm()` (`vec_addudm (psum, vrc)`) has been optimized away by the compiler. This eliminates the `xxspltib` and one `vaddudm` instruction from the final code sequence.

And we can assume that the constant load of `{ 32, 32 }` will be common-ed with other operations or hoisted out of loops. So the shift constant can be loaded early and `vrld` is not delayed. This keeps the POWER8 latency in the 19-28 cycle range.

### 7.10.3 Endian problems with doubleword operations

From the examples above we see that the construction of higher precision multiplies requires significant massaging of input and output elements. Here merge even/odd, merge high/low, swap, and splat doubleword element operations are commonly used.

PowerISA 2.06 VSX (POWER7) added the general purpose Vector Permute Doubleword Immediate (xxpermdi). The compiler generates some form of xxpermdi for the doubleword (double float, long int, bool long) merge/splat/swap operations. As xxpermdi's element selection is an immediate field, most operations require only a single instruction. All the merge/splat/swap doubleword variants are just a specific select mask value and the inputs to xxpermdi.

Which is very useful indeed for assembling, disassembling, merging, splatting, swapping, and pasting doubleword elements.

Of course it took several compiler releases to implement all the generic merge/splat/swap operations for the supported types. GCC 4.8 was the first to support vec\_xxpermdi as a built-in. GCC 4.8 also supported the generic built-ins vec\_mergeh, vec\_mergel, and vec\_splat for the vector signed/unsigned/bool long type. But endian sensitive vec\_mergeh, vec\_mergel, and vec\_splat were not supported until GCC 7. And the generic vec\_mergee, vec\_mergeo built-ins were not supported until GCC 8.

But as we have explained in [General Endian Issues](#) and [Endian problems with word operations](#) the little endian transforms applied by the compiler can cause problems for developers of multi-precision libraries. The doubleword forms of the generic merge/splat operations etc. are no exception. This is especially annoying when the endian sensitive transformd are applied between releases of the compiler.

So we need a strategy to provide endian invariant merge/splat/swap operations to be used in multi-precision arithmetic. And another set of endian sensitive operations that are mandated by the OpenPOWER ABI.

First we need a safely endian invariant version of xxpermdi to use in building other variants:

- [vec\\_permdi\(\)](#) provides the basic xxpermdi operation but nullifies the little endian transforms.

Then build the core set of endian invariant permute doubleword operations using [vec\\_permdi\(\)](#):

- Merge algebraic high/low doubleword operations [vec\\_mrgahd\(\)](#) and [vec\\_mrgald\(\)](#).
- Merge the left and right most doublewords from a double quadword operation [vec\\_pasted\(\)](#).
- Splat from the high/even or low/odd doubleword operation [vec\\_xxspltd\(\)](#).
- Swap high and low doublewords operation [vec\\_swapd\(\)](#).

We use the merge algebraic high/low doubleword operations in the implementation of [vec\\_mulhud\(\)](#), [vec\\_mulhud\(\)](#), [vec\\_vmuleud\(\)](#), and [vec\\_vmuloud\(\)](#). We use the vec\_xxspltd operation in the implementation of [vec\\_vrld\(\)](#), [vec\\_vrld\(\)](#), [vec\\_vrld\(\)](#), and [vec\\_vmuloud\(\)](#). We use the paste doubleword ([vec\\_pasted\(\)](#)) operation in the implementation of [vec\\_vsrld\(\)](#), [vec\\_vmuleud\(\)](#), and [vec\\_vmuloud\(\)](#). We use the swap doubleword operation in the implementation of [vec\\_cmpequq\(\)](#), [vec\\_cmpnequq\(\)](#), [vec\\_muludq\(\)](#), and [vec\\_mulluq\(\)](#).

Then use the compilers `__BYTE_ORDER__ = ORDER_LITTLE_ENDIAN` conditional to invert the [vec\\_permdi\(\)](#) select control for endian sensitive merge/splat doubleword operations:

- Merge even/odd doubleword operations [vec\\_mrged\(\)](#) and [vec\\_mrgod\(\)](#).
- Merge high/low doubleword operations [vec\\_mrghd\(\)](#) and [vec\\_mrgld\(\)](#).
- Splat even/odd doubleword operation [vec\\_spltd\(\)](#).

### 7.10.4 Vector Doubleword Examples

Suppose we have a requirement to convert an array of 64-bit time-interval values that need to convert to timespec format. For simplicity we will also assume that the array is nicely (Quadword) aligned and an integer multiple of 2 doublewords or 4 words.

The PowerISA provides a 64-bit TimeBase register that clocks at a constant 512MHz. The TimeBase can be read directly as either the full 64-bit value or as 32-bit upper and lower halves. For this example we assume are dealing with longer intervals (greater than  $\sim 8.38$  seconds) so the full 64-bit TimeBase is required. TimeBase values of adjacent events are subtracted to generate the intervals stored in the array.

The timespec format is a struct of unsigned int fields for seconds and nanoseconds. So the task is to convert the 512MHz 64-bit TimeBase intervals to seconds and remaining clock ticks. Then convert the remaining (subsecond) clock ticks from 512MHz to nanoseconds. The separate seconds and nanoseconds are combined in the timespec structure.

First we need to separate the raw TimeBase into the integer seconds and (subsecond) clock-ticks. Normally scalar codes would use integer divide/modulo by 512000000. Did I mention that the PowerISA vector unit does not have a integer divide operation?

Instead we can use the multiplicative inverse which is a scaled fixed point fraction calculated from the original divisor. This works nicely if the fixed radix point is just before the 64-bit fraction and we have a multiply high ([vec\\_mulhud\(\)](#)) operation. Multiplying a 64-bit unsigned integer by a 64-bit unsigned fraction generates a 128-bit product with 64-bits above (integer) and below (fraction) the radix point. The high 64-bits of the product is the integer quotient.

It turns out that generating the multiplicative inverse can be tricky. To produce correct results over the full range requires, possible pre-scaling and post-shifting, and sometimes a corrective addition is necessary. Fortunately the mathematics are well understood and are commonly used in optimizing compilers. Even better, Henry Warren's book has a whole chapter on this topic.

#### See also

"Hacker's Delight, 2nd Edition," Henry S. Warren, Jr, Addison Wesley, 2013. Chapter 10, Integer Division by Constants.

In the chapter above;

Figure 10-2 Computing the magic number for unsigned division.

provides a sample C function for generating the magic number (actually a struct containing; the magic multiplicative inverse, "add" indicator, and the shift amount.).

For the divisor 512000000 this is { 4835703278458516699, 0 , 27 }:

- the multiplier is 4835703278458516699.
- no corrective add of the dividend is required.
- the final shift is 27-bits right.

```
// Magic numbers for multiplicative inverse to divide by 512,000,000
// are 4835703278458516699 and shift right 27 bits.
const vui64_t mul_invs_clock =
    { 4835703278458516699UL, 4835703278458516699UL };
const int shift_clock = 27;
// Need const for TB clocks/second to extract remainder.
const vui64_t tb_clock_sec =
    { 512000000, 512000000 };
vui64_t tb_v, tmp, tb_clocks, seconds, nseconds;
vui64_t timespec1, timespec2;

// extract integer seconds from timebase vector.
tmp = vec_mulhud (tb_v, mul_invs_clock);
seconds = vec_srdi (tmp, shift_clock);
// Extract the remainder in tb clock ticks.
tmp = vec_muludm (seconds, tb_clock_sec);
tb_clocks = vec_sub (tb_v, tmp);
```

Next we need to convert the subseconds from TimeBase clock-ticks to nanoseconds. The subsecond remainder is now small enough (compared to a doubleword) that we can perform the conversion *in place*. The nanosecond conversion is  $((tb\_clocks * 1000000000) / 512000000)$ . And we can reduce this to  $((tb\_clocks * 1000) / 512)$ . We still have a small multiply but the divide can be converted to shift right of 9-bits.

```
const int shift_512 = 9;
const vui64_t nano_512 =
    { 1000, 1000 };

// Convert 512MHz timebase to nanoseconds.
// nseconds = tb_clocks * 1000000000 / 512000000
// reduces to (tb_clocks * 1000) >> 9
tmp = vec_muludm (tb_clocks, nano_512);
nseconds = vec_srdi (tmp, shift_512);
```

Finally we need to merge the vectors of seconds and nanoseconds into vectors of timespec. So far we have been working with 64-bit integers but the timespec is a struct of 32-bit (word) integers. Here 32-bit seconds and nanosecond provided sufficient range and precision. So the final step *packs* a pair of 64-bit timespec values into a vector of two 32-bit timespec values, each containing 2 32-bit (second, nanosecond) values.

```
timespec1 = vec_mergeh (seconds, nseconds);
timespec2 = vec_mergel (seconds, nseconds);
// seconds and nanoseconds fit int 32-bits after conversion.
// So pack the results and store the timespec.
*timespec++ = vec_vp kudum (timespec1, timespec2);
```

#### Note

`vec_sub()`, `vec_mergeh()`, and `vec_mergel()` are existing `altivec.h` generic built-ins. `vec_vp kudum()` is an existing `altivec.h` built-in that is only defined for **\_ARCH\_PWR8** and later. This header insures that `vec_vp kudum` is defined for older compilers and provides an functional equivalent implementation for POWER7.

#### 7.10.4.1 Vectorized 64-bit TimeBase conversion example

Here is the complete vectorized 64-bit TimeBase to timespec conversion example:

```
void
example_dw_convert_timebase (vui64_t *tb, vui32_t *timespec, int n)
{
    // Magic numbers for multiplicative inverse to divide by 512,000,000
    // are 4835703278458516699 and shift right 27 bits.
    const vui64_t mul_invs_clock =
        { 4835703278458516699UL, 4835703278458516699UL };
    const int shift_clock = 27;
    // Need const for TB clocks/second to extract remainder.
    const vui64_t tb_clock_sec =
        { 512000000, 512000000 };
    vui64_t tb_v, tmp, tb_clocks, seconds, nseconds;
    vui64_t timespec1, timespec2;

    // extract integer seconds from timebase vector.
    tmp = vec_mulhud (tb_v, mul_invs_clock);
    seconds = vec_srdi (tmp, shift_clock);
    // Extract the remainder in tb clock ticks.
    tmp = vec_muludm (seconds, tb_clock_sec);
    tb_clocks = vec_sub (tb_v, tmp);

    // Convert 512MHz timebase to nanoseconds.
    // nseconds = tb_clocks * 1000000000 / 512000000
    // reduces to (tb_clocks * 1000) >> 9
    tmp = vec_muludm (tb_clocks, nano_512);
    nseconds = vec_srdi (tmp, shift_512);

    timespec1 = vec_mergeh (seconds, nseconds);
    timespec2 = vec_mergel (seconds, nseconds);
    // seconds and nanoseconds fit int 32-bits after conversion.
    // So pack the results and store the timespec.
    *timespec++ = vec_vp kudum (timespec1, timespec2);
}
```

```

const int shift_512 = 9;
const vui64_t nano_512 =
{ 1000, 1000};
vui64_t tb_v, tmp, tb_clocks, seconds, nseconds;
vui64_t timespec1, timespec2;
int i;

for (i = 0; i < n; i++)
{
    tb_v = *tb++;
    // extract integer seconds from timebase vector.
    tmp = vec_mulhud (tb_v, mul_invs_clock);
    seconds = vec_srdi (tmp, shift_clock);
    // Extract remainder in tb clock ticks.
    tmp = vec_muludm (seconds, tb_clock_sec);
    tb_clocks = vec_sub (tb_v, tmp);
    // Convert 512MHz timebase to nanoseconds.
    // nseconds = tb_clocks * 1000000000 / 512000000
    // reduces to (tb_clocks * 1000) >> 9
    tmp = vec_muludm (tb_clocks, nano_512);
    nseconds = vec_srdi (tmp, shift_512);
    // Use merge high/low to interleave seconds and nseconds
    // into timespec.
    timespec1 = vec_mergeh (seconds, nseconds);
    timespec2 = vec_mergel (seconds, nseconds);
    // seconds and nanoseconds fit int 32-bits after conversion.
    // So pack the results and store the timespec.
    *timespec++ = vec_vp kudum (timespec1, timespec2);
}

```

### 7.10.5 Performance data.

High level performance estimates are provided as an aid to function selection when evaluating algorithms. For background on how *Latency* and *Throughput* are derived see: [Performance data](#).

### 7.10.6 Function Documentation

#### 7.10.6.1 static vui64\_t vec\_absdud ( vui64\_t vra, vui64\_t vrb ) [inline],[static]

Vector Absolute Difference Unsigned doubleword.

Compute the absolute difference for each doubleword. For each unsigned doubleword, subtract VRB[i] from VRA[i] and return the absolute value of the difference.

processor	Latency	Throughput
power8	4	1/cycle
power9	5	1/cycle

#### Parameters

<i>vra</i>	vector of 2 x unsigned doublewords
<i>vrb</i>	vector of 2 x unsigned doublewords

#### Returns

vector of the absolute differences.



### 7.10.6.2 static vui64\_t vec\_addudm ( vui64\_t a, vui64\_t b ) [inline],[static]

Vector Add Unsigned Doubleword Modulo.

Add two vector long int values and return modulo 64-bits result.

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

#### Parameters

<i>a</i>	128-bit vector long int.
<i>b</i>	128-bit vector long int.

#### Returns

vector long int sums of a and b.

### 7.10.6.3 static vui64\_t vec\_clzd ( vui64\_t vra ) [inline],[static]

Count leading zeros for a vector unsigned long int.

Count leading zeros for a vector \_\_int128 and return the count in a vector suitable for use with vector shift (left|right) and vector shift (left|right) by octet instructions.

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

#### Parameters

<i>vra</i>	a 128-bit vector treated a __int128.
------------	--------------------------------------

#### Returns

a 128-bit vector with bits 121:127 containing the count of leading zeros.

### 7.10.6.4 static vb64\_t vec\_cmpeqsd ( vi64\_t a, vi64\_t b ) [inline],[static]

Vector Compare Equal Signed Doubleword.

Compare each signed long (64-bit) integer and return all '1's, if a[i] == b[i], otherwise all '0's.

For POWER8 (PowerISA 2.07B) or later, use the Vector Compare Equal Unsigned DoubleWord (**vcmpequd**) instruction. Otherwise use boolean logic using word compares.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

#### Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

#### Returns

128-bit vector with each dword boolean reflecting compare equal result for each element.

**7.10.6.5** `static vb64_t vec_cmpequd ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare Equal Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return all '1's, if  $a[i] == b[i]$ , otherwise all '0's.

For POWER8 (PowerISA 2.07B) or later, use the Vector Compare Equal Unsigned DoubleWord (**vcampequd**) instruction. Otherwise use boolean logic using word compares.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

#### Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

#### Returns

128-bit vector with each dword boolean reflecting compare equal result for each element.

**7.10.6.6** `static vb64_t vec_cmpgesd ( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare Greater Than or Equal Signed Doubleword.

Compare each signed long (64-bit) integer and return all '1's, if  $a[i] \geq b[i]$ , otherwise all '0's. Use `vec_cmpgtsd` with parameters reversed to implement `vec_cmpltud`, then return the logical inverse.

processor	Latency	Throughput
power8	4	2/cycle
power9	5	2/cycle

## Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

## Returns

128-bit vector with each dword boolean reflecting compare greater then or equal result for each element.

7.10.6.7 `static vb64_t vec_cmpgeud ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare Greater Than or Equal Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return all '1's, if  $a[i] \geq b[i]$ , otherwise all '0's. Use `vec_cmpgtud` with parameters reversed to implement `vec_cmpltud`, then return the logical inverse.

processor	Latency	Throughput
power8	4	2/cycle
power9	5	2/cycle

## Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

## Returns

128-bit vector with each dword boolean reflecting compare greater then or equal result for each element.

7.10.6.8 `static vb64_t vec_cmpgtsd ( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare Greater Than Signed Doubleword.

Compare each signed long (64-bit) integer and return all '1's, if  $a[i] > b[i]$ , otherwise all '0's.

For POWER8 (PowerISA 2.07B) or later use the Vector Compare Greater Than Unsigned DoubleWord (**vcmpgtud**) instruction. Otherwise use boolean logic using word compares.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

## Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

**Returns**

128-bit vector with each dword boolean reflecting compare greater result for each element.

**7.10.6.9** `static vb64_t vec_cmpgtud ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare Greater Than Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return all '1's, if  $a[i] > b[i]$ , otherwise all '0's.

For POWER8 (PowerISA 2.07B) or later use the Vector Compare Greater Than Unsigned DoubleWord (**vcmpgtud**) instruction. Otherwise use boolean logic using word compares.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

**Returns**

128-bit vector with each dword boolean reflecting compare greater result for each element.

**7.10.6.10** `static vb64_t vec_cmpletd ( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare Less Than Equal Signed Doubleword.

Compare each signed long (64-bit) integer and return all '1's, if  $a[i] > b[i]$ , otherwise all '0's. Use `vec_cmpgtsd` with parameters reversed to implement `vec_cmpletd` then return the logical inverse.

processor	Latency	Throughput
power8	4	2/cycle
power9	5	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

**Returns**

128-bit vector with each dword boolean reflecting compare greater result for each element.

7.10.6.11 `static vb64_t vec_cmpleud ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare Less Than Equal Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return all '1's, if  $a[i] > b[i]$ , otherwise all '0's. Use `vec_cmpltud` with parameters reversed to implement `vec_cmpleud`. Use `vec_cmpltud` then return the logical inverse.

processor	Latency	Throughput
power8	4	2/cycle
power9	5	2/cycle

#### Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

#### Returns

128-bit vector with each dword boolean reflecting compare greater result for each element.

7.10.6.12 `static vb64_t vec_cmpltsd ( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare less Than Signed Doubleword.

Compare each signed long (64-bit) integer and return all '1's, if  $a[i] < b[i]$ , otherwise all '0's. Use `vec_cmpltgd` with parameters reversed to implement `vec_cmpltsd`.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

#### Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

#### Returns

128-bit vector with each dword boolean reflecting compare less result for each element.

7.10.6.13 `static vb64_t vec_cmpltud ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare less Than Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return all '1's, if  $a[i] < b[i]$ , otherwise all '0's. Use `vec_cmpltgd` with parameters reversed to implement `vec_cmpltud`.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

**Returns**

128-bit vector with each dword boolean reflecting compare less result for each element.

**7.10.6.14** `static vb64_t vec_cmpnesd( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare Not Equal Signed Doubleword.

Compare each signed long (64-bit) integer and return all '1's, if  $a[i] \neq b[i]$ , otherwise all '0's. Use `vec_cmpequd` then return the logical inverse.

processor	Latency	Throughput
power8	4	2/cycle
power9	5	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

**Returns**

128-bit vector with each dword boolean reflecting compare not equal result for each element.

**7.10.6.15** `static vb64_t vec_cmpneud( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare Not Equal Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return all '1's, if  $a[i] \neq b[i]$ , otherwise all '0's. Use `vec_cmpequd` then return the logical inverse.

processor	Latency	Throughput
power8	4	2/cycle
power9	5	2/cycle

## Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

## Returns

128-bit vector with each dword boolean reflecting compare not equal result for each element.

7.10.6.16 `static int vec_cmpsd_all_eq( vi64_t a, vi64_t b ) [inline], [static]`

Vector Compare all Equal Signed Doubleword.

Compare each signed long (64-bit) integer and return true if all elements of a and b are equal.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

## Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

## Returns

boolean int for all 128-bits, true if equal, false otherwise.

7.10.6.17 `static int vec_cmpsd_all_ge( vi64_t a, vi64_t b ) [inline], [static]`

Vector Compare all Greater Than or Equal Signed Doubleword.

Compare each signed long (64-bit) integer and return true if all elements of a  $\geq$  b.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

## Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

**Returns**

boolean int for all 128-bits, true if all Greater Than, false otherwise.

**7.10.6.18** `static int vec_cmpsd_all_gt ( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare all Greater Than Signed Doubleword.

Compare each signed long (64-bit) integer and return true if all elements of  $a > b$ .

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

**Returns**

boolean int for all 128-bits, true if all Greater Than, false otherwise.

**7.10.6.19** `static int vec_cmpsd_all_le ( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare all Less than equal Signed Doubleword.

Compare each signed long (64-bit) integer and return true if all elements of  $a \leq b$ .

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

**Returns**

boolean int for all 128-bits, true if all Greater Than, false otherwise.

**7.10.6.20** `static int vec_cmpsd_all_lt ( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare all Less than Signed Doubleword.

Compare each signed long (64-bit) integer and return true if all elements of  $a < b$ .



processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

**Returns**

boolean int for all 128-bits, true if all Greater Than, false otherwise.

**7.10.6.21** `static int vec_cmpsd_all_ne ( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare all Not Equal Signed Doubleword.

Compare each signed long (64-bit) integer and return true if all elements of a and b are not equal.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

**Returns**

boolean int for all 128-bits, true if equal, false otherwise.

**7.10.6.22** `static int vec_cmpsd_any_eq ( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare any Equal Signed Doubleword.

Compare each signed long (64-bit) integer and return true if any elements of a and b are equal.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
----------	--

**Parameters**

<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
----------	--

**Returns**

boolean int for all 128-bits, true if equal, false otherwise.

7.10.6.23 `static int vec_cmpsd_any_ge ( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare any Greater Than or Equal Signed Doubleword.

Compare each signed long (64-bit) integer and return true if any elements of a >= b.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

**Returns**

boolean int for all 128-bits, true if all Greater Than, false otherwise.

7.10.6.24 `static int vec_cmpsd_any_gt ( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare any Greater Than Signed Doubleword.

Compare each signed long (64-bit) integer and return true if all elements of a > b.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

**Returns**

boolean int for all 128-bits, true if all Greater Than, false otherwise.

7.10.6.25 `static int vec_cmpsd_any_le( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare any Less than equal Signed Doubleword.

Compare each signed long (64-bit) integer and return true if any elements of  $a \leq b$ .

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

#### Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

#### Returns

boolean int for any 128-bits, true if any Greater Than, false otherwise.

7.10.6.26 `static int vec_cmpsd_any_lt( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare any Less than Signed Doubleword.

Compare each signed long (64-bit) integer and return true if any elements of  $a < b$ .

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

#### Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

#### Returns

boolean int for any 128-bits, true if any Greater Than, false otherwise.

7.10.6.27 `static int vec_cmpsd_any_ne( vi64_t a, vi64_t b ) [inline],[static]`

Vector Compare any Not Equal Signed Doubleword.

Compare each signed long (64-bit) integer and return true if any elements of  $a$  and  $b$  are not equal.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit signed long integer (dword) elements.

**Returns**

boolean int for any 128-bits, true if equal, false otherwise.

**7.10.6.28** `static int vec_cmpud_all_eq ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare all Equal Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return true if all elements of a and b are equal.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

**Returns**

boolean int for all 128-bits, true if equal, false otherwise.

**7.10.6.29** `static int vec_cmpud_all_ge ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare all Greater Than or Equal Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return true if all elements of a >= b.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

**Returns**

boolean int for all 128-bits, true if all Greater Than, false otherwise.

**7.10.6.30** `static int vec_cmpud_all_gt ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare all Greater Than Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return true if all elements of a > b.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

**Returns**

boolean int for all 128-bits, true if all Greater Than, false otherwise.

**7.10.6.31** `static int vec_cmpud_all_le ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare all Less than equal Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return true if all elements of a <= b.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

**Returns**

boolean int for all 128-bits, true if all Greater Than, false otherwise.

**7.10.6.32** `static int vec_cmpud_all_lt ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare all Less than Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return true if all elements of a < b.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

**Returns**

boolean int for all 128-bits, true if all Greater Than, false otherwise.

**7.10.6.33** `static int vec_cmpud_all_ne ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare all Not Equal Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return true if all elements of a and b are not equal.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

**Returns**

boolean int for all 128-bits, true if equal, false otherwise.

**7.10.6.34** `static int vec_cmpud_any_eq ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare any Equal Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return true if any elements of a and b are equal.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

**Parameters**

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
----------	--

## Parameters

<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
----------	--

## Returns

boolean int for all 128-bits, true if equal, false otherwise.

7.10.6.35 `static int vec_cmpud_any_ge ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare any Greater Than or Equal Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return true if any elements of a  $\geq$  b.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

## Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

## Returns

boolean int for all 128-bits, true if all Greater Than, false otherwise.

7.10.6.36 `static int vec_cmpud_any_gt ( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare any Greater Than Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return true if all elements of a  $>$  b.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

## Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

## Returns

boolean int for all 128-bits, true if all Greater Than, false otherwise.

7.10.6.37 `static int vec_cmpud_any_le( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare any Less than equal Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return true if any elements of a <= b.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

#### Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

#### Returns

boolean int for any 128-bits, true if any Greater Than, false otherwise.

7.10.6.38 `static int vec_cmpud_any_lt( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare any Less than Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return true if any elements of a < b.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle

#### Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

#### Returns

boolean int for any 128-bits, true if any Greater Than, false otherwise.

7.10.6.39 `static int vec_cmpud_any_ne( vui64_t a, vui64_t b ) [inline],[static]`

Vector Compare any Not Equal Unsigned Doubleword.

Compare each unsigned long (64-bit) integer and return true if any elements of a and b are not equal.

processor	Latency	Throughput
power8	4-9	2/cycle
power9	3	2/cycle



## Parameters

<i>a</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.
<i>b</i>	128-bit vector treated as 2 x 64-bit unsigned long integer (dword) elements.

## Returns

boolean int for any 128-bits, true if equal, false otherwise.

7.10.6.40 `static vi64_t vec_maxsd ( vi64_t vra, vi64_t vrb ) [inline], [static]`

Vector Maximum Signed Doubleword.

For each doubleword element [0|1] of *vra* and *vr**b* compare as signed integers and return the larger value in the result.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

## Parameters

<i>vra</i>	128-bit vector long int.
<i>vr</i> <i>b</i>	128-bit vector long int.

## Returns

vector long maximum of *a* and *b*.

7.10.6.41 `static vui64_t vec_maxud ( vui64_t vra, vui64_t vrb ) [inline], [static]`

Vector Maximum Unsigned Doubleword.

For each doubleword element [0|1] of *vra* and *vr**b* compare as unsigned integers and return the larger value in the result.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

## Parameters

<i>vra</i>	128-bit vector long int.
<i>vr</i> <i>b</i>	128-bit vector long int.

**Returns**

vector unsigned long maximum of a and b.

**7.10.6.42** `static vi64_t vec_minsd ( vi64_t vra, vi64_t vrb ) [inline],[static]`

Vector Minimum Signed Doubleword.

For each doubleword element [0|1] of vra and vrb compare as signed integers and return the smaller value in the result.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

**Parameters**

<i>vra</i>	128-bit vector long int.
<i>vrb</i>	128-bit vector long int.

**Returns**

vector long minimum of a and b.

**7.10.6.43** `static vui64_t vec_minud ( vui64_t vra, vui64_t vrb ) [inline],[static]`

Vector Minimum Unsigned Doubleword.

For each doubleword element [0|1] of vra and vrb compare as unsigned integers and return the smaller value in the result.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

**Parameters**

<i>vra</i>	128-bit vector unsigned long int.
<i>vrb</i>	128-bit vector unsignedlong int.

**Returns**

vector unsigned long minimum of a and b.

**7.10.6.44** `static vui64_t vec_mrgahd ( vui128_t vra, vui128_t vrb ) [inline],[static]`

Vector Merge Algebraic High Doublewords.

Merge only the high doublewords from 2 x Algebraic quadwords across vectors *vra* and *vrh*. This effectively the Vector Merge Even doubleword operation that is not modified for endian.

For example merge the high 64-bits from 2 x 128-bit products as generated by `vec_muleud/vec_muloud`. This result is effectively a vector multiply high unsigned doubleword.

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned <code>__int128</code> .
<i>vrh</i>	128-bit vector unsigned <code>__int128</code> .

#### Returns

A vector merge from only the high doublewords of the 2 x algebraic quadwords across *vra* and *vrh*.

7.10.6.45 `static vui64_t vec_mrgald ( vui128_t vra, vui128_t vrh ) [inline],[static]`

Vector Merge Algebraic Low Doublewords.

Merge only the low doublewords from 2 x Algebraic quadwords across vectors *vra* and *vrh*. This effectively the Vector Merge Odd doubleword operation that is not modified for endian.

For example merge the low 64-bits from 2 x 128-bit products as generated by `vec_muleud/vec_muloud`. This result is effectively a vector multiply low unsigned doubleword.

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

#### Parameters

<i>vra</i>	128-bit vector unsigned <code>__int128</code> .
<i>vrh</i>	128-bit vector unsigned <code>__int128</code> .

#### Returns

A vector merge from only the low doublewords of the 2 x algebraic quadwords across *vra* and *vrh*.

7.10.6.46 `static vui64_t vec_mrged ( vui64_t VA, vui64_t VB ) [inline],[static]`

Vector Merge Even Doubleword. Merge the even doubleword elements from two vectors into the high and low doubleword elements of the result. This is effectively the VSX Permute doubleword Immediate operation modified for endian.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

## Parameters

<code>__VA</code>	a 128-bit vector as the source of the results even doubleword.
<code>__VB</code>	a 128-bit vector as the source of the results odd doubleword.

## Returns

A vector merge from only the even doublewords of the 2 x quadwords across `__VA` and `__VB`.

7.10.6.47 `static vui64_t vec_mrghd ( vui64_t __VA, vui64_t __VB ) [inline],[static]`

Vector Merge High Doubleword. Merge the high doubleword elements from two vectors into the high and low doubleword elements of the result. This is effectively the VSX Permute doubleword Immediate operation modified for endian.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

## Parameters

<code>__VA</code>	a 128-bit vector as the source of the results even doubleword.
<code>__VB</code>	a 128-bit vector as the source of the results odd doubleword.

## Returns

A vector merge from only the high doublewords of the 2 x quadwords across `__VA` and `__VB`.

7.10.6.48 `static vui64_t vec_mrgld ( vui64_t __VA, vui64_t __VB ) [inline],[static]`

Vector Merge Low Doubleword. Merge the low doubleword elements from two vectors into the high and low doubleword elements of the result. This is effectively the VSX Permute doubleword Immediate operation modified for endian.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

## Parameters

<code>__VA</code>	a 128-bit vector as the source of the results even doubleword.
<code>__VB</code>	a 128-bit vector as the source of the results odd doubleword.

**Returns**

A vector merge from only the low doublewords of the 2 x quadwords across \_\_VA and \_\_VB.

**7.10.6.49** `static vui64_t vec_mrgod ( vui64_t __VA, vui64_t __VB ) [inline],[static]`

Vector Merge Odd Doubleword. Merge the odd doubleword elements from two vectors into the high and low doubleword elements of the result. This is effectively the VSX Permute doubleword Immediate operation modified for endian.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

**Parameters**

<code>__VA</code>	a 128-bit vector as the source of the results even doubleword.
<code>__VB</code>	a 128-bit vector as the source of the results odd doubleword.

**Returns**

A vector merge from only the odd doublewords of the 2 x quadwords across \_\_VA and \_\_VB.

**7.10.6.50** `static vui64_t vec_pasted ( vui64_t __VH, vui64_t __VL ) [inline],[static]`

Vector doubleword paste. Concatenate the high doubleword of the 1st vector with the low double word of the 2nd vector.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

**Parameters**

<code>__VH</code>	a 128-bit vector as the source of the high order doubleword.
<code>__VL</code>	a 128-bit vector as the source of the low order doubleword.

**Returns**

The combined 128-bit vector composed of the high order doubleword of \_\_VH and the low order doubleword of \_\_VL.

**7.10.6.51** `static vui64_t vec_permdi ( vui64_t vra, vui64_t vrb, const int ctl ) [inline],[static]`

Vector Permute Doubleword Immediate. Combine a doubleword selected from the 1st (vra) vector with a doubleword selected from the 2nd (vrb) vector.

**Note**

This function implements the operation of a VSX Permute Doubleword Immediate instruction. This implementation is NOT Endian sensitive and the function is stable across BE/LE implementations.

The 2-bit control operand (ctl) selects which doubleword from the 1st and 2nd vector operands are transferred to the result vector. Control table:

ctl	vrt[0:63]	vrt[64:127]
0	vra[0:63]	vr[0:63]
1	vra[0:63]	vr[64:127]
2	vra[64:127]	vr[0:63]
3	vra[64:127]	vr[64:127]

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

**Parameters**

<i>vra</i>	a 128-bit vector as the source of the high order doubleword of the result.
<i>vr</i>	a 128-bit vector as the source of the low order doubleword of the result.
<i>ctl</i>	const integer where the low order 2 bits control the selection of doublewords from input vector <i>vra</i> and <i>vr</i> .

**Returns**

The combined 128-bit vector composed of the high order doubleword of *vra* and the low order doubleword of *vr*.

**7.10.6.52** `static vui64_t vec_popcntd ( vui64_t vra ) [inline],[static]`

Vector Population Count doubleword.

Count the number of '1' bits (0-64) within each doubleword element of a 128-bit vector.

processor	Latency	Throughput
power8	4	2/2 cycles
power9	3	2/cycle

For POWER8 (PowerISA 2.07B) or later use the Vector Population Count DoubleWord (**vpopcntd**) instruction. Otherwise use the pveclib `vec_popcntw` to count each word then sum across with Vector Sum across Half Signed Word Saturate (**vsum2sws**).

**Parameters**

<i>vra</i>	128-bit vector treated as 2 x 64-bit integer (dwords) elements.
------------	---

**Returns**

128-bit vector with the population count for each dword element.

**7.10.6.53** `static vui64_t vec_revbd ( vui64_t vra ) [inline],[static]`

byte reverse each doubleword for a vector unsigned long int.

For each doubleword of the input vector, reverse the order of bytes / octets within the doubleword.

processor	Latency	Throughput
power8	2-11	2/cycle
power9	3	2/cycle

**Parameters**

<i>vra</i>	a 128-bit vector unsigned long int.
------------	-------------------------------------

**Returns**

a 128-bit vector with the bytes of each doubleword reversed.

**7.10.6.54** `static vui64_t vec_rldi ( vui64_t vra, const unsigned int shb ) [inline],[static]`

Vector Rotate left Doubleword Immediate.

Rotate left each doubleword element [0-1], 0-63 bits, as specified by an immediate value. The rotate amount is a const unsigned int in the range 0-63. A rotate count of 0 returns the original value of vra. Shift counts greater than 63 bits handled modulo 64.

processor	Latency	Throughput
power8	2-4	2/cycle
power9	2-5	2/cycle

**Parameters**

<i>vra</i>	a 128-bit vector treated as a vector unsigned long int.
<i>shb</i>	rotate amount in the range 0-63.

**Returns**

128-bit vector unsigned long int, shifted left shb bits.

**7.10.6.55** `static vui64_t vec_sldi ( vui64_t vra, const unsigned int shb ) [inline],[static]`

Vector Shift left Doubleword Immediate.

Shift left each doubleword element [0-1], 0-63 bits, as specified by an immediate value. The shift amount is a const unsigned long int in the range 0-63. A shift count of 0 returns the original value of *vra*. Shift counts greater than 63 bits return zero.

processor	Latency	Throughput
power8	2-4	2/cycle
power9	2-5	2/cycle

#### Parameters

<i>vra</i>	a 128-bit vector treated as a vector unsigned long int.
<i>shb</i>	shift amount in the range 0-63.

#### Returns

128-bit vector unsigned long int, shifted left *shb* bits.

**7.10.6.56** `static vui64_t vec_splatd ( vui64_t vra, const int ctl ) [inline],[static]`

Vector splat doubleword. Duplicate the selected doubleword element across the doubleword elements of the result. This is effectively the VSX Merge doubleword operation modified for endian.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

The 1-bit control operand (*ctl*) selects which (0:1) doubleword element, from the vector operand, is replicated to both doublewords of the result vector. Control table:

ctl	vrt[0]	vrt[1]
0	<i>vra</i> [0]	<i>vra</i> [0]
1	<i>vra</i> [1]	<i>vra</i> [1]

#### Parameters

<i>vra</i>	a 128-bit vector.
<i>ctl</i>	a const integer encoding the source doubleword.

#### Returns

The original vector with the doubleword elements swapped.

**7.10.6.57** `static vui64_t vec_spltd ( vui64_t vra, const int ctl ) [inline],[static]`

**Deprecated** Vector splat doubleword. Duplicate the selected doubleword element across the doubleword elements of the result.



processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

The 1-bit control operand (*ctl*) selects which (0:1) doubleword element, from the vector operand, is replicated to both doublewords of the result vector. Control table:

ctl	vrt[0:63]	vrt[64:127]
0	vra[0:63]	vra[0:63]
1	vra[64:127]	vra[64:127]

#### Parameters

<i>vra</i>	a 128-bit vector.
<i>ctl</i>	a const integer encoding the source doubleword.

#### Returns

The original vector with the doubleword elements swapped.

**7.10.6.58** `static vi64_t vec_sradi ( vi64_t vra, const unsigned int shb )` `[inline], [static]`

Vector Shift Right Algebraic Doubleword Immediate.

Shift Right Algebraic each doubleword element [0-1], 0-63 bits, as specified by an immediate value. The shift amount is a const unsigned int in the range 0-63. A shift count of 0 returns the original value of *vra*. Shift counts greater than 63 bits return the sign bit propagated to each bit of each element.

processor	Latency	Throughput
power8	2-4	2/cycle
power9	2-5	2/cycle

#### Parameters

<i>vra</i>	a 128-bit vector treated as a vector signed long int.
<i>shb</i>	shift amount in the range 0-63.

#### Returns

128-bit vector signed long int, shifted right *shb* bits.

**7.10.6.59** `static vui64_t vec_srdi ( vui64_t vra, const unsigned int shb )` `[inline], [static]`

Vector Shift Right Doubleword Immediate.

Shift Right each doubleword element [0-1], 0-63 bits, as specified by an immediate value. The shift amount is a const unsigned int in the range 0-63. A shift count of 0 returns the original value of *vra*. Shift counts greater than 63 bits return zero.

processor	Latency	Throughput
power8	2-4	2/cycle
power9	2-5	2/cycle

## Parameters

<i>vra</i>	a 128-bit vector treated as a vector unsigned long int.
<i>shb</i>	shift amount in the range 0-63.

## Returns

128-bit vector unsigned long int, shifted right shb bits.

7.10.6.60 `static vui64_t vec_subudm ( vui64_t a, vui64_t b )` `[inline]`, `[static]`

Vector Subtract Unsigned Doubleword Modulo.

For each unsigned long (64-bit) integer element  $c[i] = a[i] + \text{NOT}(b[i]) + 1$ .

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

For POWER8 (PowerISA 2.07B) or later use the Vector Subtract Unsigned Doubleword Modulo (**vsubudm**) instruction. Otherwise use vector add word modulo forms and propagate the carry bits.

## Parameters

<i>a</i>	128-bit vector treated as 2 X unsigned long int.
<i>b</i>	128-bit vector treated as 2 X unsigned long int.

## Returns

vector unsigned long int sum of  $a[0] + \text{NOT}(b[0]) + 1$  and  $a[1] + \text{NOT}(b[1]) + 1$ .

7.10.6.61 `static vui64_t vec_swapd ( vui64_t vra )` `[inline]`, `[static]`

Vector doubleword swap. Exchange the high and low doubleword elements of a vector.

processor	Latency	Throughput
power8	2	2/cycle
power9	3	2/cycle

## Parameters

<i>vra</i>	a 128-bit vector.
------------	-------------------

## Returns

The original vector with the doubleword elements swapped.

**7.10.6.62** `static vui64_t vec_vmsumuwmm ( vui32_t vra, vui32_t vrb, vui64_t vrc )` `[inline],[static]`

Vector Multiply-Sum Unsigned Word Modulo.

Multiply the unsigned word elements of *vra* and *vrb*, internally generating doubleword products. Then generate three way sum of adjacent doubleword product pairs, plus the doubleword elements from *vrc*. The final summation is modulo 64-bits.

## Note

This function implements the operation of a Vector Multiply-Sum Unsigned Word Modulo instruction, if the PowerISA included such an instruction. This implementation is NOT endian sensitive and the function is stable across BE/LE implementations.

processor	Latency	Throughput
power8	11	1/cycle
power9	11	1/cycle

## Parameters

<i>vra</i>	128-bit vector unsigned int.
<i>vrb</i>	128-bit vector unsigned int.
<i>vrc</i>	128-bit vector unsigned long.

## Returns

vector of doubleword elements where each is the sum of the even and odd product of the *vra* and *vrb*, plus the doubleword elements of *vrc*.

**7.10.6.63** `static vui32_t vec_vpkudum ( vui64_t vra, vui64_t vrb )` `[inline],[static]`

Vector Pack Unsigned Doubleword Unsigned Modulo.

The doubleword source is the concatenation of *vra* and *vrb*. For each integer word from 0 to 3, of the result vector, do the following. The contents of bits 32:63 of the corresponding doubleword source element [*i*] placed into word element [*i*] of the result.

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

**Note**

Use `vec_vpkudum` naming but only if the compiler does not define it in `<altivec.h>`.

**Parameters**

<i>vra</i>	a 128-bit vector treated as 2 x unsigned long integers.
<i>vrb</i>	a 128-bit vector treated as 2 x unsigned long integers.

**Returns**

128-bit vector treated as 4 x unsigned integers.

**7.10.6.64** `static vui64_t vec_vrld ( vui64_t vra, vui64_t vrb )` `[inline], [static]`

Vector Rotate Left Doubleword.

Vector Rotate Left Doubleword 0-63 bits. The shift amount is from bits 58-63 and 122-127 of *vrb*.

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

**Note**

Use `vec_vrld` naming but only if the compiler does not define it in `<altivec.h>`.

**Parameters**

<i>vra</i>	a 128-bit vector treated as 2 x unsigned long integers.
<i>vrb</i>	shift amount in bits 58:63 and 122:127.

**Returns**

Left shifted vector unsigned long.

**7.10.6.65** `static vui64_t vec_vslid ( vui64_t vra, vui64_t vrb )` `[inline], [static]`

Vector Shift Left Doubleword.

Vector Shift Left Doubleword 0-63 bits. The shift amount is from bits 58-63 and 122-127 of *vrb*.

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

**Note**

Can not use `vec_sld` naming here as that would conflict with the generic Shift Left Double Vector. Use `vec_vsl` but only if the compiler does not define it in `<altivec.h>`.

**Parameters**

<i>vra</i>	a 128-bit vector treated as 2 x unsigned long integers.
<i>vr</i> <i>b</i>	shift amount in bits 58:63 and 122:127.

**Returns**

Left shifted vector unsigned long.

**7.10.6.66** `static vi64_t vec_vsr ( vi64_t vra, vui64_t vr )` `[inline],[static]`

Vector Shift Right Algebraic Doubleword.

Vector Shift Right Algebraic Doubleword 0-63 bits. The shift amount is from bits 58-63 and 122-127 of *vr**b*.

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

**Note**

Use the `vec_vsr` for consistency with `vec_vsl` above. Define `vec_vsr` only if the compiler does not define it in `<altivec.h>`.

**Parameters**

<i>vra</i>	a 128-bit vector treated as 2 x unsigned long integers.
<i>vr</i> <i>b</i>	shift amount in bits 58:63 and 122:127.

**Returns**

Right shifted vector unsigned long.

**7.10.6.67** `static vui64_t vec_vsr ( vui64_t vra, vui64_t vr )` `[inline],[static]`

Vector Shift Right Doubleword.

Vector Shift Right Doubleword 0-63 bits. The shift amount is from bits 58-63 and 122-127 of *vr**b*.

processor	Latency	Throughput
power8	2	2/cycle
power9	2	2/cycle

**Note**

Use the `vec_vsrd` for consistency with `vec_vsld` above. Define `vec_vsrd` only if the compiler does not define it in `<altivec.h>`.

**Parameters**

<i>vra</i>	a 128-bit vector treated as 2 x unsigned long integers.
<i>vrb</i>	shift amount in bits 58:63 and 122:127.

**Returns**

Right shifted vector unsigned long.

**7.10.6.68** `static vui64_t vec_xxspltd ( vui64_t vra, const int ctl )` `[inline], [static]`

Vector splat doubleword. Duplicate the selected doubleword element across the doubleword elements of the result.

**Note**

This function implements the operation of a VSX Splat Doubleword Immediate instruction. This implementation is NOT Endian sensitive and the function is stable across BE/LE implementations.

The 1-bit control operand (*ctl*) selects which (0:1) doubleword element, from the vector operand, is replicated to both doublewords of the result vector. Control table:

<b>ctl</b>	<b>vrt[0:63]</b>	<b>vrt[64:127]</b>
0	<i>vra</i> [0:63]	<i>vra</i> [0:63]
1	<i>vra</i> [64:127]	<i>vra</i> [64:127]

<b>processor</b>	<b>Latency</b>	<b>Throughput</b>
power8	2	2/cycle
power9	3	2/cycle

**Parameters**

<i>vra</i>	a 128-bit vector.
<i>ctl</i>	a const integer encoding the source doubleword.

**Returns**

The original vector with the doubleword elements swapped.

# Index

[\\_\\_VEC\\_U\\_128](#), [31](#)

[\\_\\_VF\\_128](#), [32](#)

[bcd\\_qauntize0](#)

[vec\\_bcd\\_ppc.h](#), [34](#)

[src/vec\\_bcd\\_ppc.h](#), [33](#)

[src/vec\\_char\\_ppc.h](#), [36](#)

[src/vec\\_common\\_ppc.h](#), [50](#)

[src/vec\\_f128\\_ppc.h](#), [55](#)

[src/vec\\_f32\\_ppc.h](#), [72](#)

[src/vec\\_f64\\_ppc.h](#), [85](#)

[src/vec\\_int128\\_ppc.h](#), [98](#)

[src/vec\\_int16\\_ppc.h](#), [132](#)

[src/vec\\_int32\\_ppc.h](#), [148](#)

[src/vec\\_int64\\_ppc.h](#), [163](#)

[vec\\_BCD2DFP](#)

[vec\\_bcd\\_ppc.h](#), [34](#)

[vec\\_DFP2BCD](#)

[vec\\_bcd\\_ppc.h](#), [36](#)

[vec\\_absdub](#)

[vec\\_char\\_ppc.h](#), [39](#)

[vec\\_absdud](#)

[vec\\_int64\\_ppc.h](#), [178](#)

[vec\\_absduh](#)

[vec\\_int16\\_ppc.h](#), [140](#)

[vec\\_absduw](#)

[vec\\_int32\\_ppc.h](#), [154](#)

[vec\\_absf128](#)

[vec\\_f128\\_ppc.h](#), [59](#)

[vec\\_absf32](#)

[vec\\_f32\\_ppc.h](#), [75](#)

[vec\\_absf64](#)

[vec\\_f64\\_ppc.h](#), [88](#)

[vec\\_addcq](#)

[vec\\_int128\\_ppc.h](#), [102](#)

[vec\\_addcuq](#)

[vec\\_int128\\_ppc.h](#), [103](#)

[vec\\_addecuq](#)

[vec\\_int128\\_ppc.h](#), [103](#)

[vec\\_addeq](#)

[vec\\_int128\\_ppc.h](#), [104](#)

[vec\\_addeuqm](#)

[vec\\_int128\\_ppc.h](#), [104](#)

[vec\\_addudm](#)

[vec\\_int64\\_ppc.h](#), [178](#)

[vec\\_adduqm](#)

[vec\\_int128\\_ppc.h](#), [105](#)

[vec\\_all\\_isfinitef128](#)

[vec\\_f128\\_ppc.h](#), [59](#)

[vec\\_all\\_isfinitef32](#)

[vec\\_f32\\_ppc.h](#), [75](#)

[vec\\_all\\_isfinitef64](#)

[vec\\_f64\\_ppc.h](#), [88](#)

[vec\\_all\\_isinff128](#)

[vec\\_f128\\_ppc.h](#), [60](#)

[vec\\_all\\_isinff32](#)

[vec\\_f32\\_ppc.h](#), [76](#)

[vec\\_all\\_isinff64](#)

[vec\\_f64\\_ppc.h](#), [89](#)

[vec\\_all\\_isnanf128](#)

[vec\\_f128\\_ppc.h](#), [60](#)

[vec\\_all\\_isnanf32](#)

[vec\\_f32\\_ppc.h](#), [76](#)

[vec\\_all\\_isnanf64](#)

[vec\\_f64\\_ppc.h](#), [89](#)

[vec\\_all\\_isnormalf128](#)

[vec\\_f128\\_ppc.h](#), [61](#)

[vec\\_all\\_isnormalf32](#)

[vec\\_f32\\_ppc.h](#), [77](#)

[vec\\_all\\_isnormalf64](#)

[vec\\_f64\\_ppc.h](#), [90](#)

[vec\\_all\\_issubnormalf128](#)

[vec\\_f128\\_ppc.h](#), [61](#)

[vec\\_all\\_issubnormalf32](#)

[vec\\_f32\\_ppc.h](#), [77](#)

[vec\\_all\\_issubnormalf64](#)

[vec\\_f64\\_ppc.h](#), [90](#)

[vec\\_all\\_iszerof128](#)

[vec\\_f128\\_ppc.h](#), [62](#)

[vec\\_all\\_iszerof32](#)

[vec\\_f32\\_ppc.h](#), [78](#)

[vec\\_all\\_iszerof64](#)

[vec\\_f64\\_ppc.h](#), [91](#)

[vec\\_any\\_isfinitef32](#)

[vec\\_f32\\_ppc.h](#), [78](#)

[vec\\_any\\_isfinitef64](#)

[vec\\_f64\\_ppc.h](#), [91](#)

[vec\\_any\\_isinff32](#)

[vec\\_f32\\_ppc.h](#), [79](#)

[vec\\_any\\_isinff64](#)

[vec\\_f64\\_ppc.h](#), [92](#)

[vec\\_any\\_isnanf32](#)

[vec\\_f32\\_ppc.h](#), [79](#)

[vec\\_any\\_isnanf64](#)

[vec\\_f64\\_ppc.h](#), [92](#)

[vec\\_any\\_isnormalf32](#)

[vec\\_f32\\_ppc.h](#), [80](#)

vec\_any\_isnormalf64  
     vec\_f64\_ppc.h, 93  
 vec\_any\_issubnormalf32  
     vec\_f32\_ppc.h, 80  
 vec\_any\_issubnormalf64  
     vec\_f64\_ppc.h, 93  
 vec\_any\_iszerof32  
     vec\_f32\_ppc.h, 81  
 vec\_any\_iszerof64  
     vec\_f64\_ppc.h, 94  
 vec\_bcd\_ppc.h  
     bcd\_qauntize0, 34  
     vec\_BCD2DFP, 34  
     vec\_DFP2BCD, 36  
     vec\_bcdadd, 34  
     vec\_bcdctb100s, 34  
     vec\_bcddiv, 35  
     vec\_bcdmul, 35  
     vec\_bcdsub, 36  
 vec\_bcdadd  
     vec\_bcd\_ppc.h, 34  
 vec\_bcdctb100s  
     vec\_bcd\_ppc.h, 34  
 vec\_bcddiv  
     vec\_bcd\_ppc.h, 35  
 vec\_bcdmul  
     vec\_bcd\_ppc.h, 35  
 vec\_bcdsub  
     vec\_bcd\_ppc.h, 36  
 vec\_char\_ppc.h  
     vec\_absdub, 39  
     vec\_clzb, 39  
     vec\_isalnum, 40  
     vec\_isalpha, 40  
     vec\_isdigit, 41  
     vec\_mrgahb, 41  
     vec\_mrgalb, 41  
     vec\_mrgeb, 42  
     vec\_mrgob, 42  
     vec\_mulhsb, 43  
     vec\_mulhub, 43  
     vec\_mulubm, 44  
     vec\_popcntb, 44  
     vec\_shift\_leftdo, 46  
     vec\_slbi, 46  
     vec\_srabi, 47  
     vec\_srbi, 47  
     vec\_tolower, 48  
     vec\_toupper, 48  
     vec\_vmrgeb, 48  
     vec\_vmrgeb, 49  
 vec\_clzb  
     vec\_char\_ppc.h, 39  
 vec\_clzd  
     vec\_int64\_ppc.h, 179  
 vec\_clzh  
     vec\_int16\_ppc.h, 140  
 vec\_clzq  
     vec\_int128\_ppc.h, 105  
 vec\_clzw  
     vec\_int32\_ppc.h, 155  
 vec\_cmpeqsd  
     vec\_int64\_ppc.h, 179  
 vec\_cmpeqsq  
     vec\_int128\_ppc.h, 105  
 vec\_cmpequd  
     vec\_int64\_ppc.h, 180  
 vec\_cmpequq  
     vec\_int128\_ppc.h, 106  
 vec\_cmpgesd  
     vec\_int64\_ppc.h, 180  
 vec\_cmpgesq  
     vec\_int128\_ppc.h, 106  
 vec\_cmpgeud  
     vec\_int64\_ppc.h, 181  
 vec\_cmpgeuq  
     vec\_int128\_ppc.h, 107  
 vec\_cmpgtsd  
     vec\_int64\_ppc.h, 181  
 vec\_cmpgtsq  
     vec\_int128\_ppc.h, 107  
 vec\_cmpgtud  
     vec\_int64\_ppc.h, 182  
 vec\_cmpgtuq  
     vec\_int128\_ppc.h, 108  
 vec\_cmpleud  
     vec\_int64\_ppc.h, 182  
 vec\_cmpleuq  
     vec\_int128\_ppc.h, 108  
 vec\_cmpleud  
     vec\_int64\_ppc.h, 182  
 vec\_cmpleuq  
     vec\_int128\_ppc.h, 109  
 vec\_cmpltsd  
     vec\_int64\_ppc.h, 183  
 vec\_cmpltsq  
     vec\_int128\_ppc.h, 109  
 vec\_cmpltud  
     vec\_int64\_ppc.h, 183  
 vec\_cmpltuq  
     vec\_int128\_ppc.h, 110  
 vec\_cmpnesd  
     vec\_int64\_ppc.h, 184  
 vec\_cmpnesq  
     vec\_int128\_ppc.h, 110  
 vec\_cmpneud  
     vec\_int64\_ppc.h, 184  
 vec\_cmpneuq  
     vec\_int128\_ppc.h, 111  
 vec\_cmppsd\_all\_eq  
     vec\_int64\_ppc.h, 185  
 vec\_cmppsd\_all\_ge  
     vec\_int64\_ppc.h, 185  
 vec\_cmppsd\_all\_gt  
     vec\_int64\_ppc.h, 186  
 vec\_cmppsd\_all\_le



vec\_int64\_ppc.h, 186  
vec\_cmpsd\_all\_lt  
    vec\_int64\_ppc.h, 186  
vec\_cmpsd\_all\_ne  
    vec\_int64\_ppc.h, 187  
vec\_cmpsd\_any\_eq  
    vec\_int64\_ppc.h, 187  
vec\_cmpsd\_any\_ge  
    vec\_int64\_ppc.h, 188  
vec\_cmpsd\_any\_gt  
    vec\_int64\_ppc.h, 188  
vec\_cmpsd\_any\_le  
    vec\_int64\_ppc.h, 188  
vec\_cmpsd\_any\_lt  
    vec\_int64\_ppc.h, 189  
vec\_cmpsd\_any\_ne  
    vec\_int64\_ppc.h, 189  
vec\_cmpsq\_all\_eq  
    vec\_int128\_ppc.h, 111  
vec\_cmpsq\_all\_ge  
    vec\_int128\_ppc.h, 112  
vec\_cmpsq\_all\_gt  
    vec\_int128\_ppc.h, 112  
vec\_cmpsq\_all\_le  
    vec\_int128\_ppc.h, 112  
vec\_cmpsq\_all\_lt  
    vec\_int128\_ppc.h, 113  
vec\_cmpsq\_all\_ne  
    vec\_int128\_ppc.h, 113  
vec\_cmpud\_all\_eq  
    vec\_int64\_ppc.h, 190  
vec\_cmpud\_all\_ge  
    vec\_int64\_ppc.h, 190  
vec\_cmpud\_all\_gt  
    vec\_int64\_ppc.h, 191  
vec\_cmpud\_all\_le  
    vec\_int64\_ppc.h, 191  
vec\_cmpud\_all\_lt  
    vec\_int64\_ppc.h, 191  
vec\_cmpud\_all\_ne  
    vec\_int64\_ppc.h, 192  
vec\_cmpud\_any\_eq  
    vec\_int64\_ppc.h, 192  
vec\_cmpud\_any\_ge  
    vec\_int64\_ppc.h, 193  
vec\_cmpud\_any\_gt  
    vec\_int64\_ppc.h, 193  
vec\_cmpud\_any\_le  
    vec\_int64\_ppc.h, 193  
vec\_cmpud\_any\_lt  
    vec\_int64\_ppc.h, 194  
vec\_cmpud\_any\_ne  
    vec\_int64\_ppc.h, 194  
vec\_cmpuq\_all\_eq  
    vec\_int128\_ppc.h, 114  
vec\_cmpuq\_all\_ge  
    vec\_int128\_ppc.h, 114  
vec\_int128\_ppc.h, 115  
vec\_cmpuq\_all\_le  
    vec\_int128\_ppc.h, 115  
vec\_cmpuq\_all\_lt  
    vec\_int128\_ppc.h, 115  
vec\_cmpuq\_all\_ne  
    vec\_int128\_ppc.h, 116  
vec\_cmul100cuq  
    vec\_int128\_ppc.h, 116  
vec\_cmul100ecuq  
    vec\_int128\_ppc.h, 117  
vec\_cmul10cuq  
    vec\_int128\_ppc.h, 117  
vec\_cmul10ecuq  
    vec\_int128\_ppc.h, 118  
vec\_const\_huge\_valf128  
    vec\_f128\_ppc.h, 62  
vec\_const\_inff128  
    vec\_f128\_ppc.h, 62  
vec\_const\_nanf128  
    vec\_f128\_ppc.h, 62  
vec\_const\_nansf128  
    vec\_f128\_ppc.h, 63  
vec\_copysignf128  
    vec\_f128\_ppc.h, 63  
vec\_copysignf32  
    vec\_f32\_ppc.h, 81  
vec\_copysignf64  
    vec\_f64\_ppc.h, 94  
vec\_f128\_ppc.h  
    vec\_absf128, 59  
    vec\_all\_isfinitef128, 59  
    vec\_all\_isinff128, 60  
    vec\_all\_isnanf128, 60  
    vec\_all\_isnormalf128, 61  
    vec\_all\_issubnormalf128, 61  
    vec\_all\_iszerof128, 62  
    vec\_const\_huge\_valf128, 62  
    vec\_const\_inff128, 62  
    vec\_const\_nanf128, 62  
    vec\_const\_nansf128, 63  
    vec\_copysignf128, 63  
    vec\_isfinitef128, 63  
    vec\_isinf\_signf128, 64  
    vec\_isinff128, 64  
    vec\_isnanf128, 65  
    vec\_isnormalf128, 65  
    vec\_issubnormalf128, 66  
    vec\_iszerof128, 66  
    vec\_setb\_qp, 67  
    vec\_signbitf128, 67  
    vec\_xfer\_bin128\_2\_vui128t, 67  
    vec\_xfer\_bin128\_2\_vui16t, 68  
    vec\_xfer\_bin128\_2\_vui32t, 68  
    vec\_xfer\_bin128\_2\_vui64t, 69  
    vec\_xfer\_bin128\_2\_vui8t, 69  
    vec\_xfer\_vui128t\_2\_bin128, 69  
    vec\_xfer\_vui16t\_2\_bin128, 70

- vec\_xfer\_vui32t\_2\_bin128, 70
- vec\_xfer\_vui64t\_2\_bin128, 71
- vec\_xfer\_vui8t\_2\_bin128, 71
- vec\_f32\_ppc.h
  - vec\_absf32, 75
  - vec\_all\_isfinitef32, 75
  - vec\_all\_isinff32, 76
  - vec\_all\_isnanf32, 76
  - vec\_all\_isnormalf32, 77
  - vec\_all\_issubnormalf32, 77
  - vec\_all\_iszerof32, 78
  - vec\_any\_isfinitef32, 78
  - vec\_any\_isinff32, 79
  - vec\_any\_isnanf32, 79
  - vec\_any\_isnormalf32, 80
  - vec\_any\_issubnormalf32, 80
  - vec\_any\_iszerof32, 81
  - vec\_copysignf32, 81
  - vec\_isfinitef32, 82
  - vec\_isinff32, 82
  - vec\_isnanf32, 83
  - vec\_isnormalf32, 83
  - vec\_issubnormalf32, 84
  - vec\_iszerof32, 84
- vec\_f64\_ppc.h
  - vec\_absf64, 88
  - vec\_all\_isfinitef64, 88
  - vec\_all\_isinff64, 89
  - vec\_all\_isnanf64, 89
  - vec\_all\_isnormalf64, 90
  - vec\_all\_issubnormalf64, 90
  - vec\_all\_iszerof64, 91
  - vec\_any\_isfinitef64, 91
  - vec\_any\_isinff64, 92
  - vec\_any\_isnanf64, 92
  - vec\_any\_isnormalf64, 93
  - vec\_any\_issubnormalf64, 93
  - vec\_any\_iszerof64, 94
  - vec\_copysignf64, 94
  - vec\_isfinitef64, 95
  - vec\_isinff64, 95
  - vec\_isnanf64, 96
  - vec\_isnormalf64, 96
  - vec\_issubnormalf64, 97
  - vec\_iszerof64, 97
  - vec\_pack\_longdouble, 98
  - vec\_unpack\_longdouble, 98
- vec\_int128\_ppc.h
  - vec\_addcq, 102
  - vec\_addcuq, 103
  - vec\_addecuq, 103
  - vec\_addeq, 104
  - vec\_addeuqm, 104
  - vec\_adduqm, 105
  - vec\_clzq, 105
  - vec\_cmpeqsq, 105
  - vec\_cmpequq, 106
  - vec\_cmpgesq, 106
  - vec\_cmpgeuq, 107
  - vec\_cmpgtsq, 107
  - vec\_cmpgtuq, 108
  - vec\_cmplesq, 108
  - vec\_cmpleuq, 109
  - vec\_cmpltuq, 109
  - vec\_cmpltuq, 110
  - vec\_cmpnesq, 110
  - vec\_cmpneuq, 111
  - vec\_cmpsq\_all\_eq, 111
  - vec\_cmpsq\_all\_ge, 112
  - vec\_cmpsq\_all\_gt, 112
  - vec\_cmpsq\_all\_le, 112
  - vec\_cmpsq\_all\_lt, 113
  - vec\_cmpsq\_all\_ne, 113
  - vec\_cmpuq\_all\_eq, 114
  - vec\_cmpuq\_all\_ge, 114
  - vec\_cmpuq\_all\_gt, 115
  - vec\_cmpuq\_all\_le, 115
  - vec\_cmpuq\_all\_lt, 115
  - vec\_cmpuq\_all\_ne, 116
  - vec\_cmul100cuq, 116
  - vec\_cmul100ecuq, 117
  - vec\_cmul10cuq, 117
  - vec\_cmul10ecuq, 118
  - vec\_msumudm, 118
  - vec\_mul10cuq, 118
  - vec\_mul10ecuq, 119
  - vec\_mul10euq, 119
  - vec\_mul10uq, 120
  - vec\_muleud, 120
  - vec\_mulhud, 121
  - vec\_mulluq, 121
  - vec\_muloud, 122
  - vec\_muludm, 122
  - vec\_muludq, 123
  - vec\_popcntq, 123
  - vec\_revbq, 124
  - vec\_setb\_cyq, 124
  - vec\_setb\_ncq, 124
  - vec\_setb\_sq, 125
  - vec\_sldq, 125
  - vec\_slq, 126
  - vec\_slq4, 126
  - vec\_slq5, 126
  - vec\_slqi, 127
  - vec\_srq, 127
  - vec\_srq4, 127
  - vec\_srq5, 128
  - vec\_srqi, 128
  - vec\_subcuq, 129
  - vec\_subecuq, 129
  - vec\_subeuqm, 129
  - vec\_subuqm, 131
  - vec\_vmuleud, 131
  - vec\_vmuloud, 132
- vec\_int16\_ppc.h
  - vec\_absduh, 140

- vec\_clzh, [140](#)
- vec\_mrgahh, [141](#)
- vec\_mrgalh, [141](#)
- vec\_mrgeh, [142](#)
- vec\_mrgoh, [142](#)
- vec\_mulhsh, [143](#)
- vec\_mulhuh, [143](#)
- vec\_muluhm, [144](#)
- vec\_popcnth, [144](#)
- vec\_revbh, [145](#)
- vec\_slhi, [145](#)
- vec\_srahi, [146](#)
- vec\_srhi, [146](#)
- vec\_vmrgeh, [147](#)
- vec\_vmrgoh, [147](#)
- vec\_int32\_ppc.h
  - vec\_absduw, [154](#)
  - vec\_clzw, [155](#)
  - vec\_mrgahw, [155](#)
  - vec\_mrgalw, [156](#)
  - vec\_mrgew, [156](#)
  - vec\_mrgow, [157](#)
  - vec\_mulesw, [157](#)
  - vec\_muleuw, [158](#)
  - vec\_mulhsw, [158](#)
  - vec\_mulhuw, [159](#)
  - vec\_mulosw, [159](#)
  - vec\_mulouw, [160](#)
  - vec\_muluwm, [160](#)
  - vec\_popcntw, [160](#)
  - vec\_revbw, [161](#)
  - vec\_slwi, [161](#)
  - vec\_srawi, [162](#)
  - vec\_srwi, [162](#)
- vec\_int64\_ppc.h
  - vec\_absdud, [178](#)
  - vec\_addudm, [178](#)
  - vec\_clzd, [179](#)
  - vec\_cmpeqsd, [179](#)
  - vec\_cmpequd, [180](#)
  - vec\_cmpgesd, [180](#)
  - vec\_cmpgeud, [181](#)
  - vec\_cmpgtsd, [181](#)
  - vec\_cmpgtud, [182](#)
  - vec\_cmpleud, [182](#)
  - vec\_cmpleud, [182](#)
  - vec\_cmpltsd, [183](#)
  - vec\_cmpltud, [183](#)
  - vec\_cmpnesd, [184](#)
  - vec\_cmpneud, [184](#)
  - vec\_cmpsd\_all\_eq, [185](#)
  - vec\_cmpsd\_all\_ge, [185](#)
  - vec\_cmpsd\_all\_gt, [186](#)
  - vec\_cmpsd\_all\_le, [186](#)
  - vec\_cmpsd\_all\_lt, [186](#)
  - vec\_cmpsd\_all\_ne, [187](#)
  - vec\_cmpsd\_any\_eq, [187](#)
  - vec\_cmpsd\_any\_ge, [188](#)
  - vec\_cmpsd\_any\_gt, [188](#)
  - vec\_cmpsd\_any\_le, [188](#)
  - vec\_cmpsd\_any\_lt, [189](#)
  - vec\_cmpsd\_any\_ne, [189](#)
  - vec\_cmpud\_all\_eq, [190](#)
  - vec\_cmpud\_all\_ge, [190](#)
  - vec\_cmpud\_all\_gt, [191](#)
  - vec\_cmpud\_all\_le, [191](#)
  - vec\_cmpud\_all\_lt, [191](#)
  - vec\_cmpud\_all\_ne, [192](#)
  - vec\_cmpud\_any\_eq, [192](#)
  - vec\_cmpud\_any\_ge, [193](#)
  - vec\_cmpud\_any\_gt, [193](#)
  - vec\_cmpud\_any\_le, [193](#)
  - vec\_cmpud\_any\_lt, [194](#)
  - vec\_cmpud\_any\_ne, [194](#)
  - vec\_maxsd, [195](#)
  - vec\_maxud, [195](#)
  - vec\_minsd, [196](#)
  - vec\_minud, [196](#)
  - vec\_mrgahd, [196](#)
  - vec\_mrgald, [197](#)
  - vec\_mrged, [197](#)
  - vec\_mrghd, [198](#)
  - vec\_mrgld, [198](#)
  - vec\_mrgod, [199](#)
  - vec\_pasted, [199](#)
  - vec\_permdi, [199](#)
  - vec\_popcntd, [200](#)
  - vec\_revbd, [201](#)
  - vec\_rldi, [201](#)
  - vec\_sldi, [201](#)
  - vec\_splatd, [202](#)
  - vec\_spltd, [202](#)
  - vec\_sradi, [203](#)
  - vec\_srdi, [203](#)
  - vec\_subudm, [204](#)
  - vec\_swapd, [204](#)
  - vec\_vmsumuw, [205](#)
  - vec\_vpkudum, [205](#)
  - vec\_vrld, [206](#)
  - vec\_vsld, [206](#)
  - vec\_vsrld, [207](#)
  - vec\_vsrld, [207](#)
  - vec\_xxspltd, [208](#)
- vec\_isalnum
  - vec\_char\_ppc.h, [40](#)
- vec\_isalpha
  - vec\_char\_ppc.h, [40](#)
- vec\_isdigit
  - vec\_char\_ppc.h, [41](#)
- vec\_isfinitef128
  - vec\_f128\_ppc.h, [63](#)
- vec\_isfinitef32
  - vec\_f32\_ppc.h, [82](#)
- vec\_isfinitef64
  - vec\_f64\_ppc.h, [95](#)
- vec\_isinf\_signf128

vec\_f128\_ppc.h, 64  
 vec\_isinff128  
   vec\_f128\_ppc.h, 64  
 vec\_isinff32  
   vec\_f32\_ppc.h, 82  
 vec\_isinff64  
   vec\_f64\_ppc.h, 95  
 vec\_isnanf128  
   vec\_f128\_ppc.h, 65  
 vec\_isnanf32  
   vec\_f32\_ppc.h, 83  
 vec\_isnanf64  
   vec\_f64\_ppc.h, 96  
 vec\_isnormalf128  
   vec\_f128\_ppc.h, 65  
 vec\_isnormalf32  
   vec\_f32\_ppc.h, 83  
 vec\_isnormalf64  
   vec\_f64\_ppc.h, 96  
 vec\_issubnormalf128  
   vec\_f128\_ppc.h, 66  
 vec\_issubnormalf32  
   vec\_f32\_ppc.h, 84  
 vec\_issubnormalf64  
   vec\_f64\_ppc.h, 97  
 vec\_iszerof128  
   vec\_f128\_ppc.h, 66  
 vec\_iszerof32  
   vec\_f32\_ppc.h, 84  
 vec\_iszerof64  
   vec\_f64\_ppc.h, 97  
 vec\_maxsd  
   vec\_int64\_ppc.h, 195  
 vec\_maxud  
   vec\_int64\_ppc.h, 195  
 vec\_minsd  
   vec\_int64\_ppc.h, 196  
 vec\_minud  
   vec\_int64\_ppc.h, 196  
 vec\_mrgahb  
   vec\_char\_ppc.h, 41  
 vec\_mrgahd  
   vec\_int64\_ppc.h, 196  
 vec\_mrgahh  
   vec\_int16\_ppc.h, 141  
 vec\_mrgahw  
   vec\_int32\_ppc.h, 155  
 vec\_mrgalb  
   vec\_char\_ppc.h, 41  
 vec\_mrgald  
   vec\_int64\_ppc.h, 197  
 vec\_mrgalh  
   vec\_int16\_ppc.h, 141  
 vec\_mrgalw  
   vec\_int32\_ppc.h, 156  
 vec\_mrgeb  
   vec\_char\_ppc.h, 42  
 vec\_mrged  
   vec\_int64\_ppc.h, 197  
 vec\_mrggeh  
   vec\_int16\_ppc.h, 142  
 vec\_mrgew  
   vec\_int32\_ppc.h, 156  
 vec\_mrghd  
   vec\_int64\_ppc.h, 198  
 vec\_mrgld  
   vec\_int64\_ppc.h, 198  
 vec\_mrgob  
   vec\_char\_ppc.h, 42  
 vec\_mrgod  
   vec\_int64\_ppc.h, 199  
 vec\_mrgoh  
   vec\_int16\_ppc.h, 142  
 vec\_mrgow  
   vec\_int32\_ppc.h, 157  
 vec\_msumudm  
   vec\_int128\_ppc.h, 118  
 vec\_mul10cuq  
   vec\_int128\_ppc.h, 118  
 vec\_mul10ecuq  
   vec\_int128\_ppc.h, 119  
 vec\_mul10euq  
   vec\_int128\_ppc.h, 119  
 vec\_mul10uq  
   vec\_int128\_ppc.h, 120  
 vec\_mulesw  
   vec\_int32\_ppc.h, 157  
 vec\_muleud  
   vec\_int128\_ppc.h, 120  
 vec\_muleuw  
   vec\_int32\_ppc.h, 158  
 vec\_mulhsb  
   vec\_char\_ppc.h, 43  
 vec\_mulhsh  
   vec\_int16\_ppc.h, 143  
 vec\_mulhsw  
   vec\_int32\_ppc.h, 158  
 vec\_mulhub  
   vec\_char\_ppc.h, 43  
 vec\_mulhud  
   vec\_int128\_ppc.h, 121  
 vec\_mulhuh  
   vec\_int16\_ppc.h, 143  
 vec\_mulhuw  
   vec\_int32\_ppc.h, 159  
 vec\_mulluq  
   vec\_int128\_ppc.h, 121  
 vec\_mulosw  
   vec\_int32\_ppc.h, 159  
 vec\_muloud  
   vec\_int128\_ppc.h, 122  
 vec\_mulouw  
   vec\_int32\_ppc.h, 160  
 vec\_mulubm  
   vec\_char\_ppc.h, 44  
 vec\_muludm

- vec\_int128\_ppc.h, [122](#)
- vec\_muludq
  - vec\_int128\_ppc.h, [123](#)
- vec\_muluhm
  - vec\_int16\_ppc.h, [144](#)
- vec\_muluwm
  - vec\_int32\_ppc.h, [160](#)
- vec\_pack\_longdouble
  - vec\_f64\_ppc.h, [98](#)
- vec\_pasted
  - vec\_int64\_ppc.h, [199](#)
- vec\_permdi
  - vec\_int64\_ppc.h, [199](#)
- vec\_popcntb
  - vec\_char\_ppc.h, [44](#)
- vec\_popcntd
  - vec\_int64\_ppc.h, [200](#)
- vec\_popcnth
  - vec\_int16\_ppc.h, [144](#)
- vec\_popcntq
  - vec\_int128\_ppc.h, [123](#)
- vec\_popcntw
  - vec\_int32\_ppc.h, [160](#)
- vec\_revbd
  - vec\_int64\_ppc.h, [201](#)
- vec\_revbh
  - vec\_int16\_ppc.h, [145](#)
- vec\_revbq
  - vec\_int128\_ppc.h, [124](#)
- vec\_revbw
  - vec\_int32\_ppc.h, [161](#)
- vec\_rldi
  - vec\_int64\_ppc.h, [201](#)
- vec\_setb\_cyq
  - vec\_int128\_ppc.h, [124](#)
- vec\_setb\_ncq
  - vec\_int128\_ppc.h, [124](#)
- vec\_setb\_qp
  - vec\_f128\_ppc.h, [67](#)
- vec\_setb\_sq
  - vec\_int128\_ppc.h, [125](#)
- vec\_shift\_leftdo
  - vec\_char\_ppc.h, [46](#)
- vec\_signbitf128
  - vec\_f128\_ppc.h, [67](#)
- vec\_slbi
  - vec\_char\_ppc.h, [46](#)
- vec\_sldi
  - vec\_int64\_ppc.h, [201](#)
- vec\_sldq
  - vec\_int128\_ppc.h, [125](#)
- vec\_slhi
  - vec\_int16\_ppc.h, [145](#)
- vec\_slq
  - vec\_int128\_ppc.h, [126](#)
- vec\_slq4
  - vec\_int128\_ppc.h, [126](#)
- vec\_slq5
  - vec\_int128\_ppc.h, [126](#)
- vec\_slqi
  - vec\_int128\_ppc.h, [127](#)
- vec\_slwi
  - vec\_int32\_ppc.h, [161](#)
- vec\_splatd
  - vec\_int64\_ppc.h, [202](#)
- vec\_spltd
  - vec\_int64\_ppc.h, [202](#)
- vec\_srabi
  - vec\_char\_ppc.h, [47](#)
- vec\_sradi
  - vec\_int64\_ppc.h, [203](#)
- vec\_srahi
  - vec\_int16\_ppc.h, [146](#)
- vec\_srawi
  - vec\_int32\_ppc.h, [162](#)
- vec\_srbi
  - vec\_char\_ppc.h, [47](#)
- vec\_srdi
  - vec\_int64\_ppc.h, [203](#)
- vec\_srhi
  - vec\_int16\_ppc.h, [146](#)
- vec\_srq
  - vec\_int128\_ppc.h, [127](#)
- vec\_srq4
  - vec\_int128\_ppc.h, [127](#)
- vec\_srq5
  - vec\_int128\_ppc.h, [128](#)
- vec\_srqi
  - vec\_int128\_ppc.h, [128](#)
- vec\_srwi
  - vec\_int32\_ppc.h, [162](#)
- vec\_subcuq
  - vec\_int128\_ppc.h, [129](#)
- vec\_subecuq
  - vec\_int128\_ppc.h, [129](#)
- vec\_subeuqm
  - vec\_int128\_ppc.h, [129](#)
- vec\_subudm
  - vec\_int64\_ppc.h, [204](#)
- vec\_subuqm
  - vec\_int128\_ppc.h, [131](#)
- vec\_swapd
  - vec\_int64\_ppc.h, [204](#)
- vec\_tolower
  - vec\_char\_ppc.h, [48](#)
- vec\_toupper
  - vec\_char\_ppc.h, [48](#)
- vec\_unpack\_longdouble
  - vec\_f64\_ppc.h, [98](#)
- vec\_vmrgeb
  - vec\_char\_ppc.h, [48](#)
- vec\_vmrgeh
  - vec\_int16\_ppc.h, [147](#)
- vec\_vmrgeb
  - vec\_char\_ppc.h, [49](#)
- vec\_vmrgeh

- vec\_int16\_ppc.h, [147](#)
- vec\_vmsumuwmm
  - vec\_int64\_ppc.h, [205](#)
- vec\_vmuleud
  - vec\_int128\_ppc.h, [131](#)
- vec\_vmuloud
  - vec\_int128\_ppc.h, [132](#)
- vec\_vpkudum
  - vec\_int64\_ppc.h, [205](#)
- vec\_vrld
  - vec\_int64\_ppc.h, [206](#)
- vec\_vslld
  - vec\_int64\_ppc.h, [206](#)
- vec\_vsradd
  - vec\_int64\_ppc.h, [207](#)
- vec\_vsrdd
  - vec\_int64\_ppc.h, [207](#)
- vec\_xfer\_bin128\_2\_vui128t
  - vec\_f128\_ppc.h, [67](#)
- vec\_xfer\_bin128\_2\_vui16t
  - vec\_f128\_ppc.h, [68](#)
- vec\_xfer\_bin128\_2\_vui32t
  - vec\_f128\_ppc.h, [68](#)
- vec\_xfer\_bin128\_2\_vui64t
  - vec\_f128\_ppc.h, [69](#)
- vec\_xfer\_bin128\_2\_vui8t
  - vec\_f128\_ppc.h, [69](#)
- vec\_xfer\_vui128t\_2\_bin128
  - vec\_f128\_ppc.h, [69](#)
- vec\_xfer\_vui16t\_2\_bin128
  - vec\_f128\_ppc.h, [70](#)
- vec\_xfer\_vui32t\_2\_bin128
  - vec\_f128\_ppc.h, [70](#)
- vec\_xfer\_vui64t\_2\_bin128
  - vec\_f128\_ppc.h, [71](#)
- vec\_xfer\_vui8t\_2\_bin128
  - vec\_f128\_ppc.h, [71](#)
- vec\_xxspltd
  - vec\_int64\_ppc.h, [208](#)